

PURDUE CS47100

INTRODUCTION TO AI

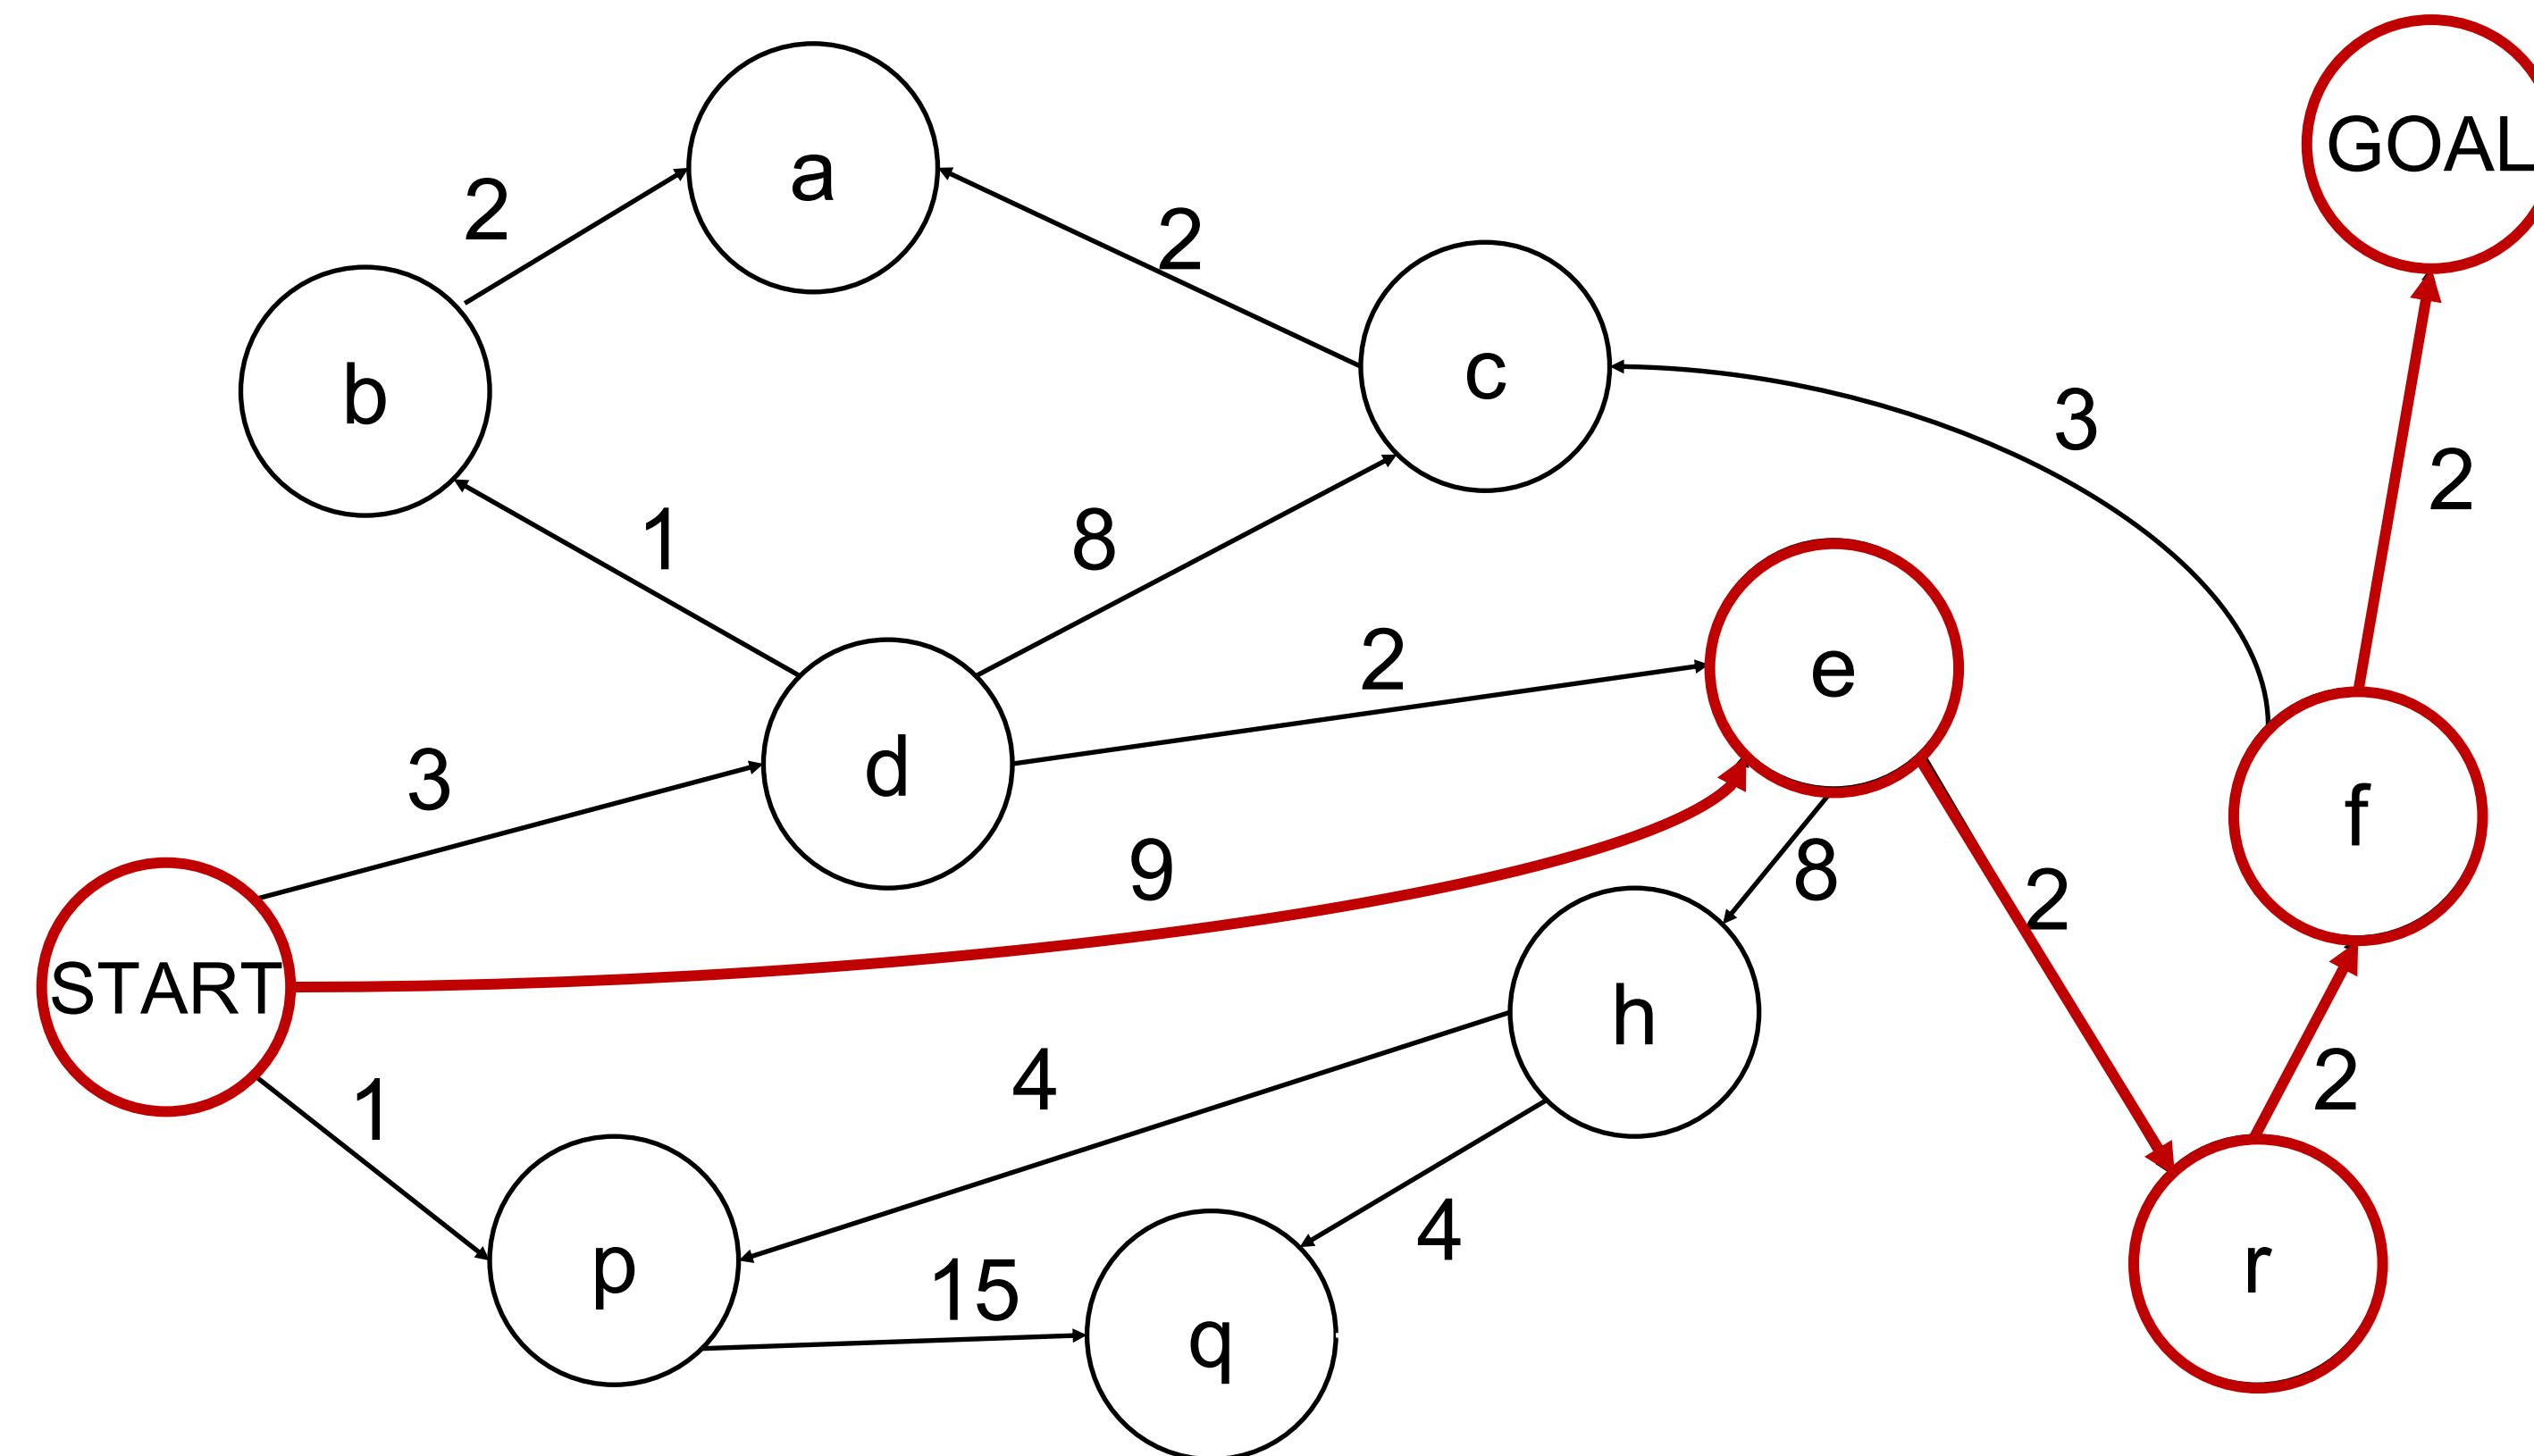
ANNOUNCEMENTS

- ▶ Assignment 1 is released on Brightspace
- ▶ Submit to GradeScope!
- ▶ Due date: 11:59pm, September 23 (Friday)
- ▶ 100 points + 28 bonus points
 - ▶ It can contributes up to $128/100*15\% = 19.2\%$ of your final grades
 - ▶ If you'd like to use extension days on this assignment, specify it clearly in your written pdf submission!

RECAP: UNINFORMED SEARCH STRATEGIES

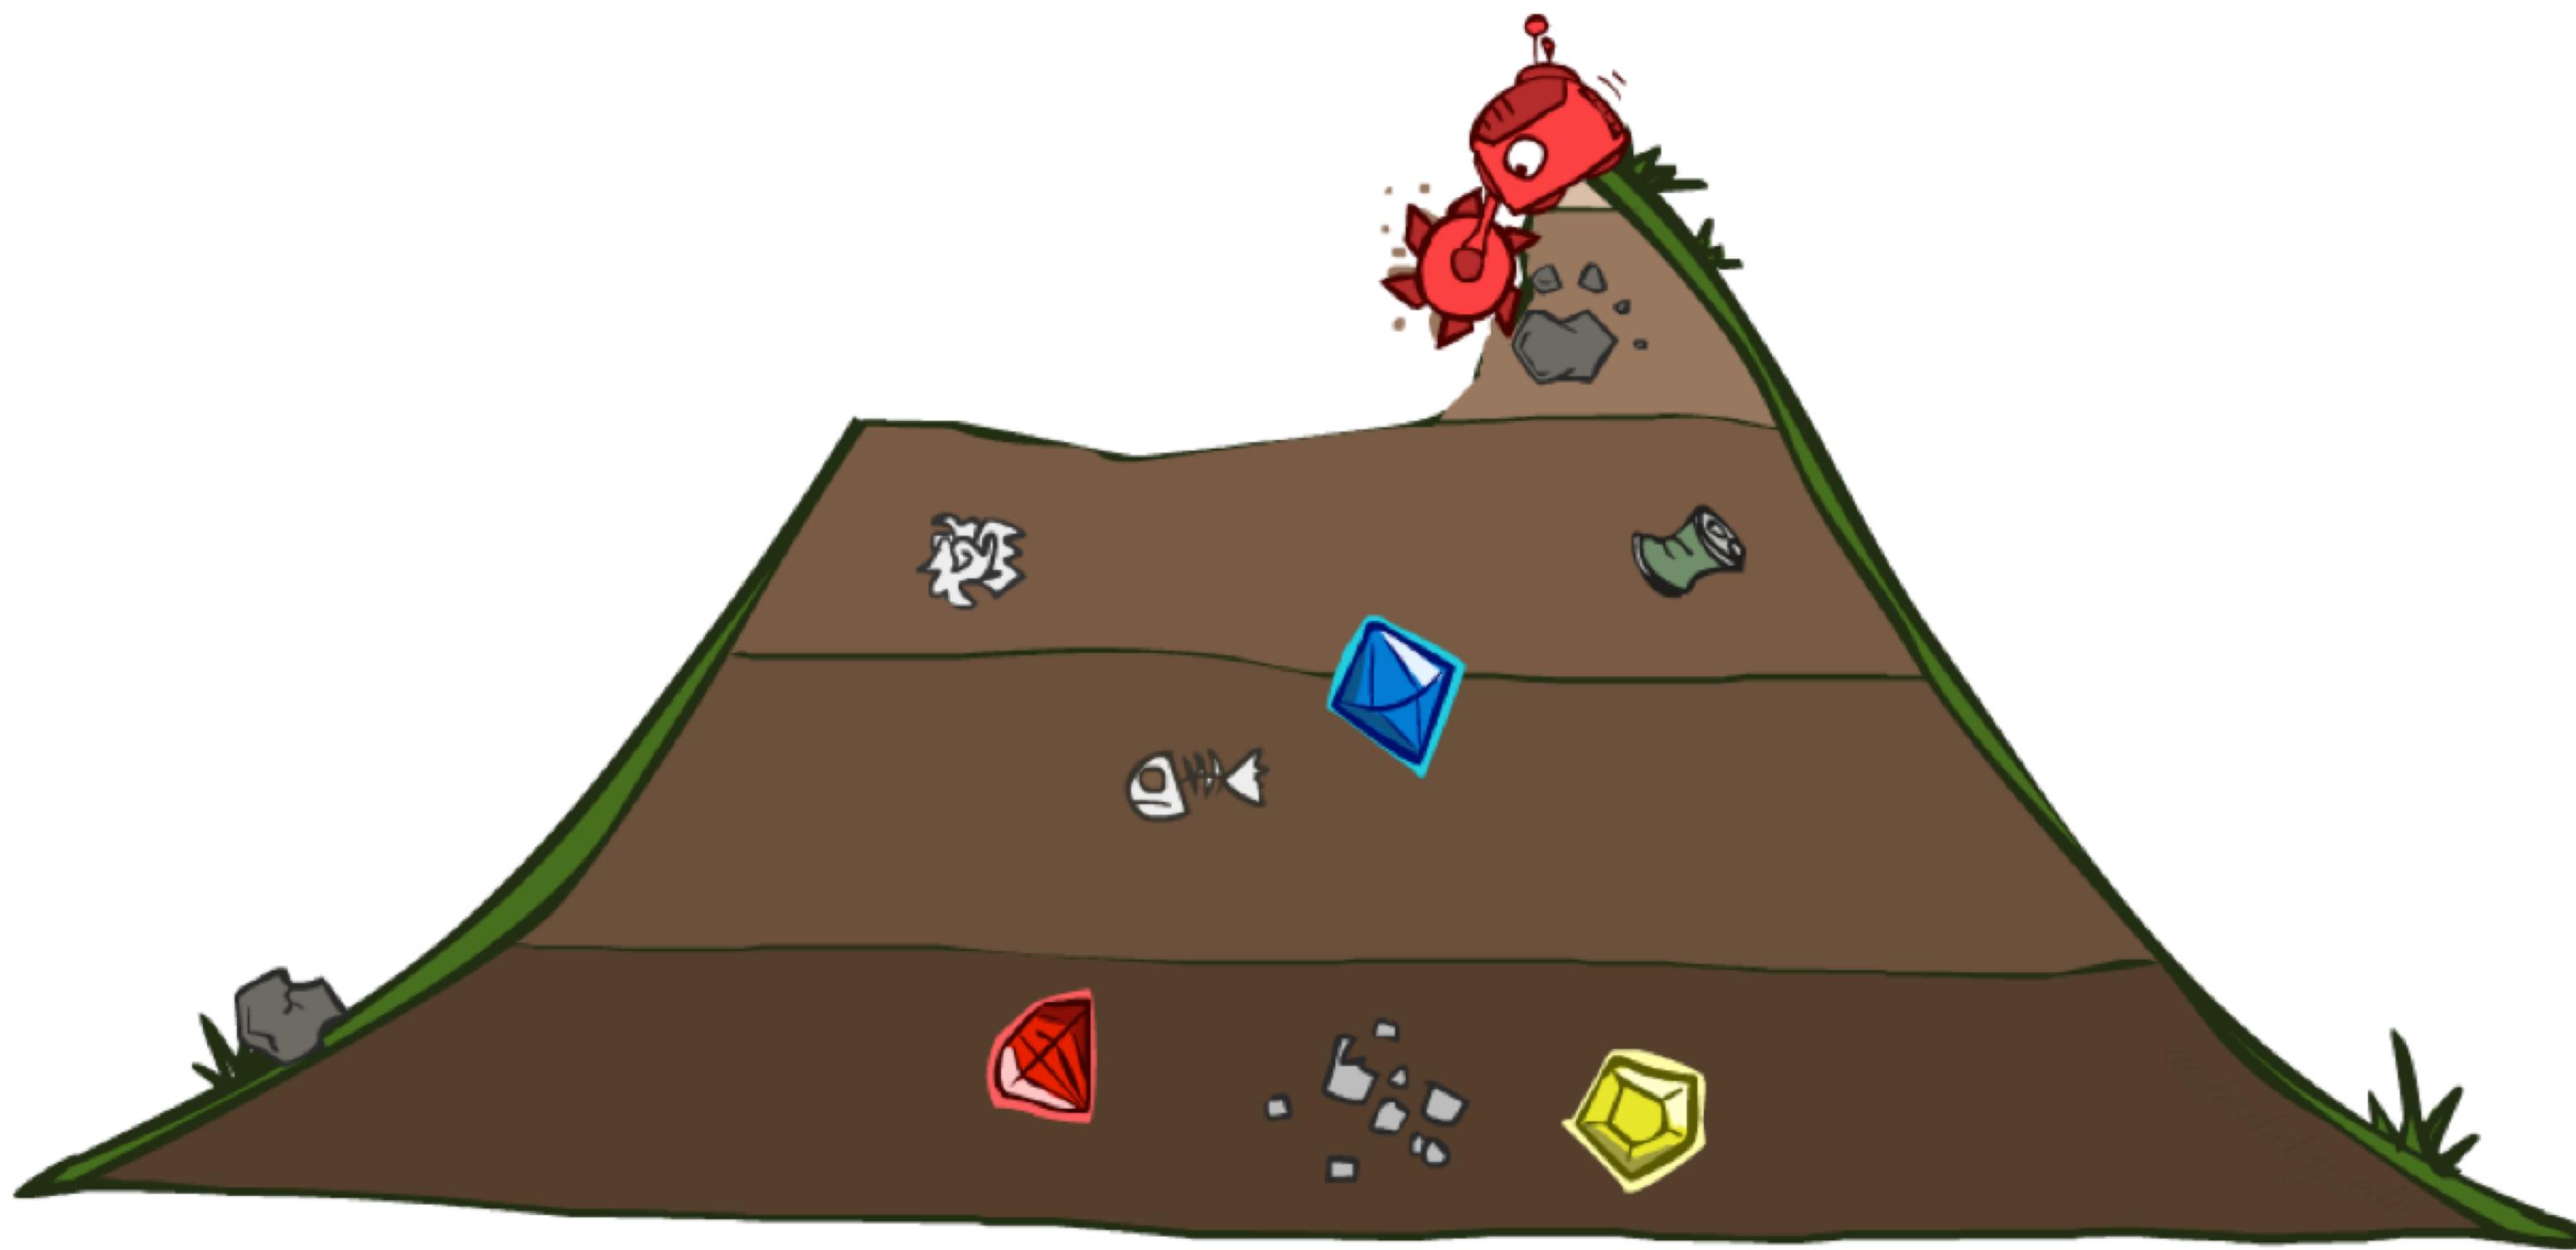
- ▶ The search algorithm has no additional information about states beyond what has been provided in the problem formulation
 - ▶ Depth-first search
 - ▶ Breadth-first search
 - ▶ Depth-limited search
 - ▶ Iterative deepening depth-first search
 - ▶ Bi-directional search

COST-SENSITIVE SEARCH



BFS finds the shortest path in terms of number of actions.
It does not find the least-cost path.

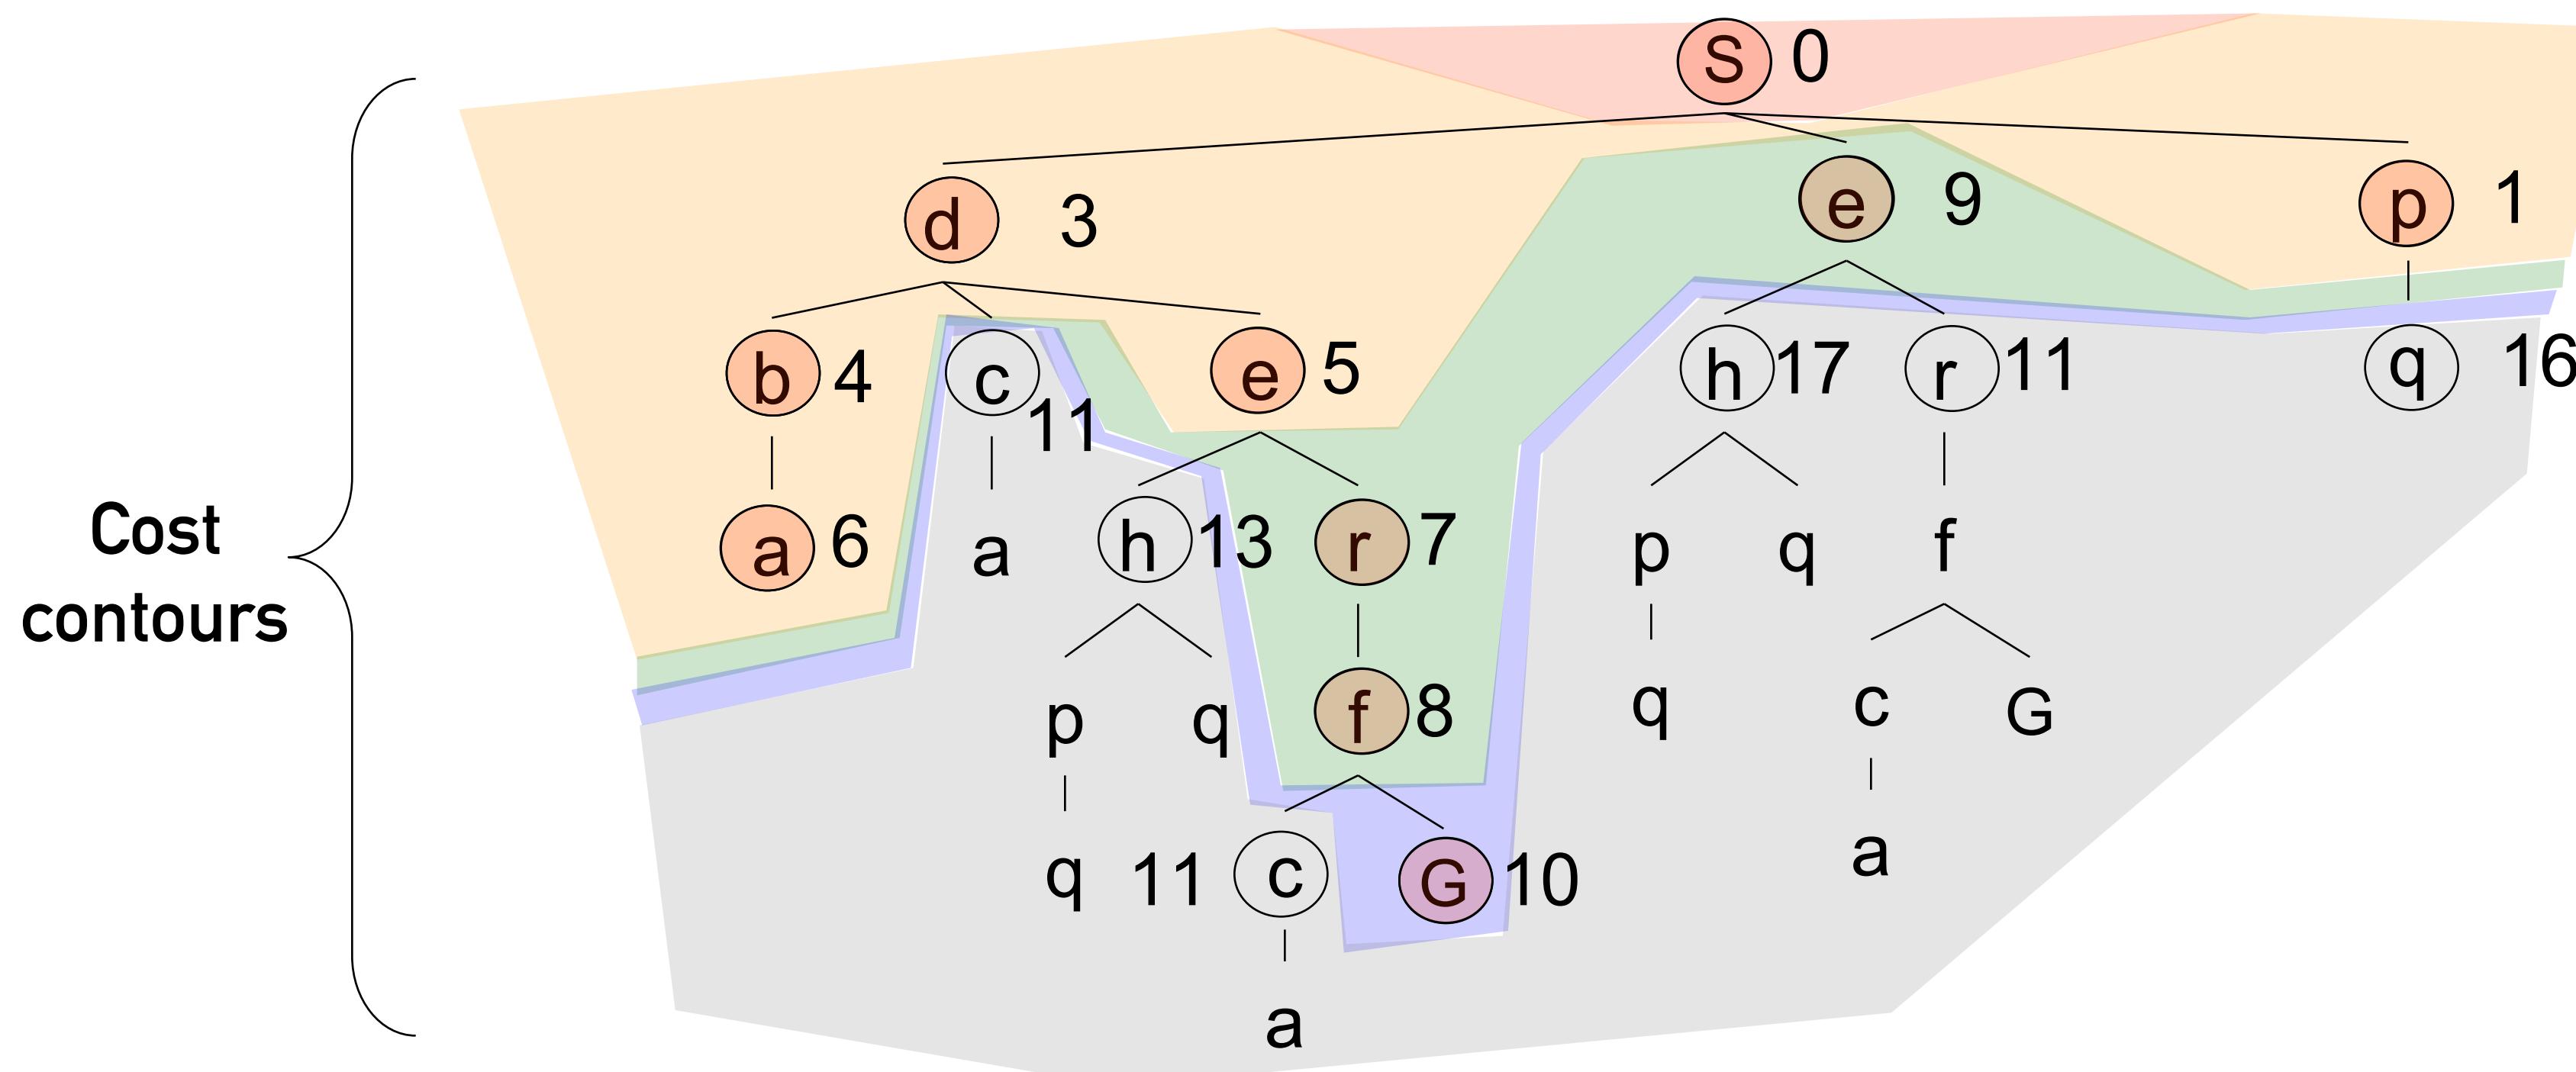
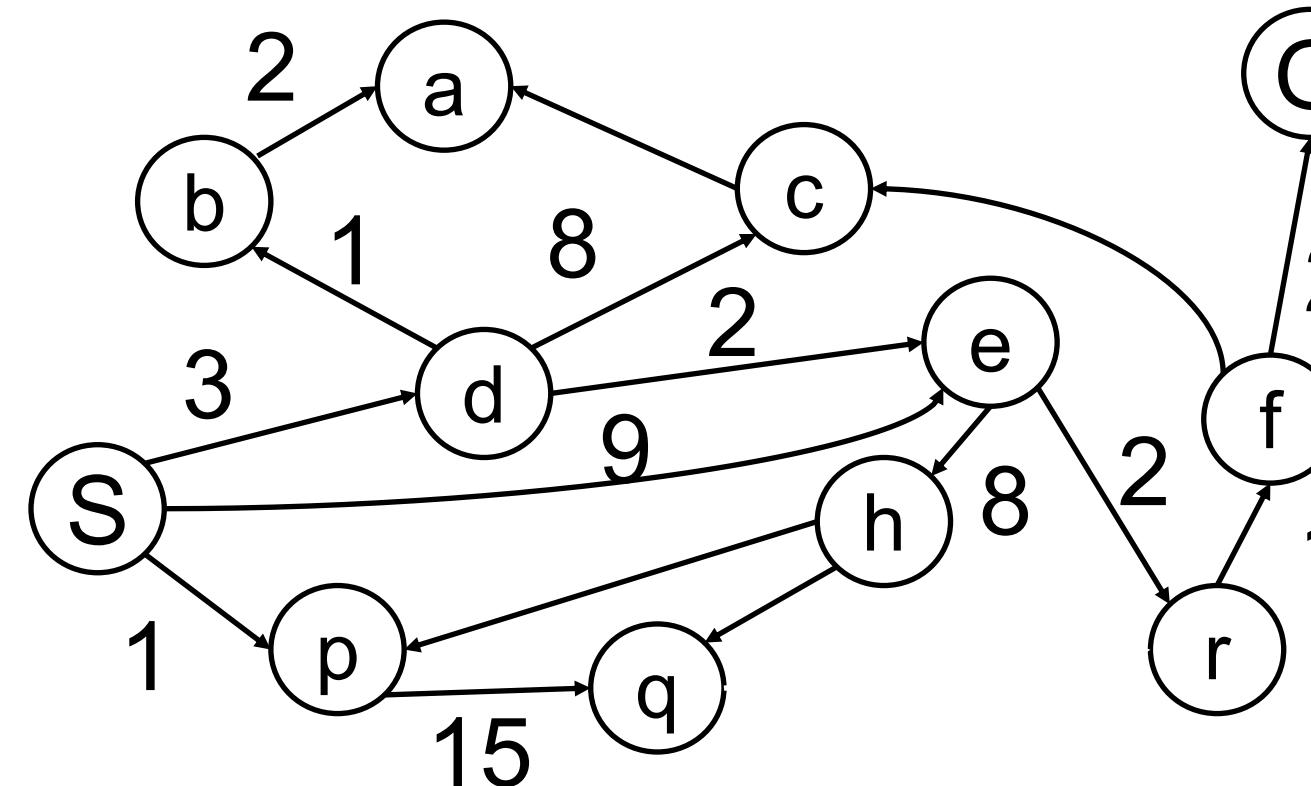
UNIFORM COST SEARCH



UNIFORM COST SEARCH

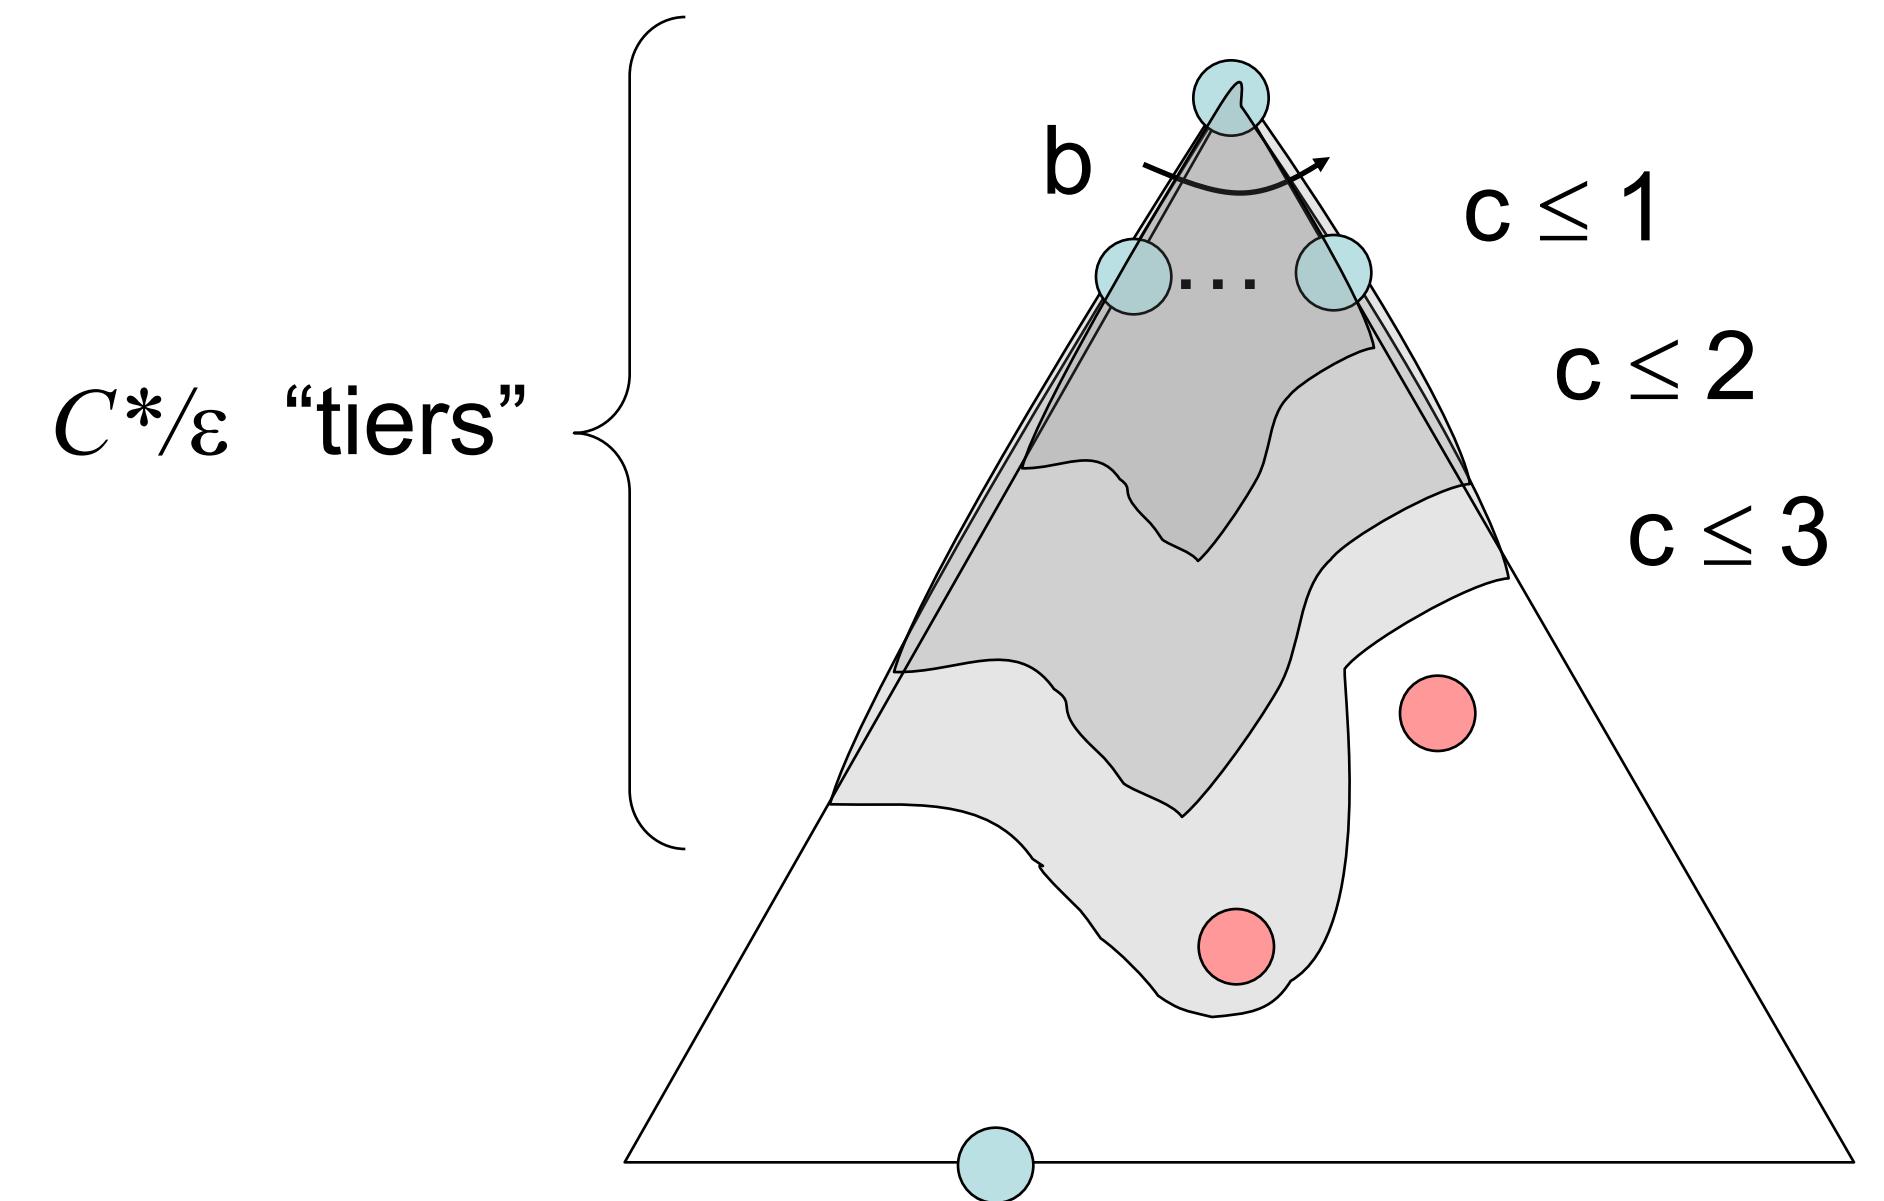
Strategy: expand cheapest node first:

Fringe is a priority queue
(priority: cumulative cost)



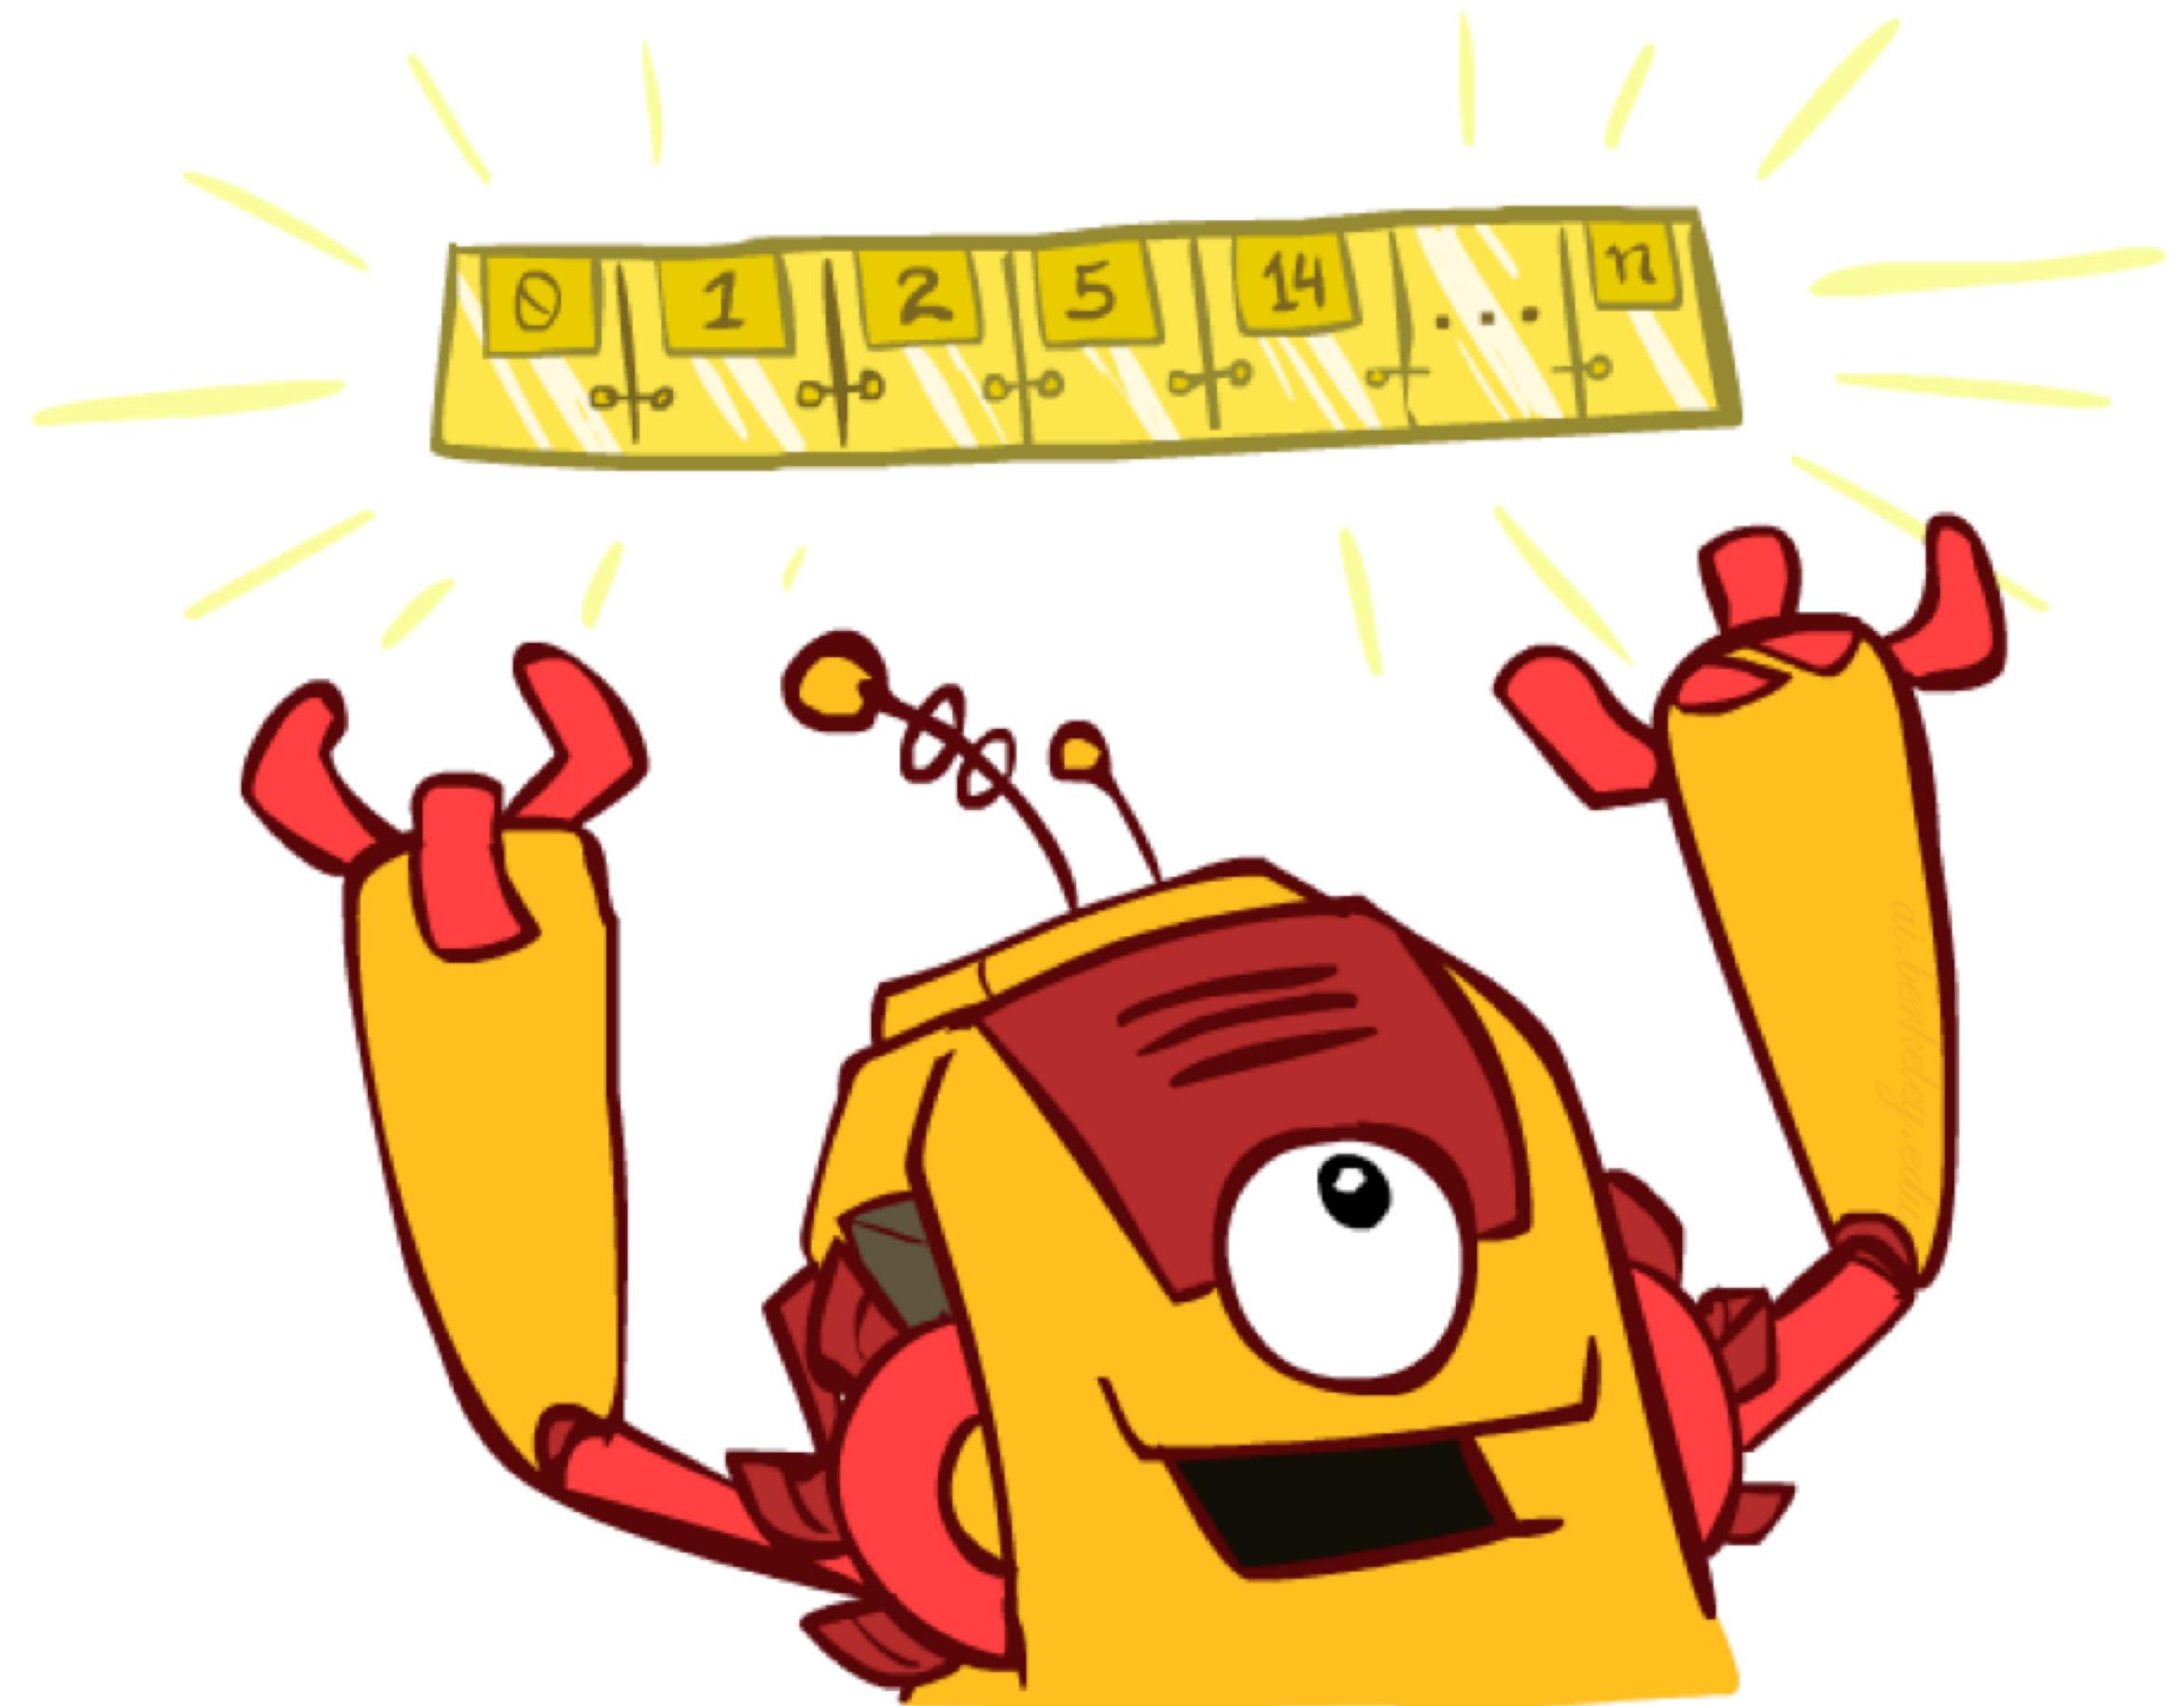
UNIFORM COST SEARCH (UCS) PROPERTIES

- ▶ Time complexity
 - ▶ Processes all nodes with cost less than cheapest solution
 - ▶ If that solution costs C^* and actions cost at least ε , then the “effective depth” is roughly C^*/ε
 - ▶ Takes time $O(b^{1+C^*/\varepsilon})$
- ▶ Space complexity
 - ▶ Has roughly the successors of the last tier, so $O(b^{1+C^*/\varepsilon})$
- ▶ Is it complete?
 - ▶ Yes, assuming best solution has finite cost and min action cost is positive
- ▶ Is it optimal?
 - ▶ Yes



ONE SEARCH QUEUE

- ▶ All these search algorithms are the same except for fringe strategies
 - ▶ Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
 - ▶ Practically, for DFS and BFS, you can avoid the $\log(n)$ overhead from an actual priority queue, by using stacks and queues
 - ▶ Can even code one implementation that takes a variable queuing object



SUMMARY – UNINFORMED SEARCH STRATEGIES

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^{d+1})$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^{d+1})$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

Figure 3.21 Evaluation of search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

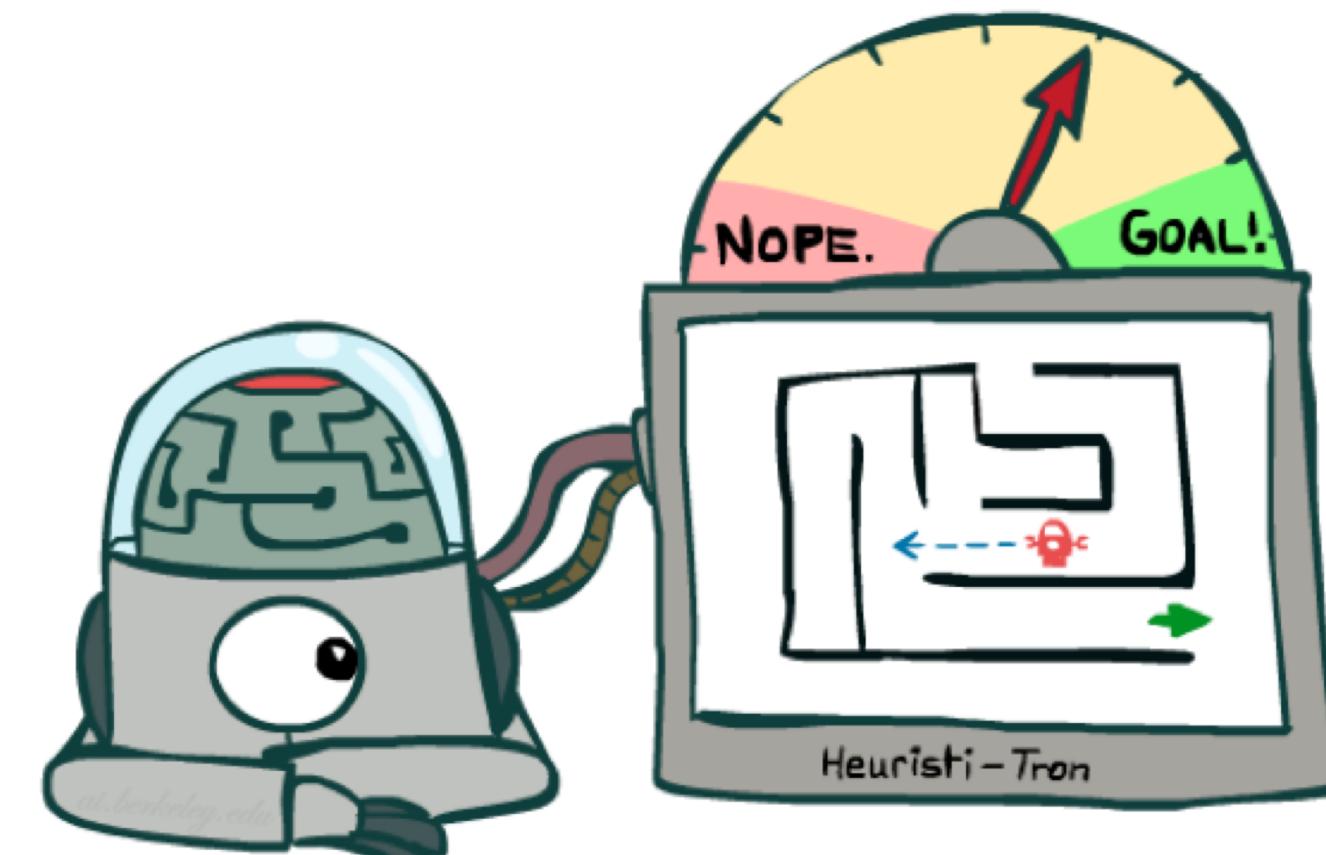
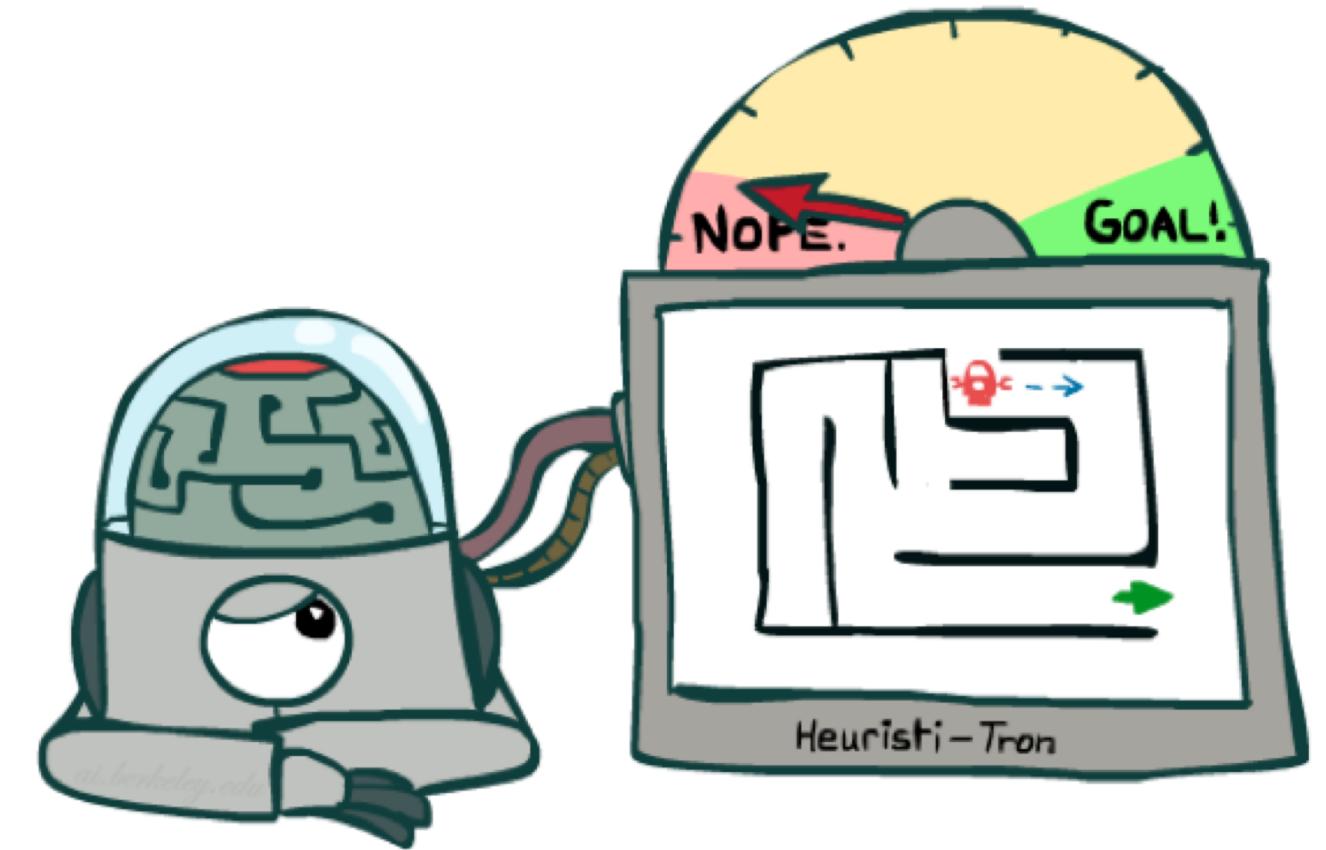
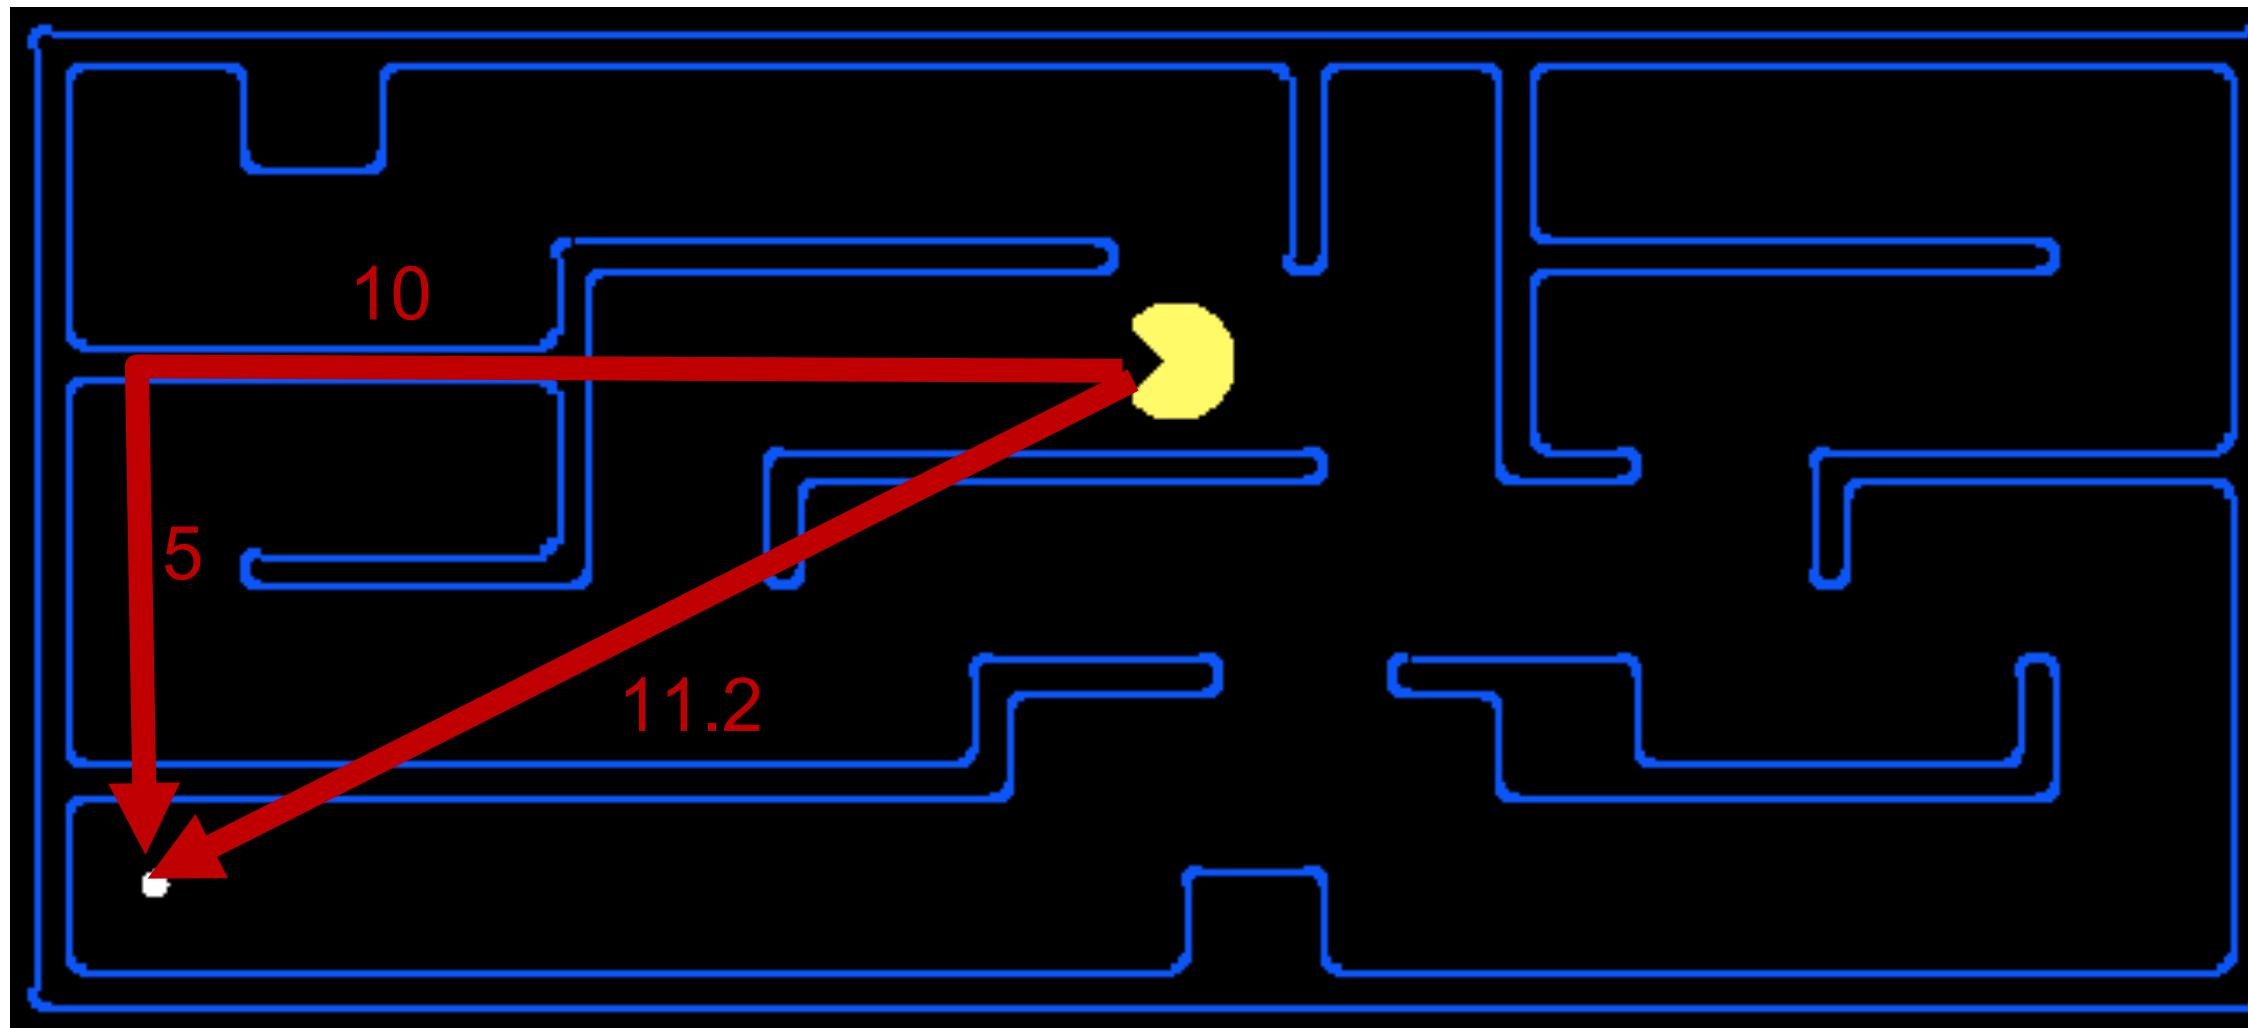
INFORMED SEARCH

- ▶ Determine how to prioritize the expansion of different states based on information beyond what has been provided in the problem formulation
- ▶ **Key idea:** Uses an evaluation function $f(n)$, which estimate the “cost” of node n
 - ▶ $f(n)$ usually contains a heuristic function $h(n)$, which is the estimated cost of the cheapest path from the state at node n to a goal state.
 - ▶ **Best-first search:** Select node with minimum $f(n)$ for expansion

HEURISTIC FUNCTION

- ▶ A **heuristic function** is:

 - ▶ A **non-negative** function that estimates how close a state is to a goal, designed for a particular search problem
 - ▶ $h(n) = 0$ if n is a goal state
 - ▶ Examples: Manhattan distance, Euclidean distance for pathing

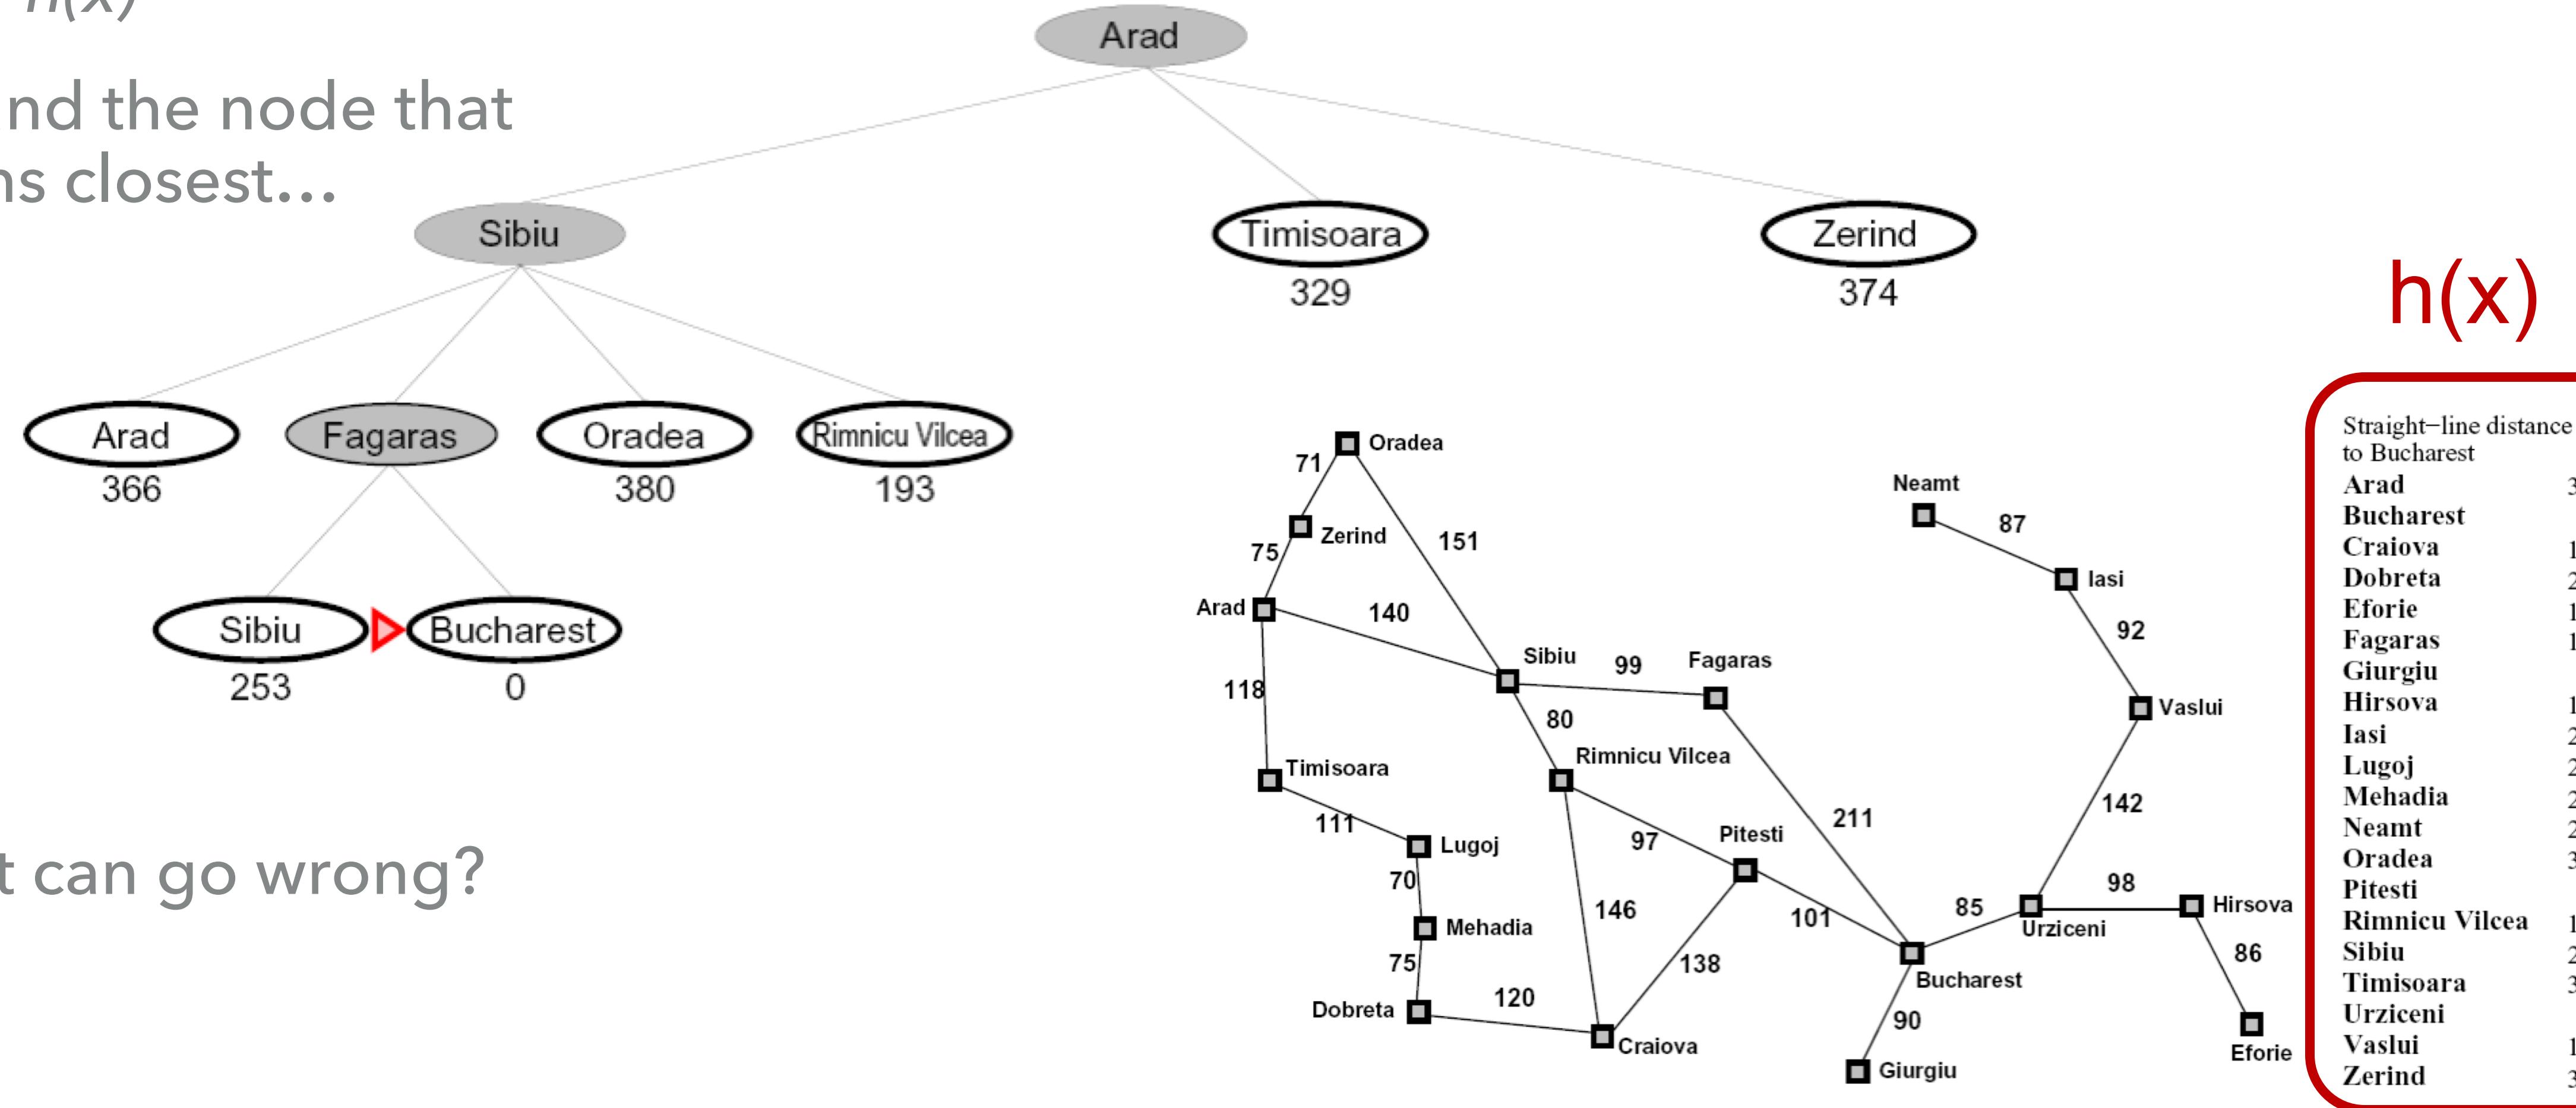


GREEDY SEARCH



GREEDY SEARCH

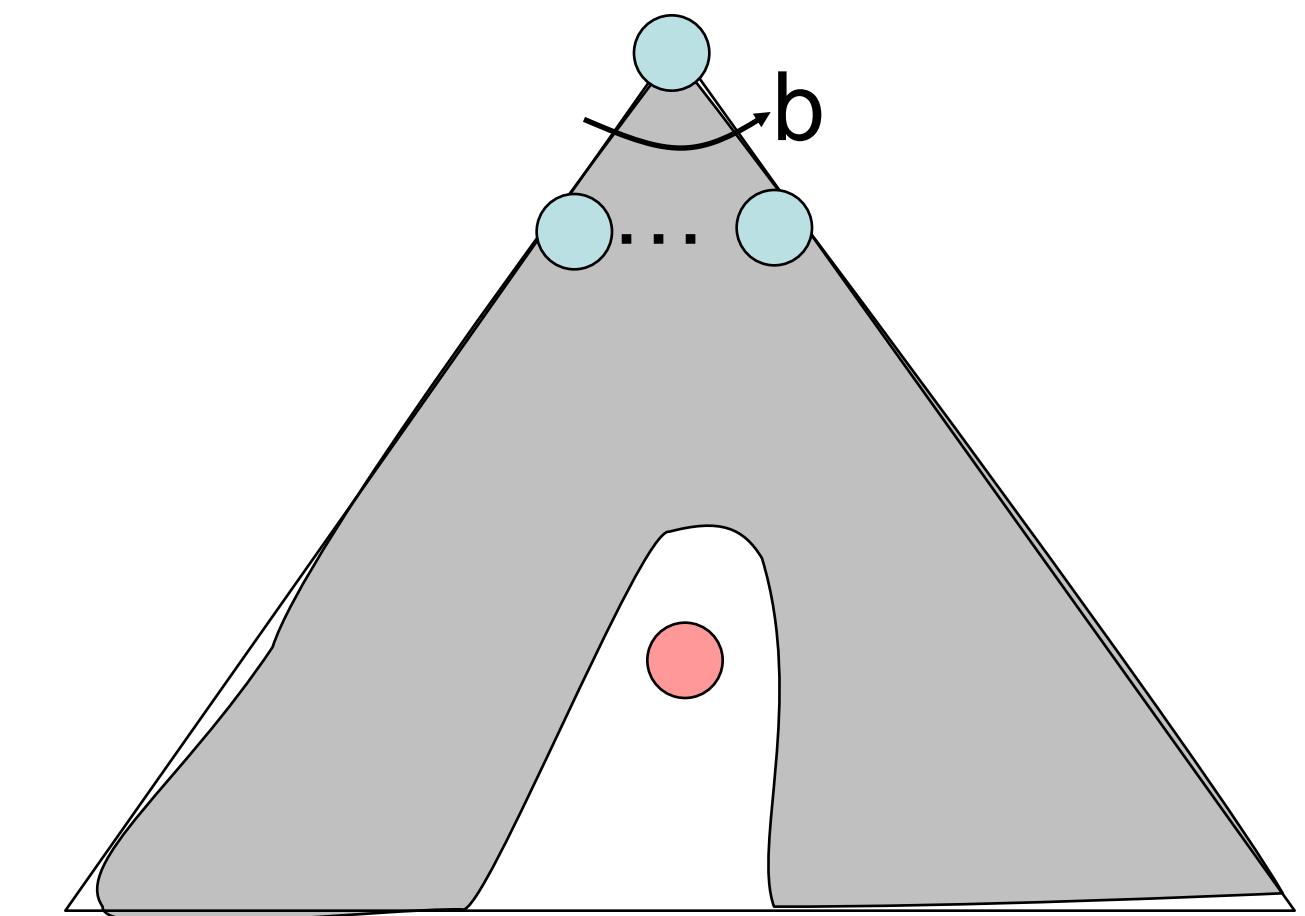
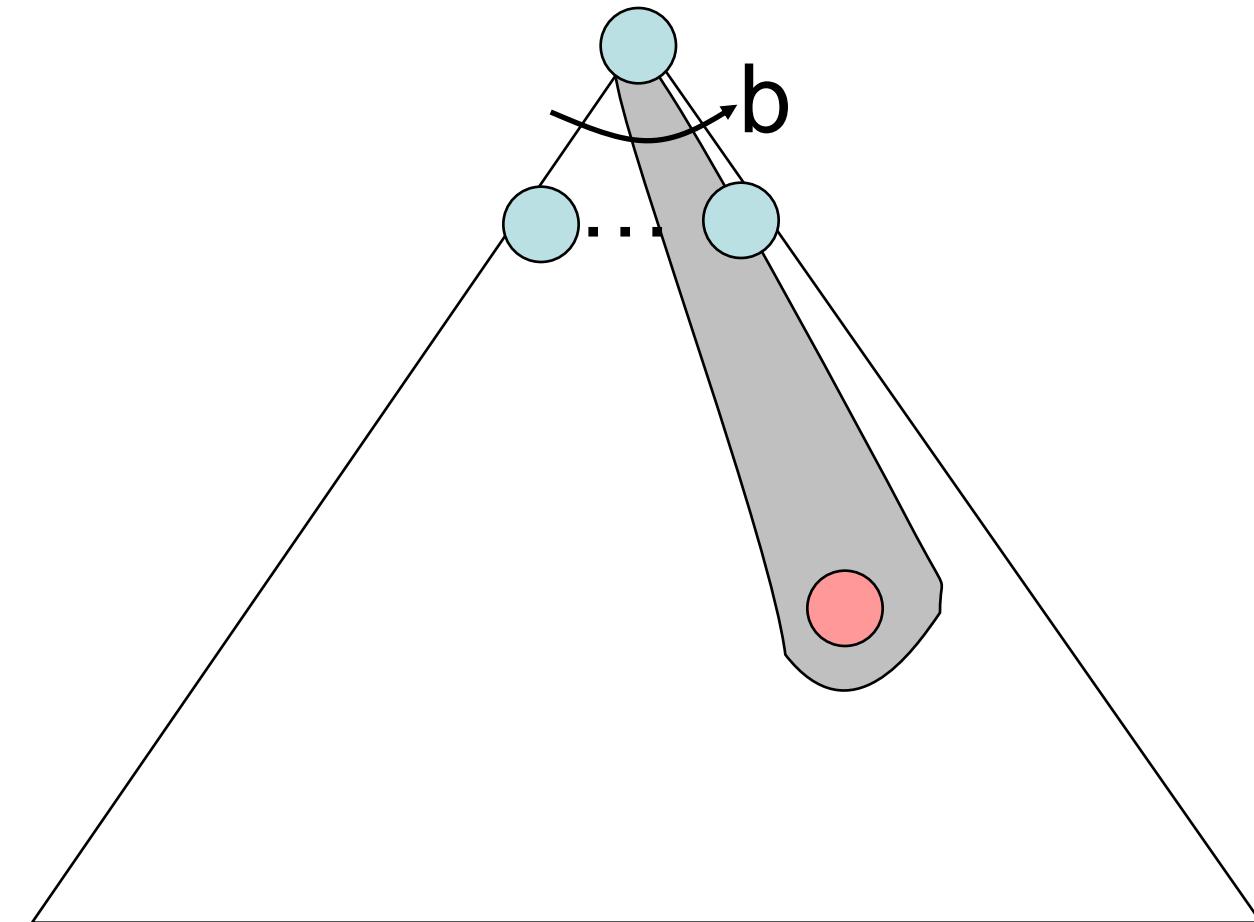
- ▶ $f(x) = h(x)$
- ▶ Expand the node that seems closest...



- ▶ What can go wrong?

GREEDY SEARCH

- ▶ Strategy: expand a node that you think is closest to a goal state
- ▶ Heuristic: estimate of distance to nearest goal for each state
- ▶ Common case: Best-first takes you straight to the (wrong) goal
- ▶ Worst-case: like a badly-guided DFS



A* SEARCH

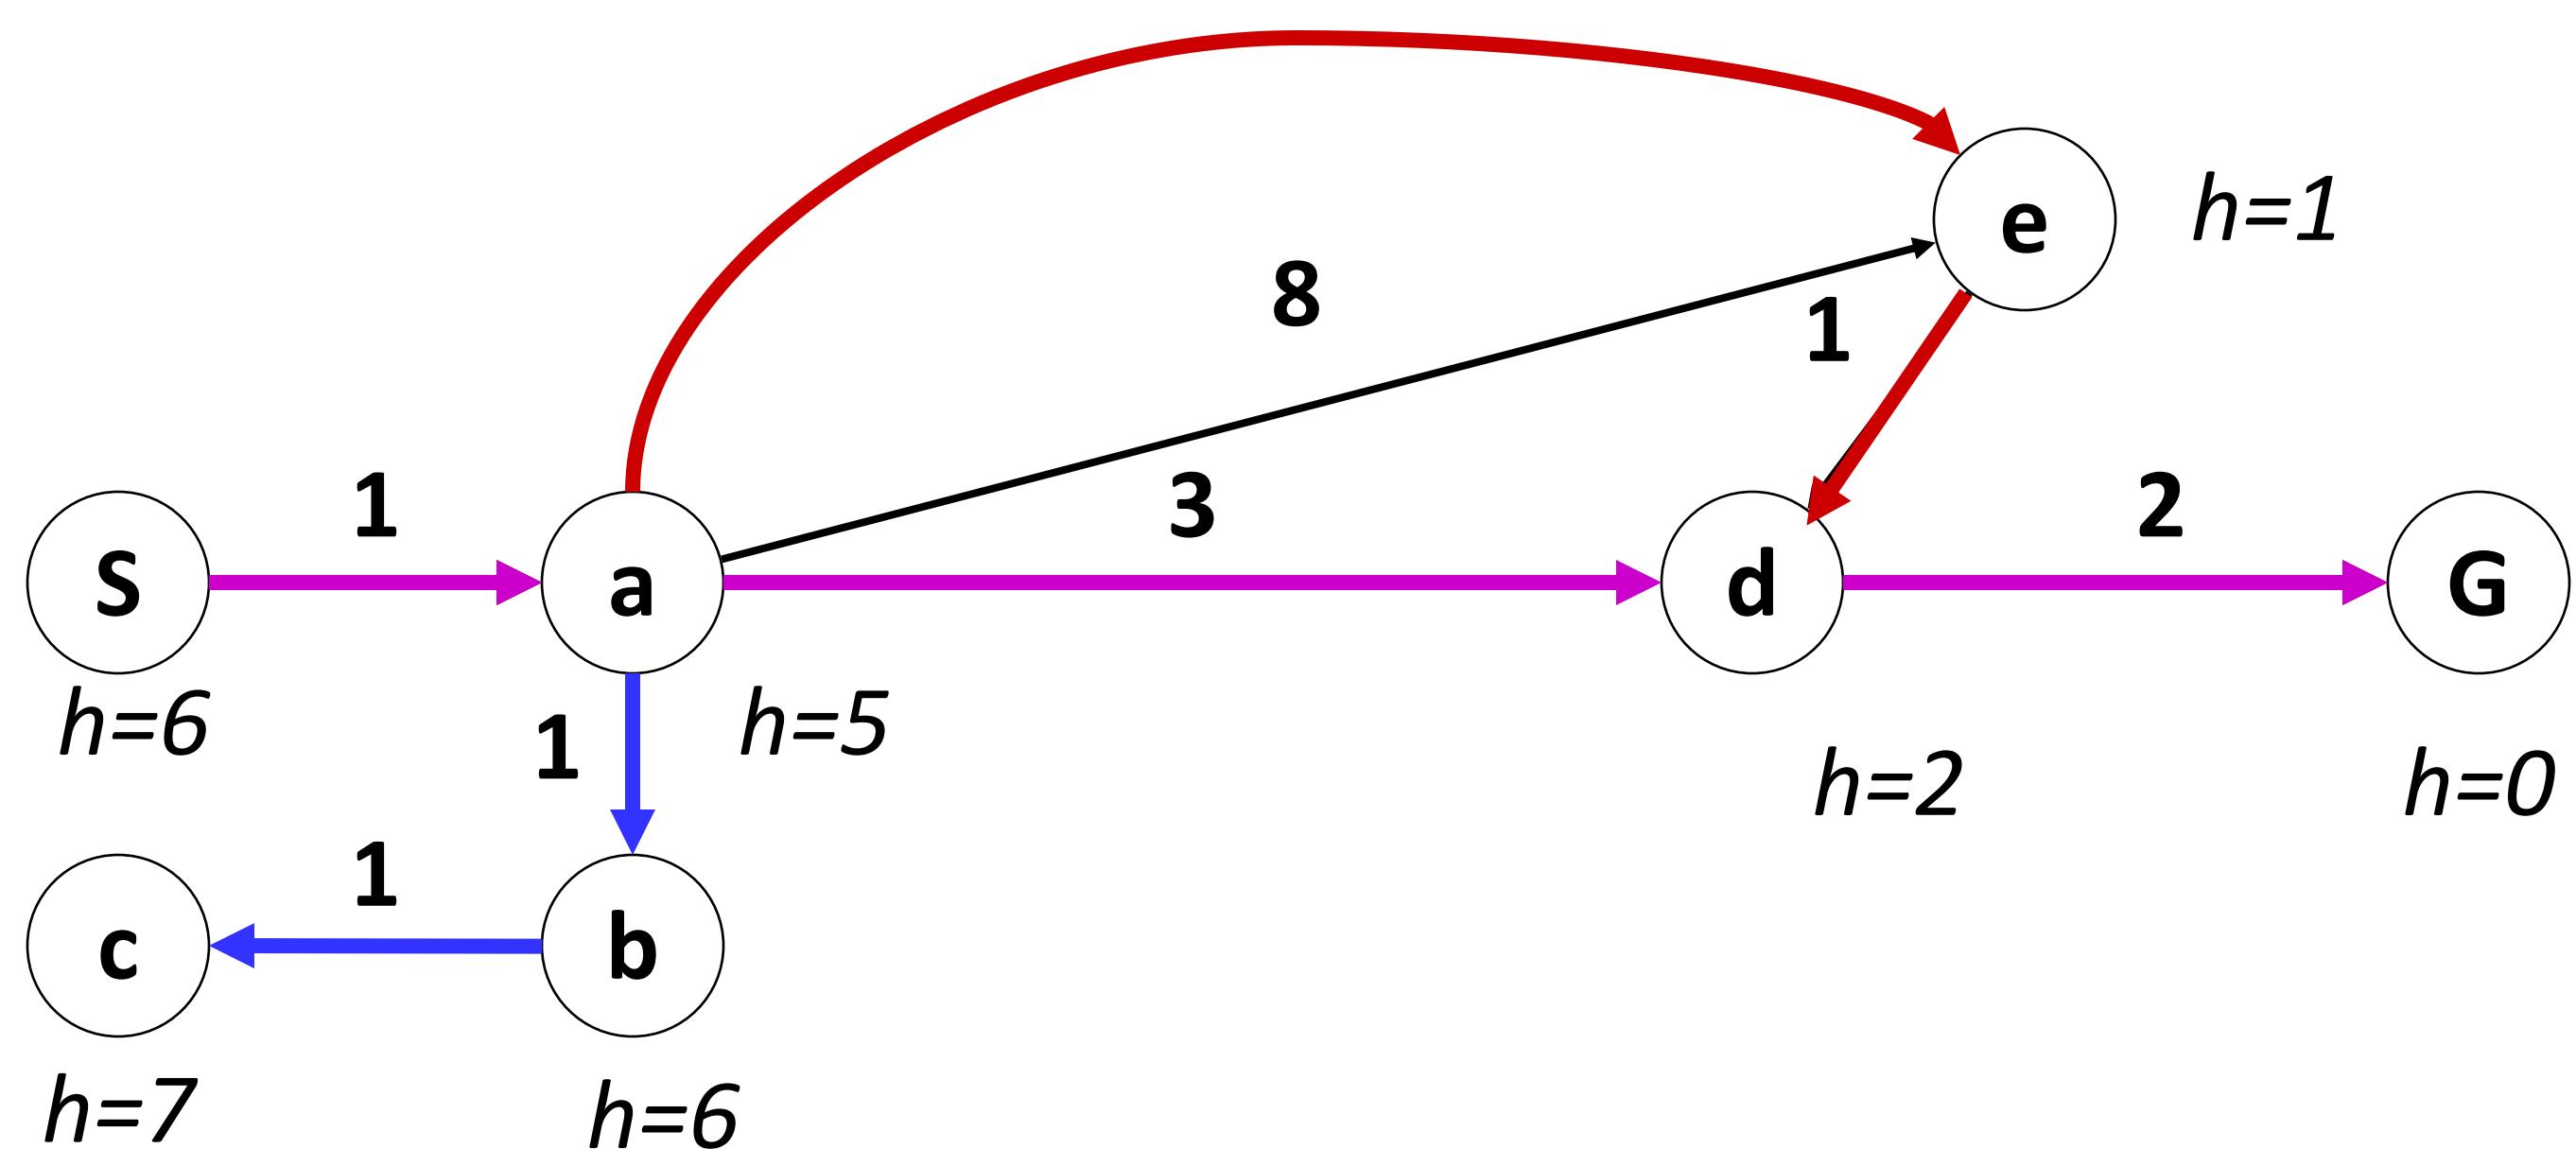


A* SEARCH

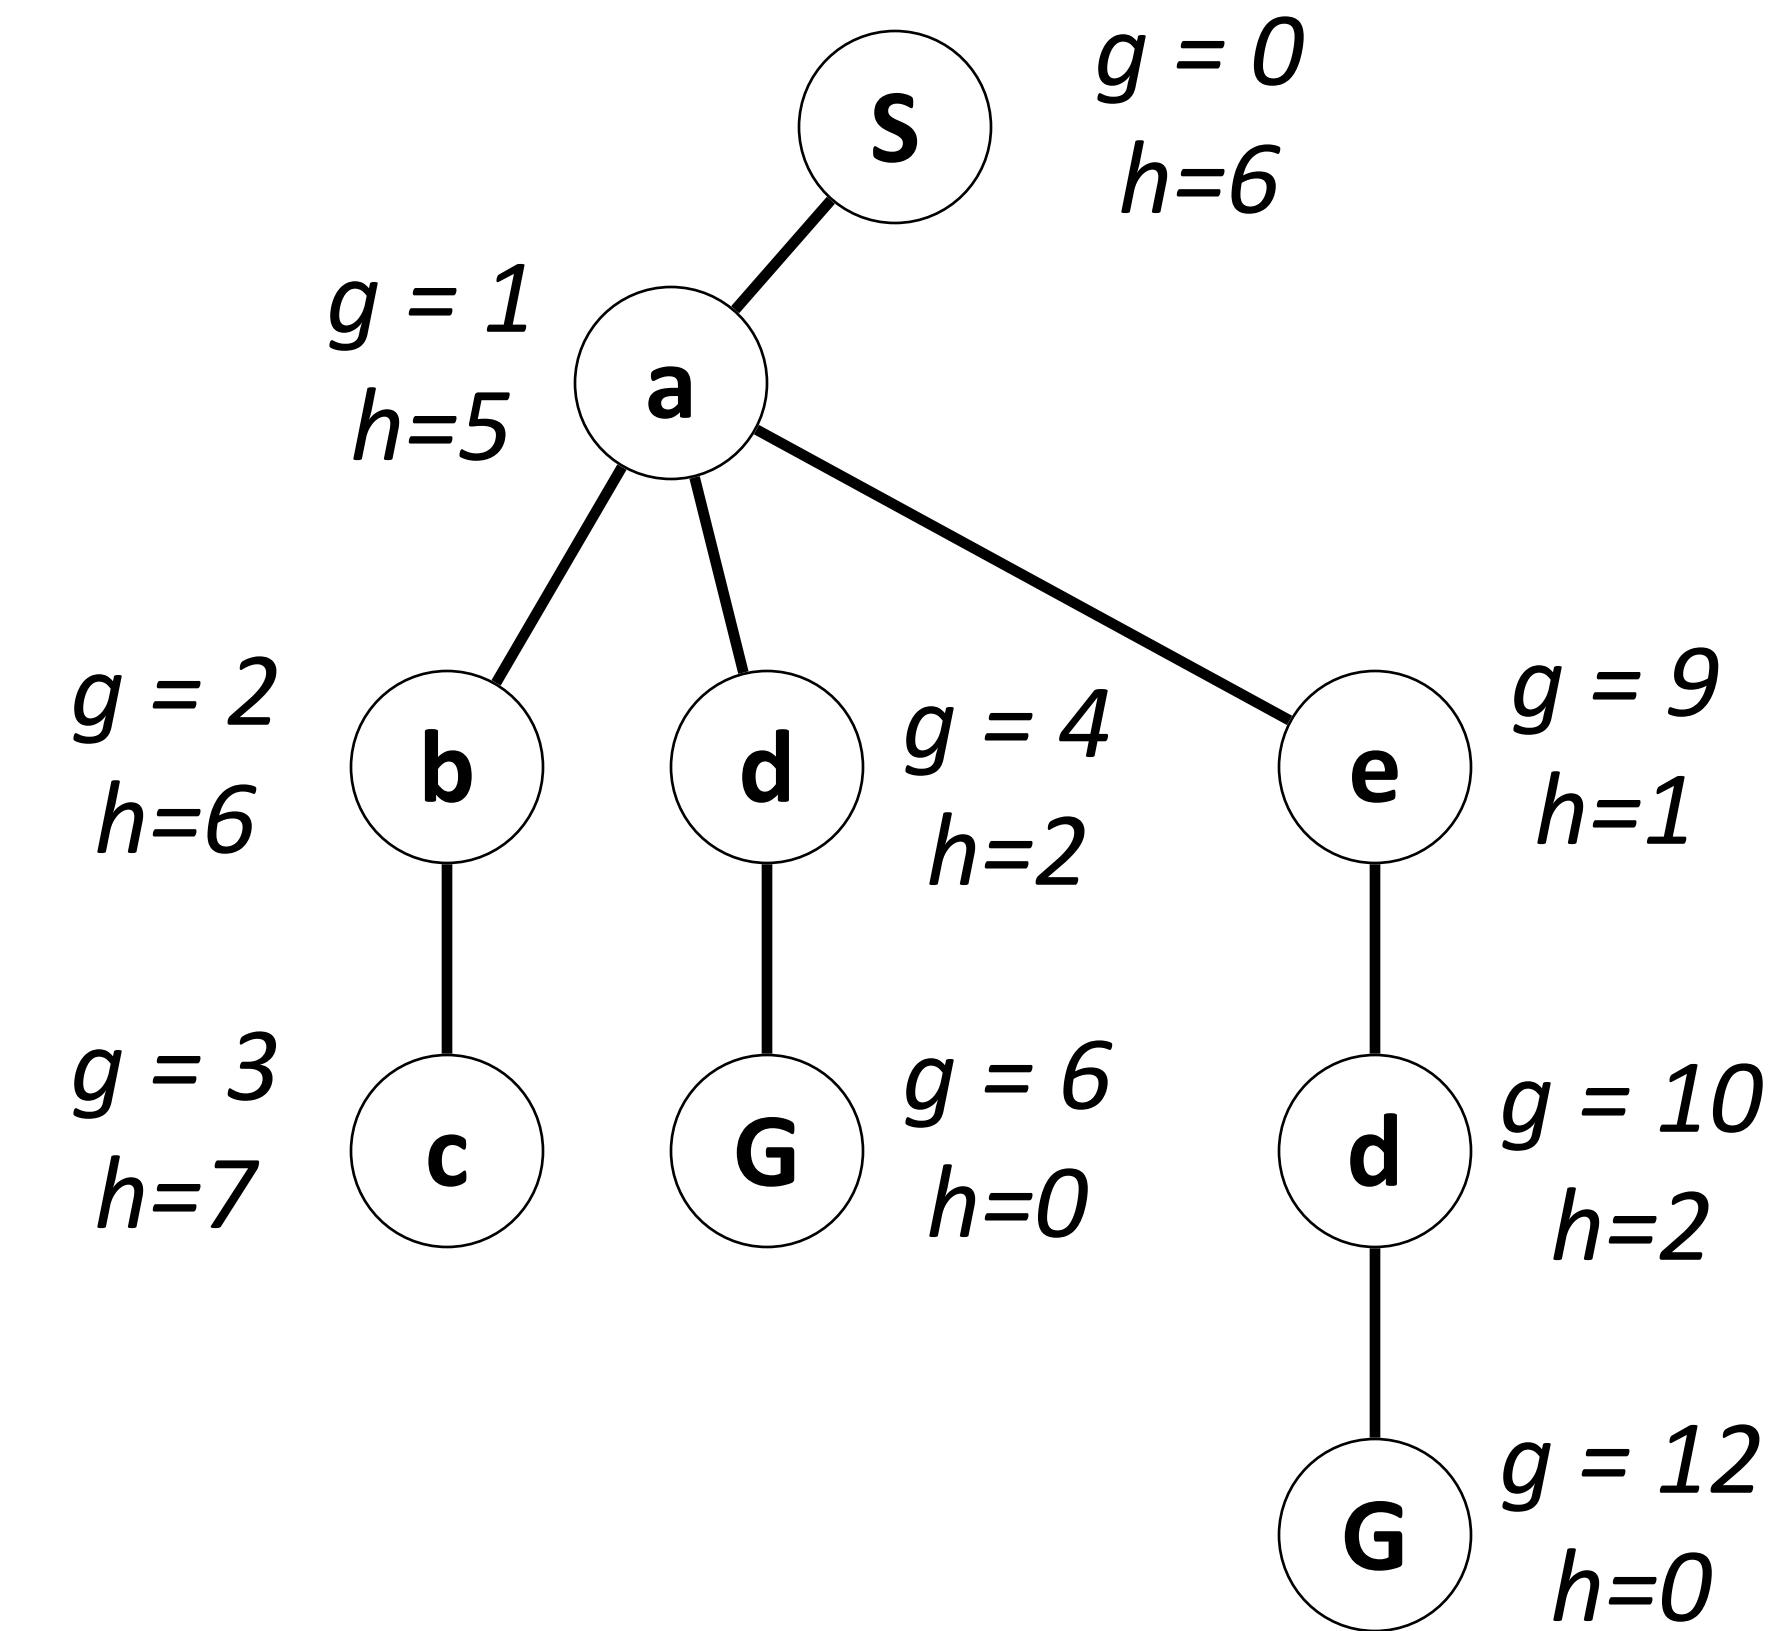
- ▶ $f(n) = g(n) + h(n)$
 - ▶ $g(n)$ = path cost from start node to current node n **Uniform cost search**
 - ▶ $h(n)$ = estimate of path cost from n to the goal **Greedy best-first search**
- ▶ Expands the node n if the estimated cost of the solution passing through it is the lowest
- ▶ Combines uniform cost search and greedy search
- ▶ A* search is optimal if $h(n)$ is an **underestimate** of the actual cost from n to the goal

COMBINING UCS AND GREEDY

- ▶ Uniform-cost orders by path cost, or *backward cost* $g(n)$
- ▶ Greedy orders by goal proximity, or *forward cost* $h(n)$



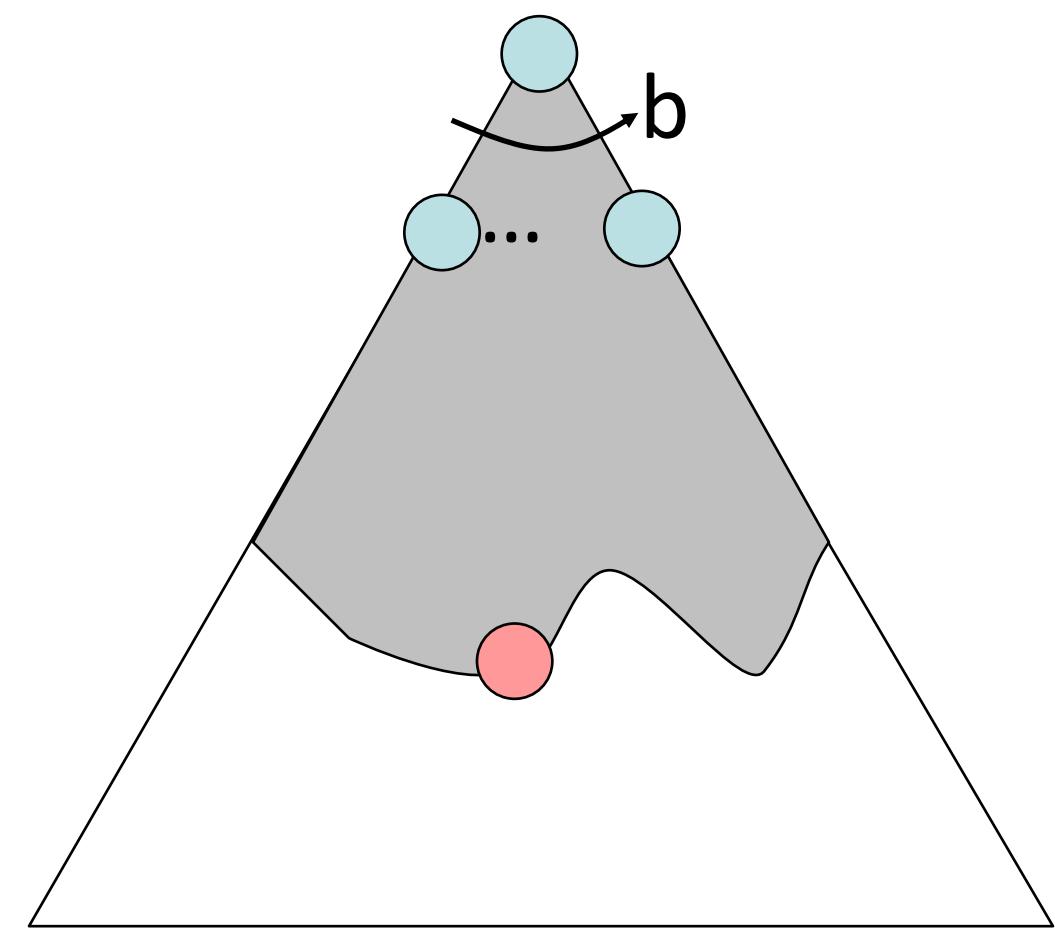
- ▶ A* Search orders by the sum: $f(n) = g(n) + h(n)$



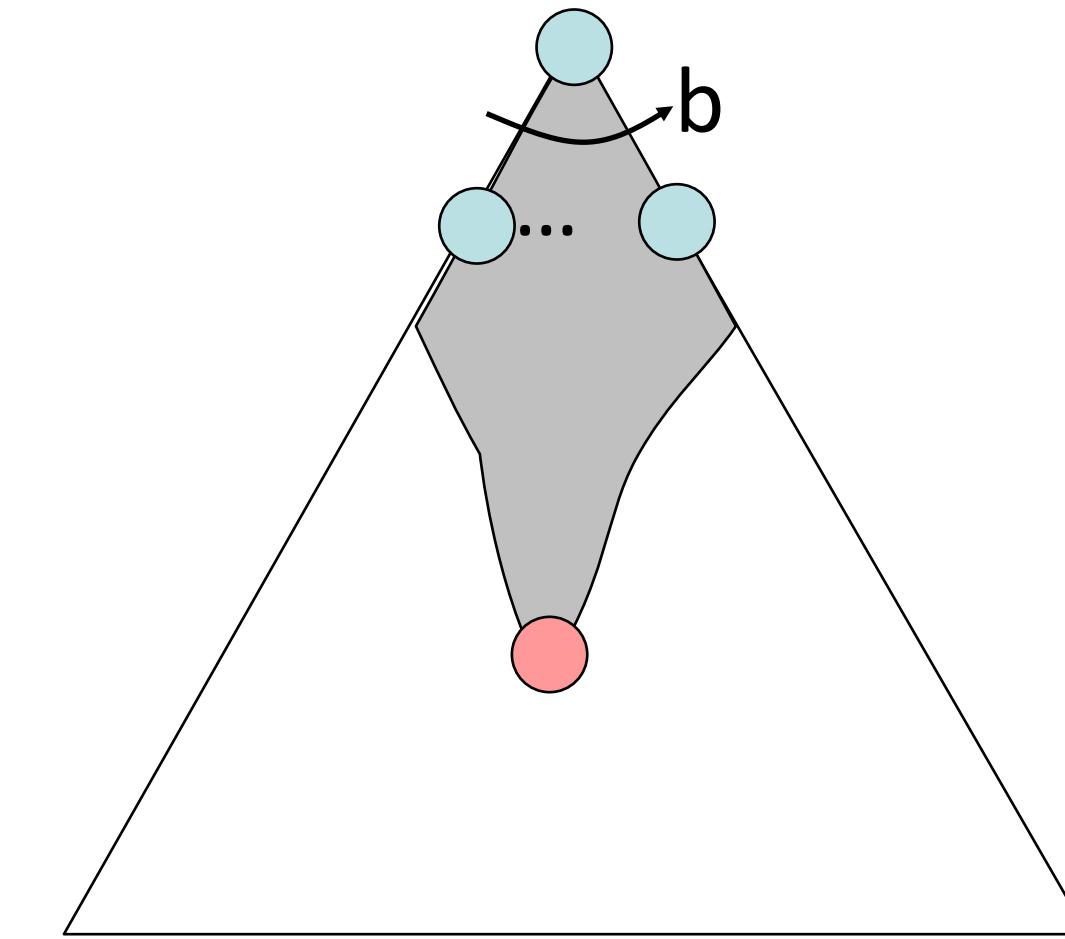
Example: Teg Grenager

UCS VS A*

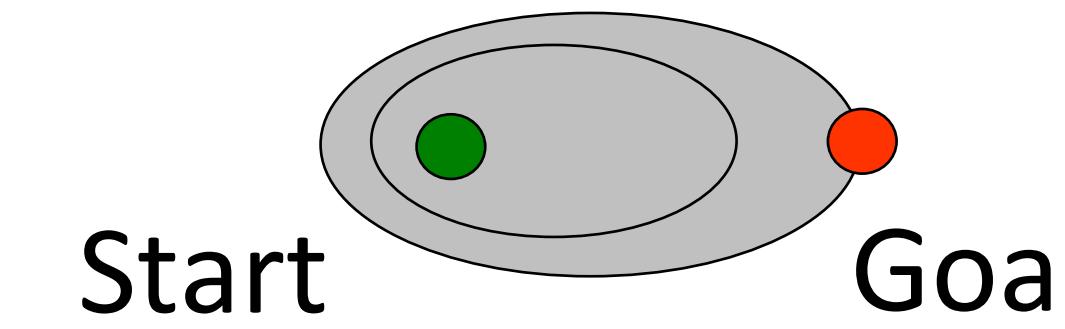
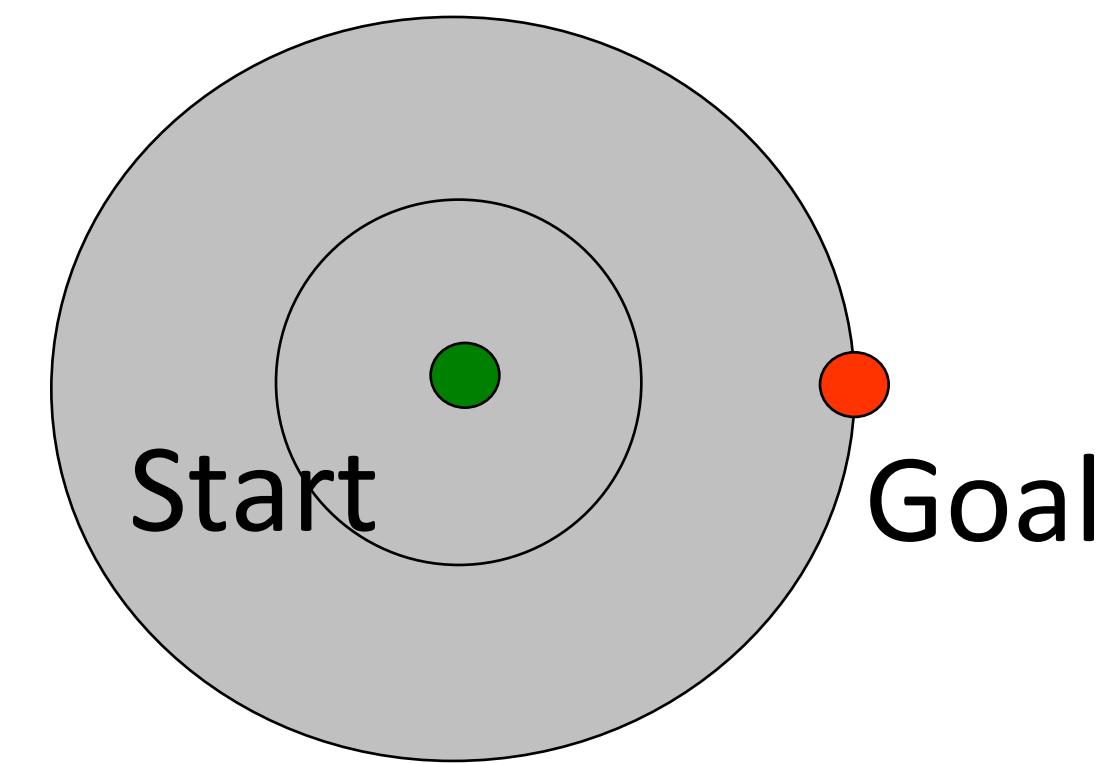
Uniform-Cost



A*



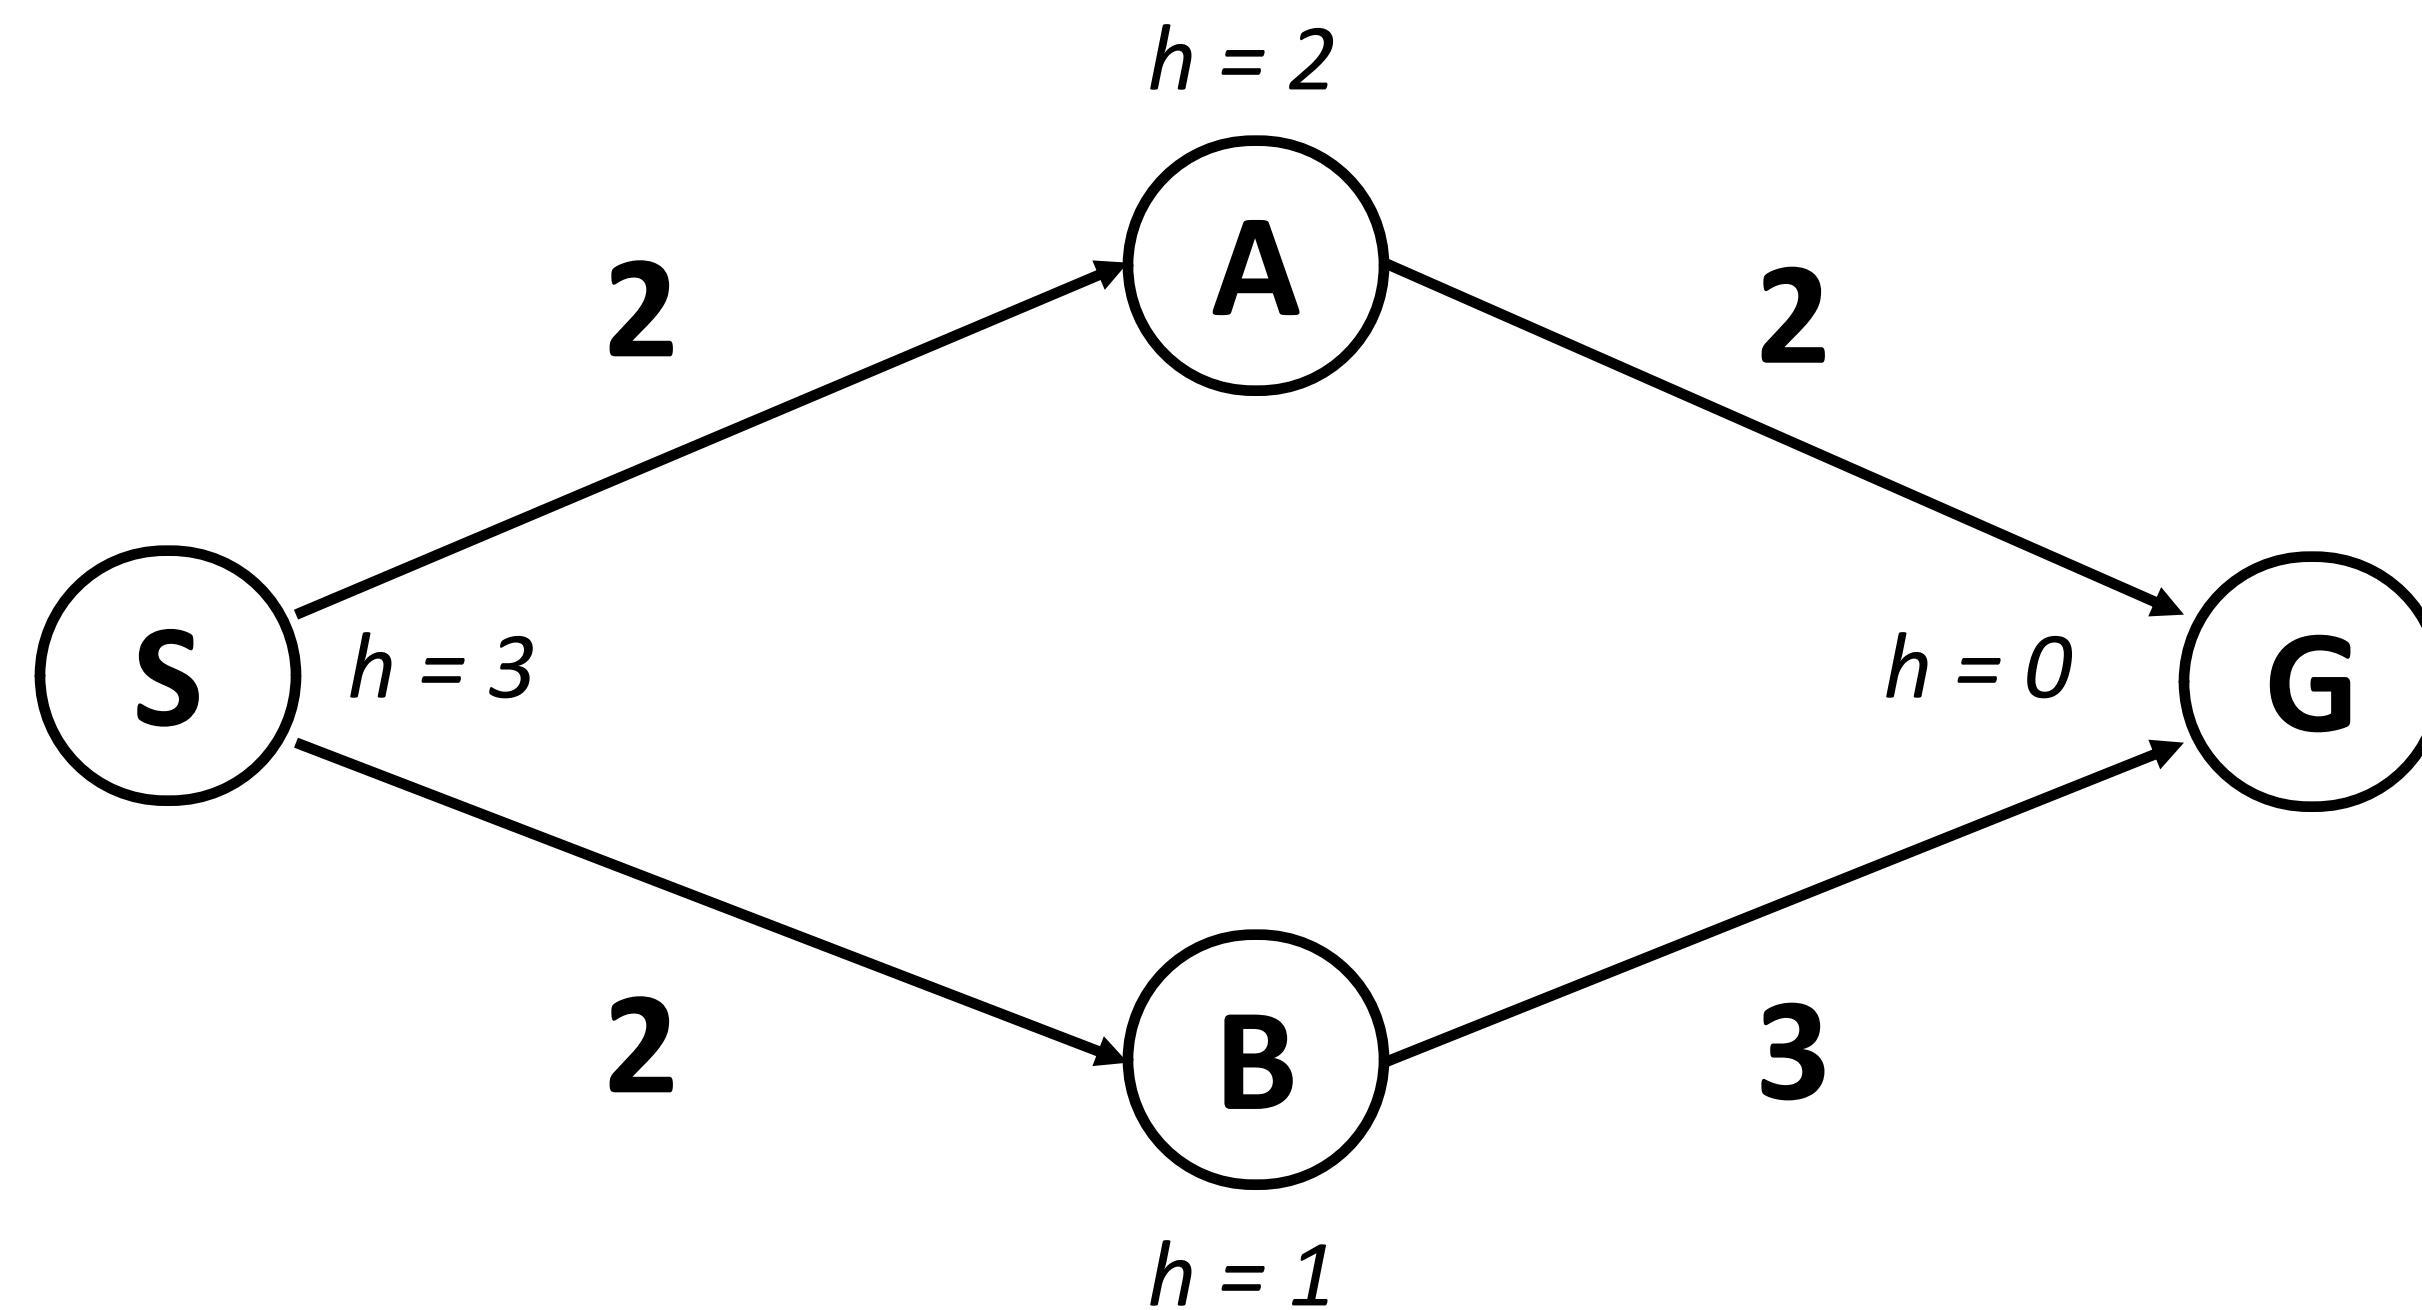
Uniform-cost expands
equally in all "directions"



A* expands mainly
toward the goal, but
does hedge its bets to
ensure optimality

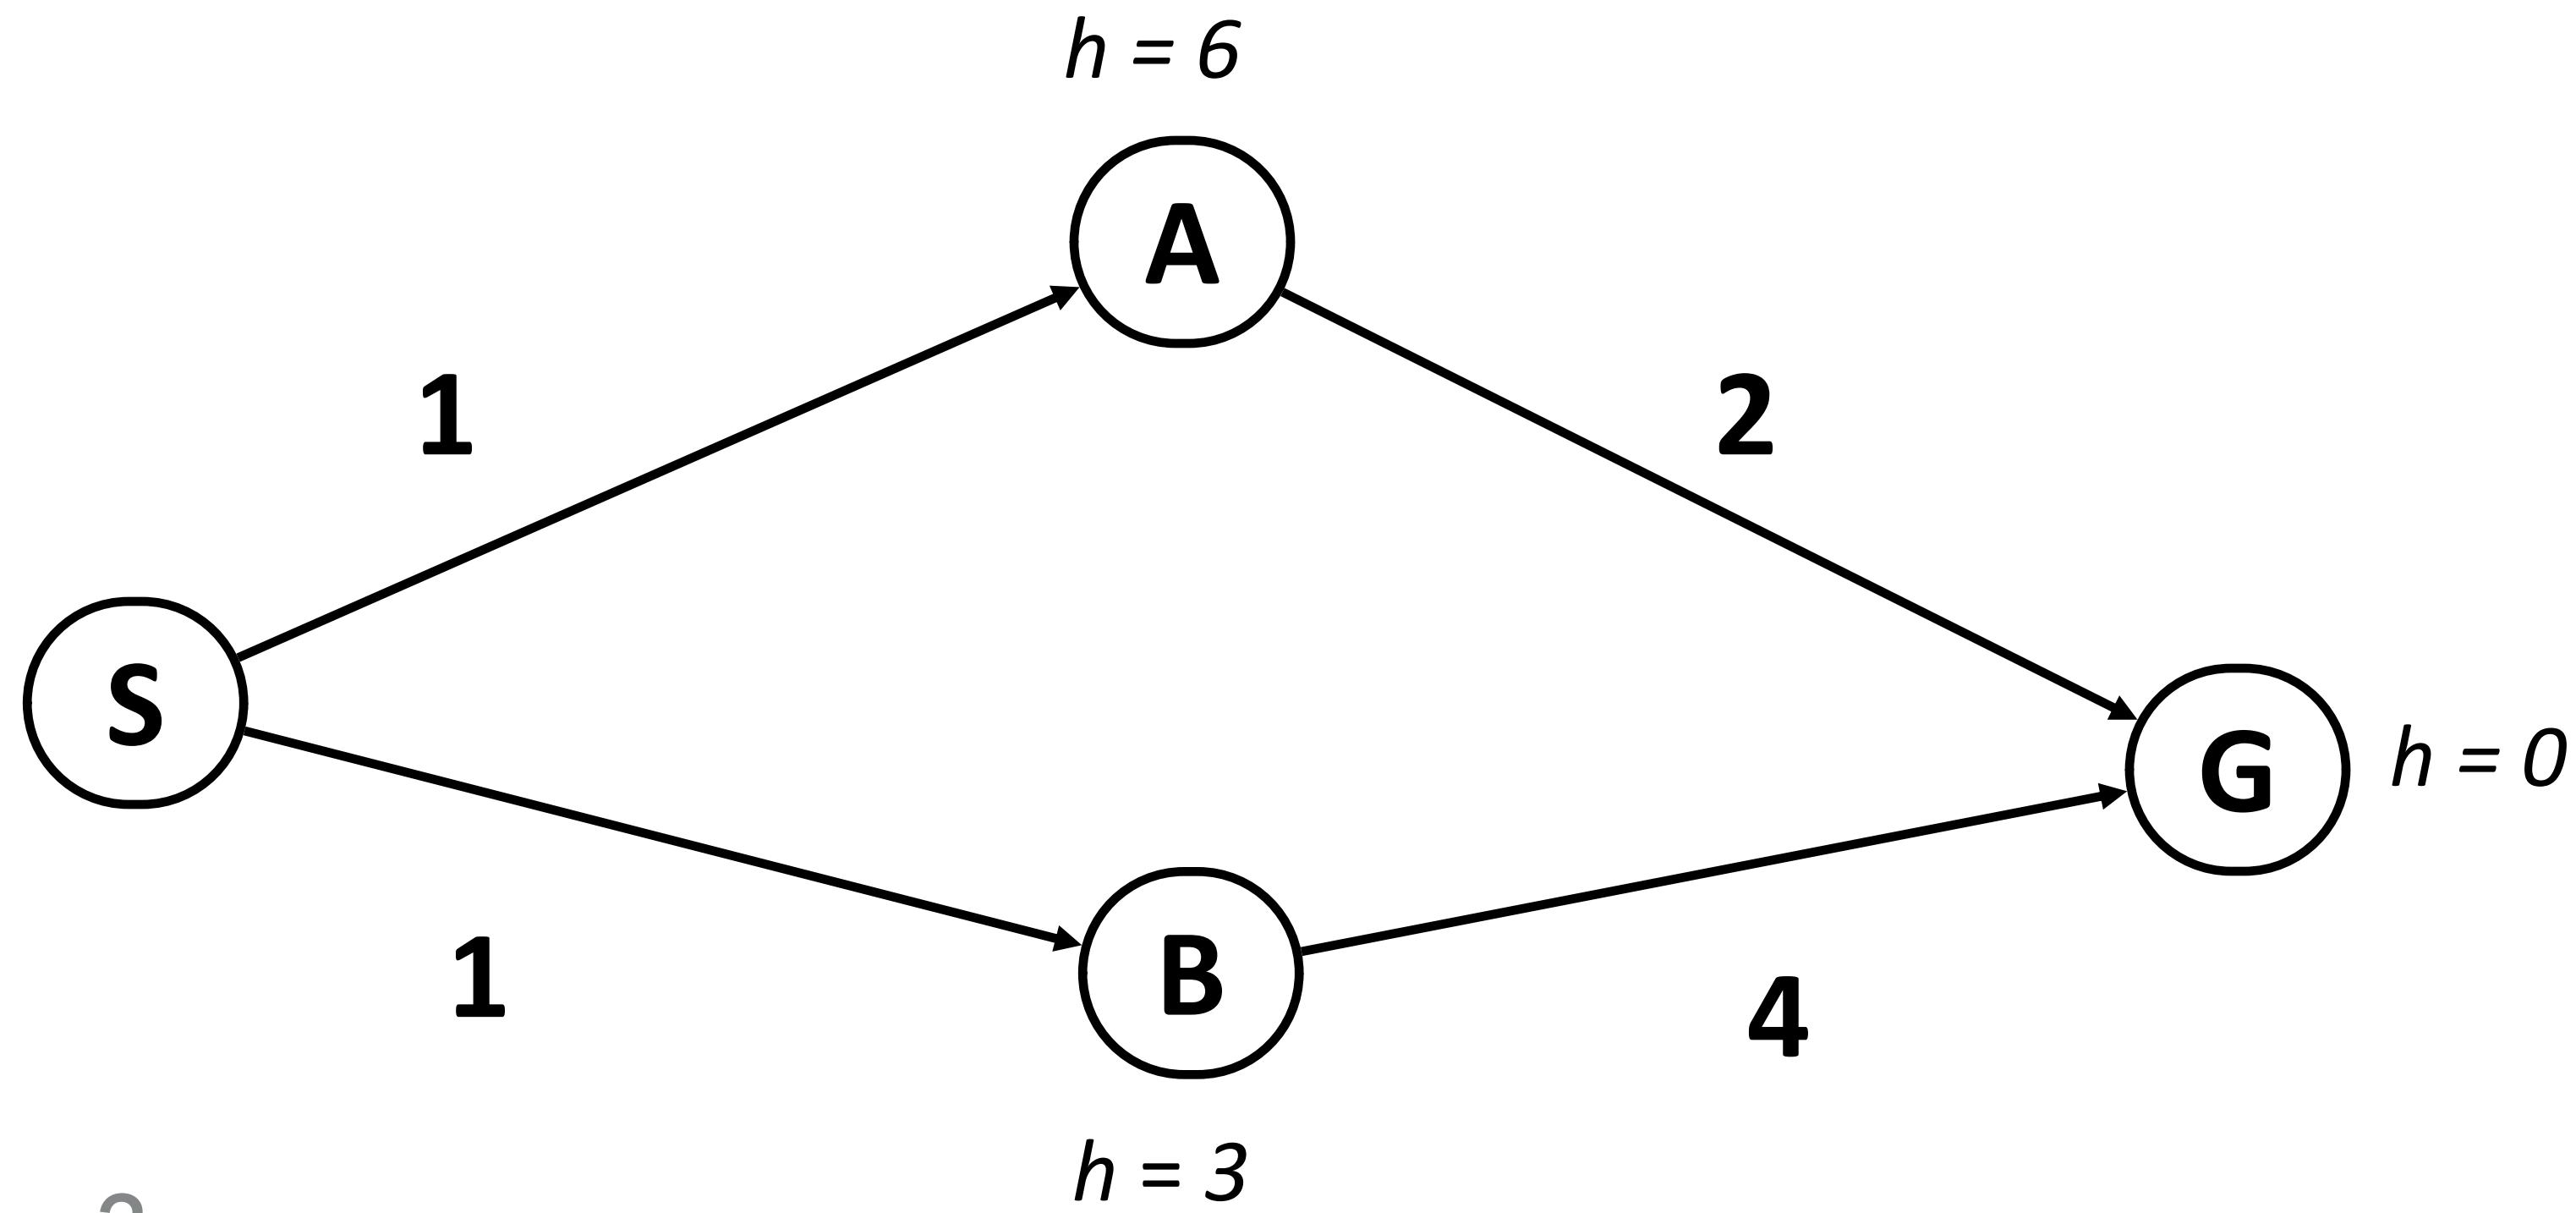
WHEN SHOULD A* TERMINATE?

- ▶ Should we stop when a goal state enters the fringe?



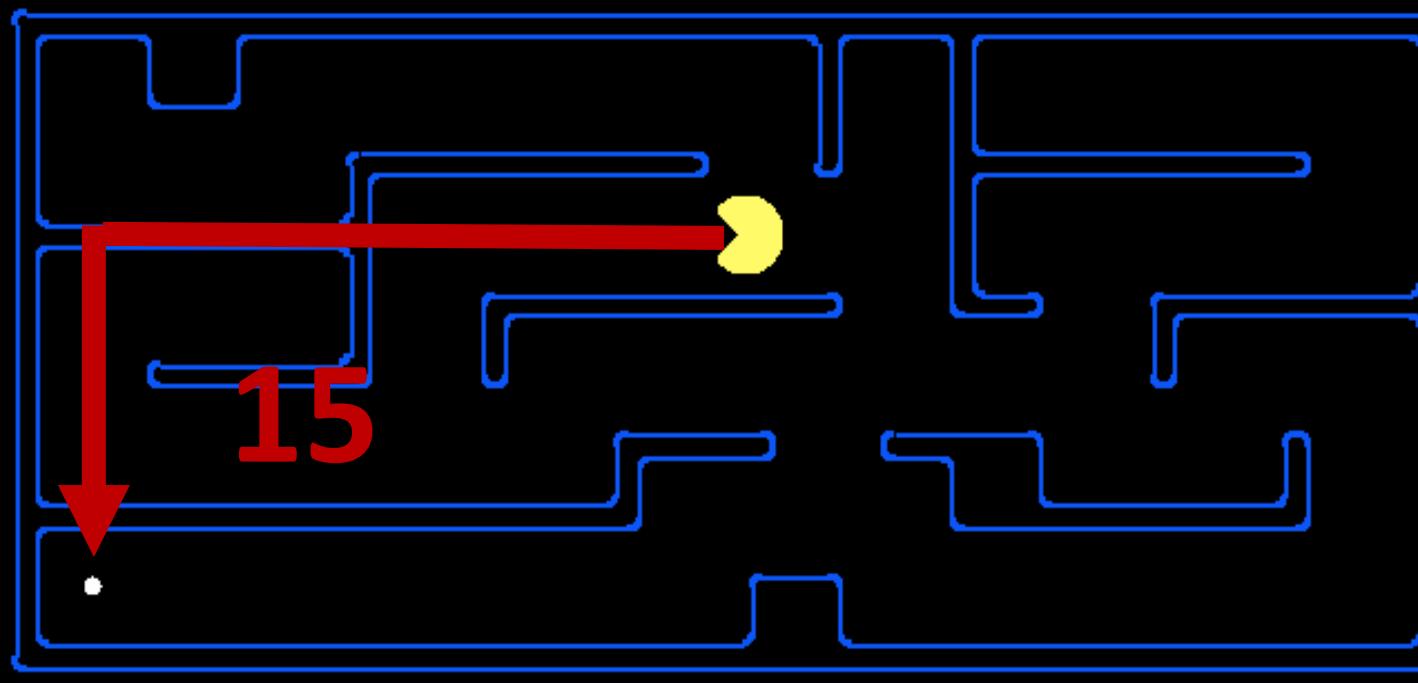
- ▶ No: only stop when we a goal state is expanded (i.e., exits the fringe)

IS A* OPTIMAL?



- ▶ What went wrong?
- ▶ We need **estimated costs $h(n)$ to be less than actual costs**

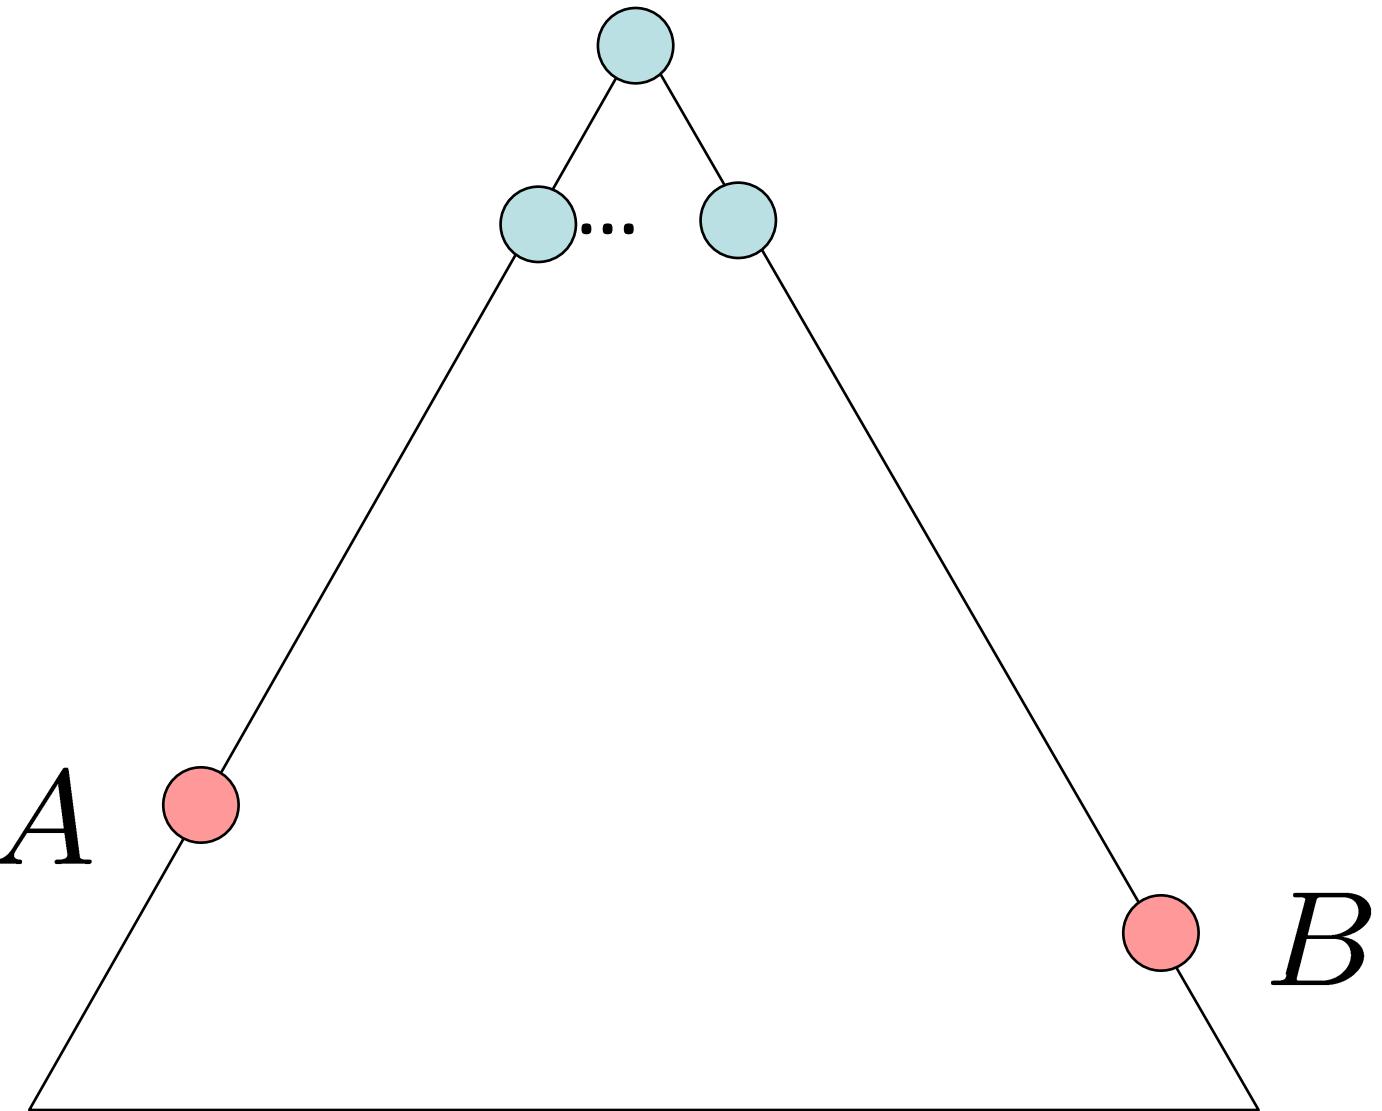
ADMISSIBLE HEURISTICS

- ▶ A heuristic h is **admissible** (optimistic) if $0 \leq h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to a nearest goal.
- ▶ Example:

- ▶ Coming up with admissible heuristics is most of what's involved in using A* in practice

OPTIMALITY OF A* TREE SEARCH

Assume:

- ▶ A is an optimal goal node
- ▶ B is a suboptimal goal node
- ▶ h is admissible



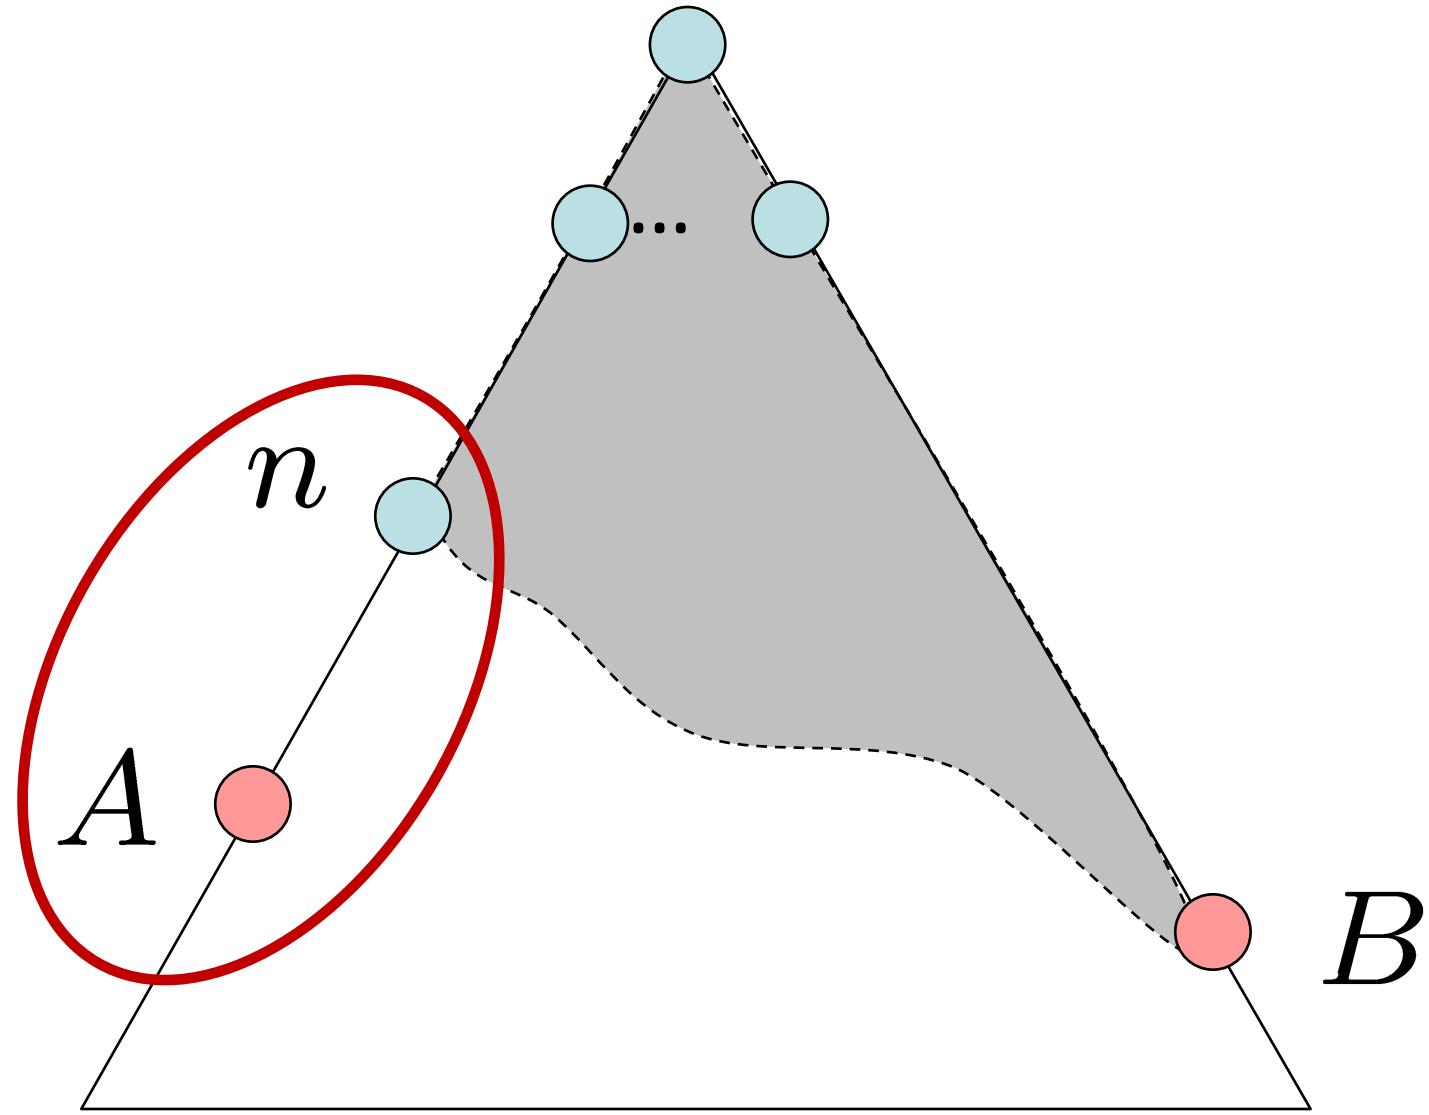
Claim:

- ▶ A will exit the fringe before B

OPTIMALITY OF A* TREE SEARCH

Proof:

- ▶ Imagine B is on the fringe
 - ▶ Some ancestor n of A is on the fringe
 - ▶ Claim: n will be expanded before B
1. $f(n)$ is less or equal to $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

$$g(A) = f(A)$$

Definition of f-cost

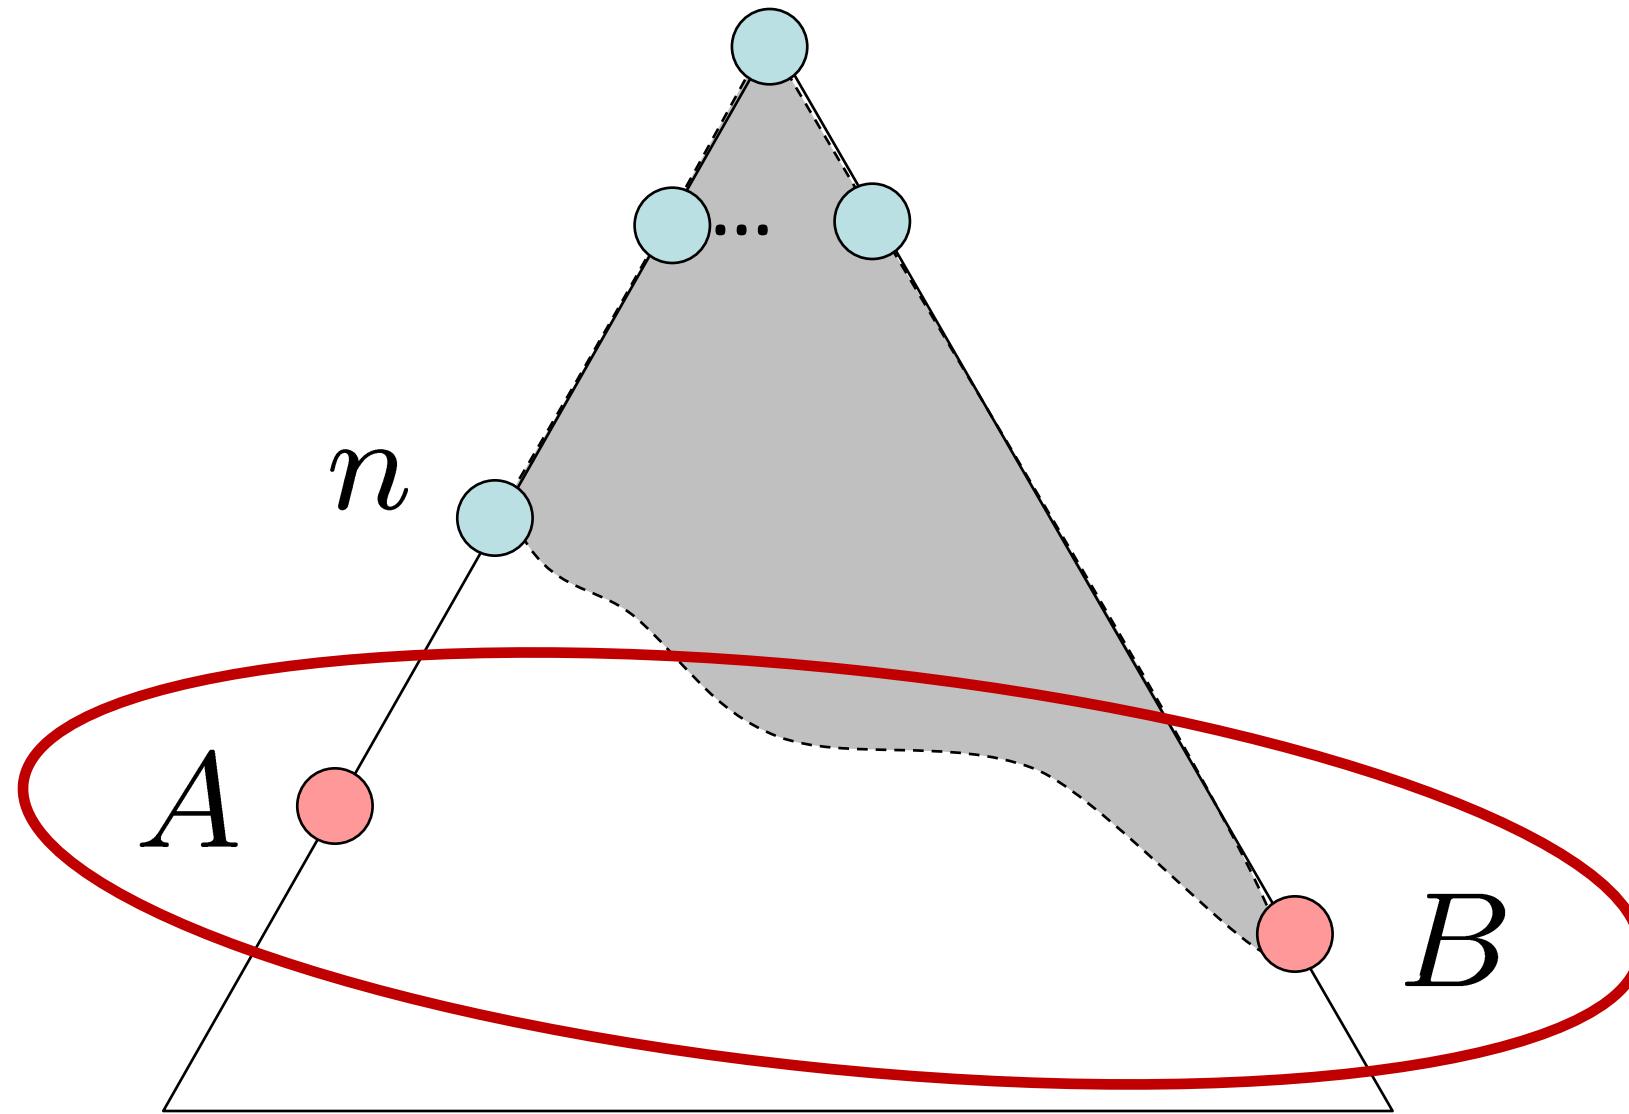
Admissibility of h

$h = 0$ at a goal

OPTIMALITY OF A* TREE SEARCH

Proof:

- ▶ Imagine B is on the fringe
 - ▶ Some ancestor n of A is on the fringe
 - ▶ Claim: n will be expanded before B
1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$



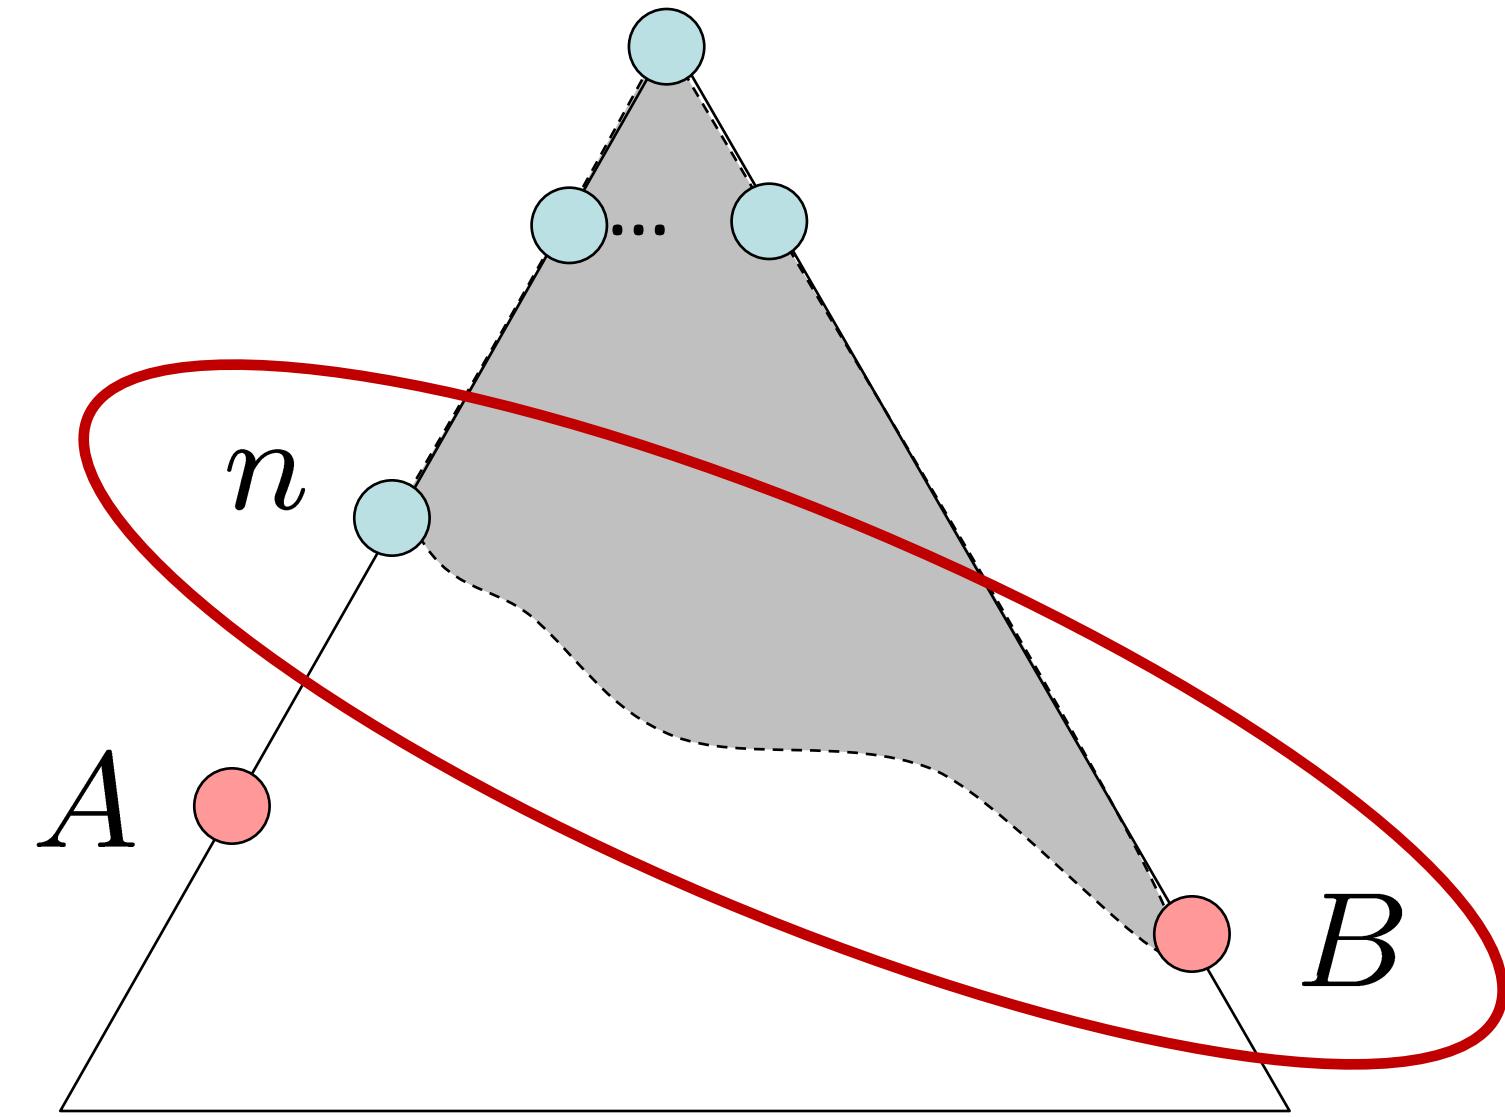
$$\begin{aligned} g(A) &< g(B) \\ f(A) &< f(B) \end{aligned}$$

B is suboptimal
 $h = 0$ at a goal

OPTIMALITY OF A* TREE SEARCH

Proof:

- ▶ Imagine B is on the fringe
- ▶ Some ancestor n of A is on the fringe
- ▶ Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$
 3. n expands before B

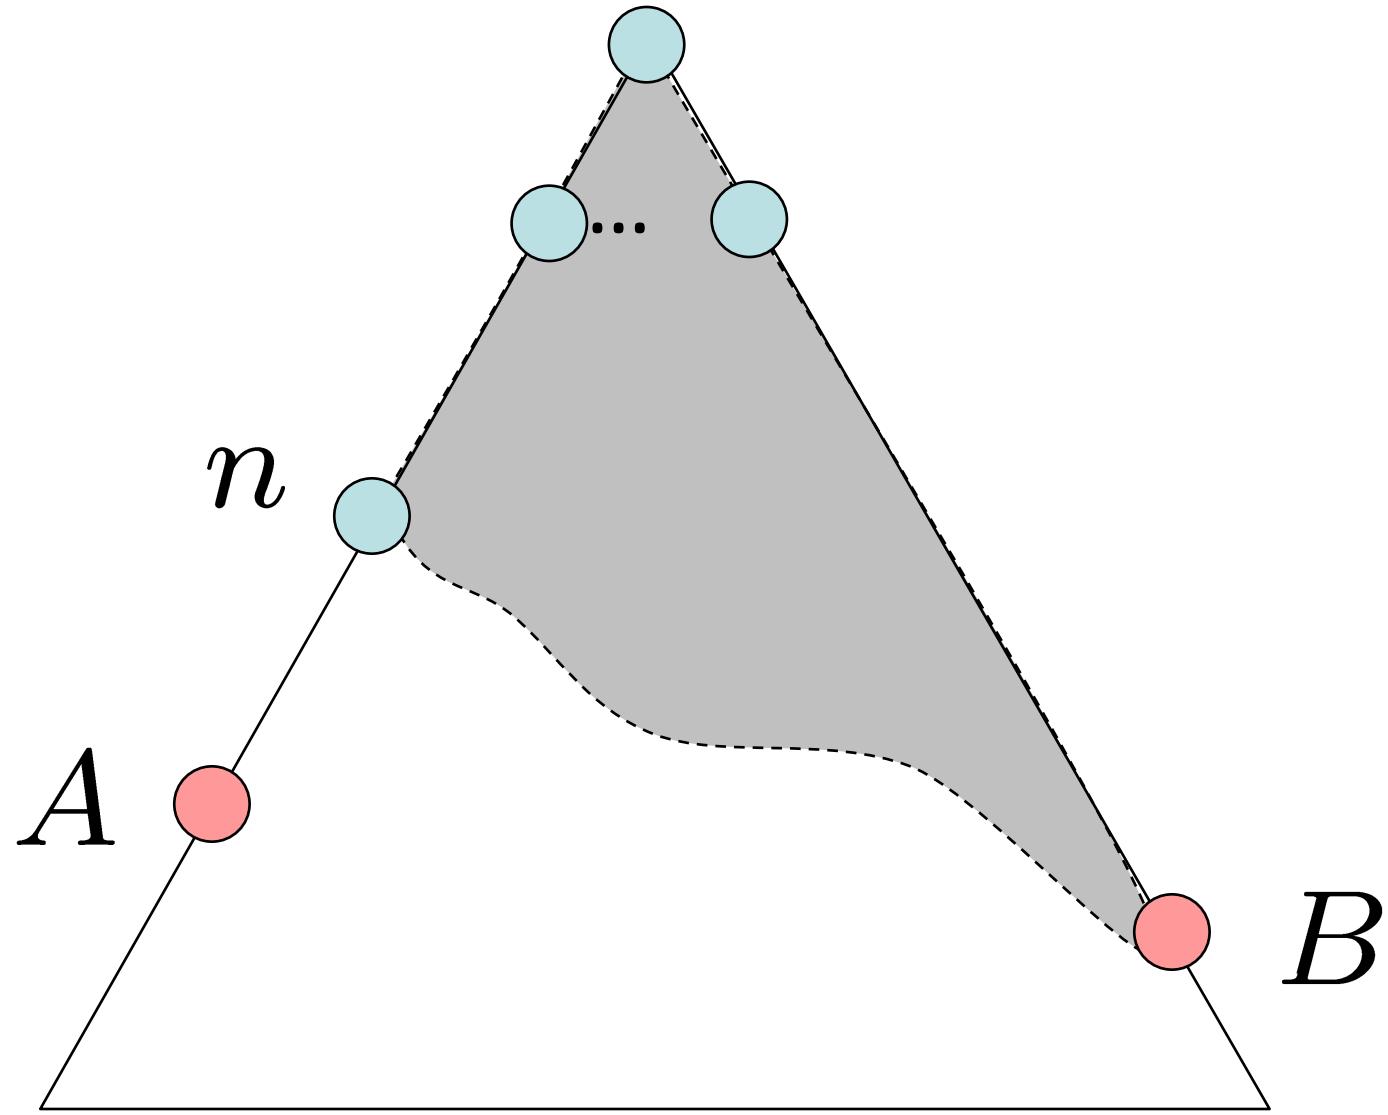


$$f(n) \leq f(A) < f(B)$$

OPTIMALITY OF A* TREE SEARCH

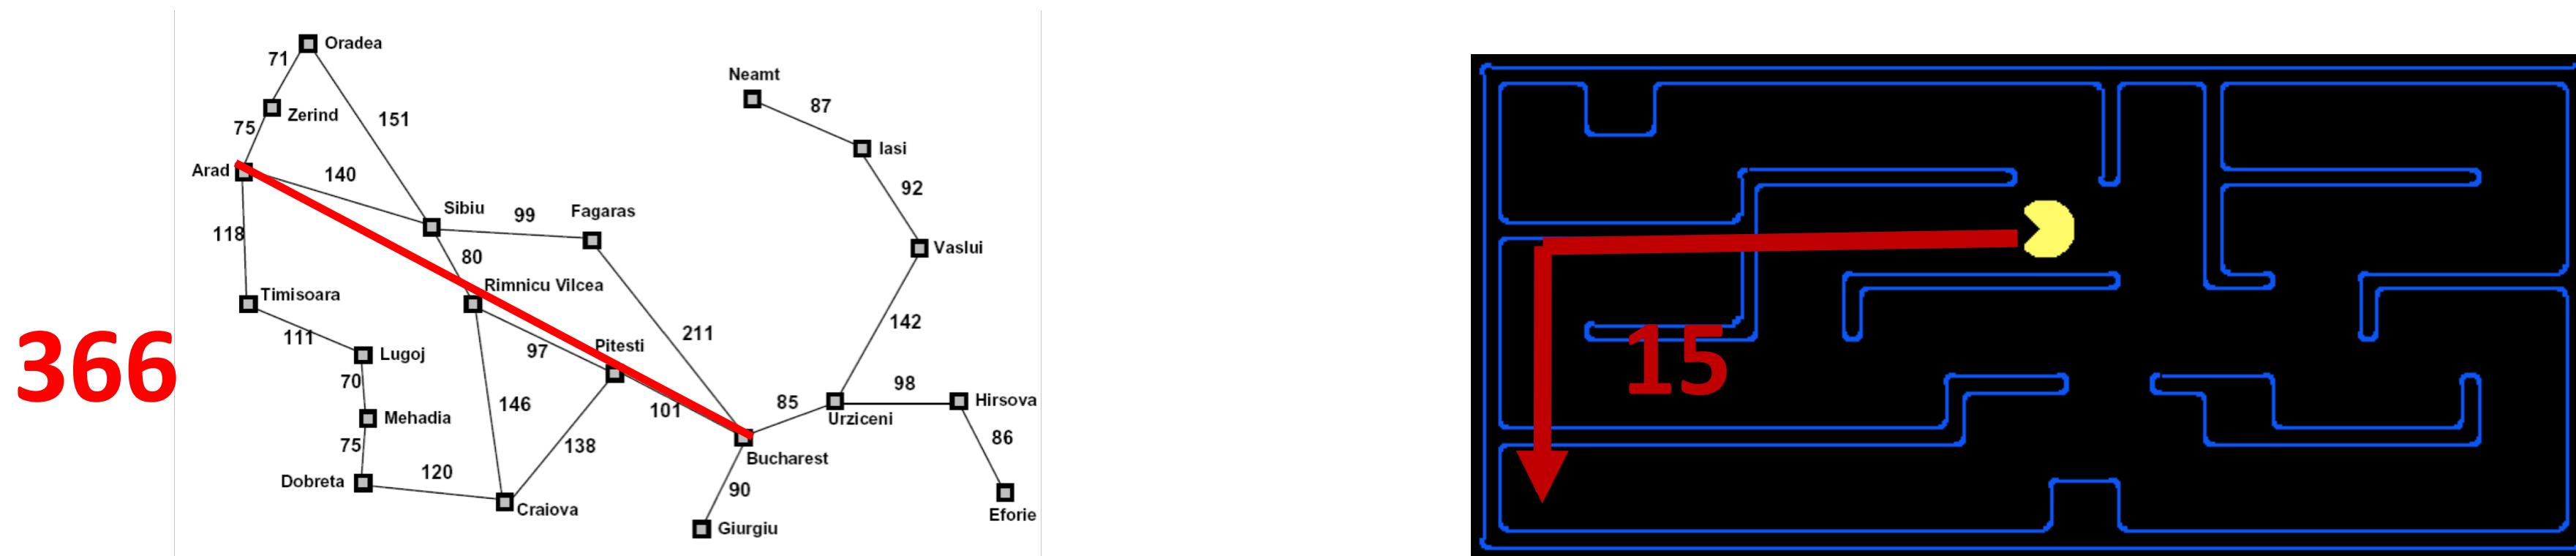
Proof:

- ▶ Imagine B is on the fringe; some ancestor n of A is on the fringe
- ▶ Claim: n will be expanded before B
 - 1. $f(n)$ is less or equal to $f(A)$
 - 2. $f(A)$ is less than $f(B)$
 - 3. n expands before B
- ▶ All ancestors of A expand before B
- ▶ A expands before B
- ▶ A* search is optimal



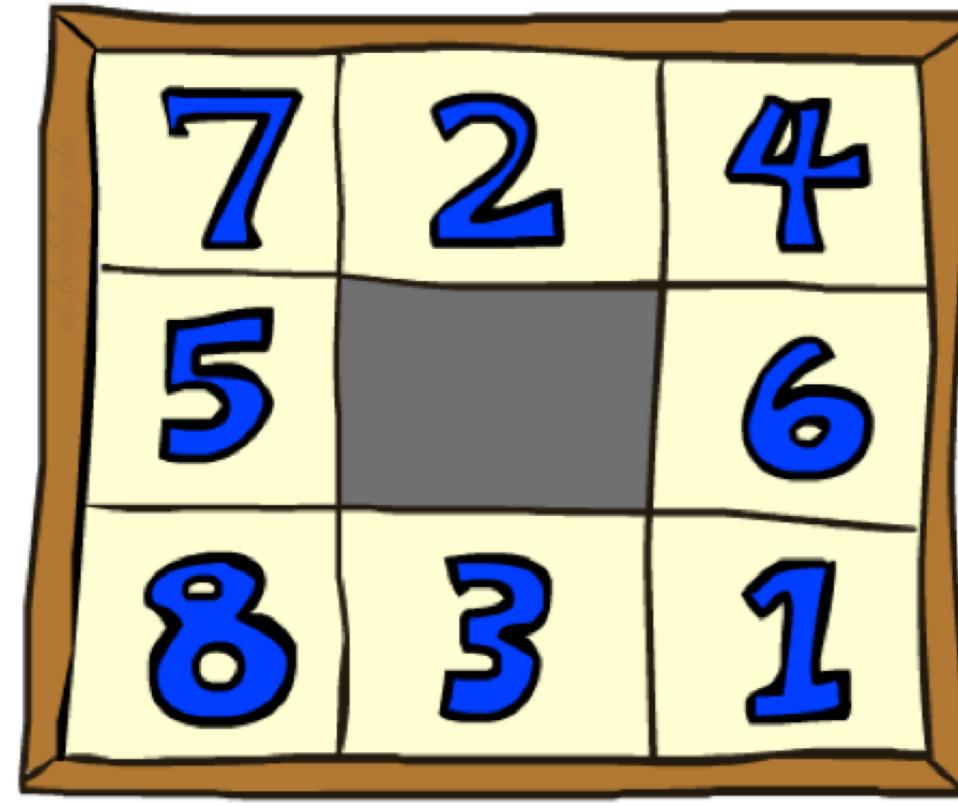
CREATING ADMISSIBLE HEURISTICS

- ▶ Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- ▶ Often, admissible heuristics are solutions to *relaxed problems*, where new actions are available

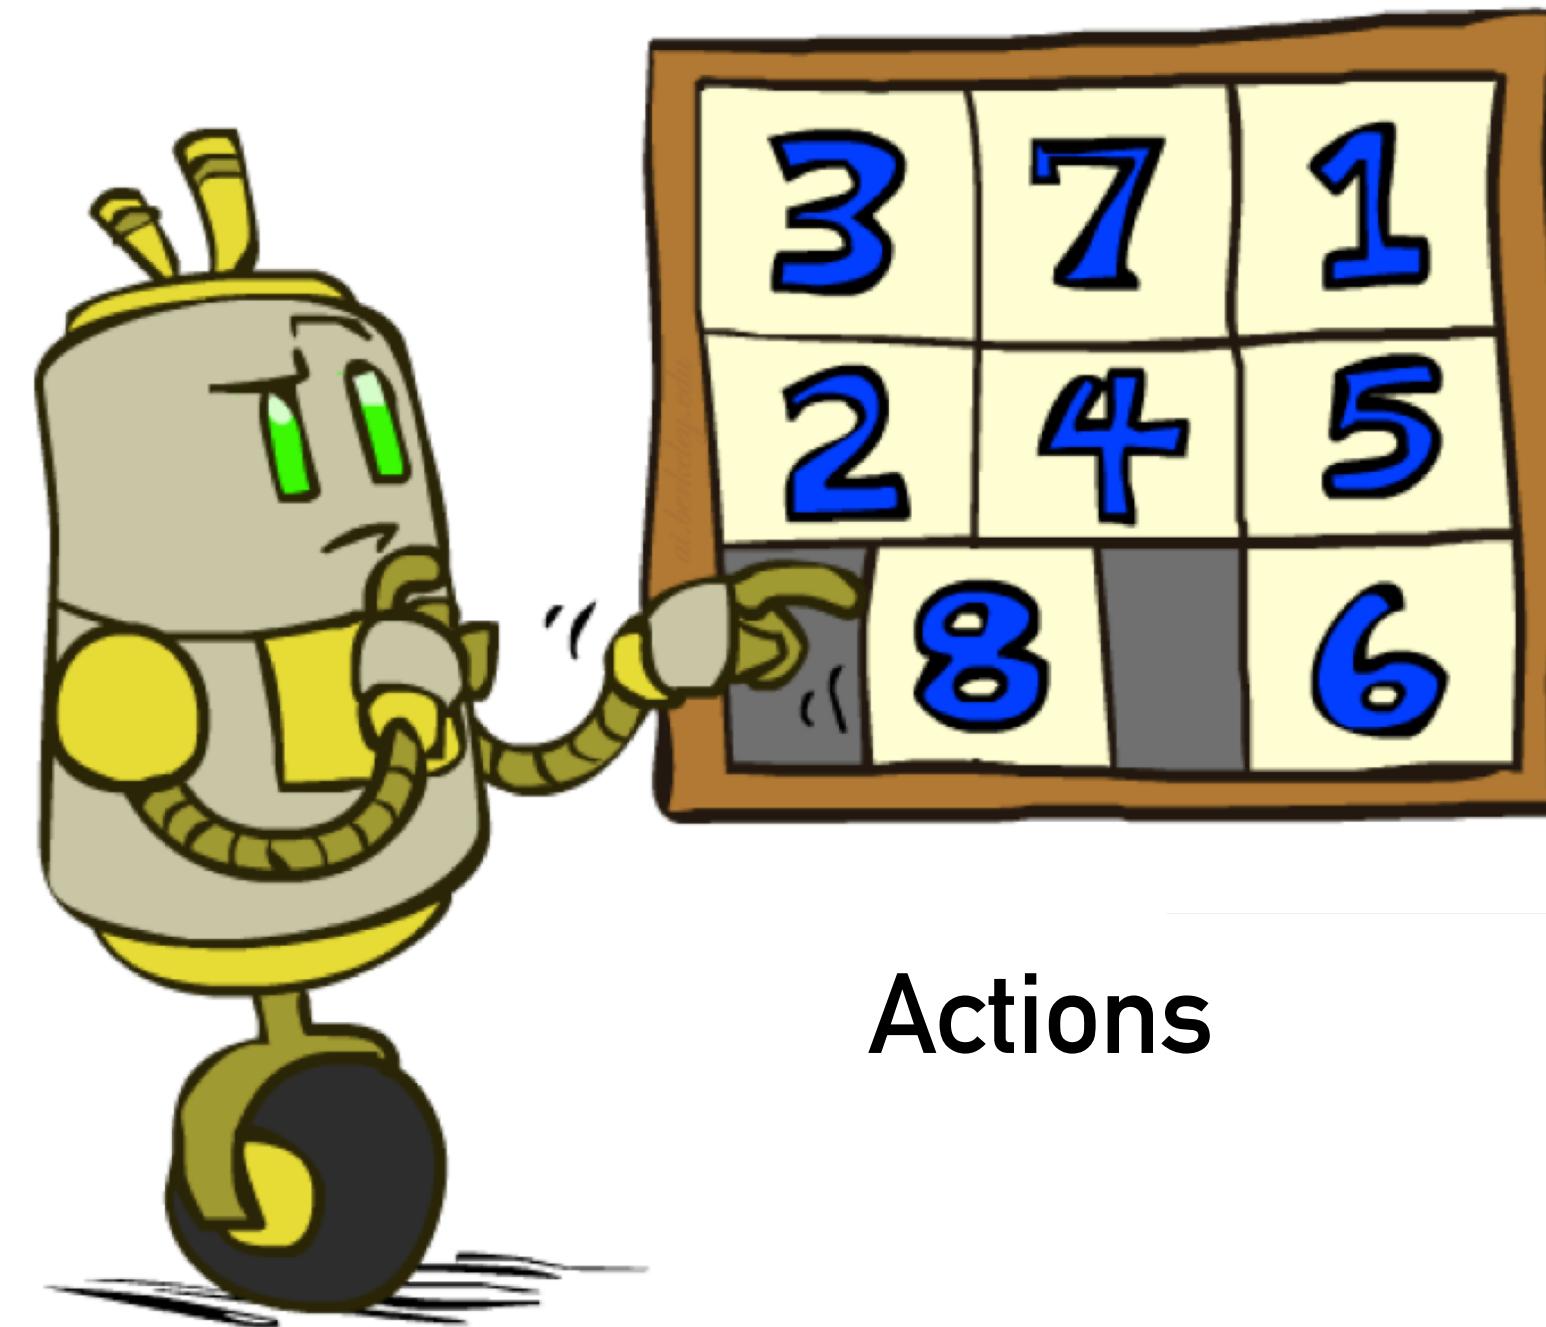


- ▶ Inadmissible heuristics are often useful too

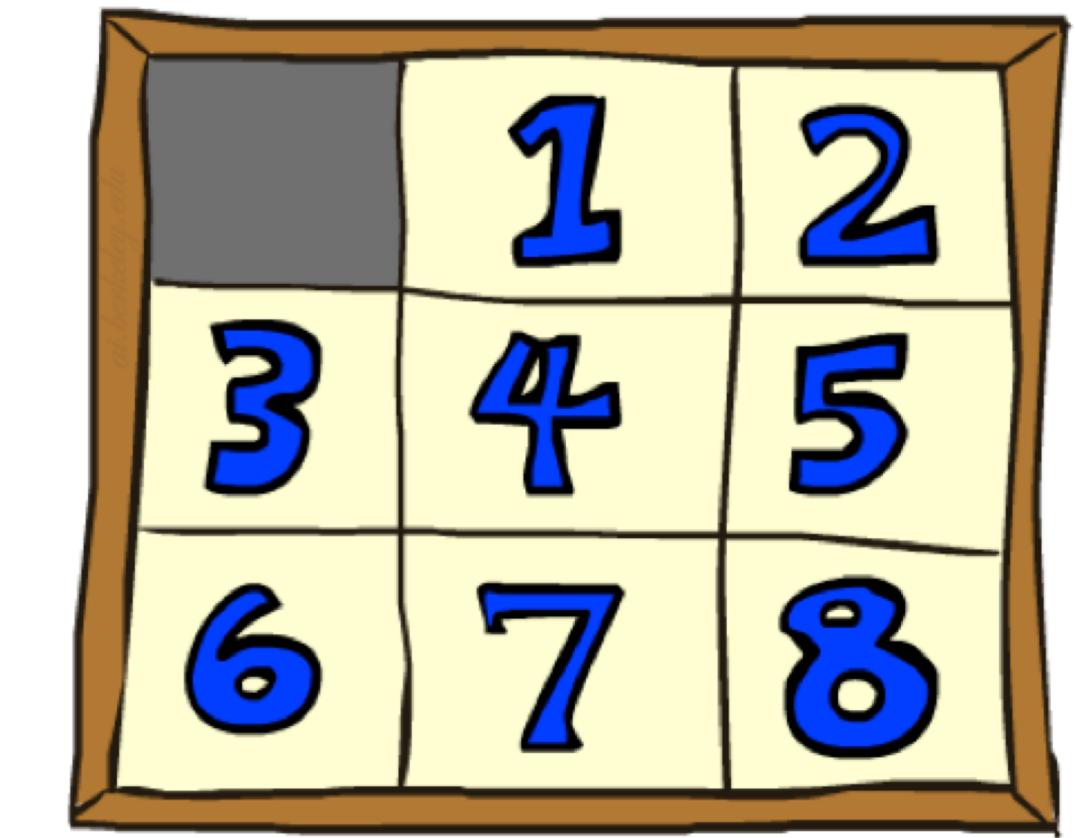
ADMISSIBLE HEURISTICS EXAMPLE: 8 PUZZLE



Start State



Actions



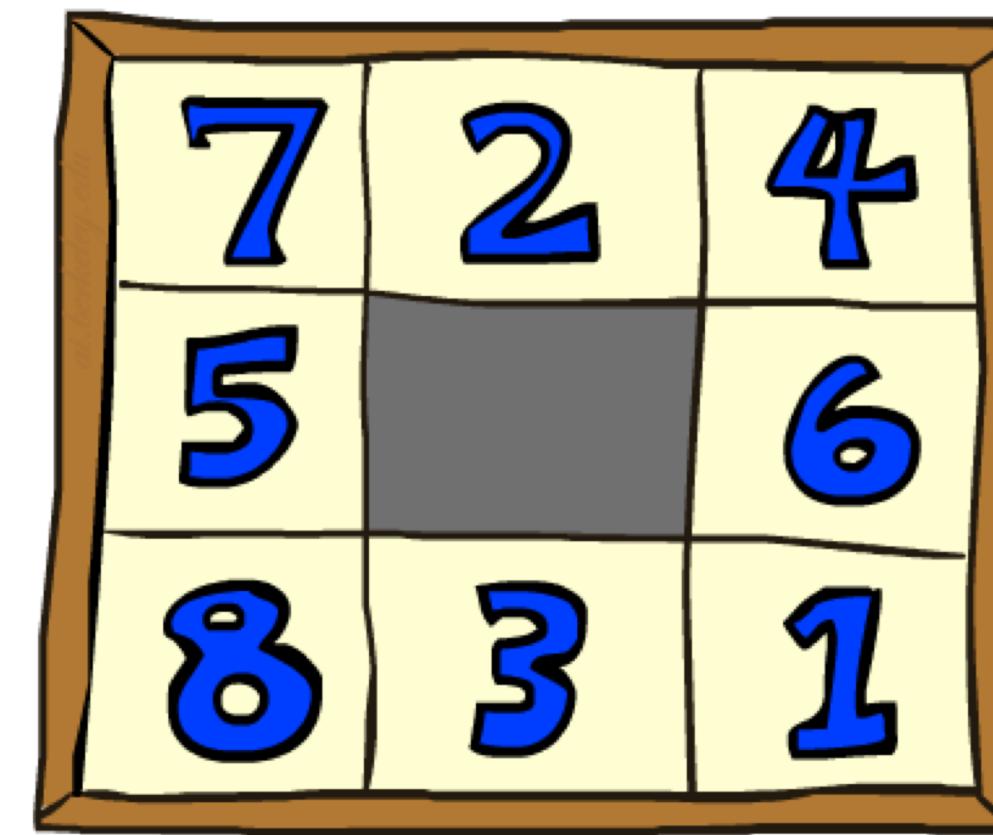
Goal State

What could be an admissible heuristic function for 8-puzzle problems?

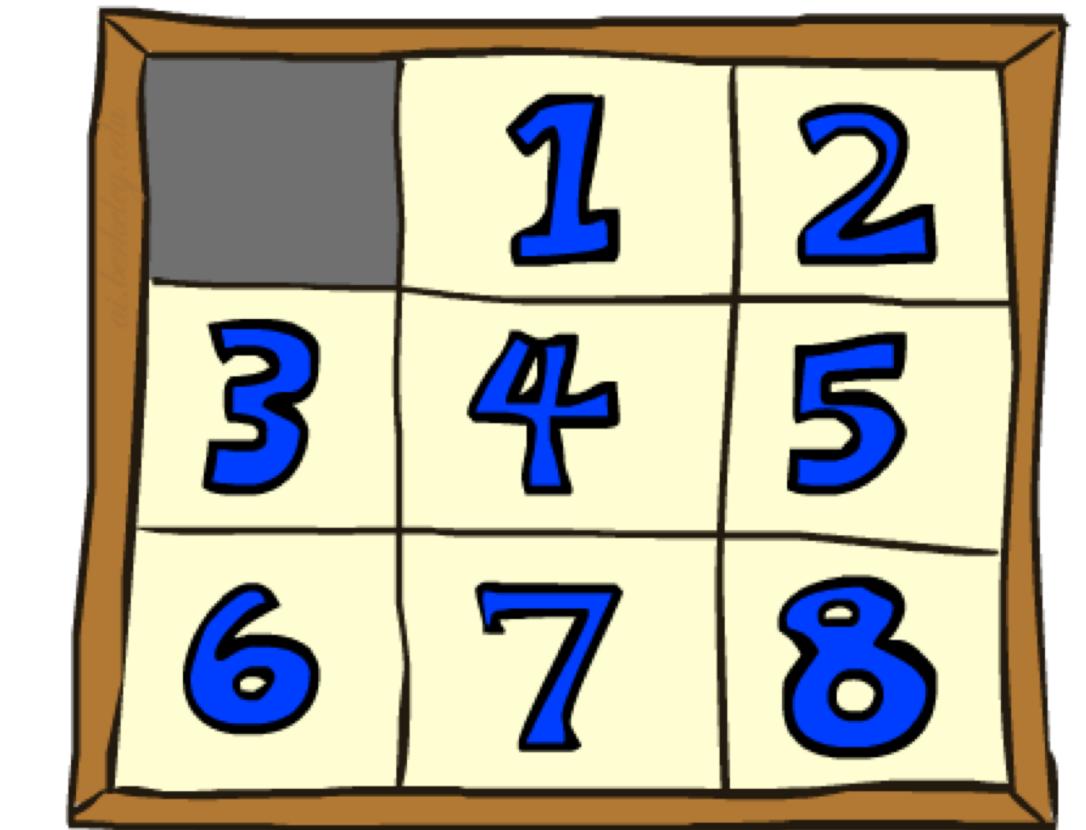
Hint: A tile can move from Location A to B if A and B are adjacent and B is blank...Relax it?

8 PUZZLE HEURISTIC I

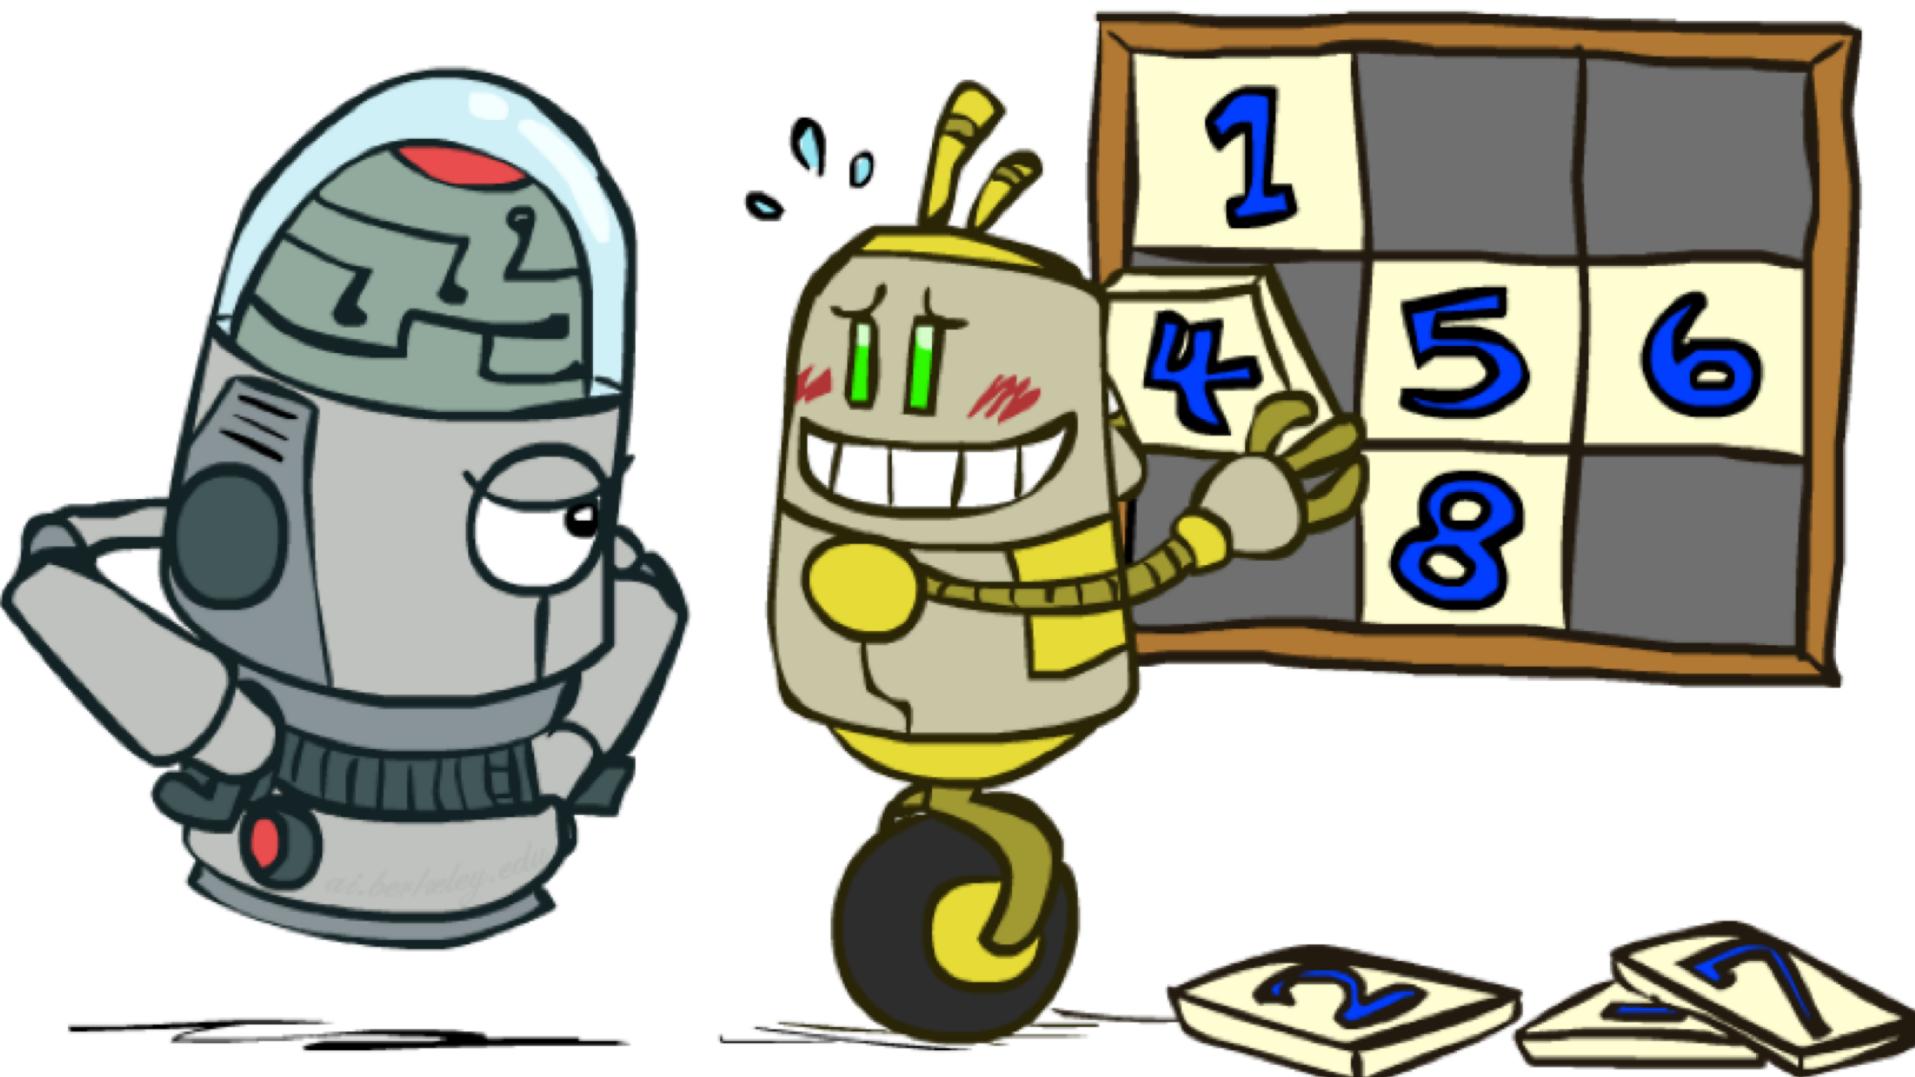
- ▶ Heuristic: Number of misplaced tiles
- ▶ Why is it admissible?
- ▶ $h(\text{start}) = 8$
- ▶ This is a *relaxed-problem* heuristic



Start State



Goal State



Average nodes expanded when the optimal path has...

	...4 steps	...8 steps	...12 steps
UCS	112	6,300	3.6×10^6
TILES	13	39	227