

PURDUE CS47100

---

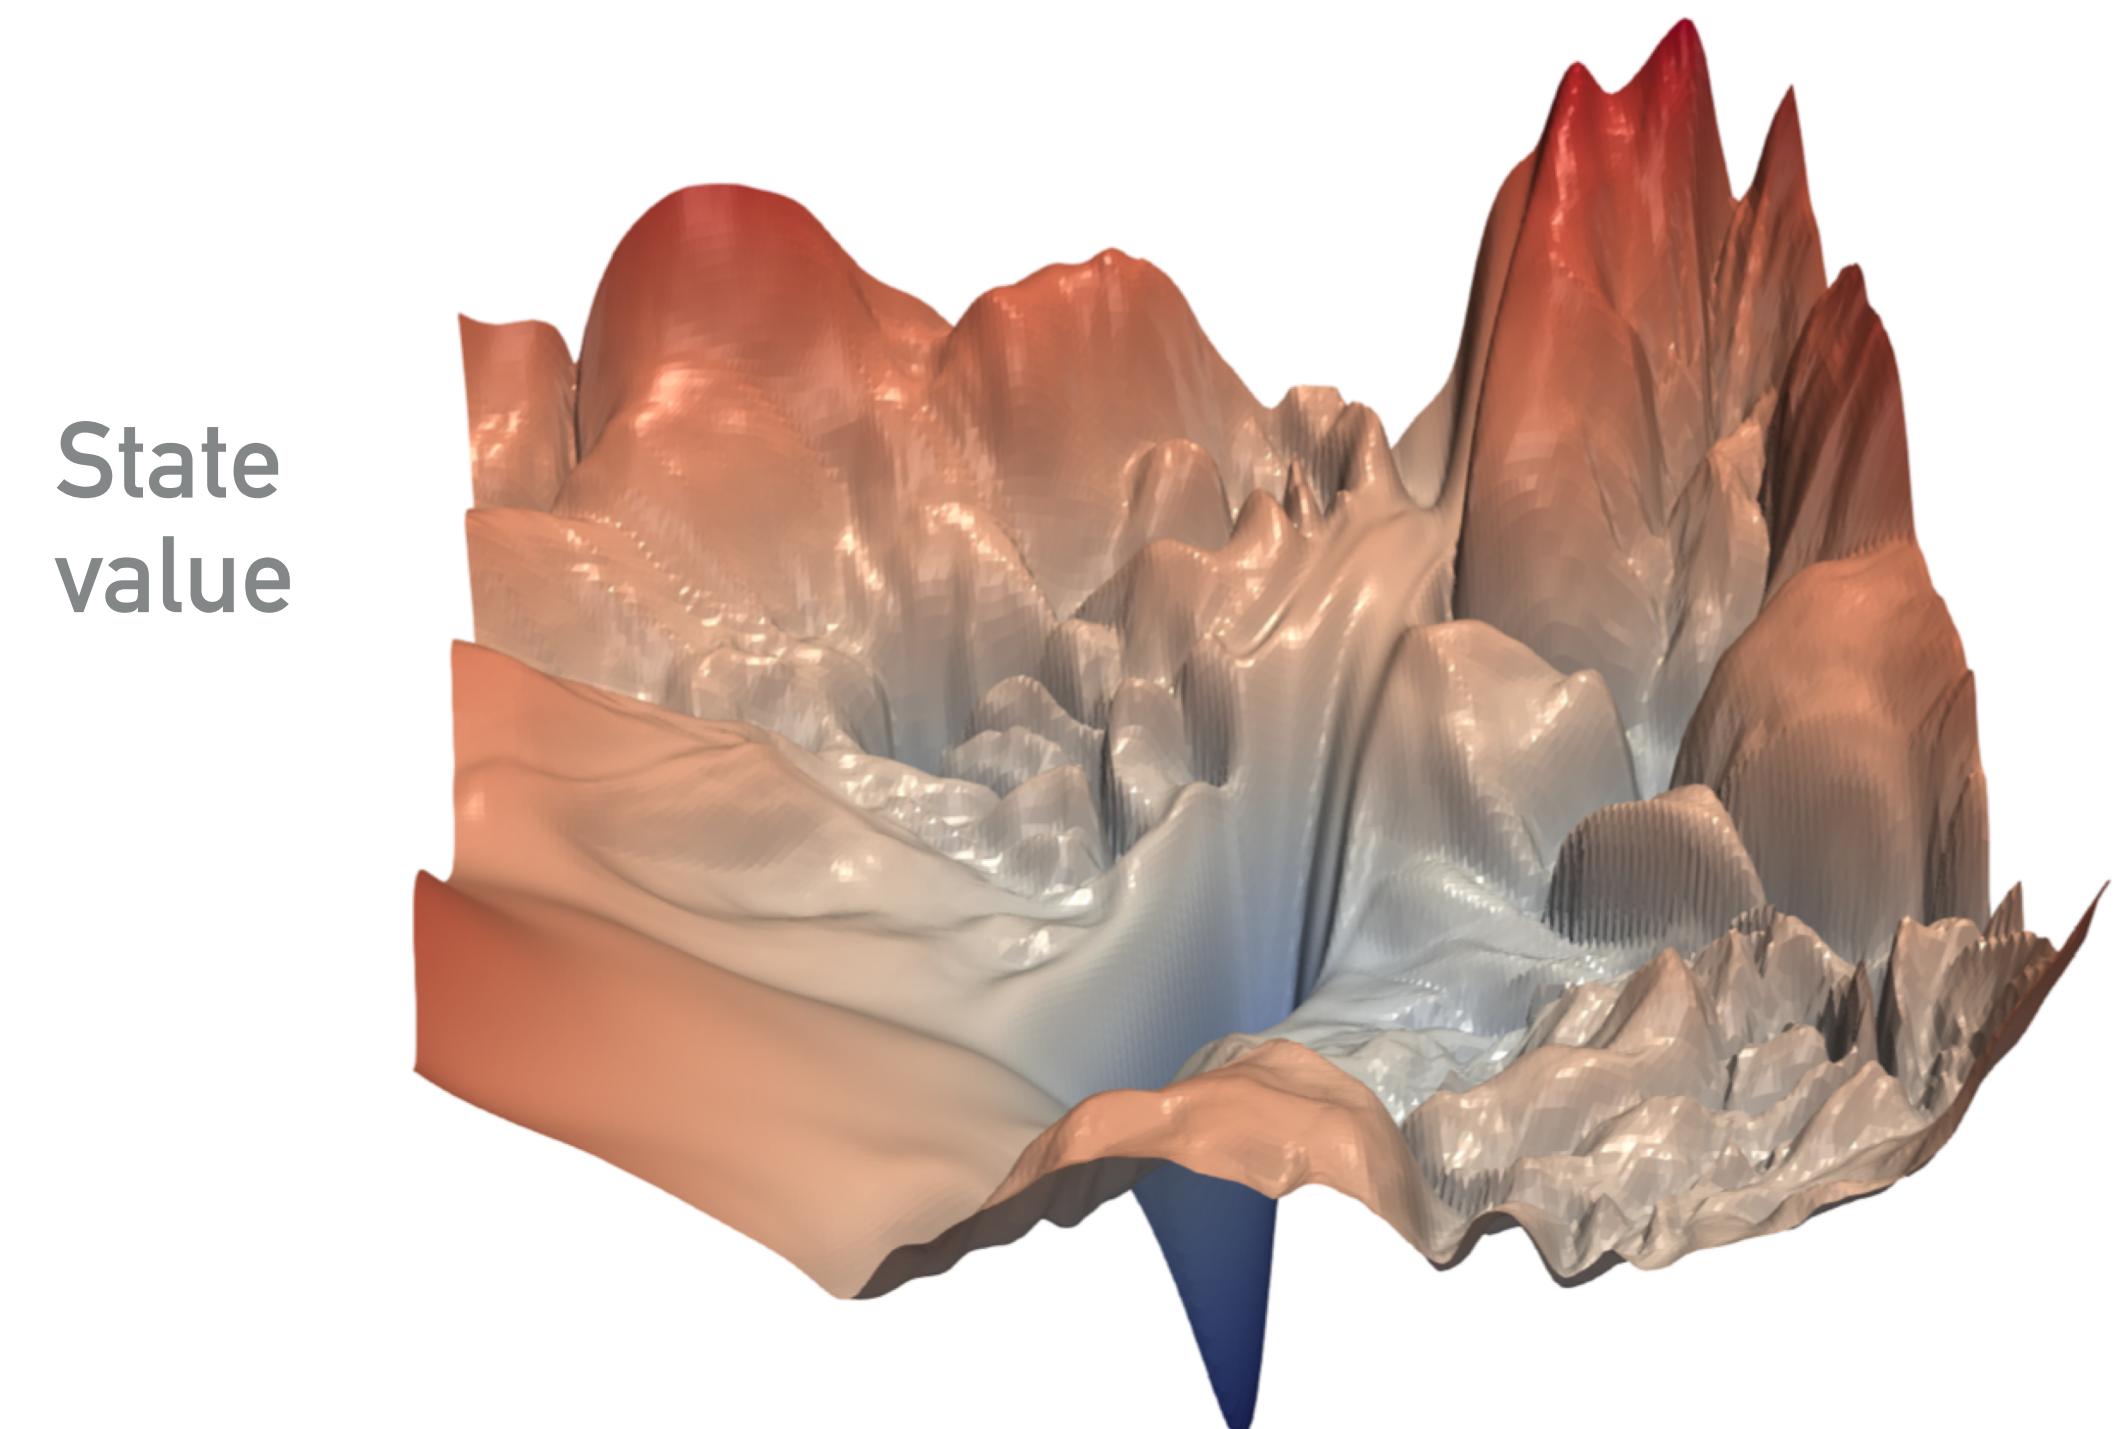
# INTRODUCTION TO AI

---

## RECAP: GRAPH SEARCH & LOCAL SEARCH

- ▶ Graph search
  - ▶ Never expand the same state twice!
  - ▶ A\* needs an consistent heuristic for the graph search algorithm to be optimal
- ▶ Local search algorithm
  - ▶ Useful when path to goal state does not matter / solve pure optimization problem
  - ▶ Hill-climbing search: always move to a neighboring state with the optimal objective function value

# STATE SPACE LANDSCAPES IN PRACTICE (FOR NEURAL NETWORKS)



State location

## VARIANTS OF LOCAL HILL-CLIMBING

- ▶ Stochastic hill-climbing
  - ▶ Select randomly from all moves that improve the value of the objective function
- ▶ Random-restart hill-climbing
  - ▶ Conducts a series of hill-climbing searches, starting from random positions
  - ▶ Very frequently used general method in AI

## SIMULATED ANNEALING

- ▶ Inspired by statistical physics/metallurgy
- ▶ Key idea: mostly goes “uphill”, but occasionally travels “downhill” to escape local optimum
- ▶ The likelihood to go downhill is controlled by a “temperature schedule”
- ▶ More and more “conservative” (i.e., less likely to go downhill) as the search progress

---

# SIMULATED ANNEALING SEARCH

```
function SIMULATED_ANNEALING (problem, schedule) return a solution state
    current ← initial_state;
    for t = 1 to ∞:
        T ← schedule(t)
        if T = 0 then return current
        next ← a randomly selected neighbor of current
        ΔE = next.value - current.value
        if ΔE > 0:
            current ← next
        else:
            current ← next with probability  $e^{\Delta E / T}$ 
```

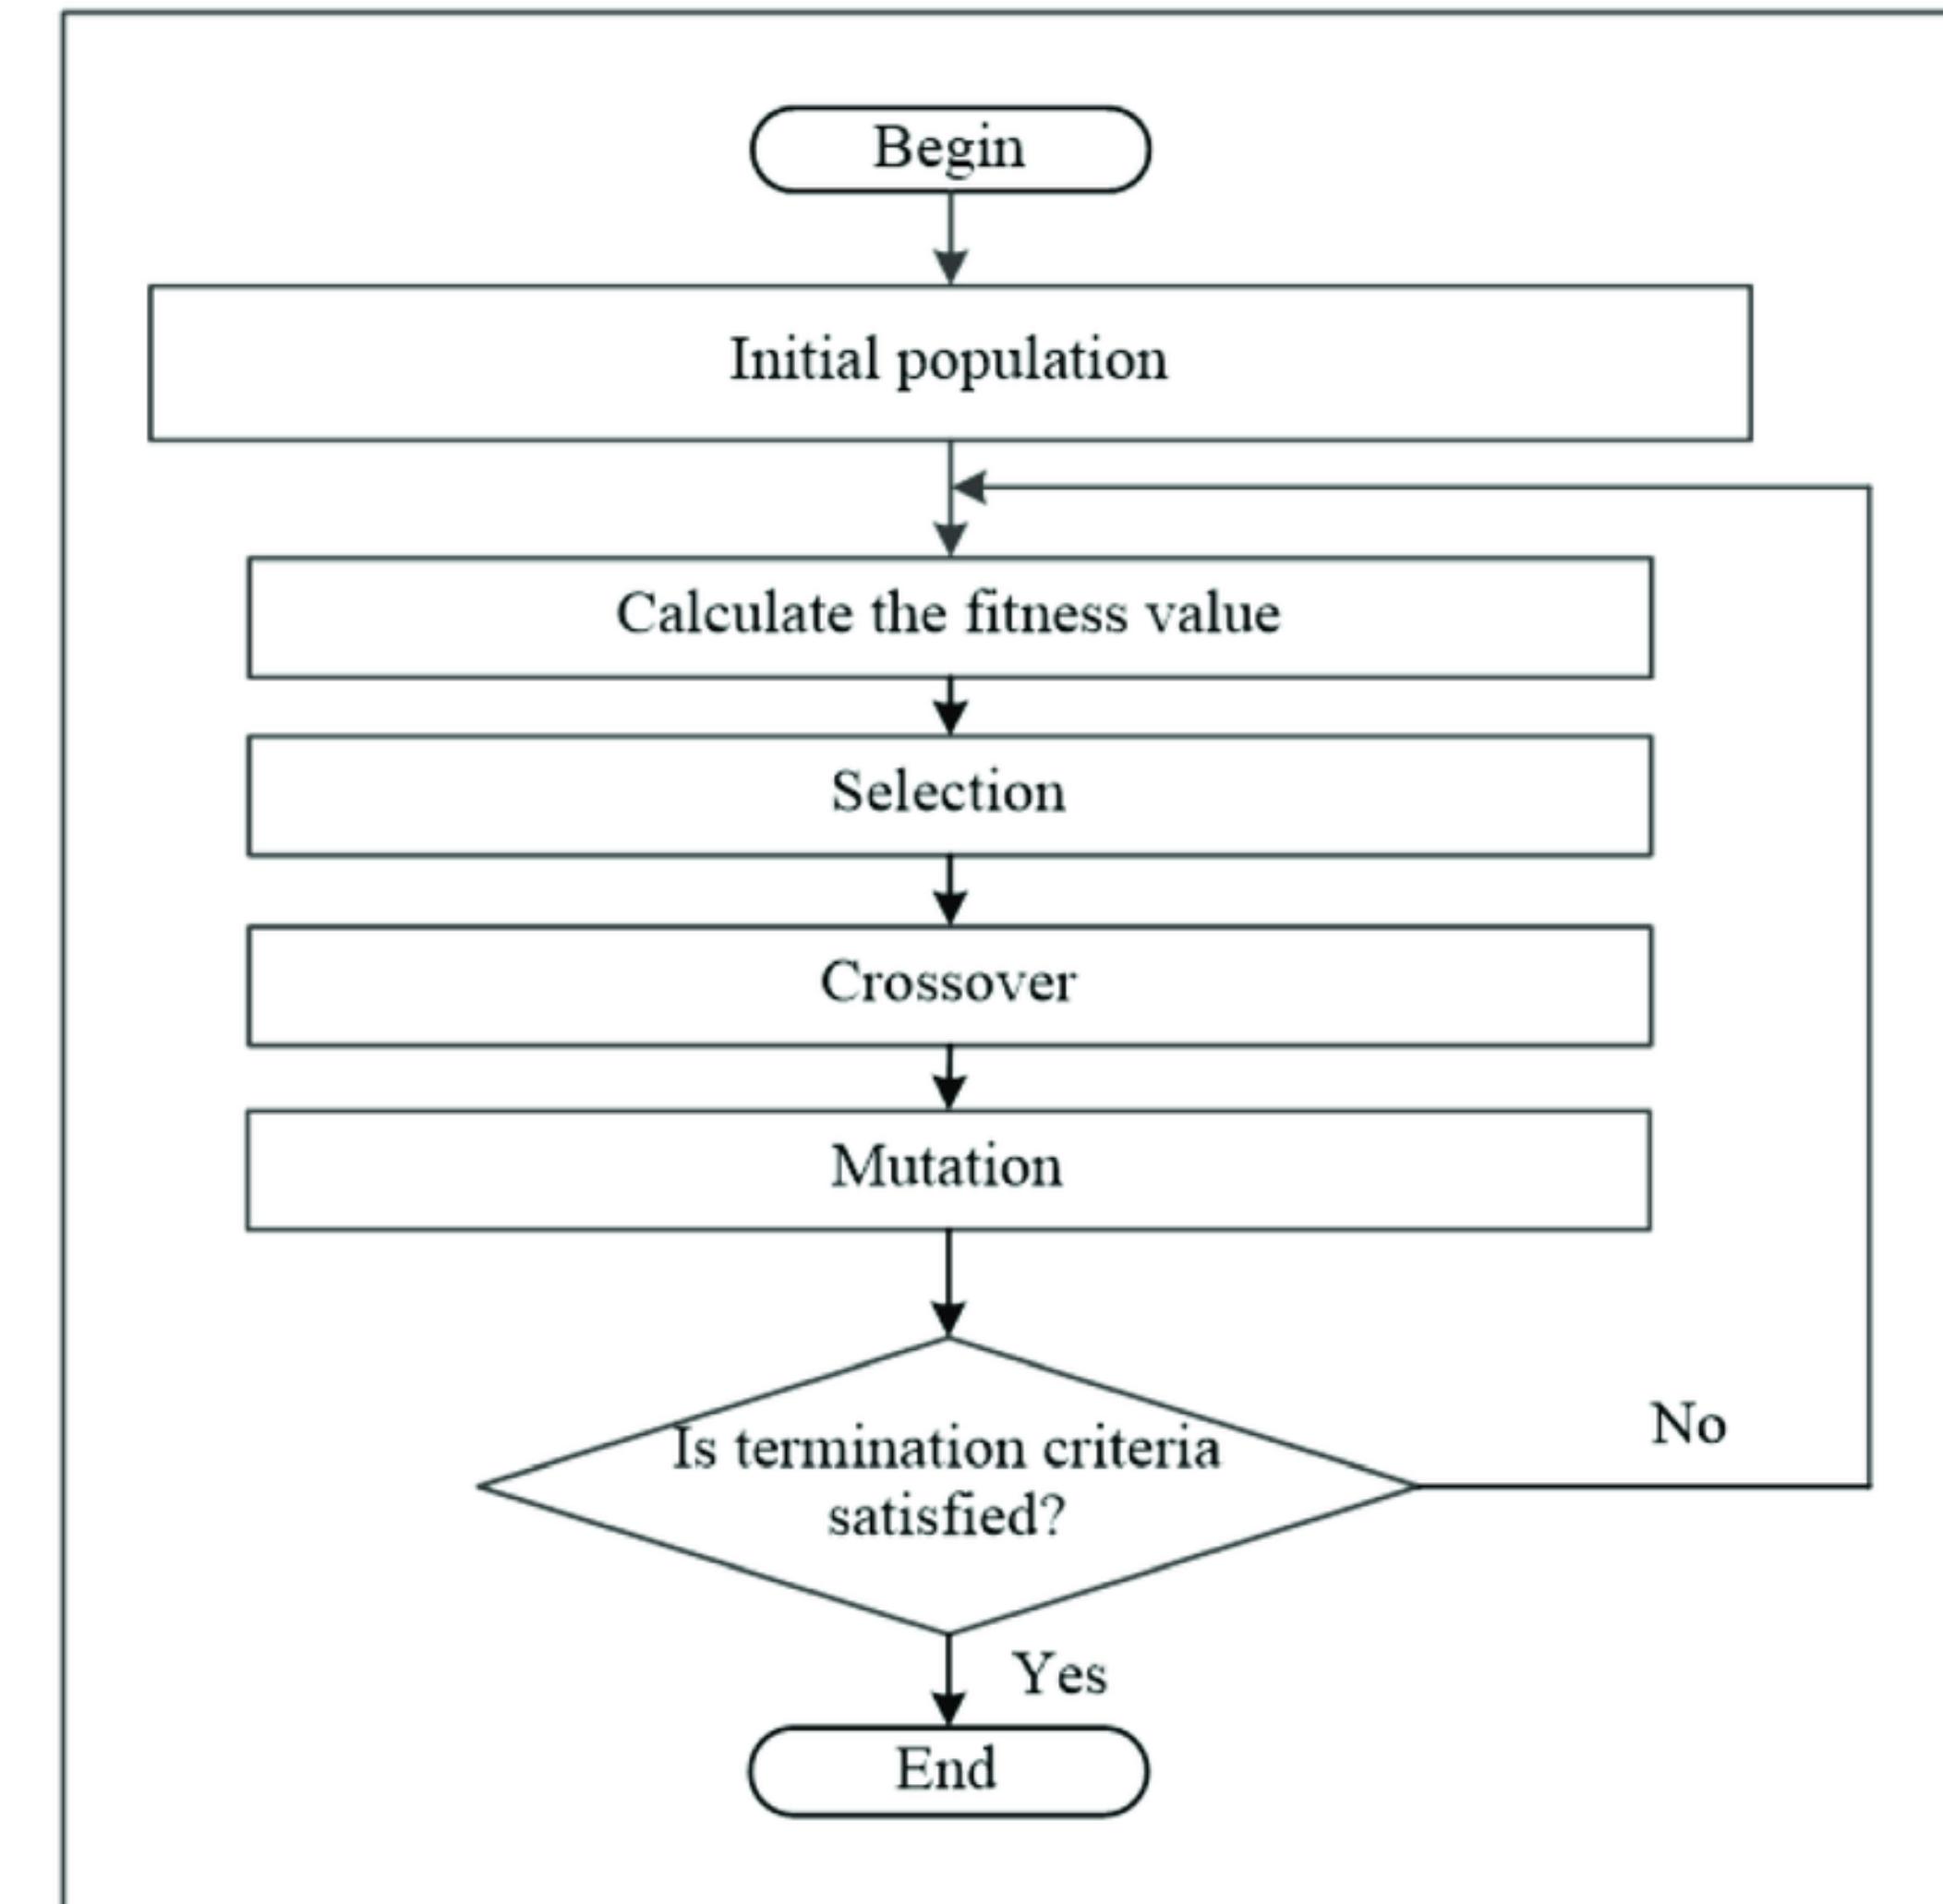
---

## LOCAL BEAM SEARCH

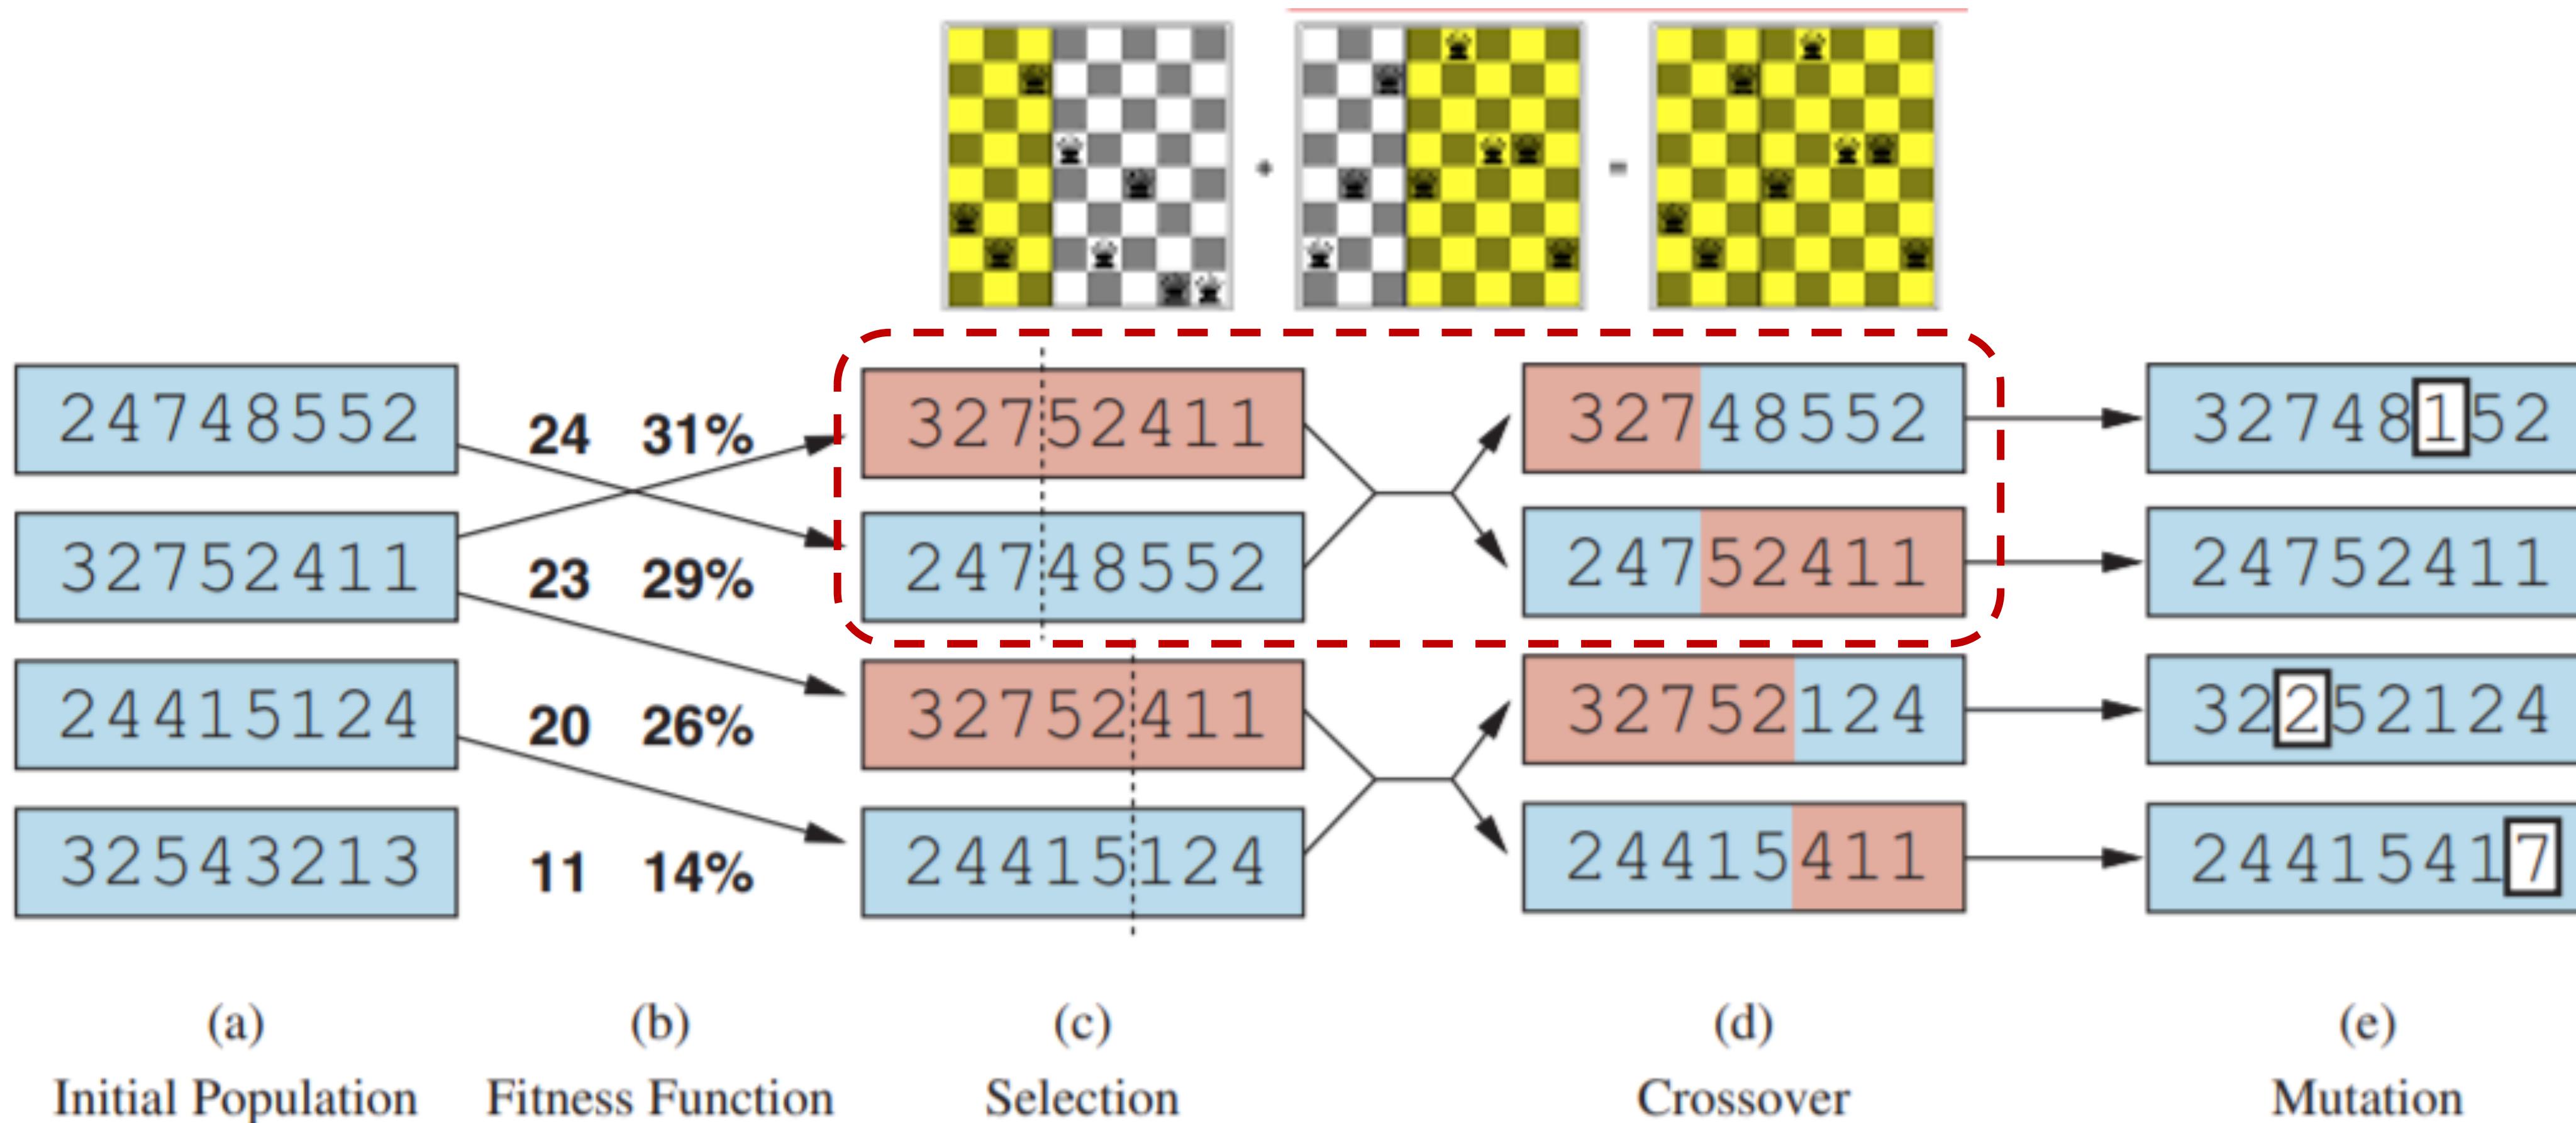
- ▶ Similar to hill-climbing, but...
  - ▶ Keeps track of  $k$  current states rather than just a single current state
  - ▶ Select the  $k$  best neighbors among all neighbors of the  $k$  current states
- ▶ Analogy
  - ▶ Hill-climbing: "...trying to find the top of Mt. Everest in a thick fog while suffering amnesia"
  - ▶ Local beam search: "Doing this with several friends, each of whom has a short-range radio and an altimeter"
  - ▶ Stochastic beam search – Select successors at random weighted by value

# GENETIC ALGORITHM

- ▶ Inspired by evolutionary biology
- ▶ Mimic the evolution of a population under natural selection: the fittest (i.e., the ones with better objective function values) survives



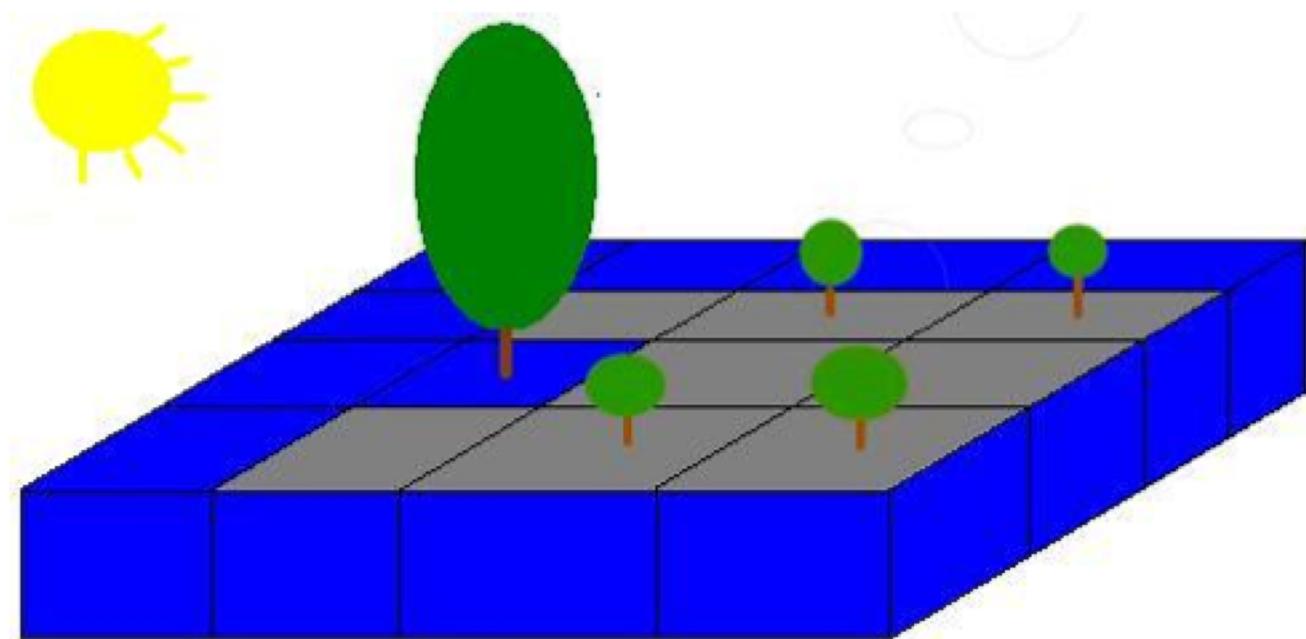
# GENETIC ALGORITHM EXAMPLE: EIGHT-QUEENS PROBLEM



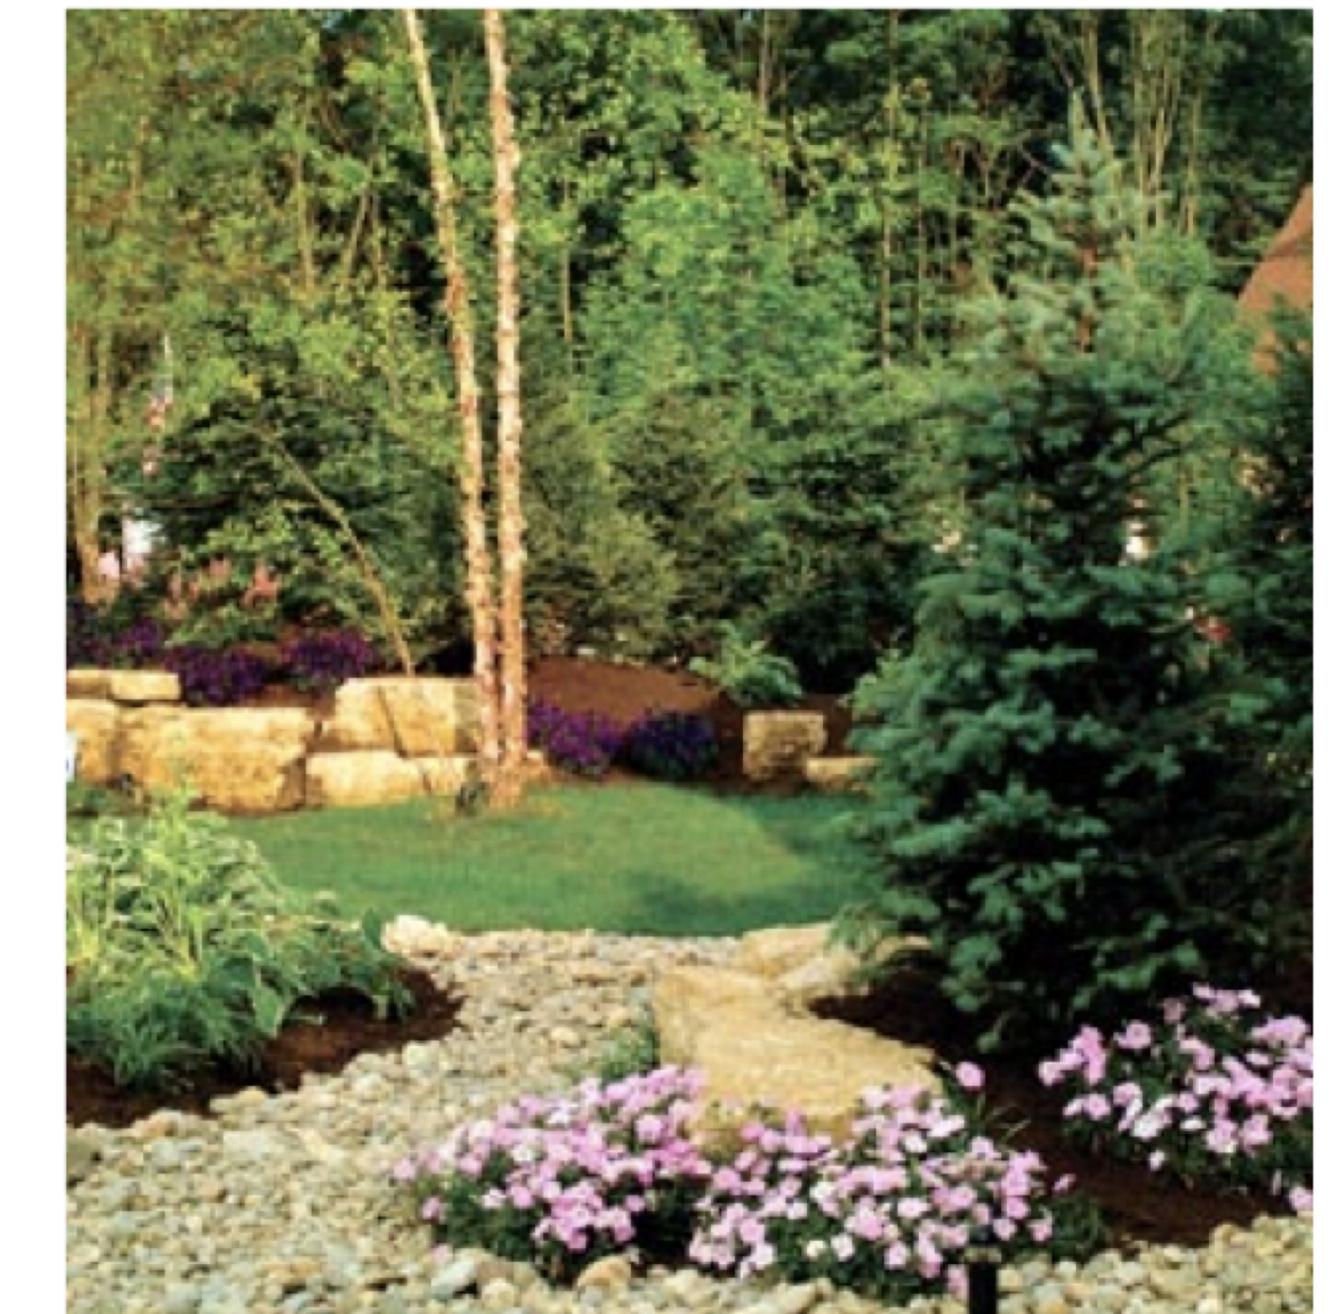
**Fitness function:** number of pairs of  
non-attacking queens

# APPLICATION: OPTIMIZING PLANT PLACEMENT TO SAVE WATER

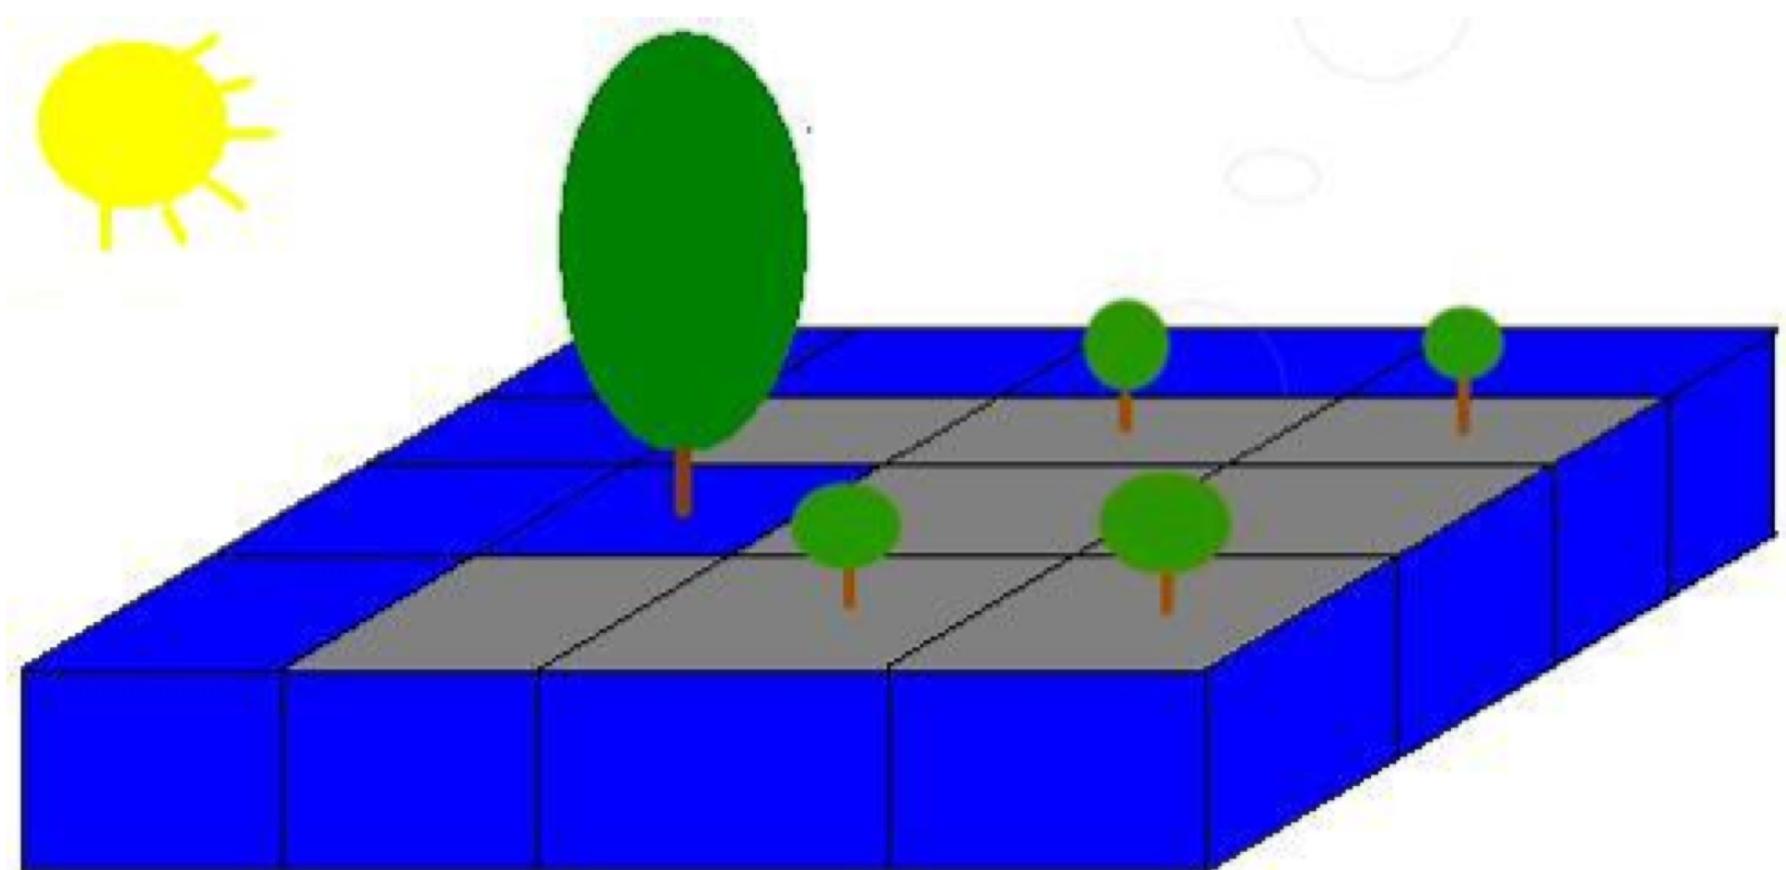
- ▶ Species have differences in size, light requirement, and water use
- ▶ Complex interactions between species, e.g., shading changes water requirements



Facilitation: when a “nurse plant” improves conditions for other plants’ growth



## APPLICATION: OPTIMIZING PLANT PLACEMENT TO SAVE WATER



$$\text{landscape\_fitness} = \frac{1}{n} \sum_{i=1}^n \text{fitness}_i$$

$$\text{fitness}_i = \alpha * \text{growth}_i + \beta * \text{water\_efficiency}_i$$

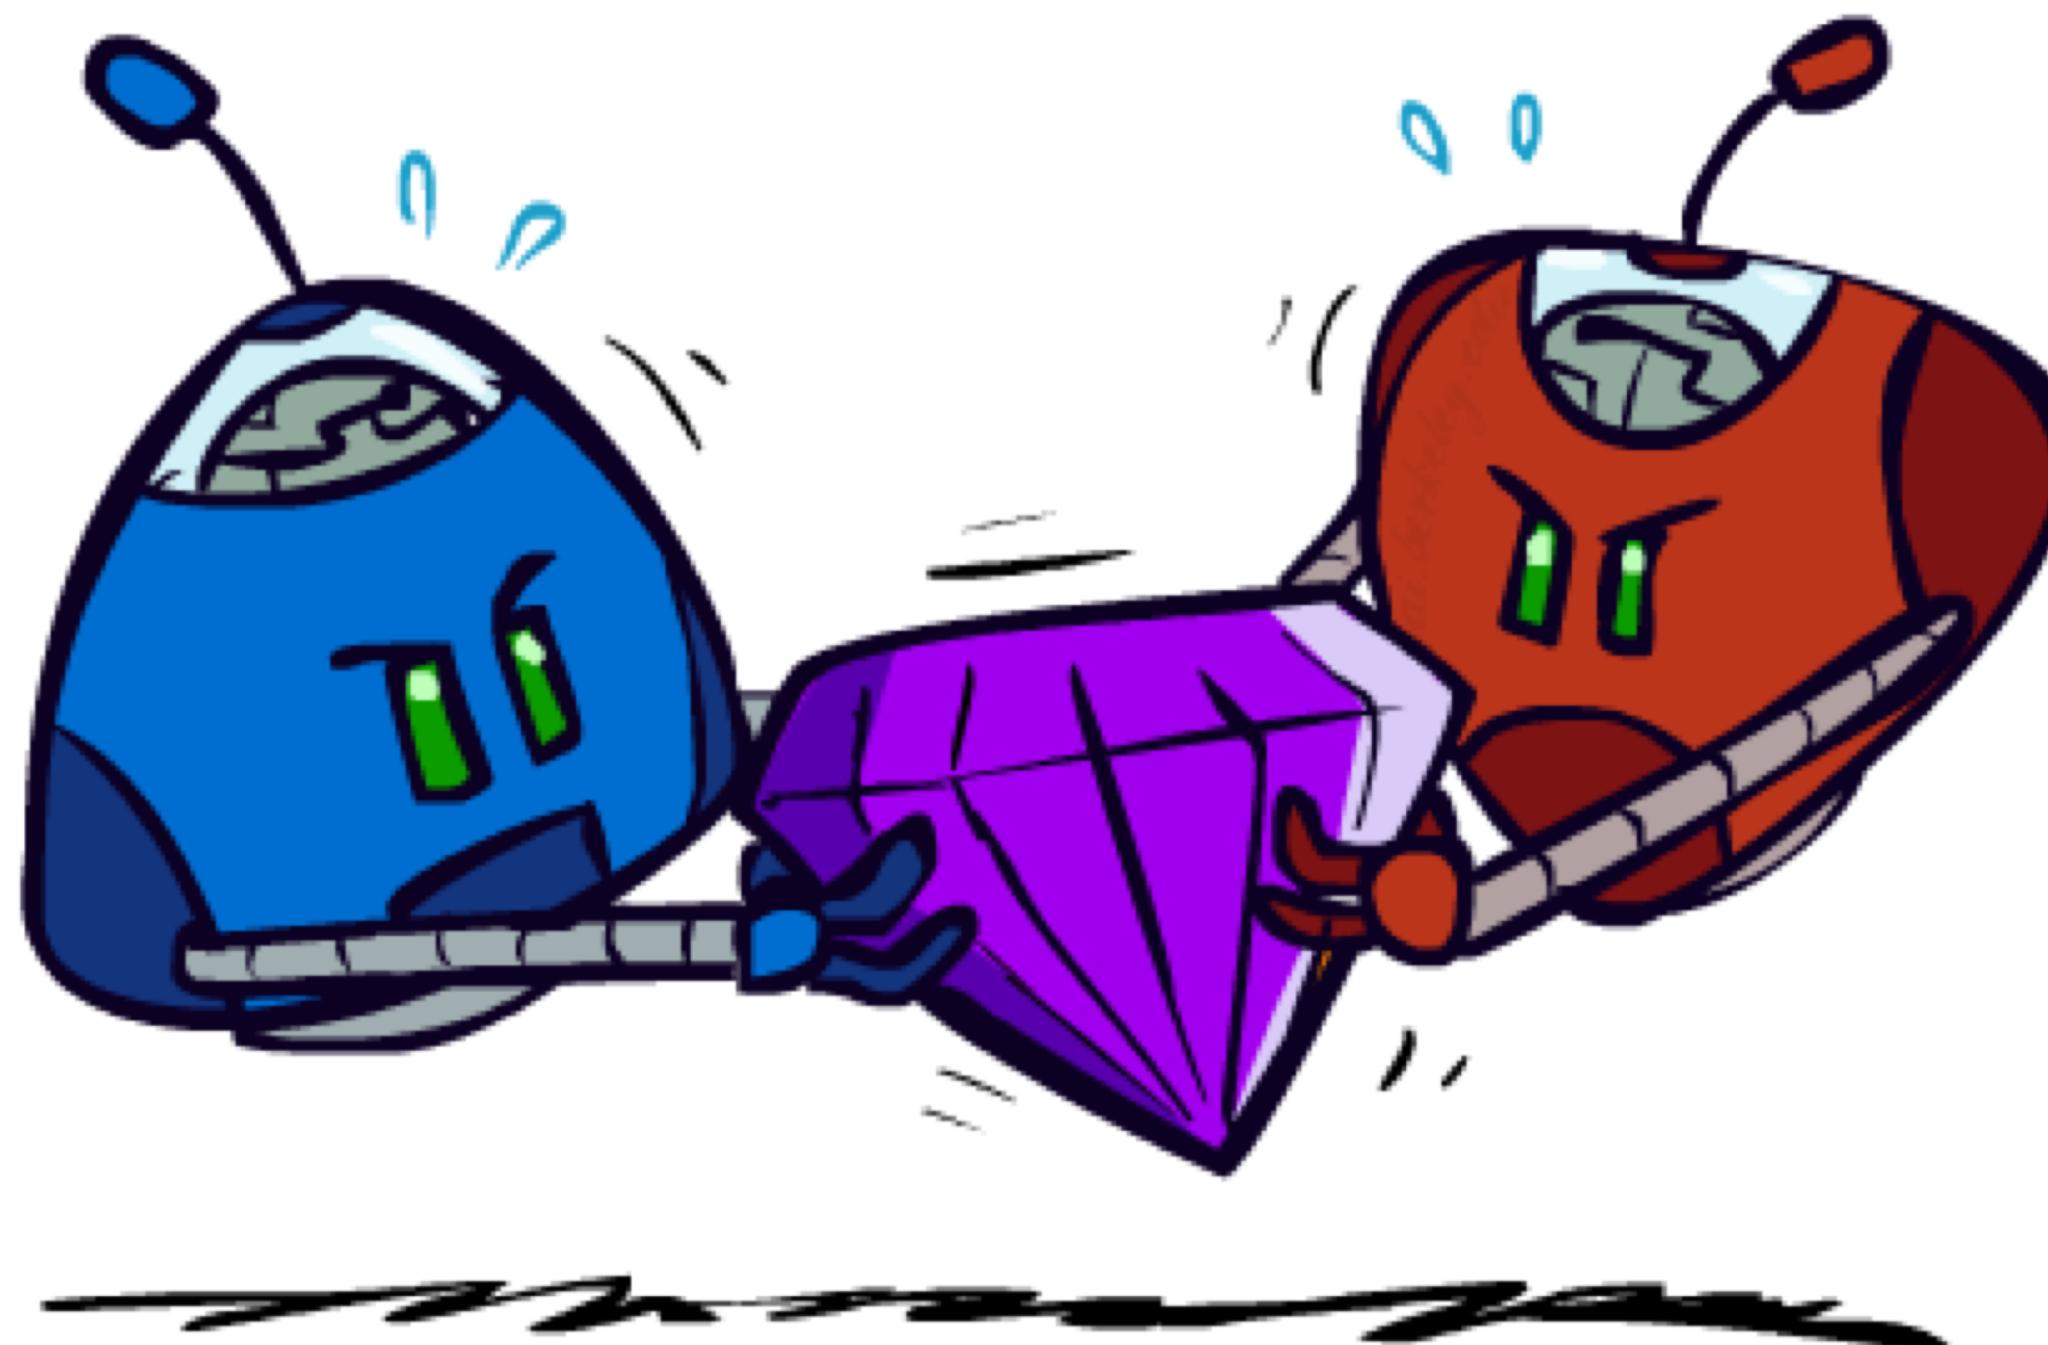
$$\text{growth}_i = f(\text{lightReq}_i, \text{waterReq}_i, \text{light}_{(x,y)}, \text{water}_{(x,y)})$$

$$\text{water\_efficiency} = g\left(\sum_{x,y} \text{water}_{(x,y)}\right)$$

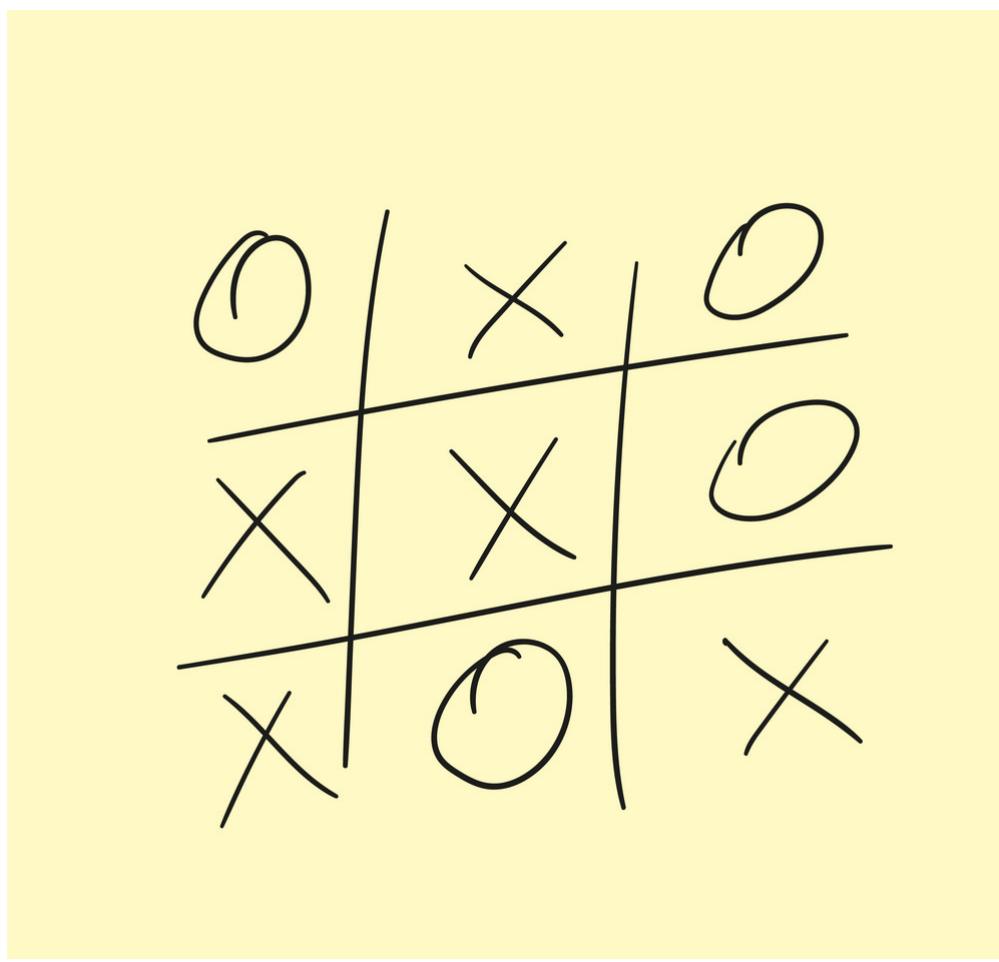
Simulated annealing: Grow  $k$  plants in a  $N \times M$  landscape while determine a watering plan to maximize the landscape fitness score

---

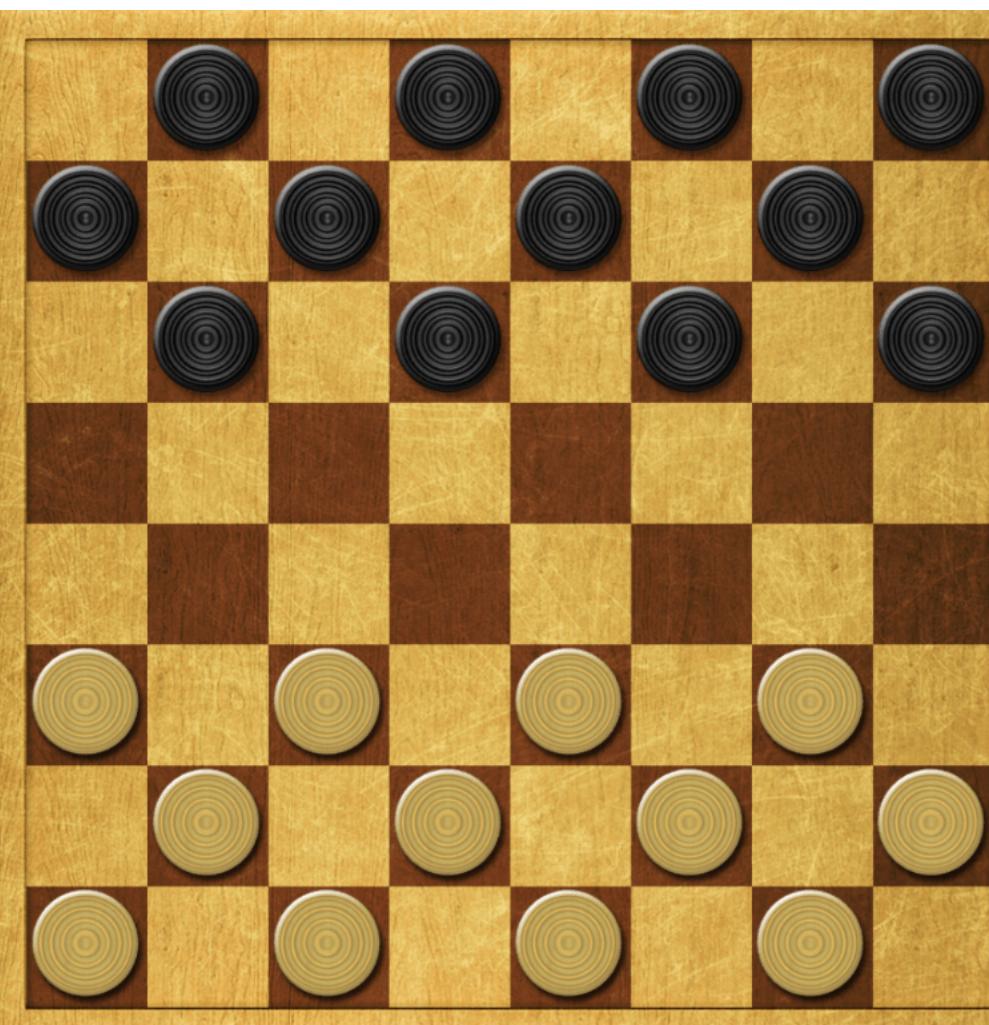
# ADVERSARIAL GAMES



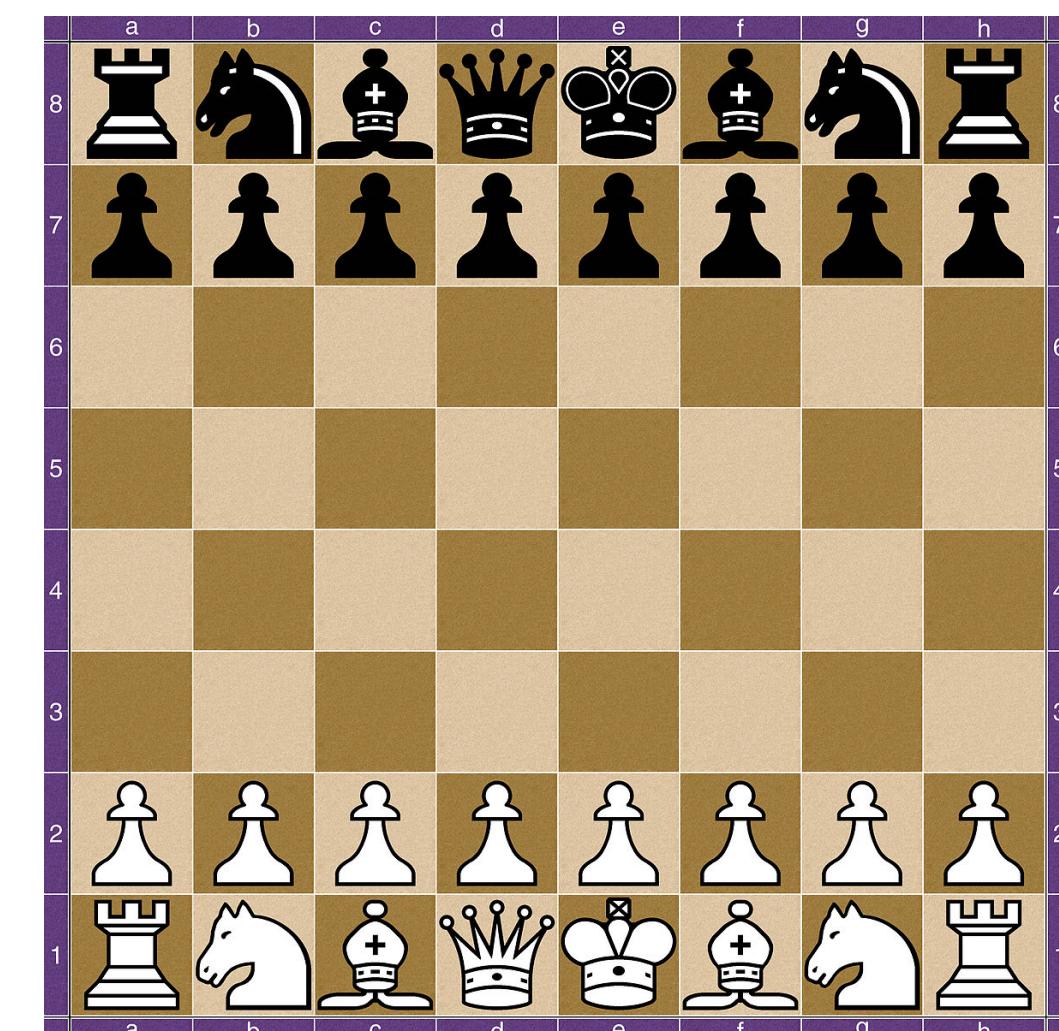
# GAMES ARE EVERYWHERE



Tic-Tac-Toe



Checkers

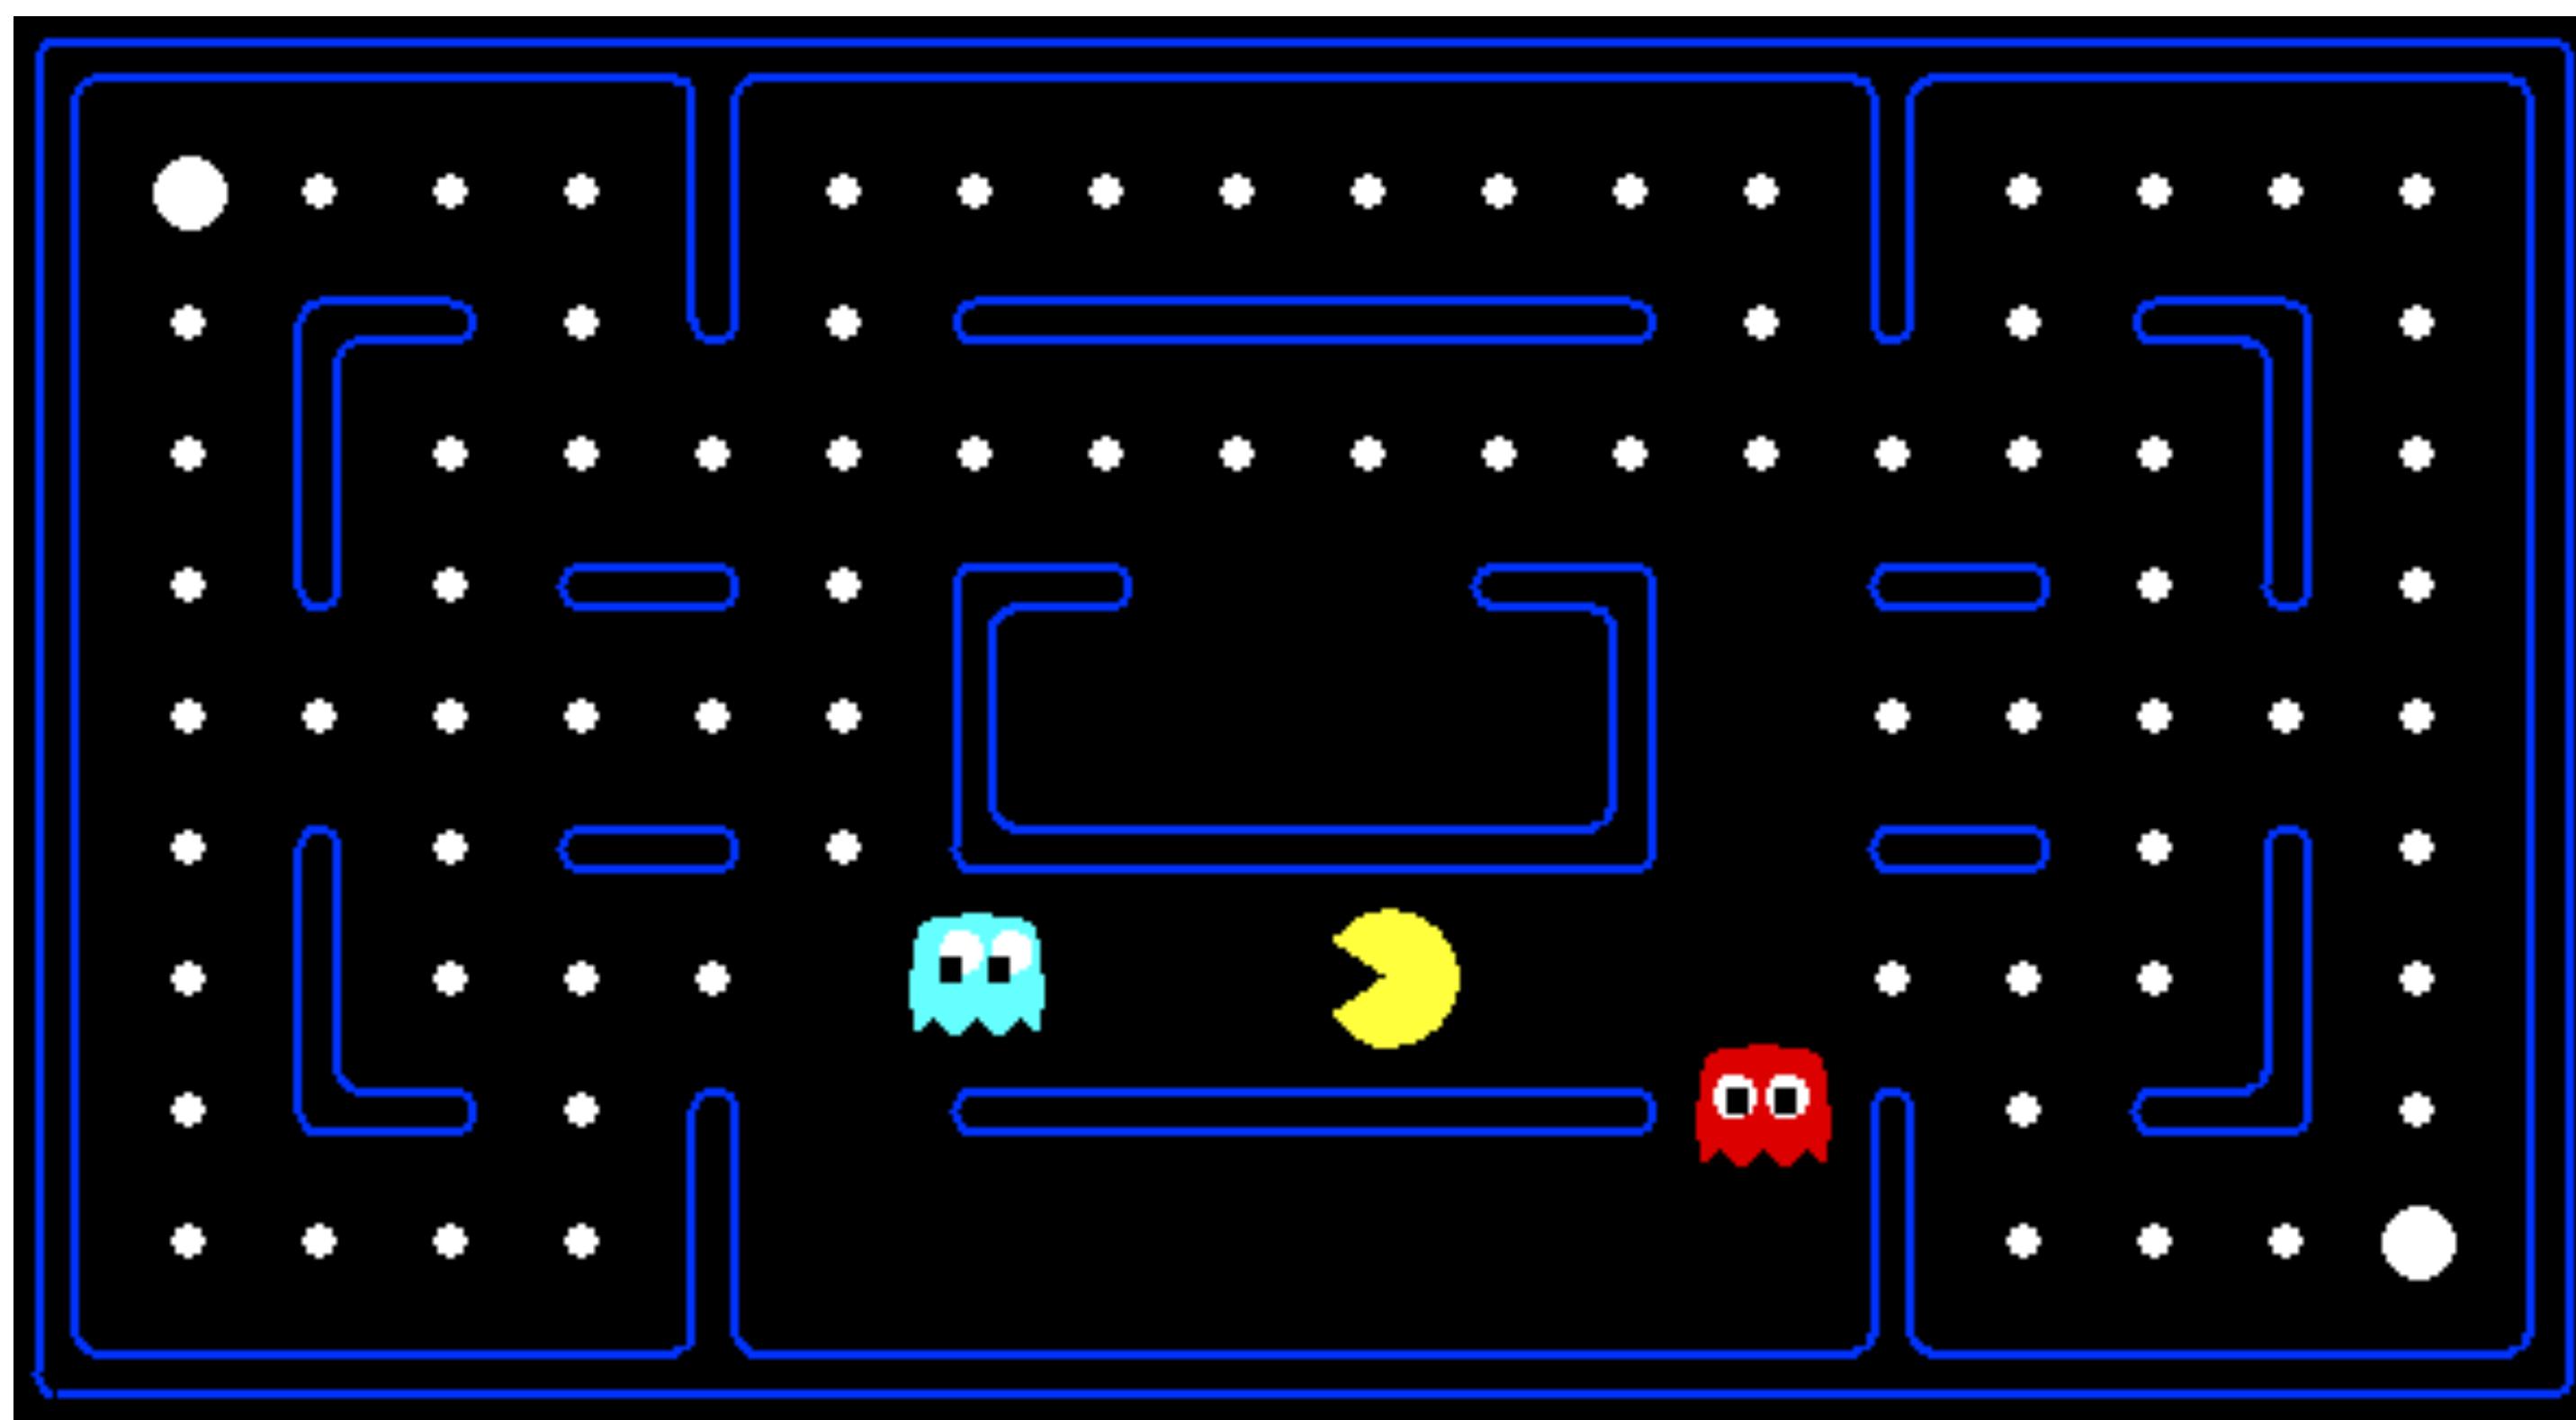


Chess



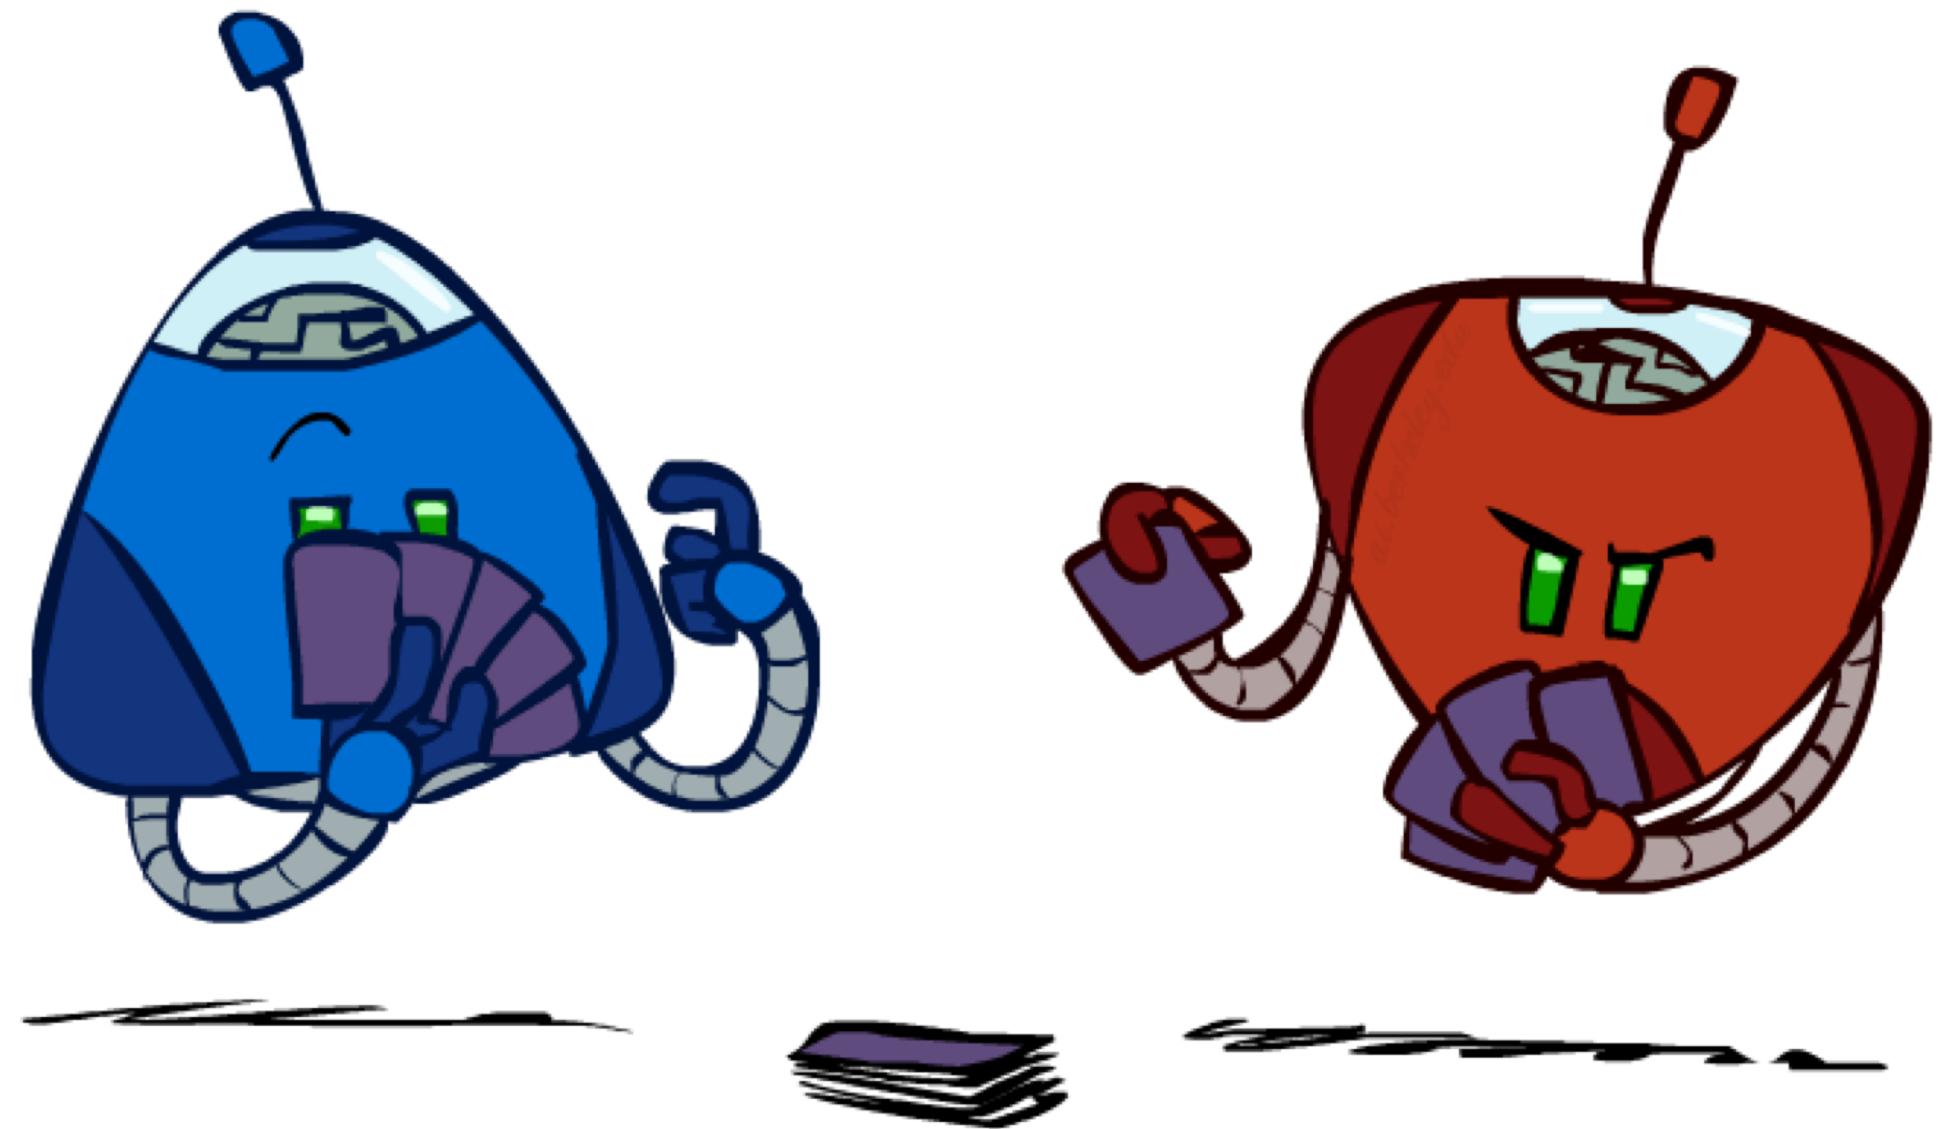
Go

# BEHAVIOR FROM COMPUTATION



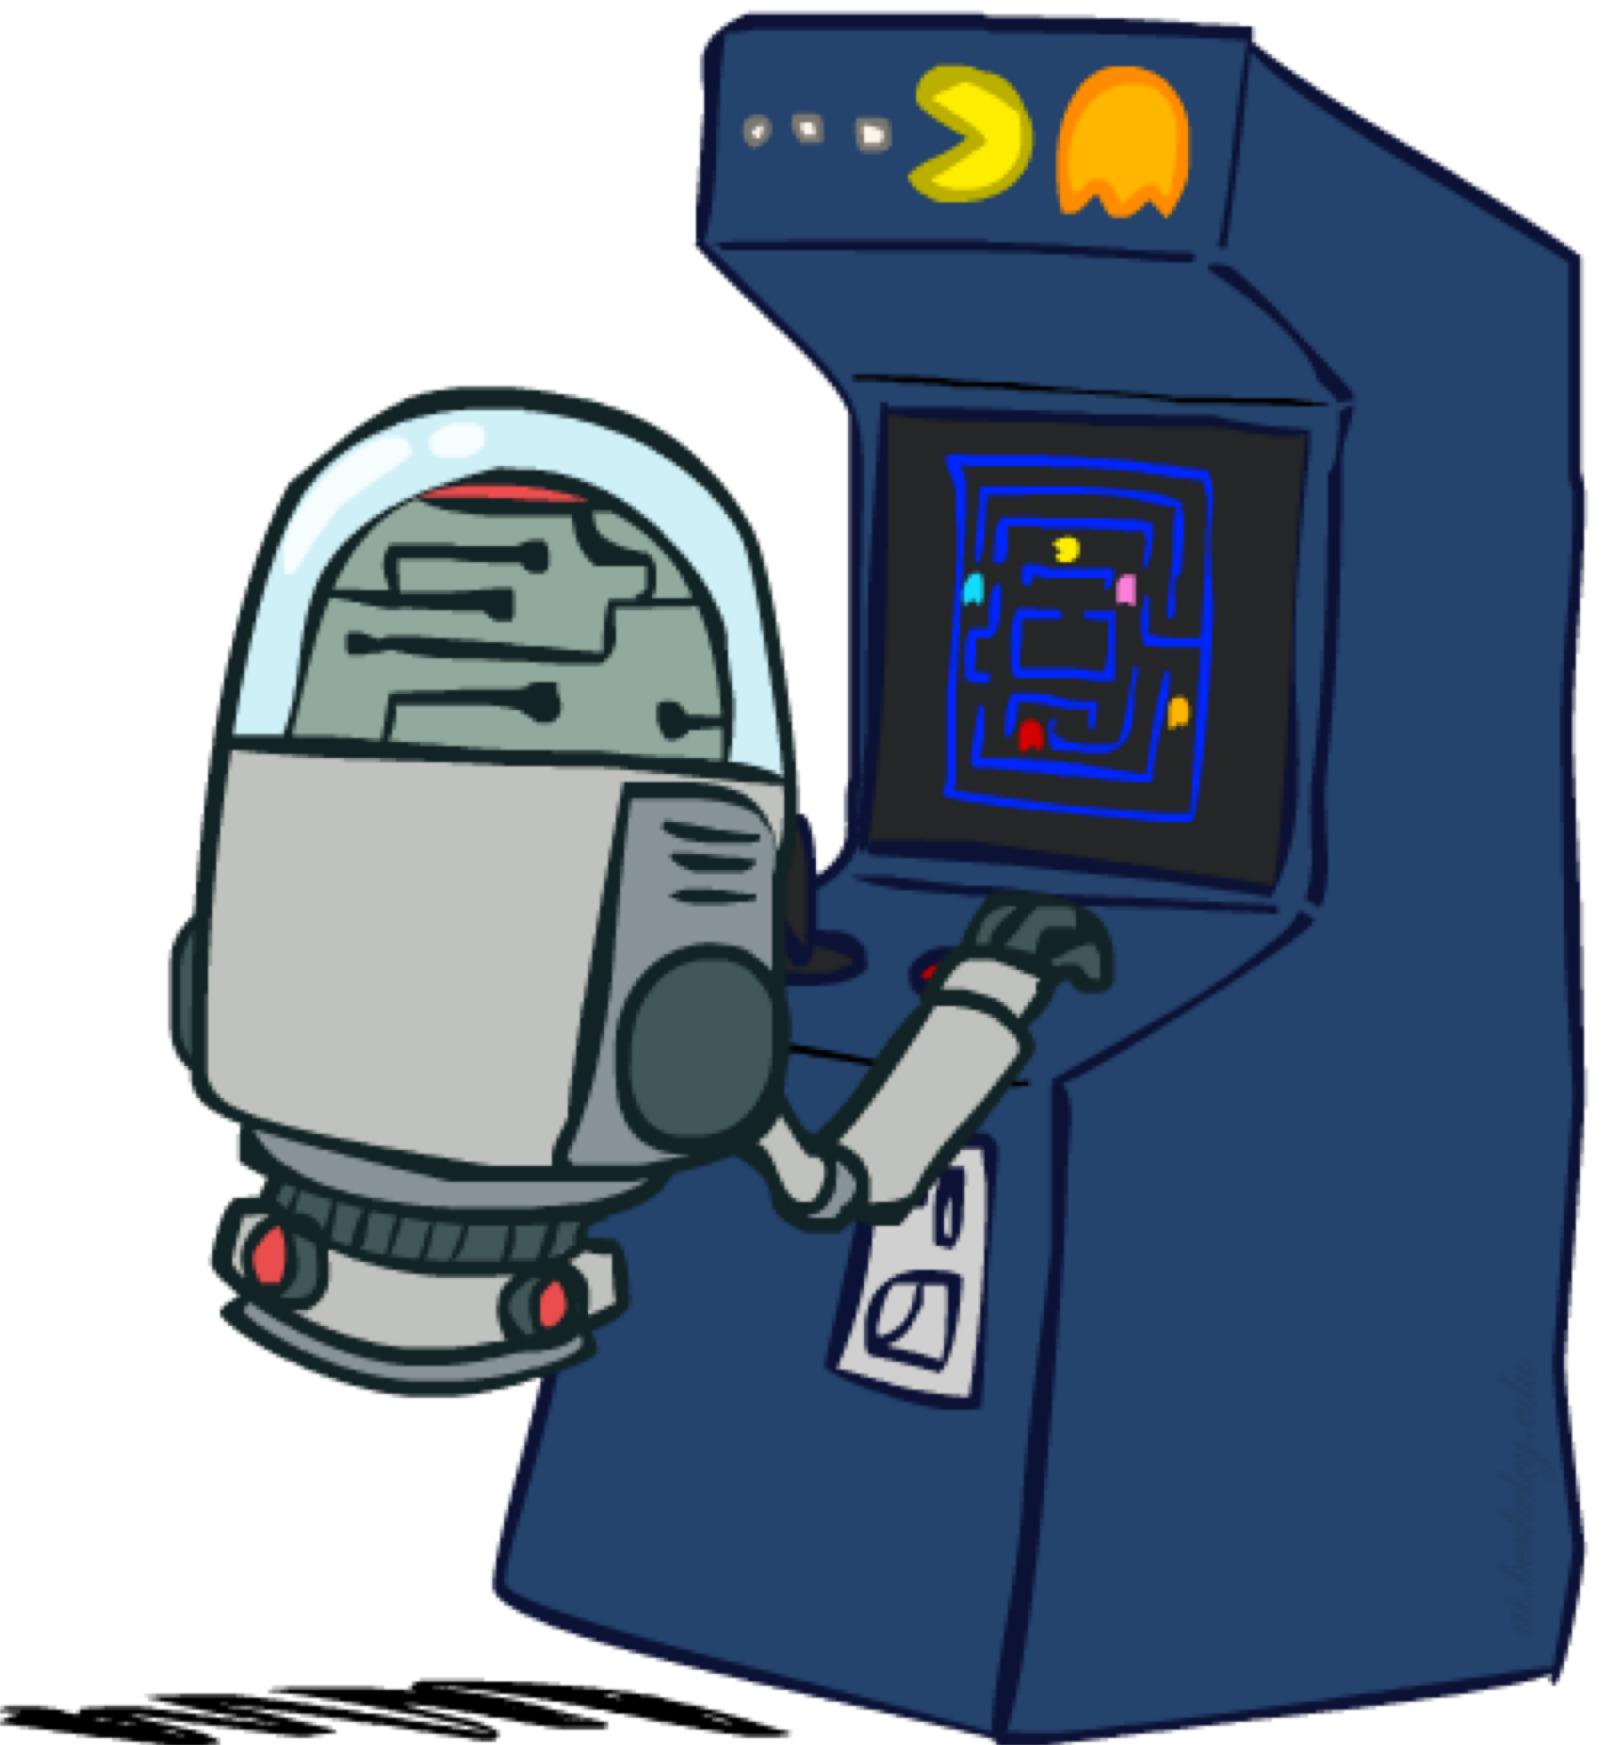
## TYPES OF GAMES

- ▶ Many different kinds of games
- ▶ Axes:
  - ▶ Deterministic vs. stochastic
  - ▶ One, two or more players
  - ▶ Zero sum vs. general sum
  - ▶ Perfect information (can you see the state) vs. partial information
- ▶ Algorithms need to calculate a strategy (**policy**) which recommends a move from each state

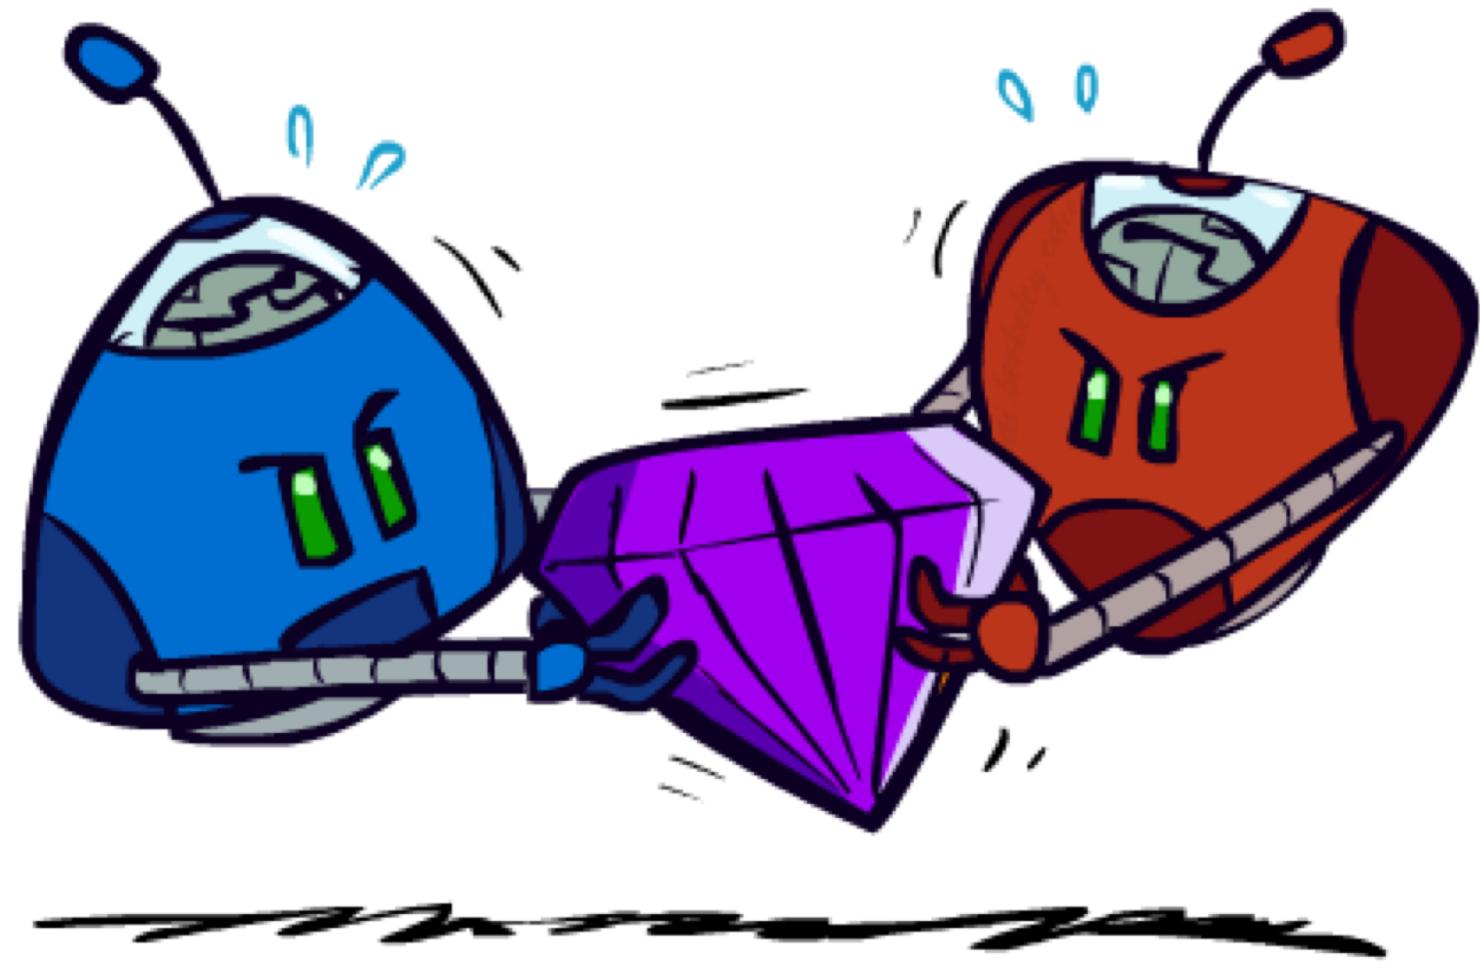


# DETERMINISTIC GAMES

- ▶ Problem formulation:
  - ▶ States:  $S$  (start at  $s_0$ )
  - ▶ Players:  $P=\{1\dots N\}$  (usually take turns)
  - ▶ Actions:  $A$  (may depend on player / state)
  - ▶ Transition Function:  $S \times A \rightarrow S$
  - ▶ Terminal Test:  $S \rightarrow \{t,f\}$
  - ▶ Terminal Utilities:  $S \times P \rightarrow R$
  - ▶ Solution for a player is a **policy**:  $S \rightarrow A$



# ZERO-SUM VS. GENERAL GAMES



## ► Zero-Sum Games

- ▶ Agents have opposite *utilities* (values on outcomes)
- ▶ Can then think of outcome as a single value that one maximizes and the other minimizes
- ▶ Adversarial, pure competition

## ► General Games

- ▶ Agents have independent *utilities* (values on outcomes)
- ▶ Cooperation, indifference, competition, and more are all possible
- ▶ More later on non-zero-sum games