

PURDUE CS47100

---

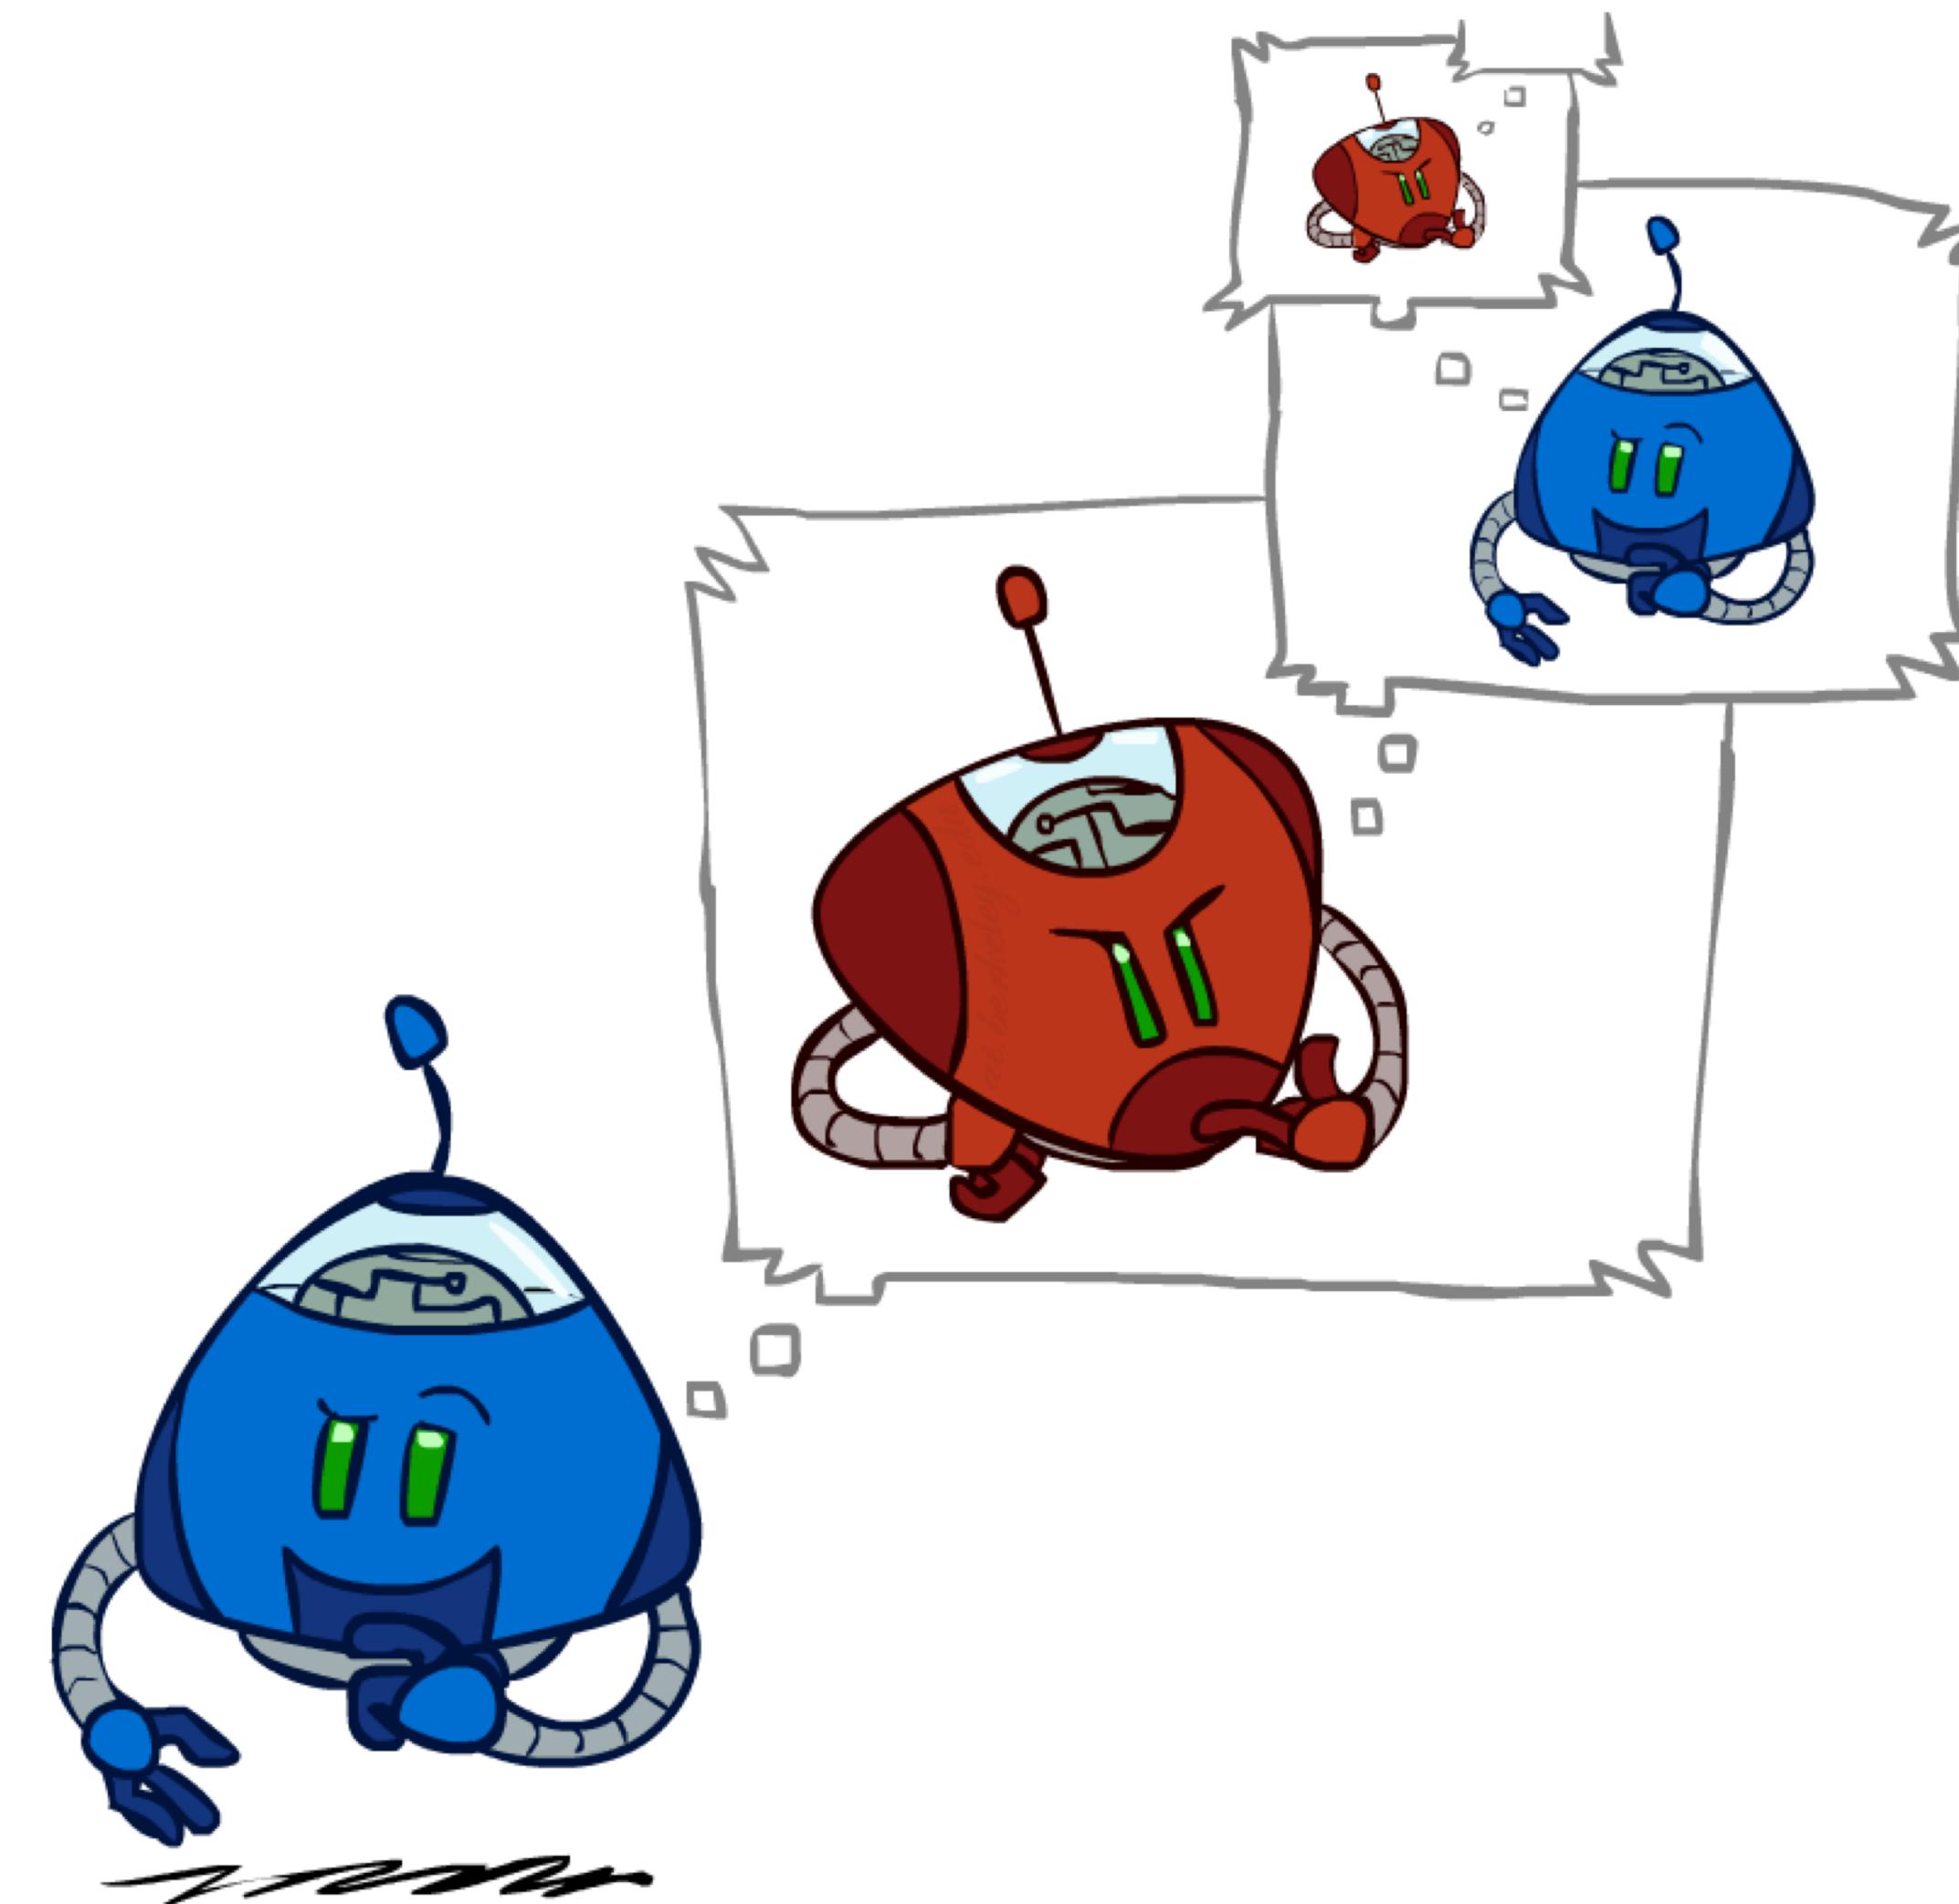
# INTRODUCTION TO AI

---

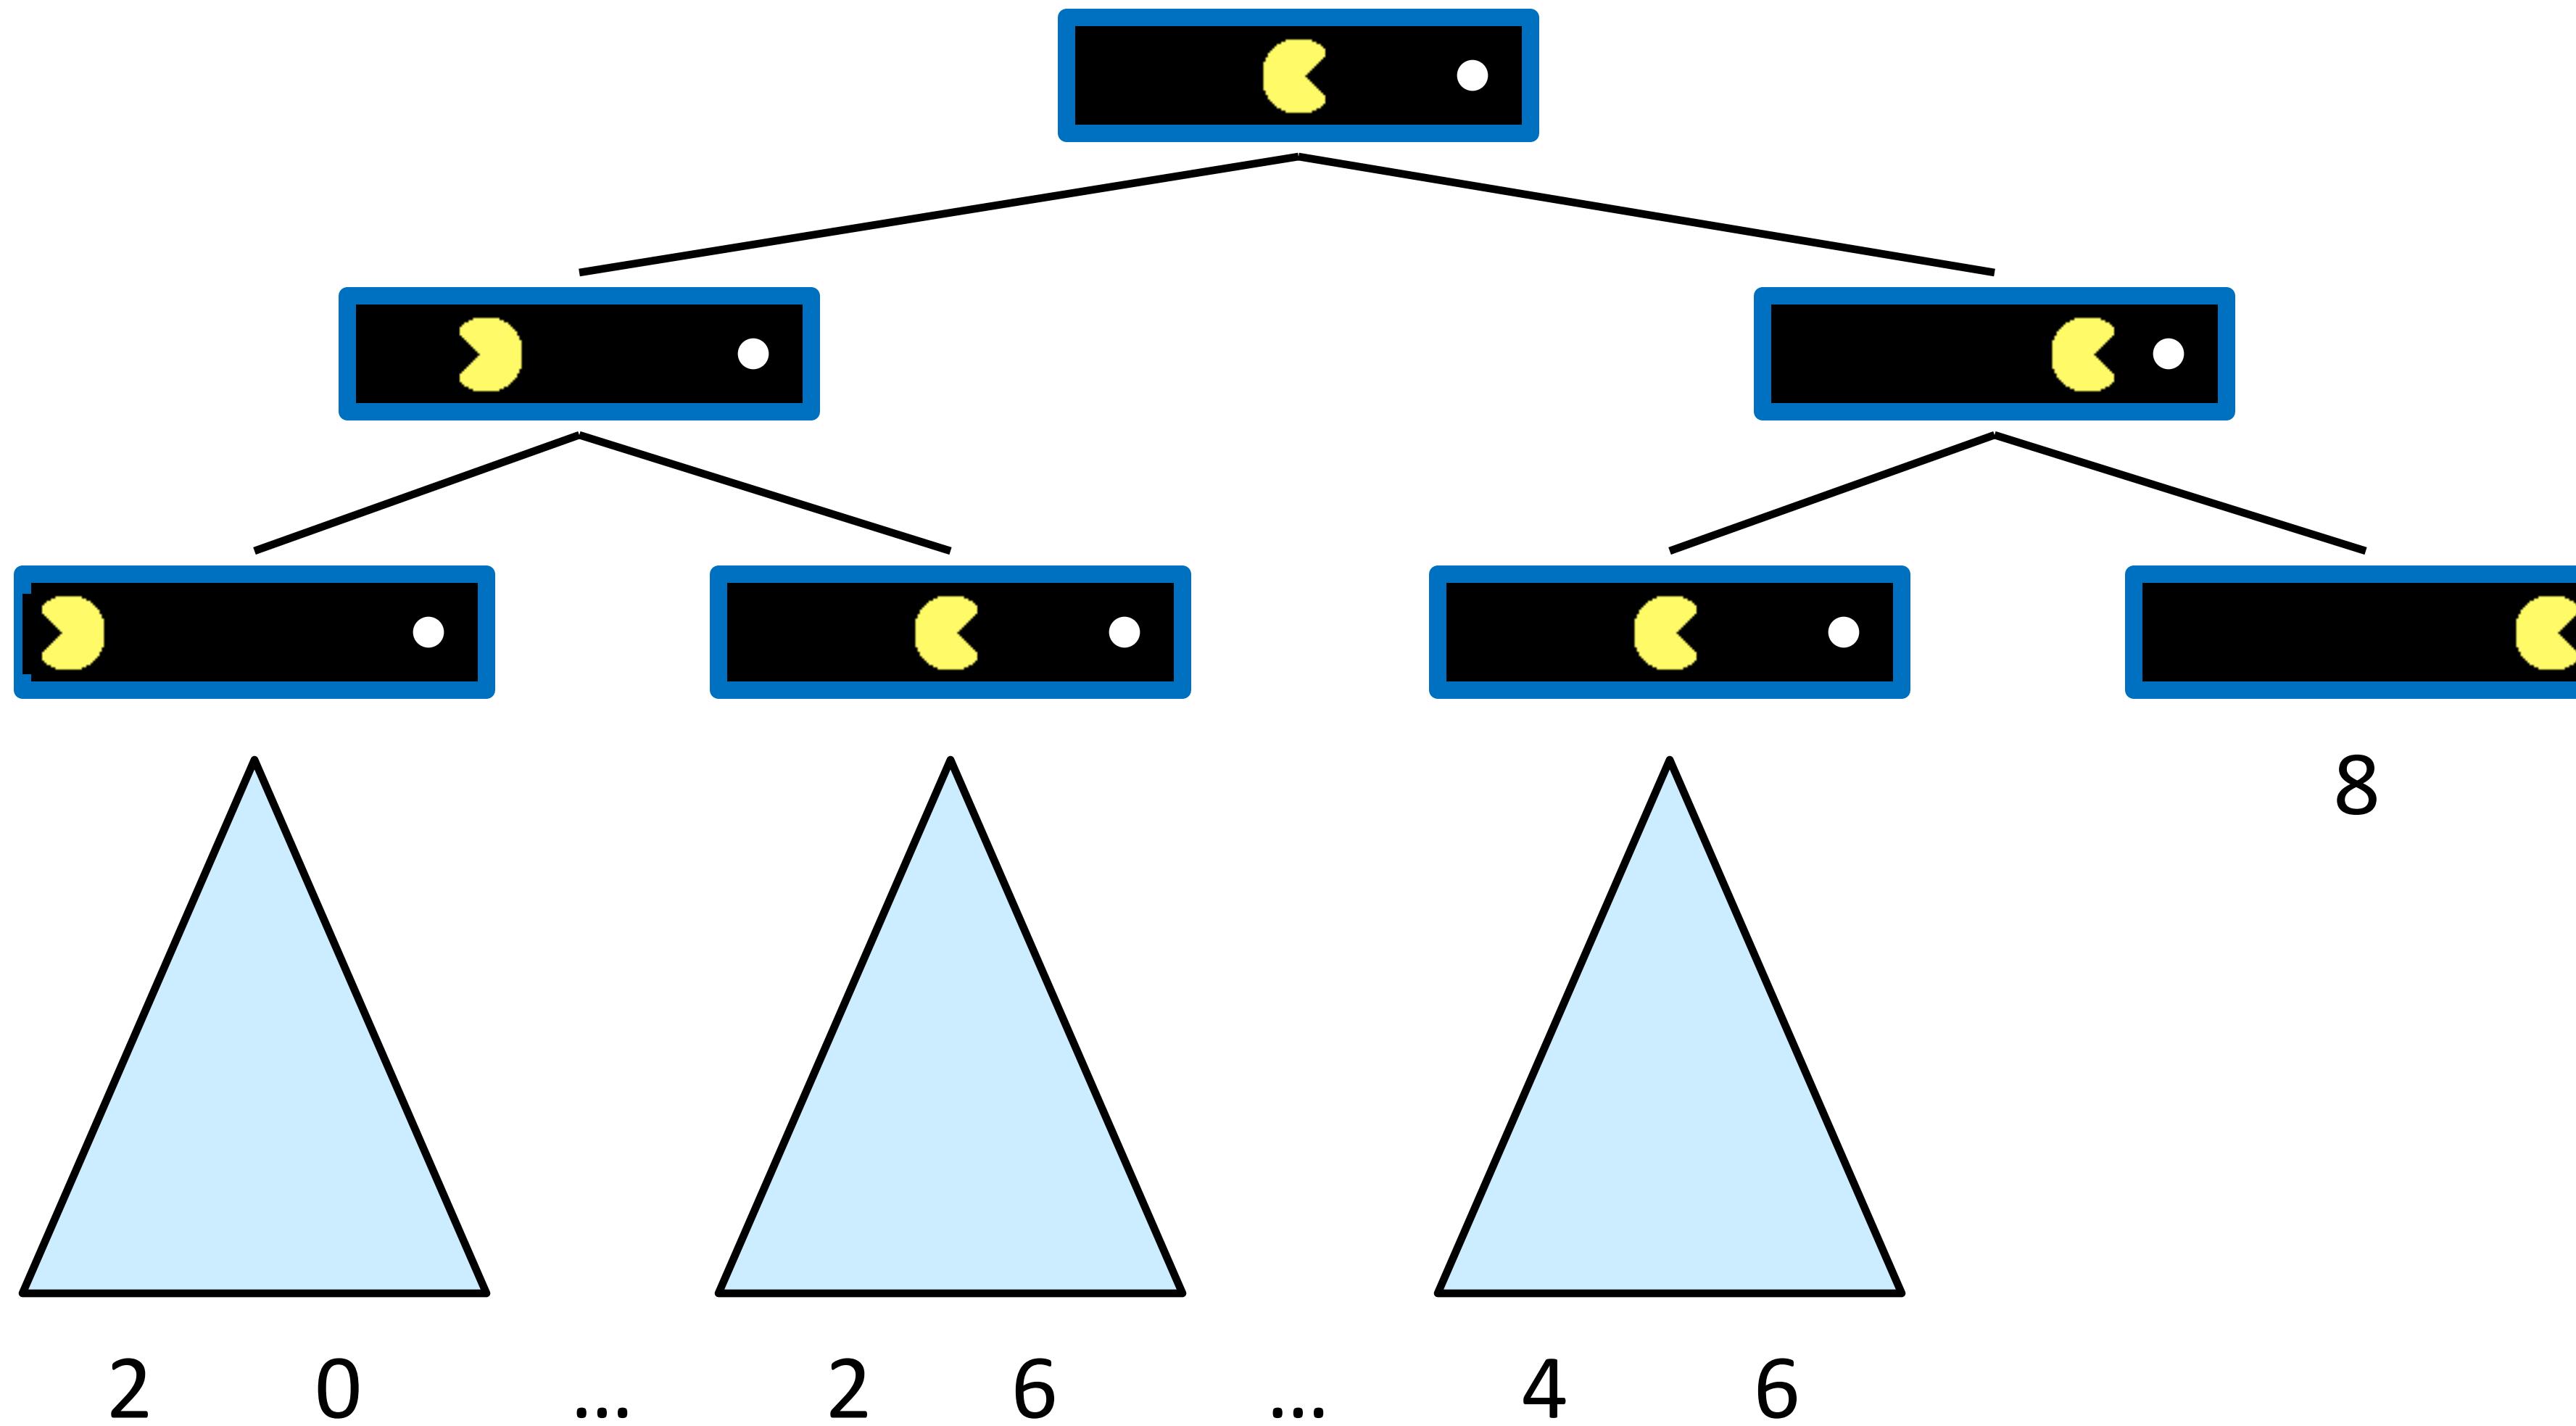
# RECAP: LOCAL SEARCH & ADVERSARIAL SEARCH

- ▶ Local search
  - ▶ Hill climbing, simulated annealing, beam search, genetic algorithm
- ▶ Adversarial search
  - ▶ States:  $S$  (start at  $s_0$ )
  - ▶ Players:  $P=\{1\dots N\}$  (usually take turns)
  - ▶ Actions:  $A$  (may depend on player / state)
  - ▶ Transition Function:  $S \times A \rightarrow S$
  - ▶ Terminal Test:  $S \rightarrow \{t,f\}$
  - ▶ Terminal Utilities:  $S \times P \rightarrow R$
  - ▶ Solution for a player is a **policy**:  $S \rightarrow A$

# ADVERSARIAL SEARCH

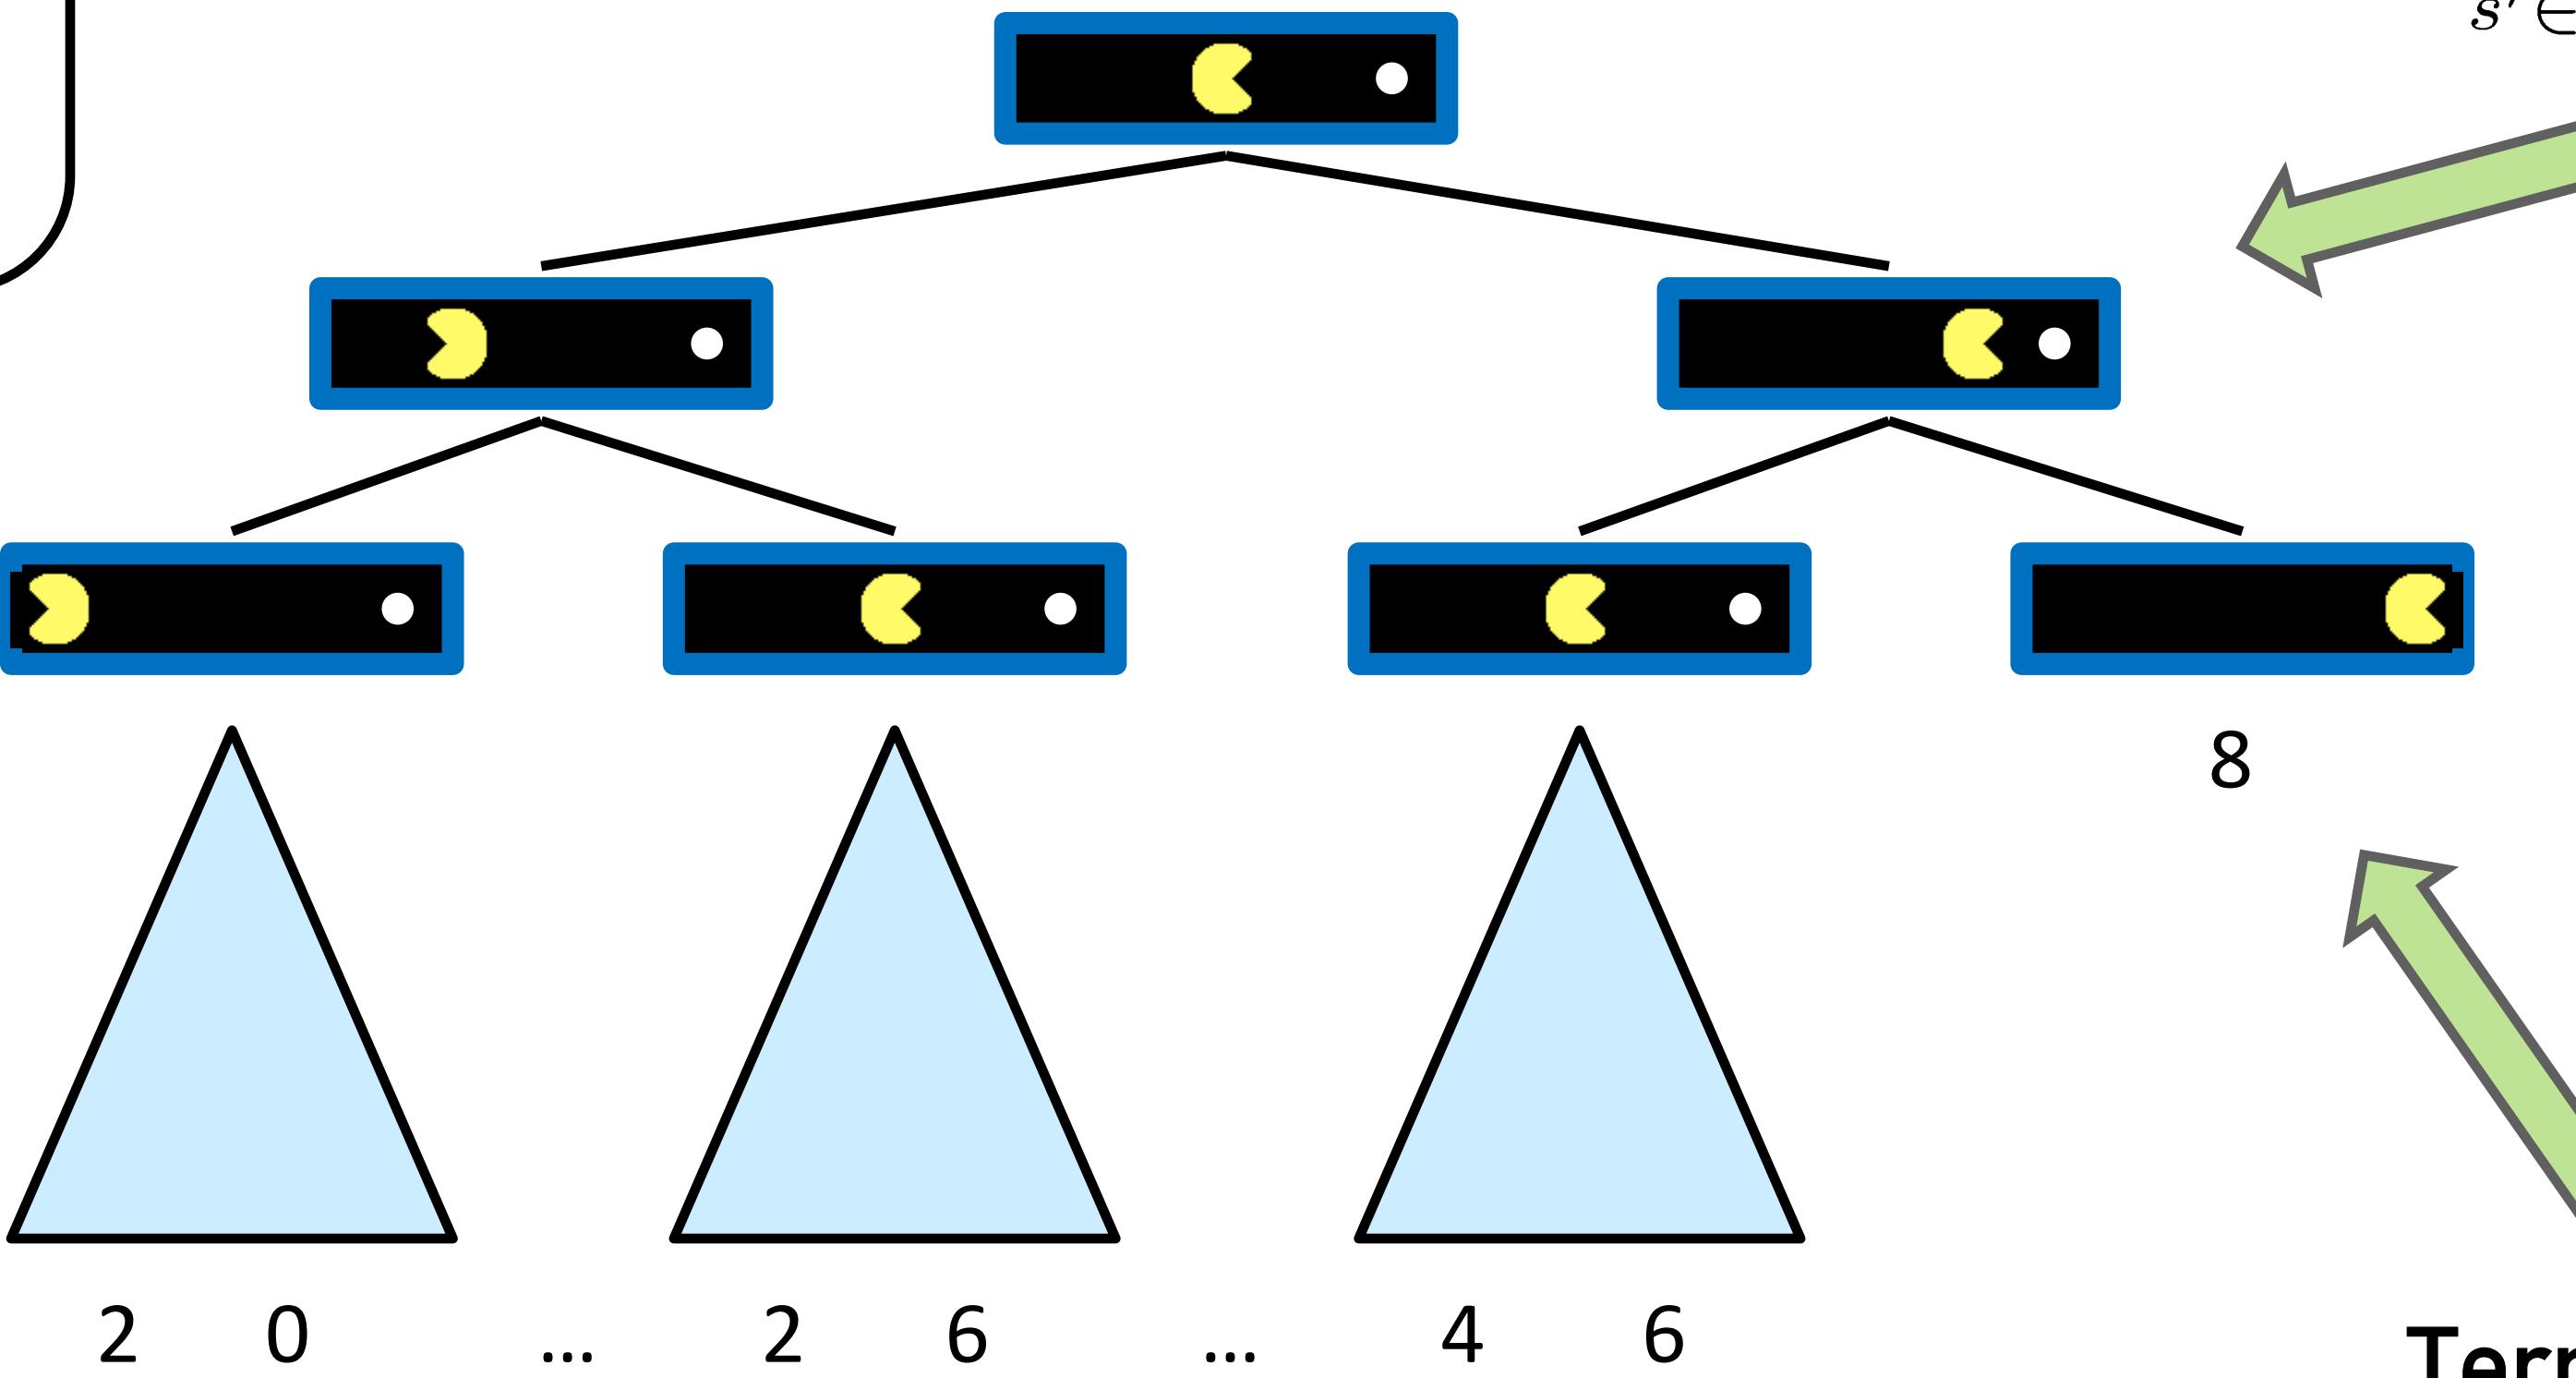


# SINGLE-AGENT TREES



# VALUE OF A STATE

**Value of a state:**  
The best achievable outcome (utility) from that state



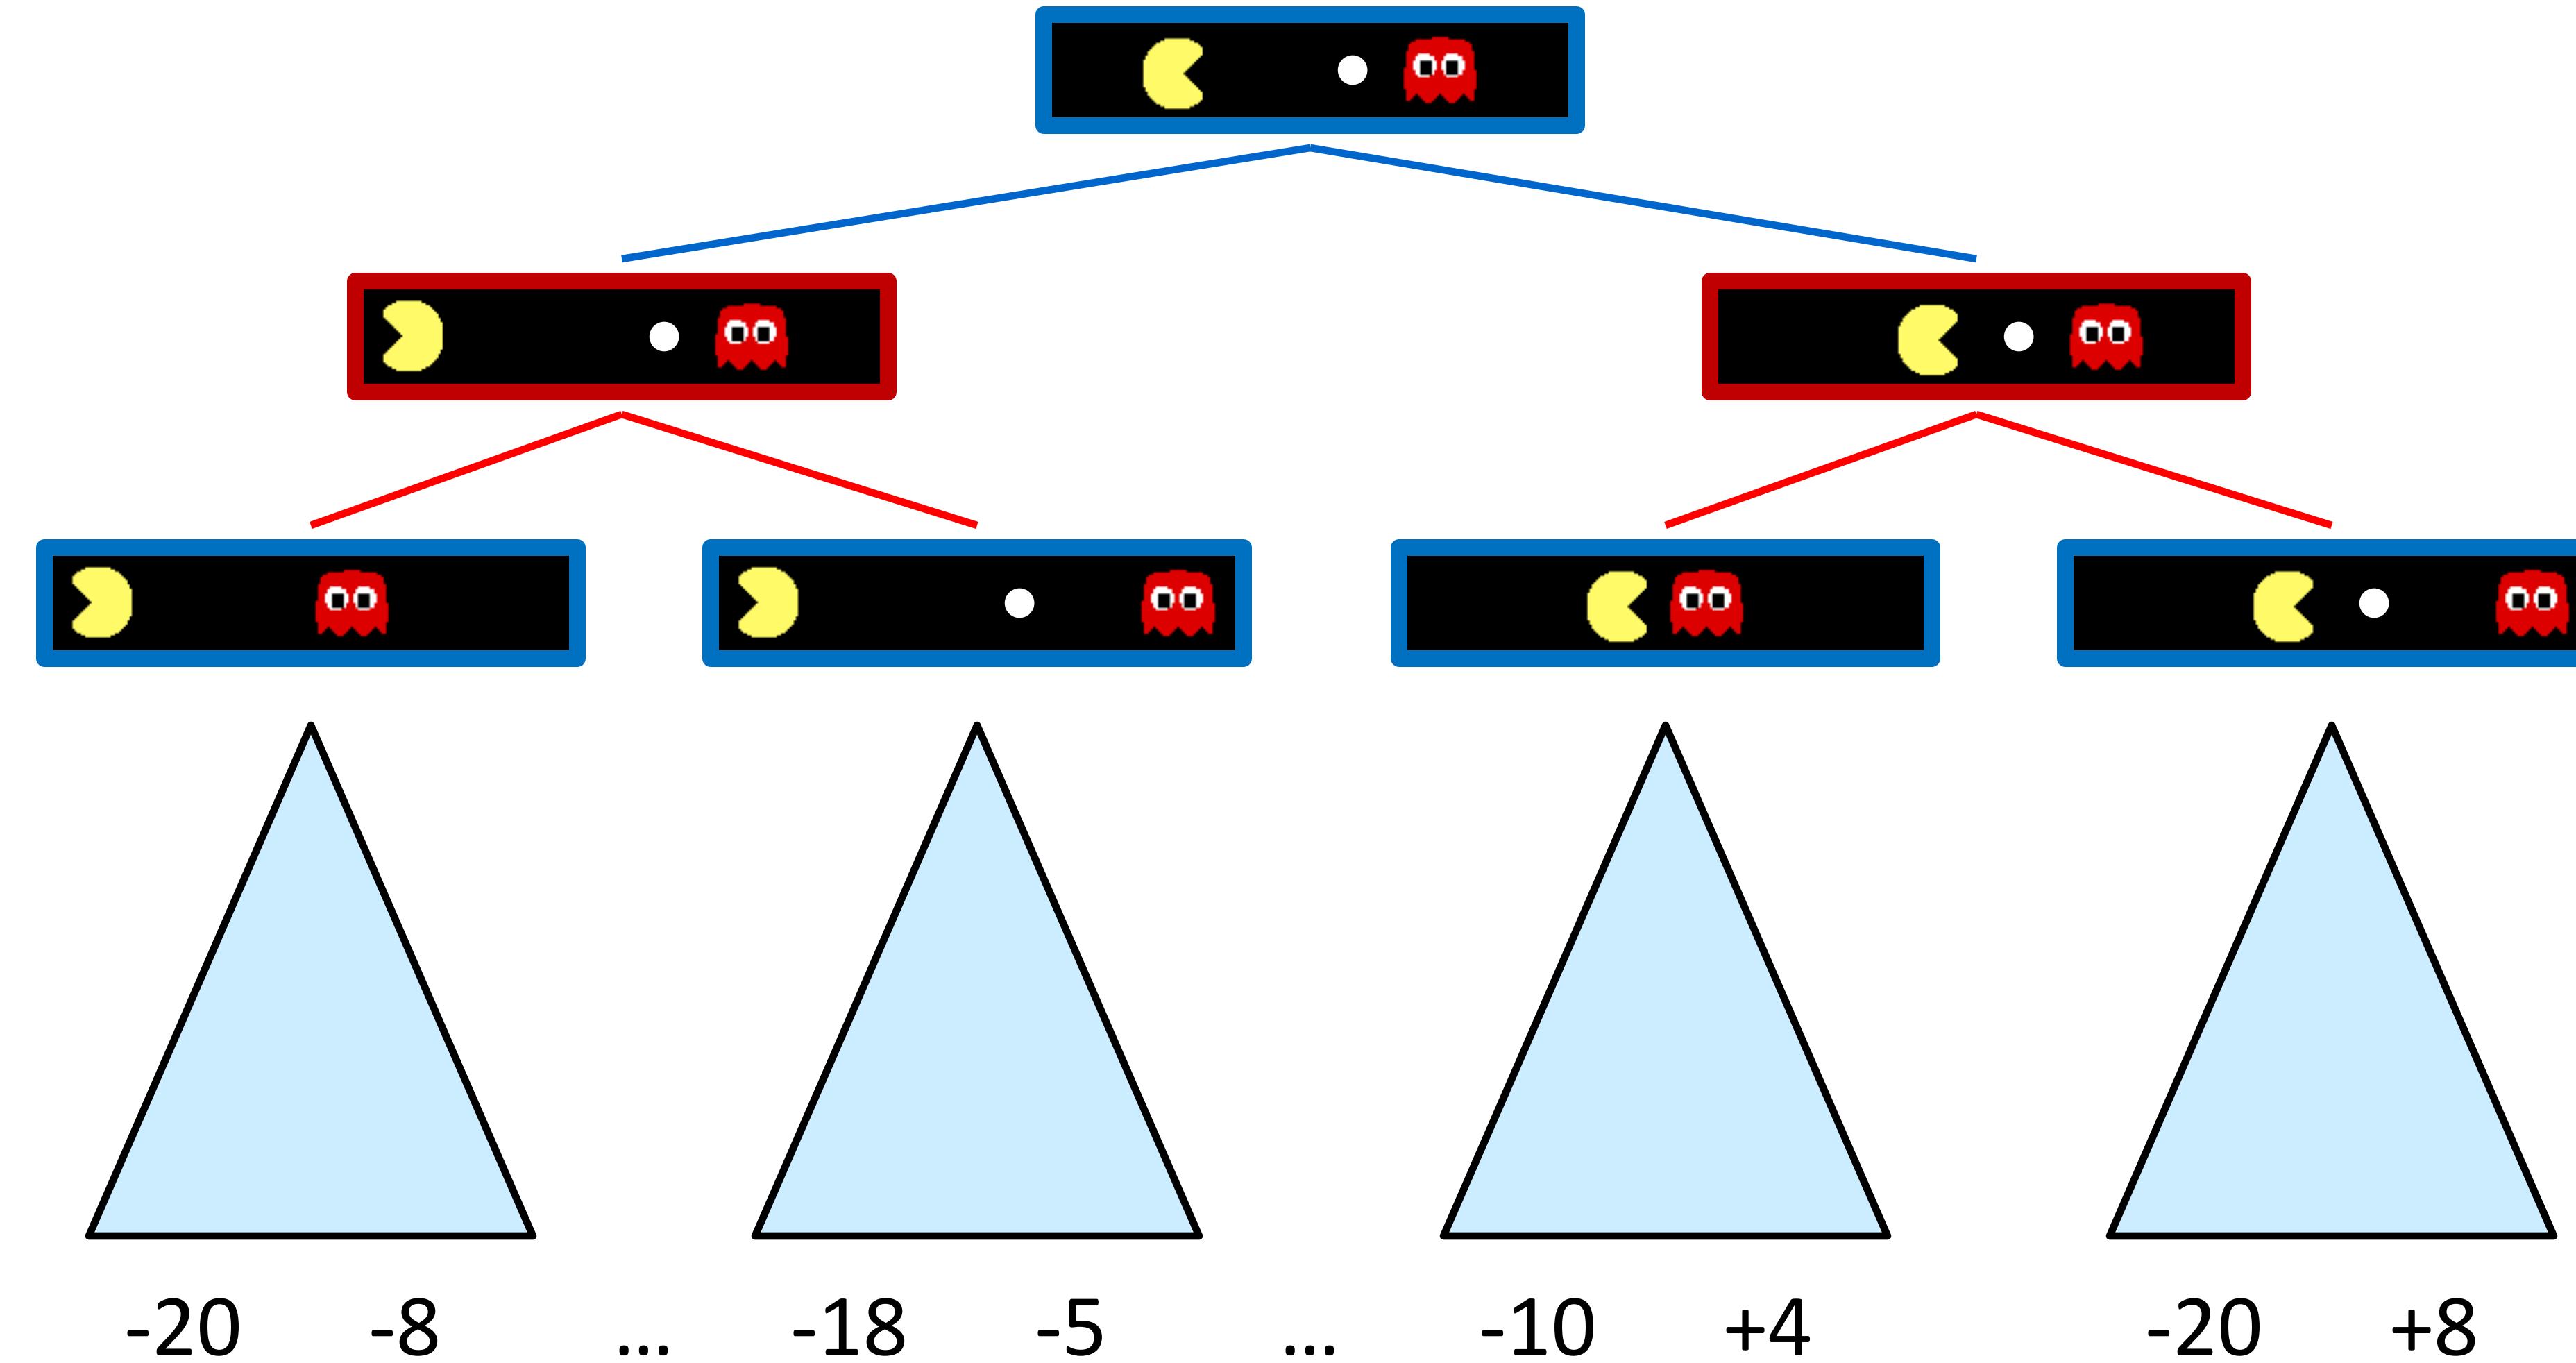
**Non-Terminal States:**

$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

**Terminal States:**

$$V(s) = \text{known}$$

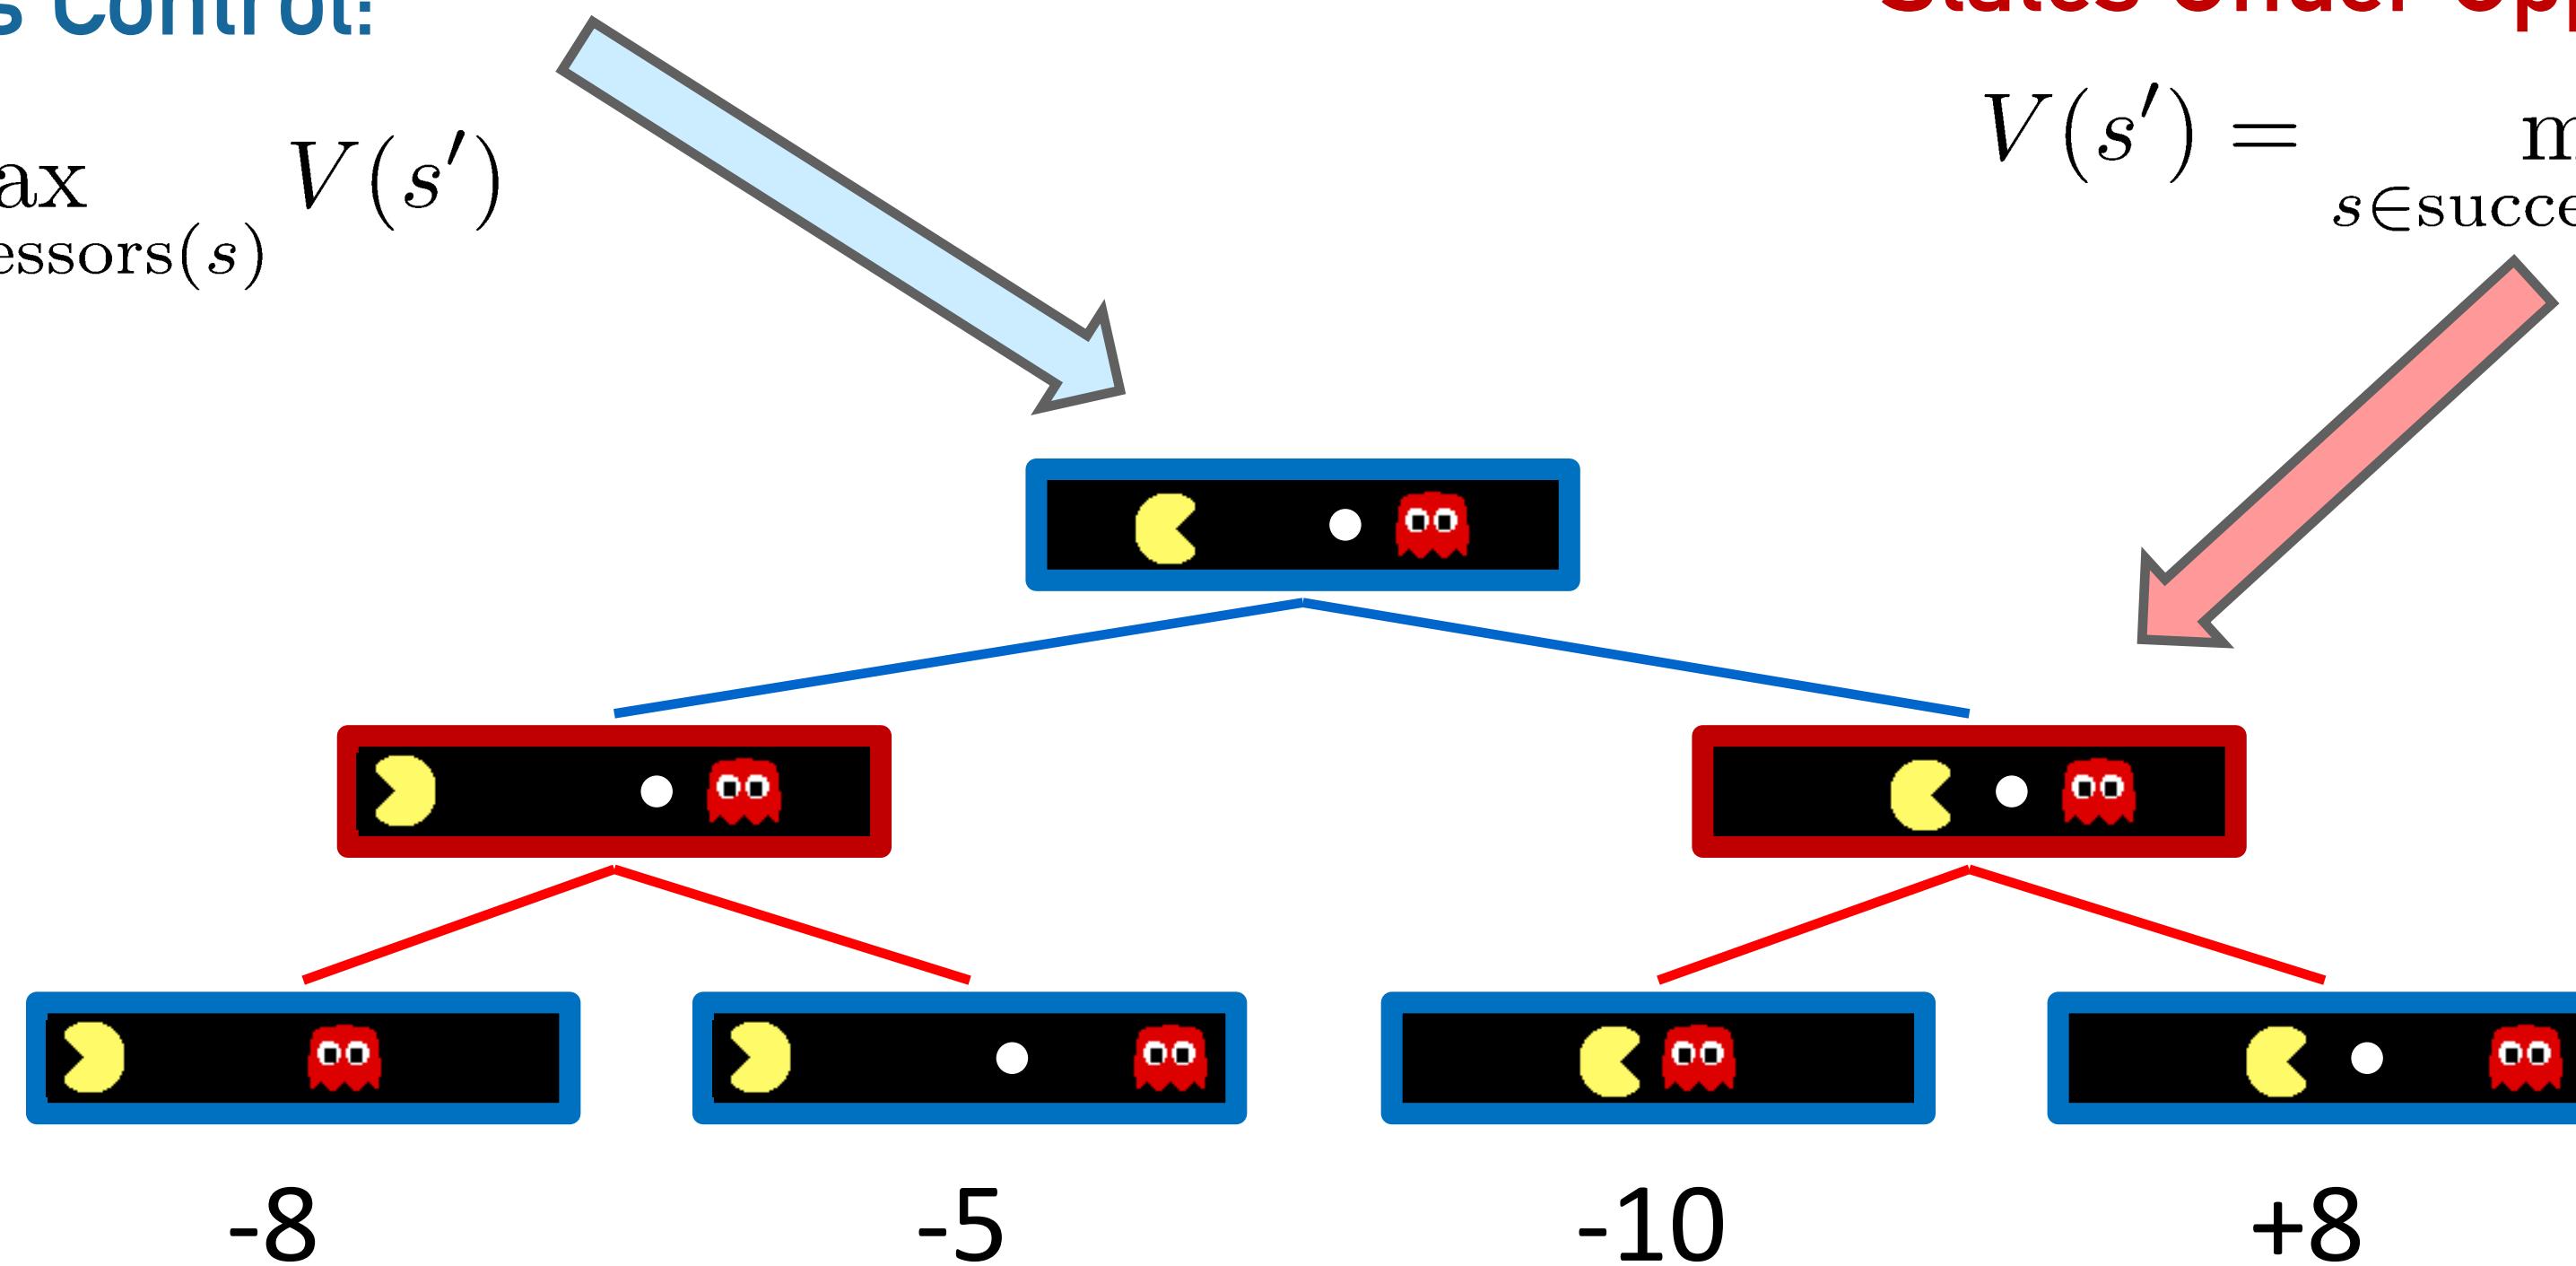
# ADVERSARIAL GAME TREES



# MINIMAX VALUES

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$



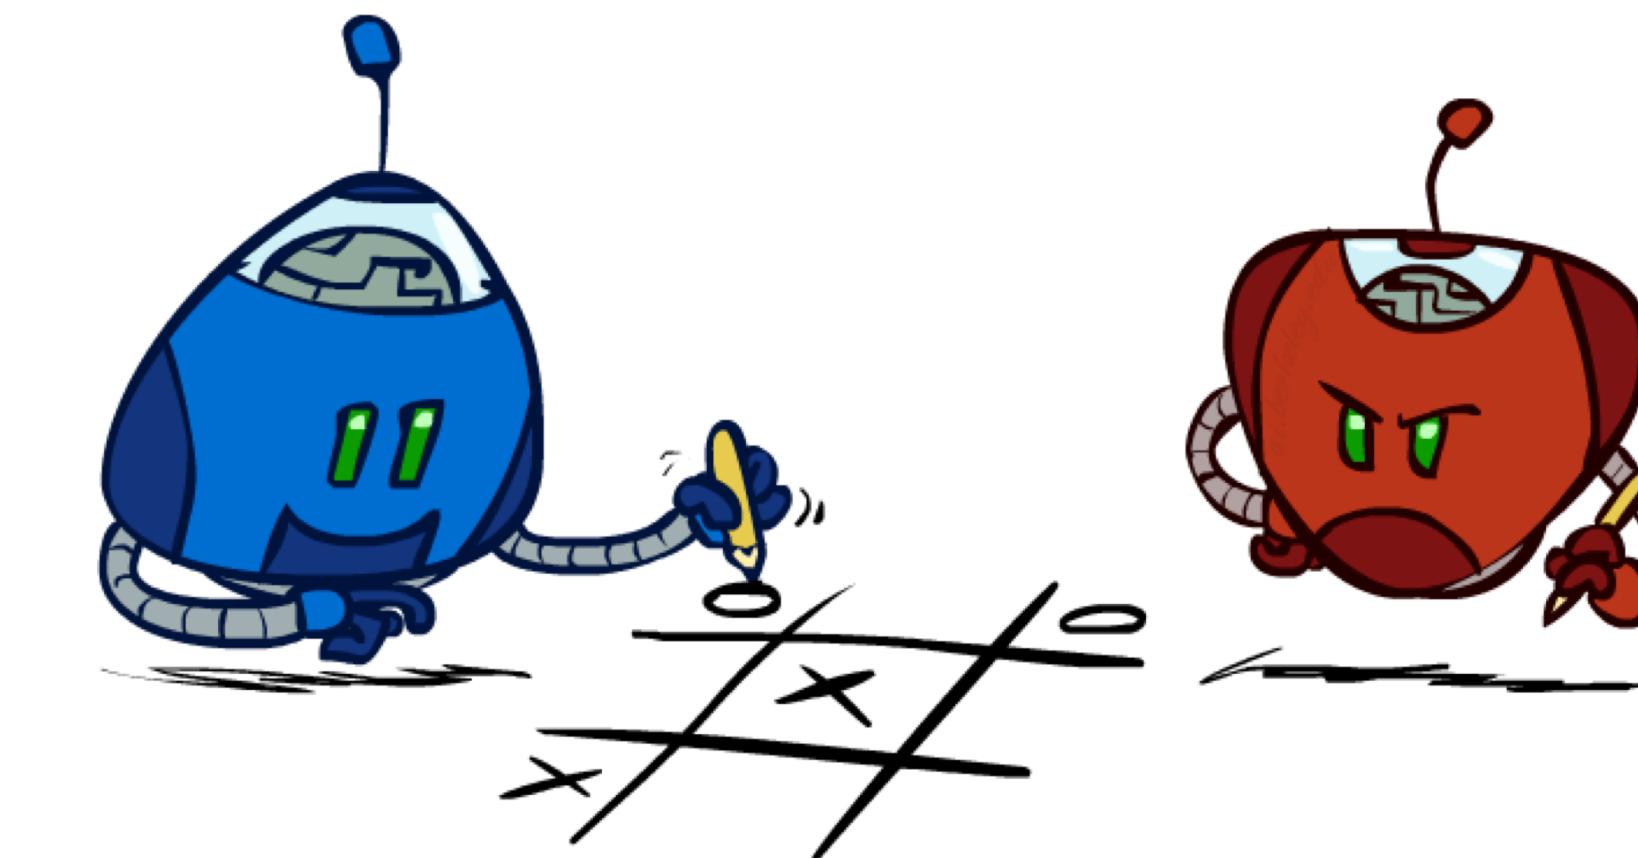
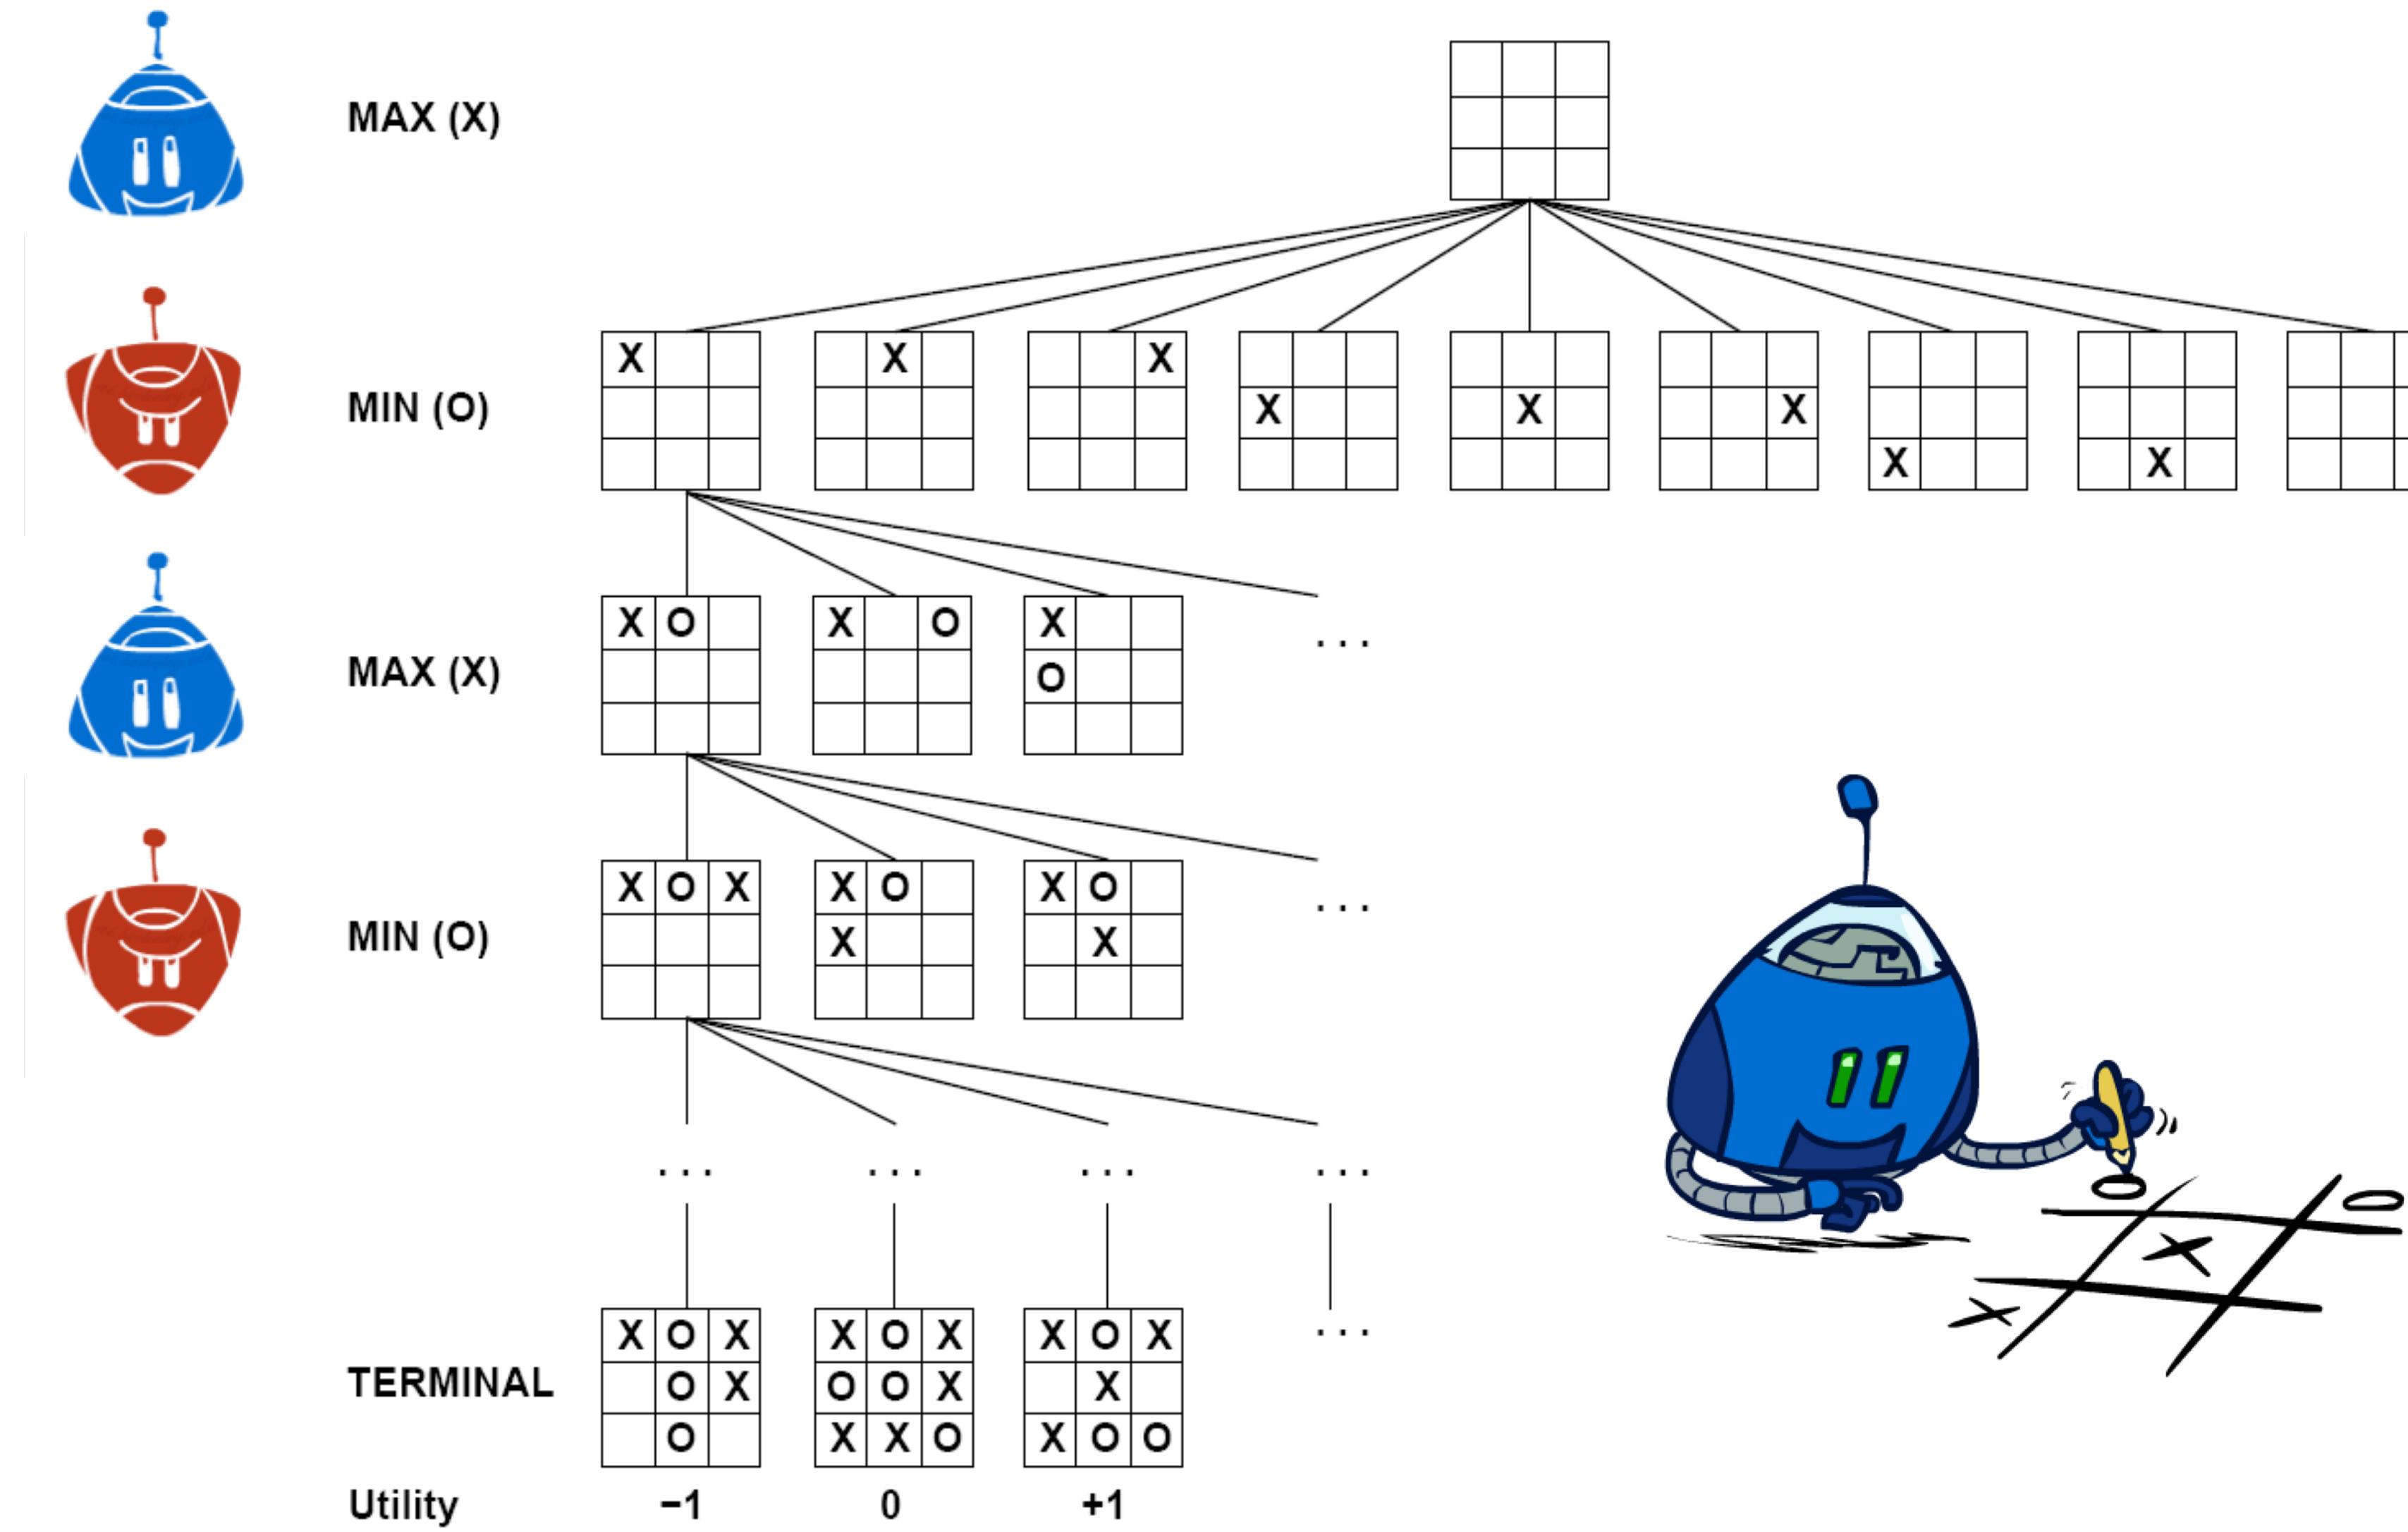
States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Terminal States:

$$V(s) = \text{known}$$

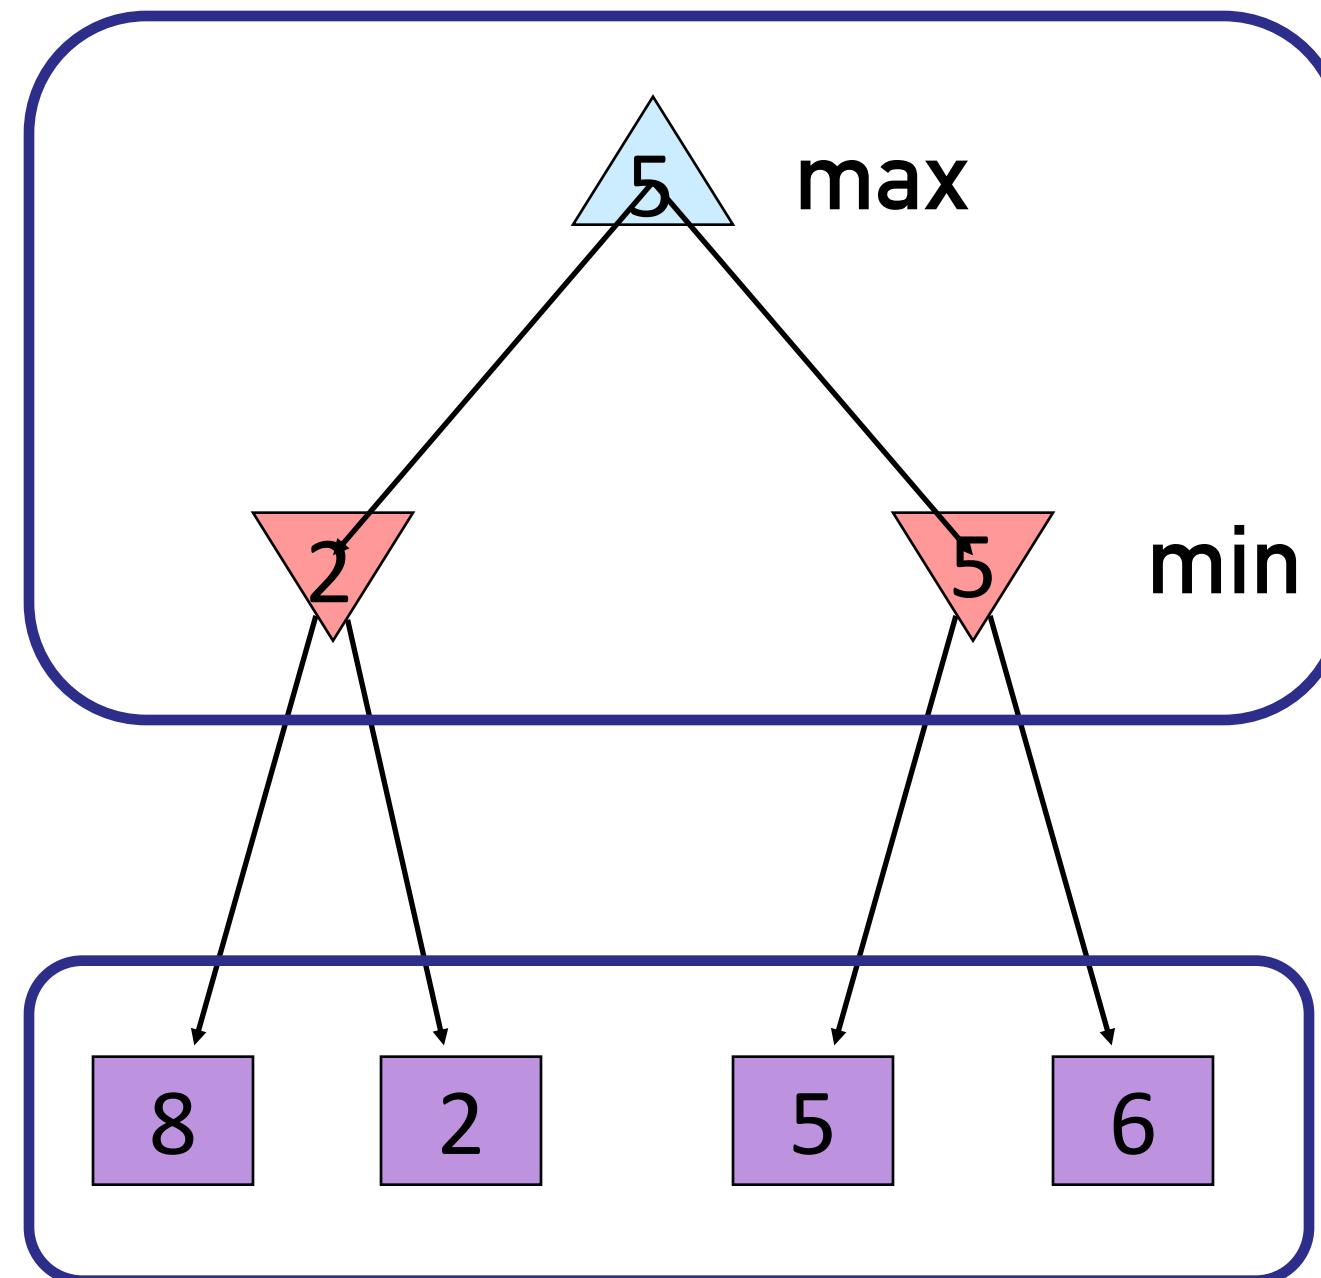
# TIC-TAC-TOE GAME TREE



# ADVERSARIAL SEARCH (MINIMAX)

- ▶ Deterministic, zero-sum games:
  - ▶ Tic-tac-toe, chess, checkers
  - ▶ One player maximizes result
  - ▶ The other minimizes result
  
- ▶ Minimax search:
  - ▶ A state-space search tree
  - ▶ Players alternate turns
  - ▶ Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary

Minimax values:  
computed recursively

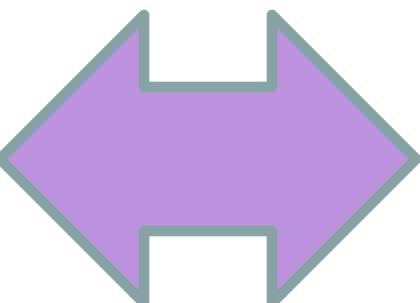


Terminal values:  
part of the game

## MINIMAX IMPLEMENTATION

```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, min-value(successor))
    return v
```

```
def min-value(state):
    initialize v = +∞
    for each successor of state:
        v = min(v, max-value(successor))
    return v
```



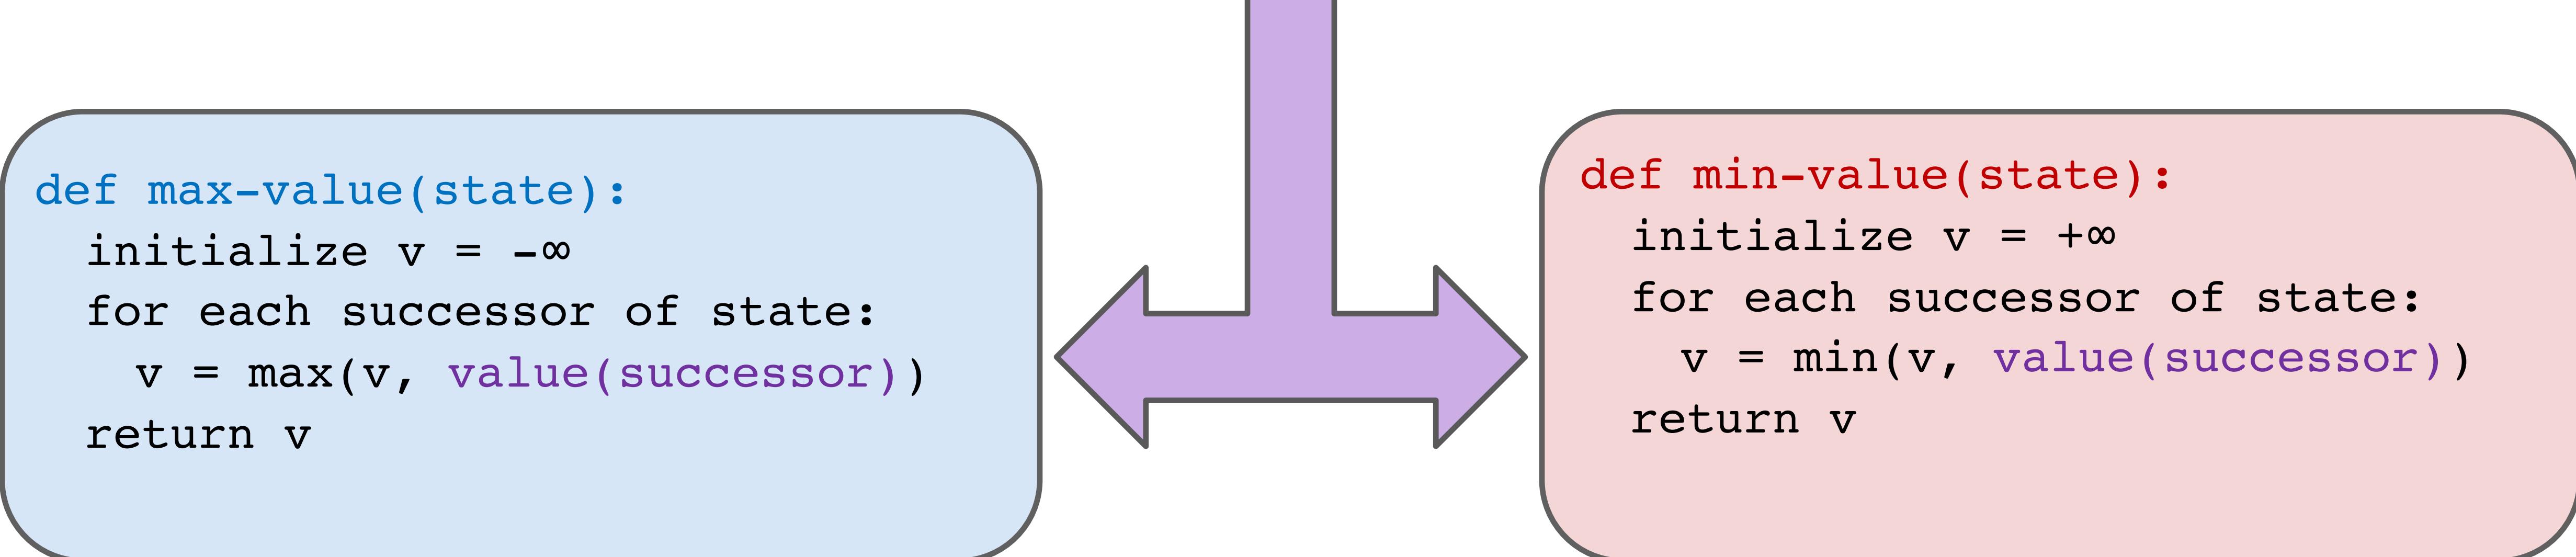
$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

## MINIMAX IMPLEMENTATION (DISPATCH)

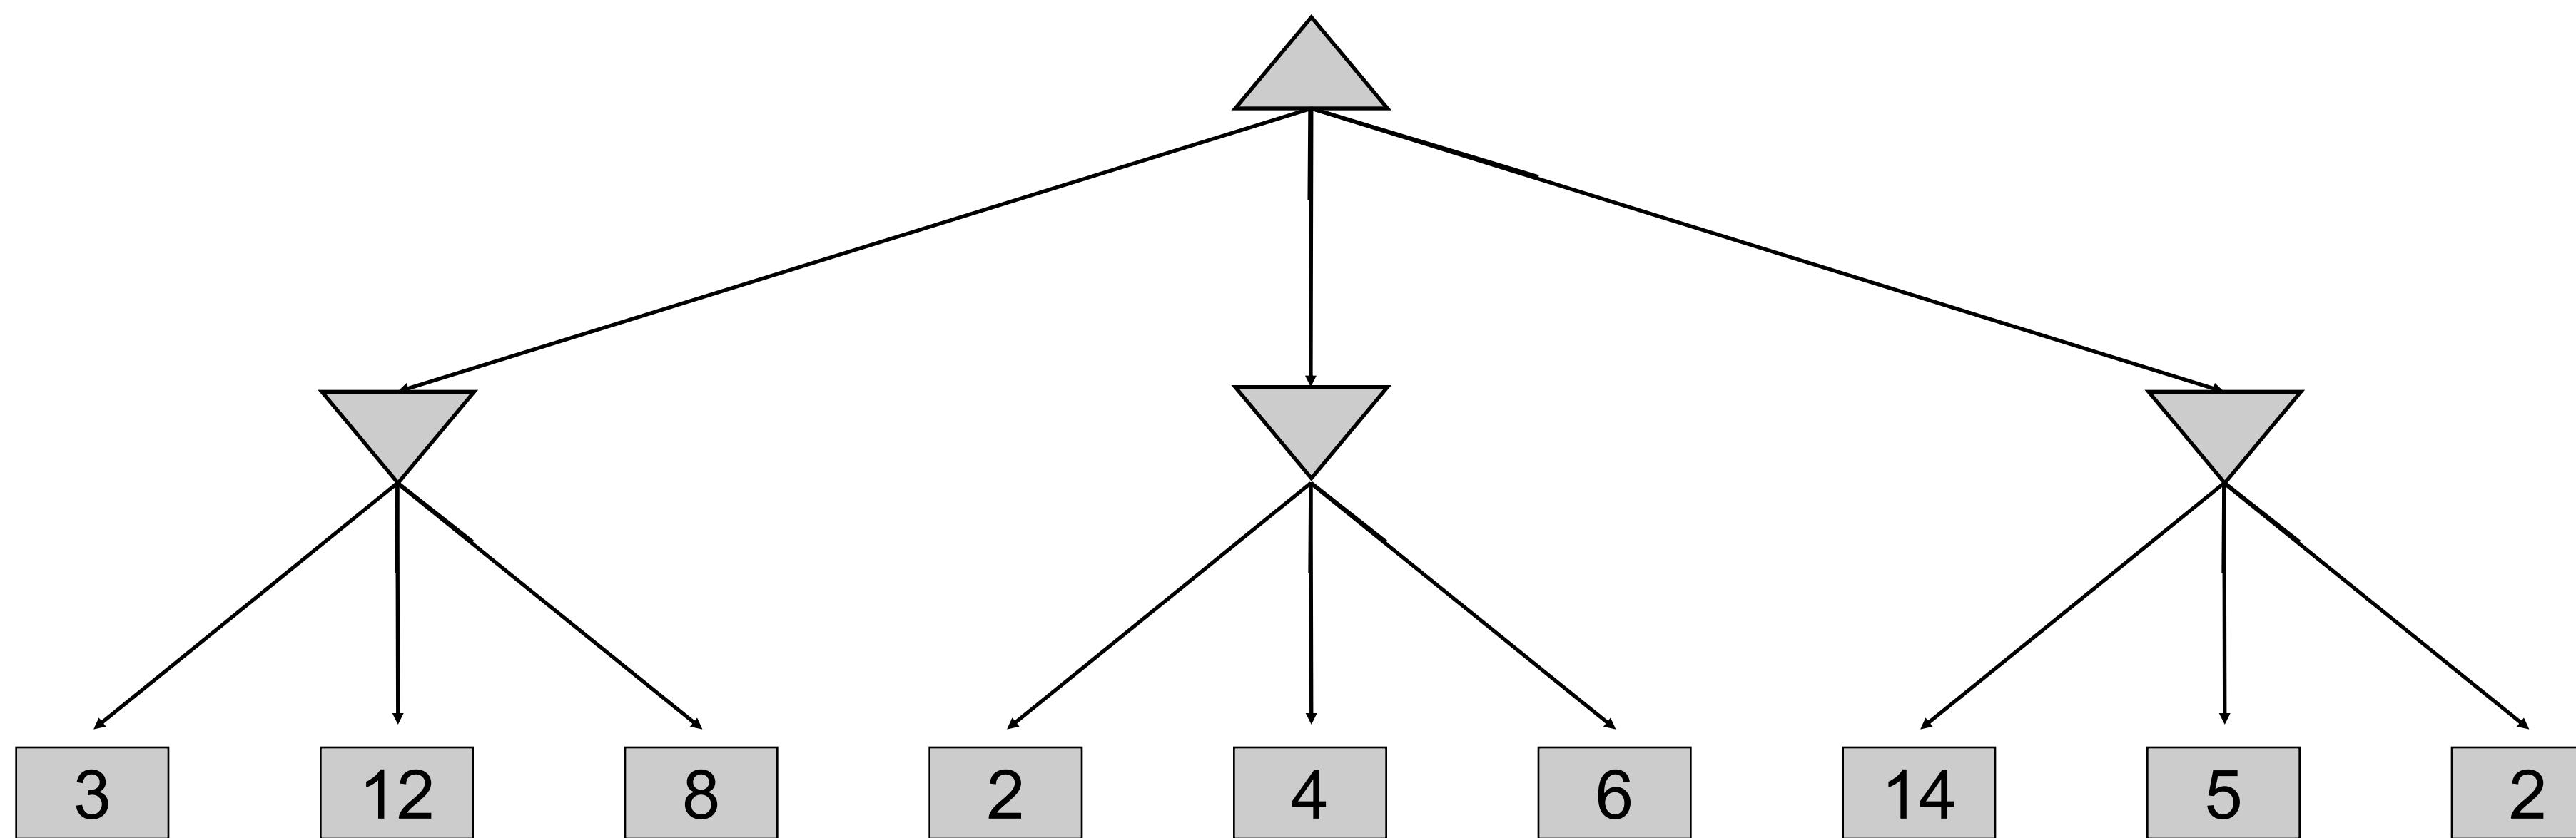
```
def value(state):
    if the state is a terminal state: return the state's utility
    if the state's agent is MAX: return max-value(state)
    if the state's agent is MIN: return min-value(state)
```

```
def max-value(state):
    initialize v = -∞
    for each successor of state:
        v = max(v, value(successor))
    return v
```

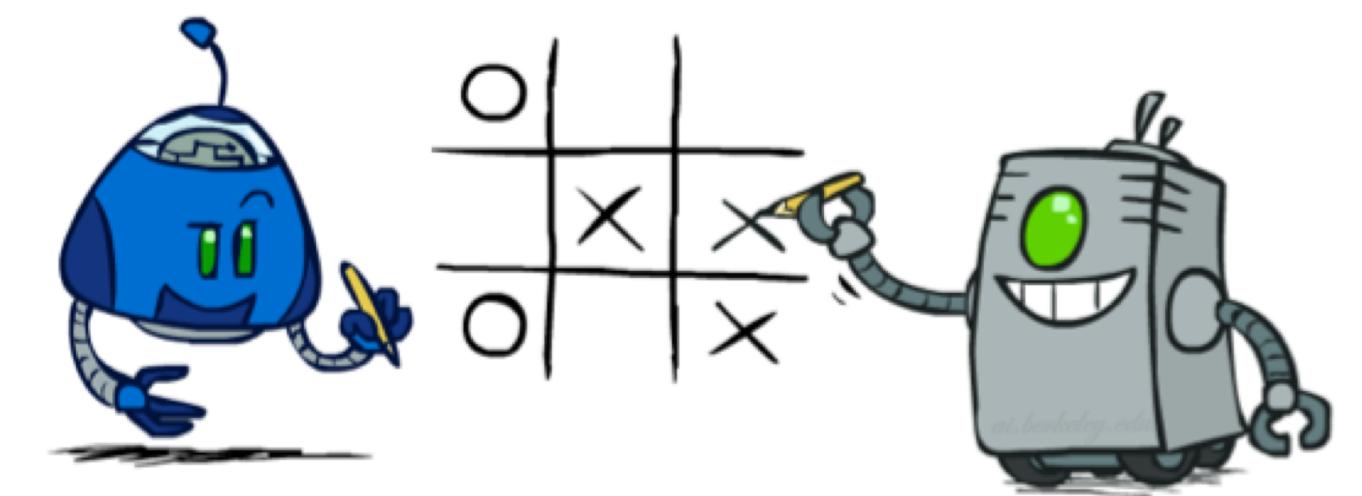
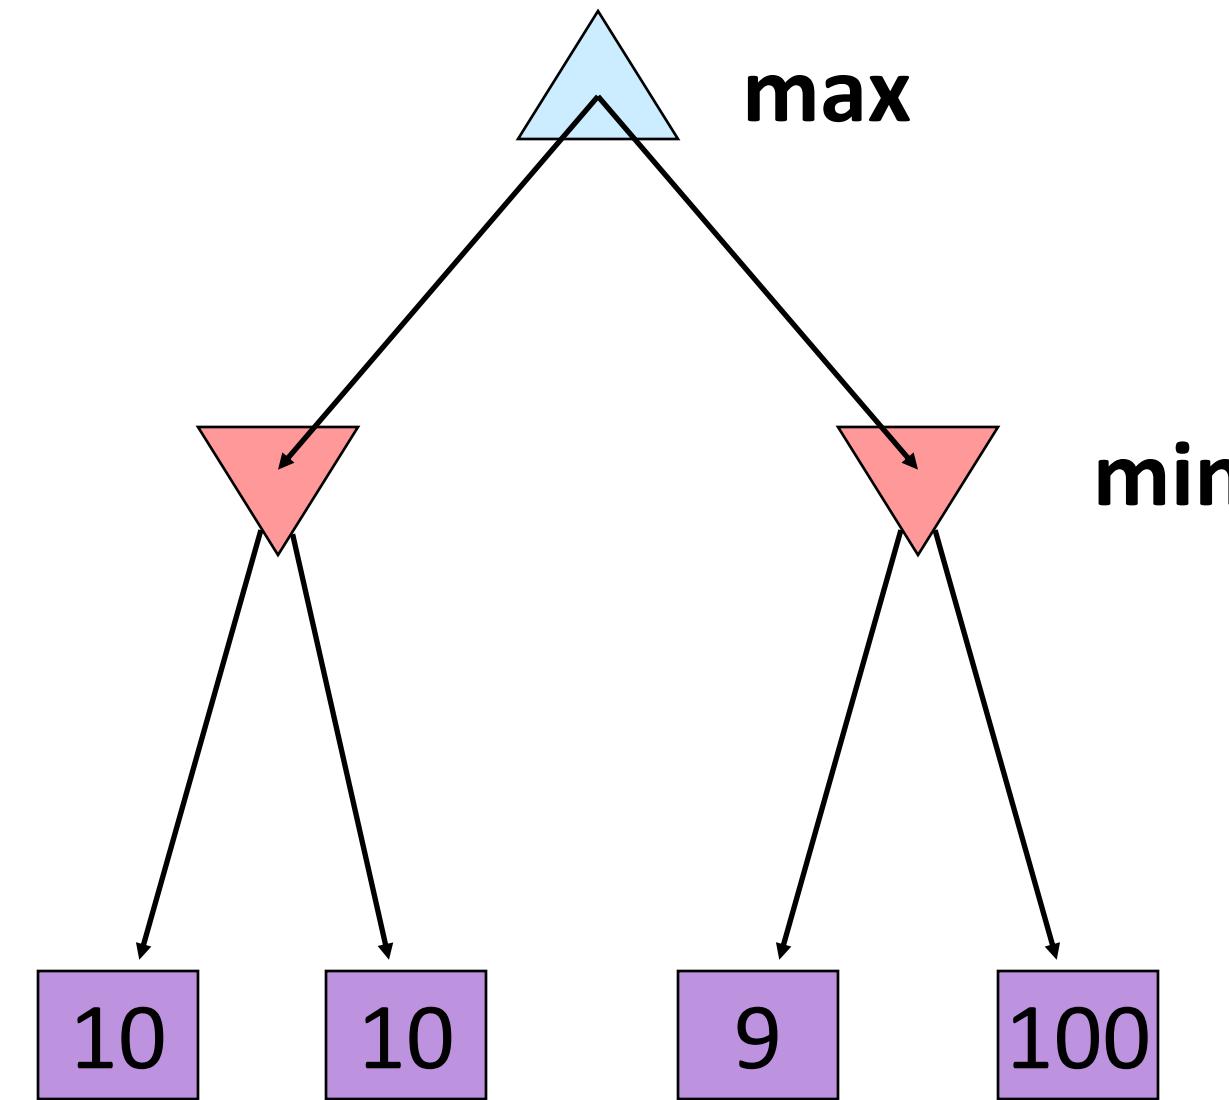
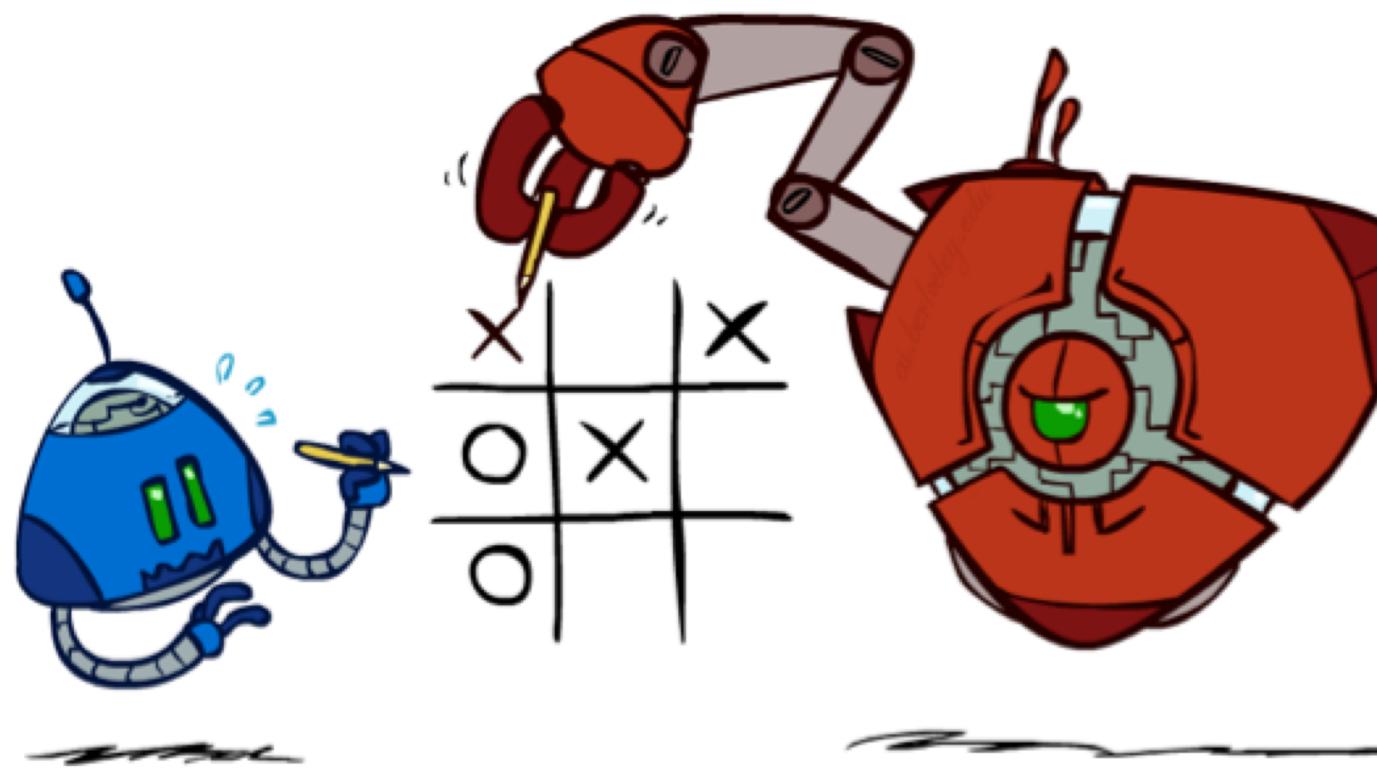


```
def min-value(state):
    initialize v = +∞
    for each successor of state:
        v = min(v, value(successor))
    return v
```

## MINIMAX EXAMPLE



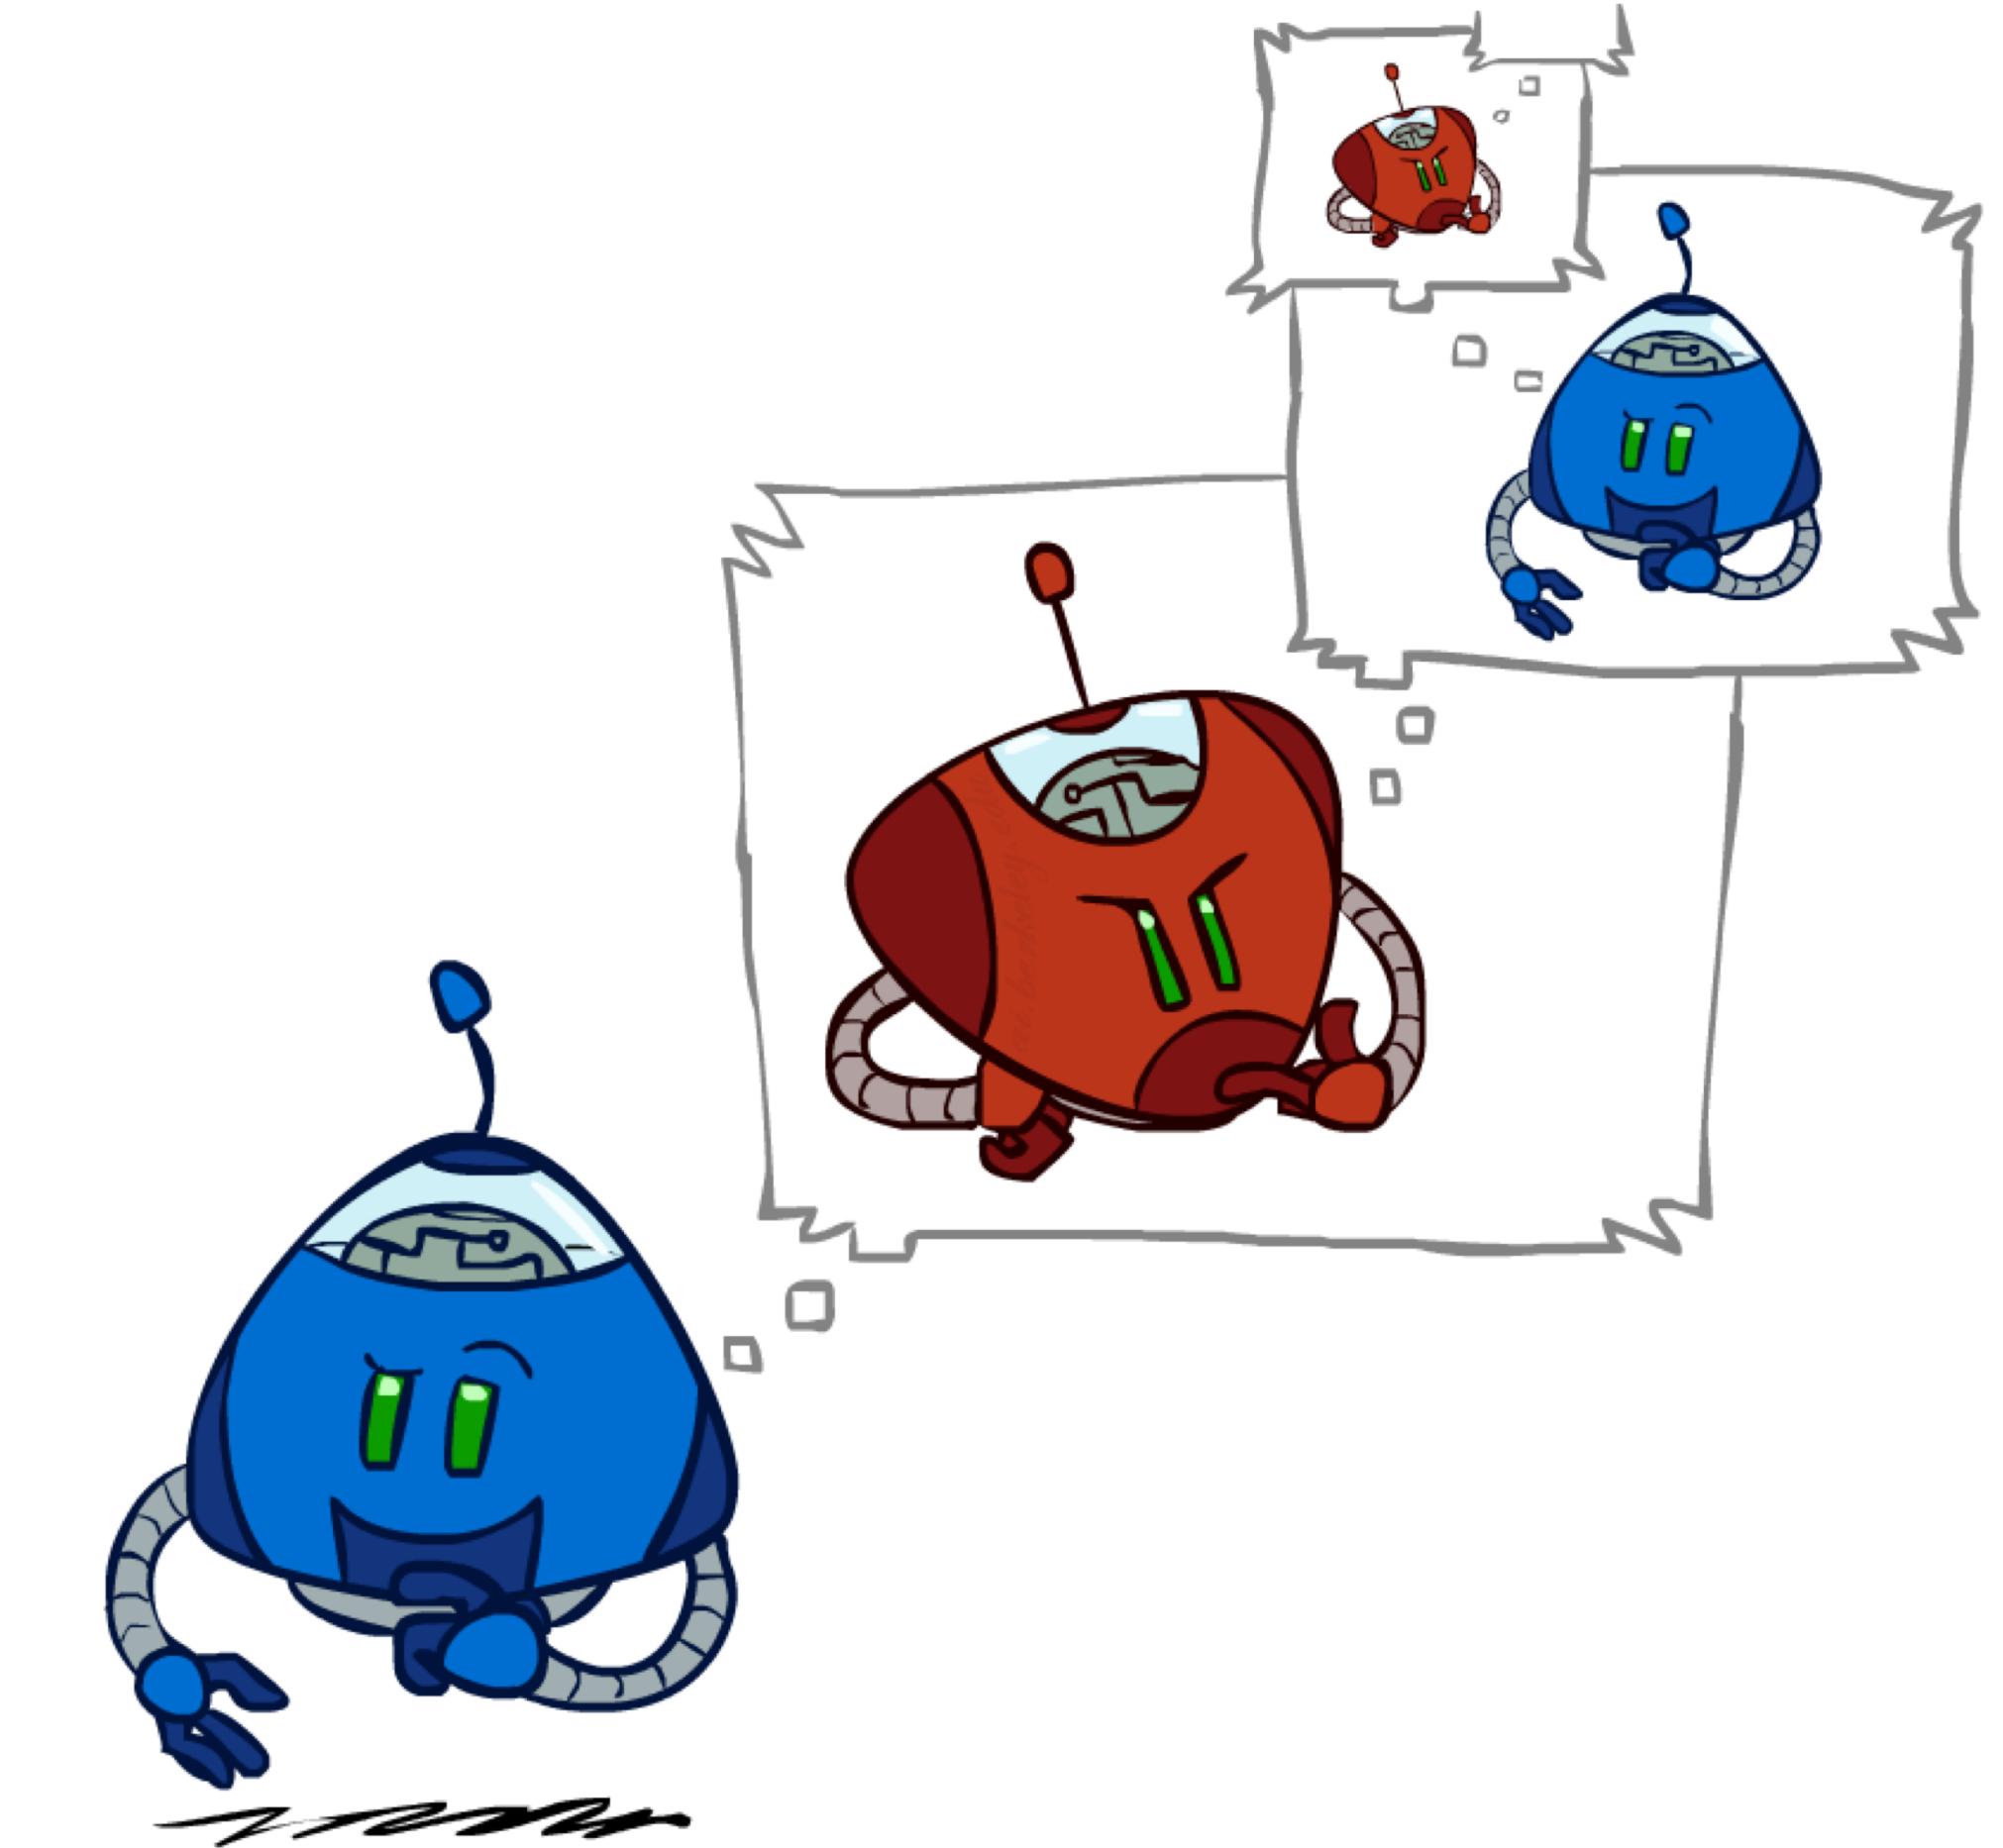
## MINIMAX PROPERTIES



Optimal against a rational player. Otherwise?

# MINIMAX EFFICIENCY

- ▶ Efficiency of minimax search
  - ▶ Just like (exhaustive) DFS
  - ▶ Time:  $O(b^m)$
  - ▶ Space:  $O(bm)$
- ▶ Example: For chess,  $b \approx 35$ ,  $m \approx 100$ 
  - ▶ Exact solution is completely infeasible



## GAME TREE SIZES

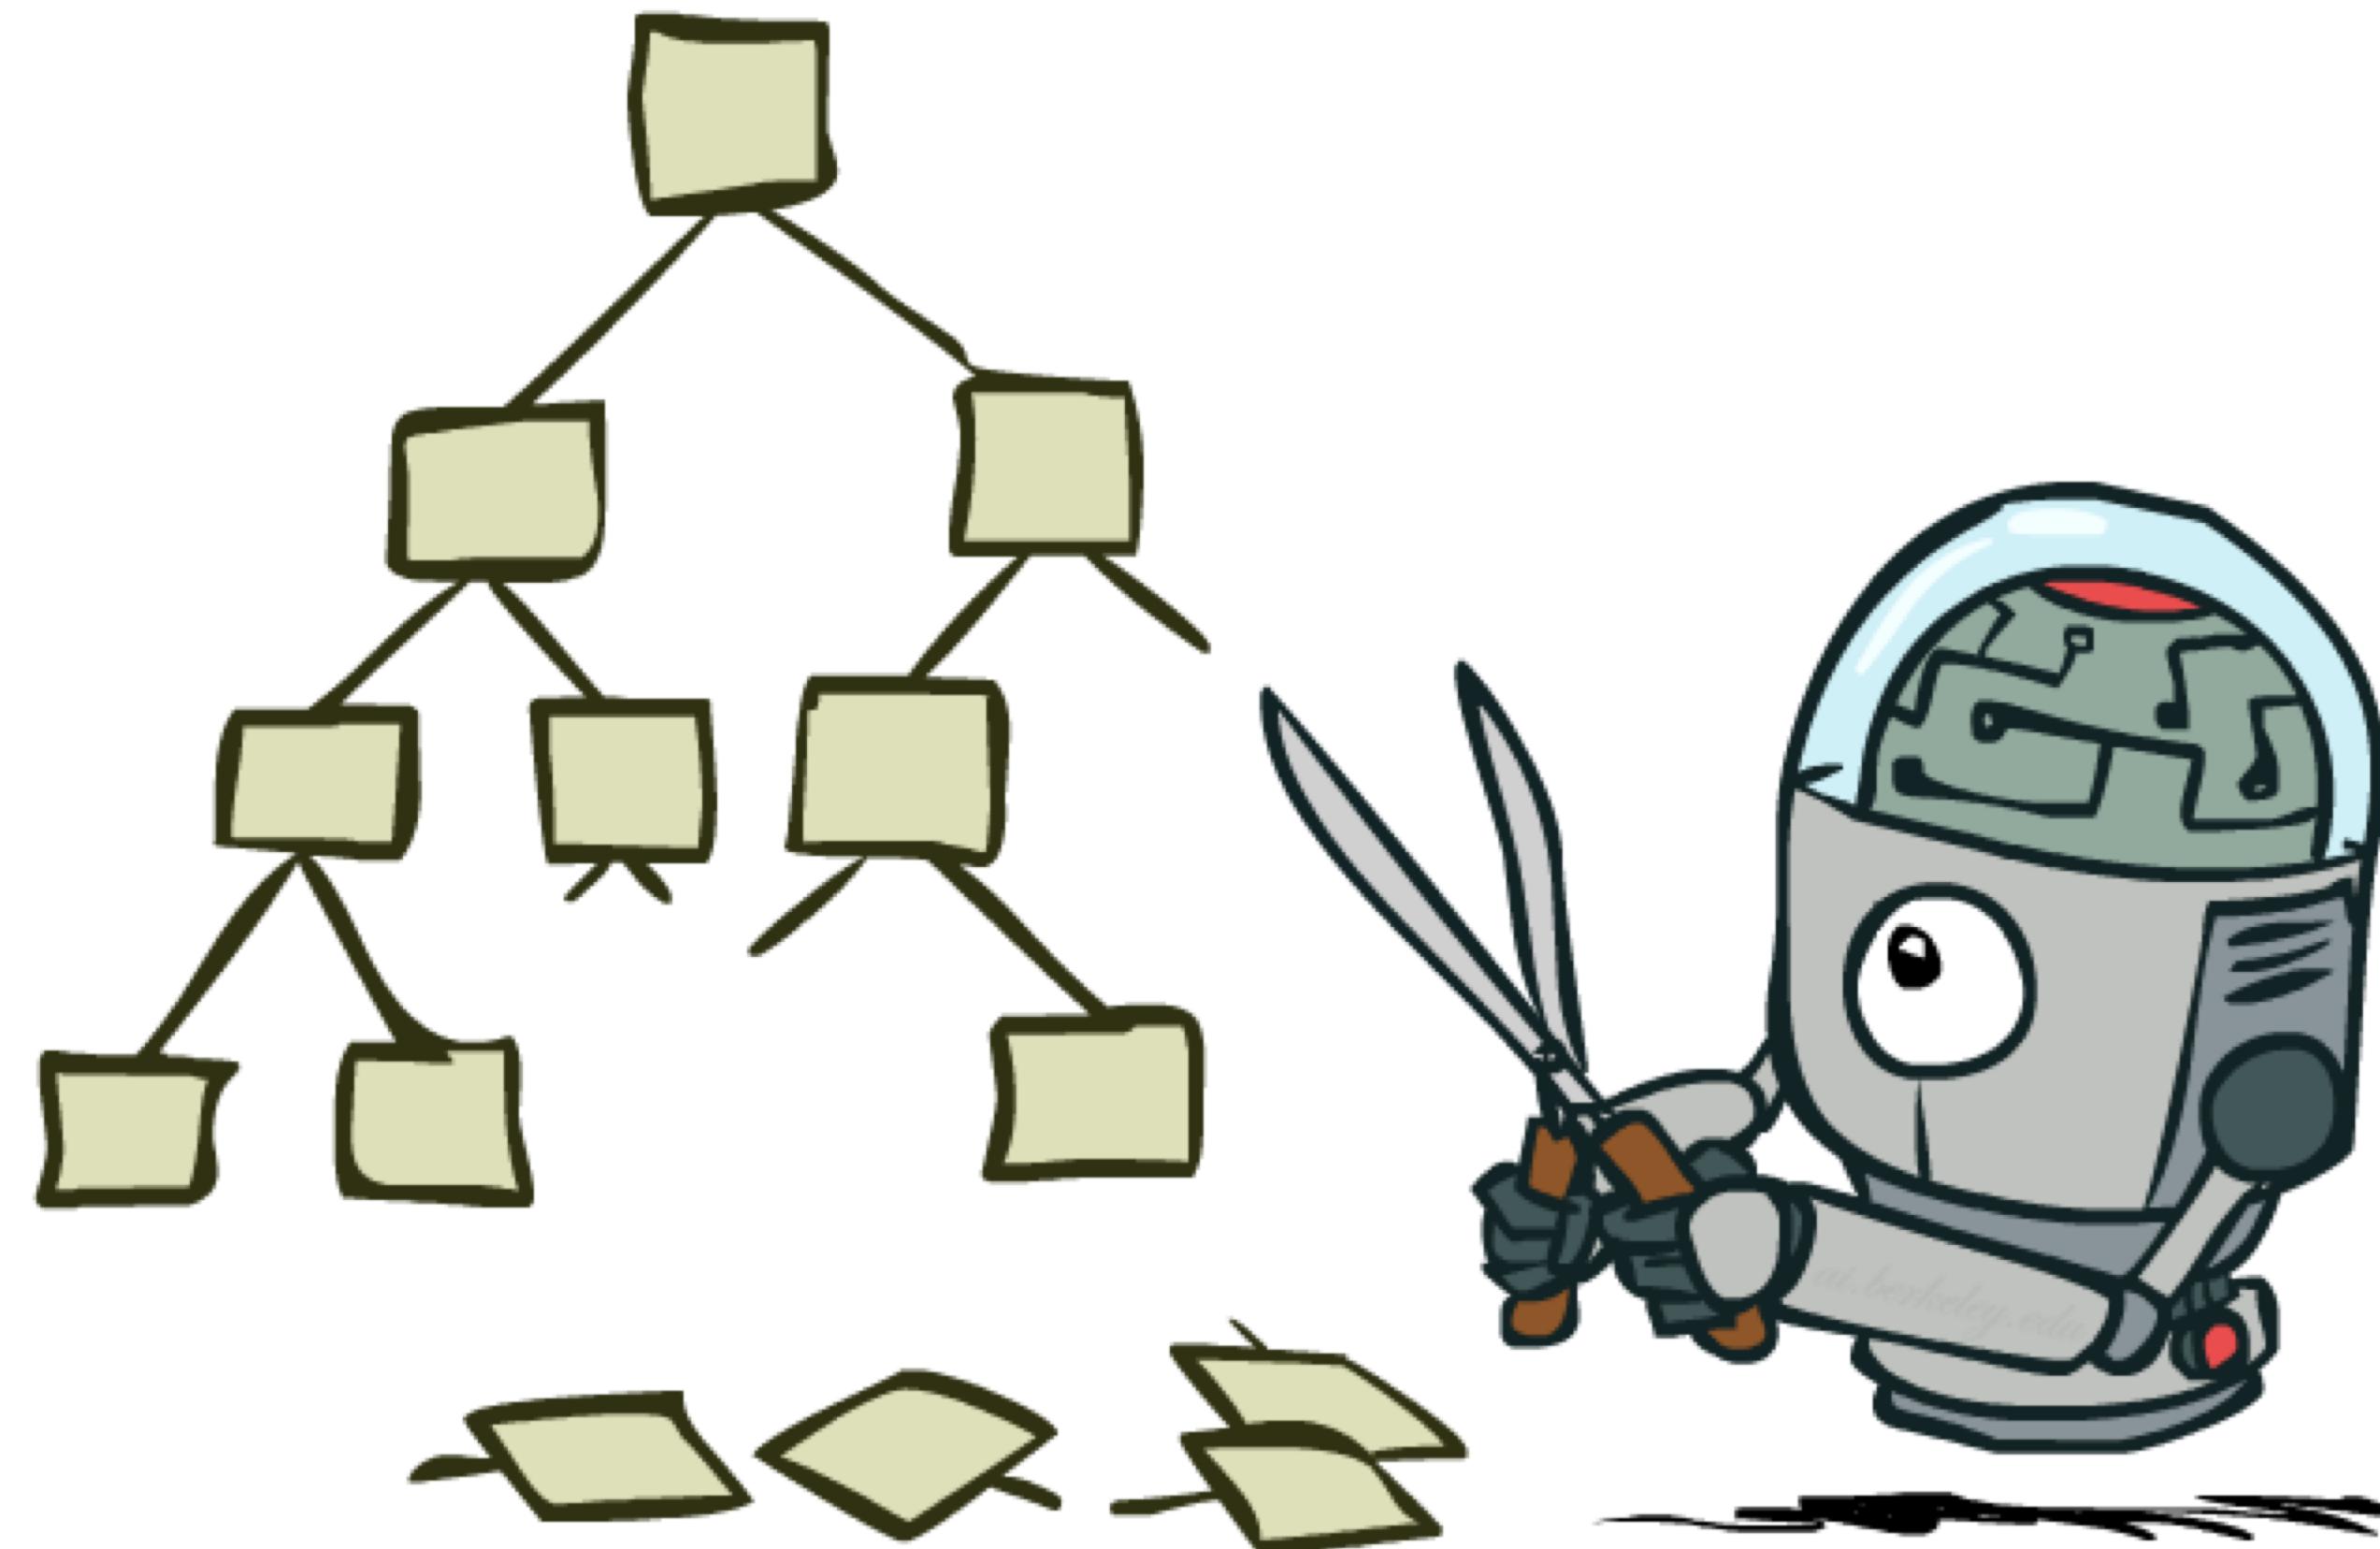
- ▶ Tic-tac-toe:  $10^5$
- ▶ Checkers:  $10^{31}$
- ▶ Chess:  $10^{123}$
- ▶ Backgammon:  $10^{144}$
- ▶ Go:  $10^{360}$

Assume that a computer can evaluate 1 million board configurations per second.

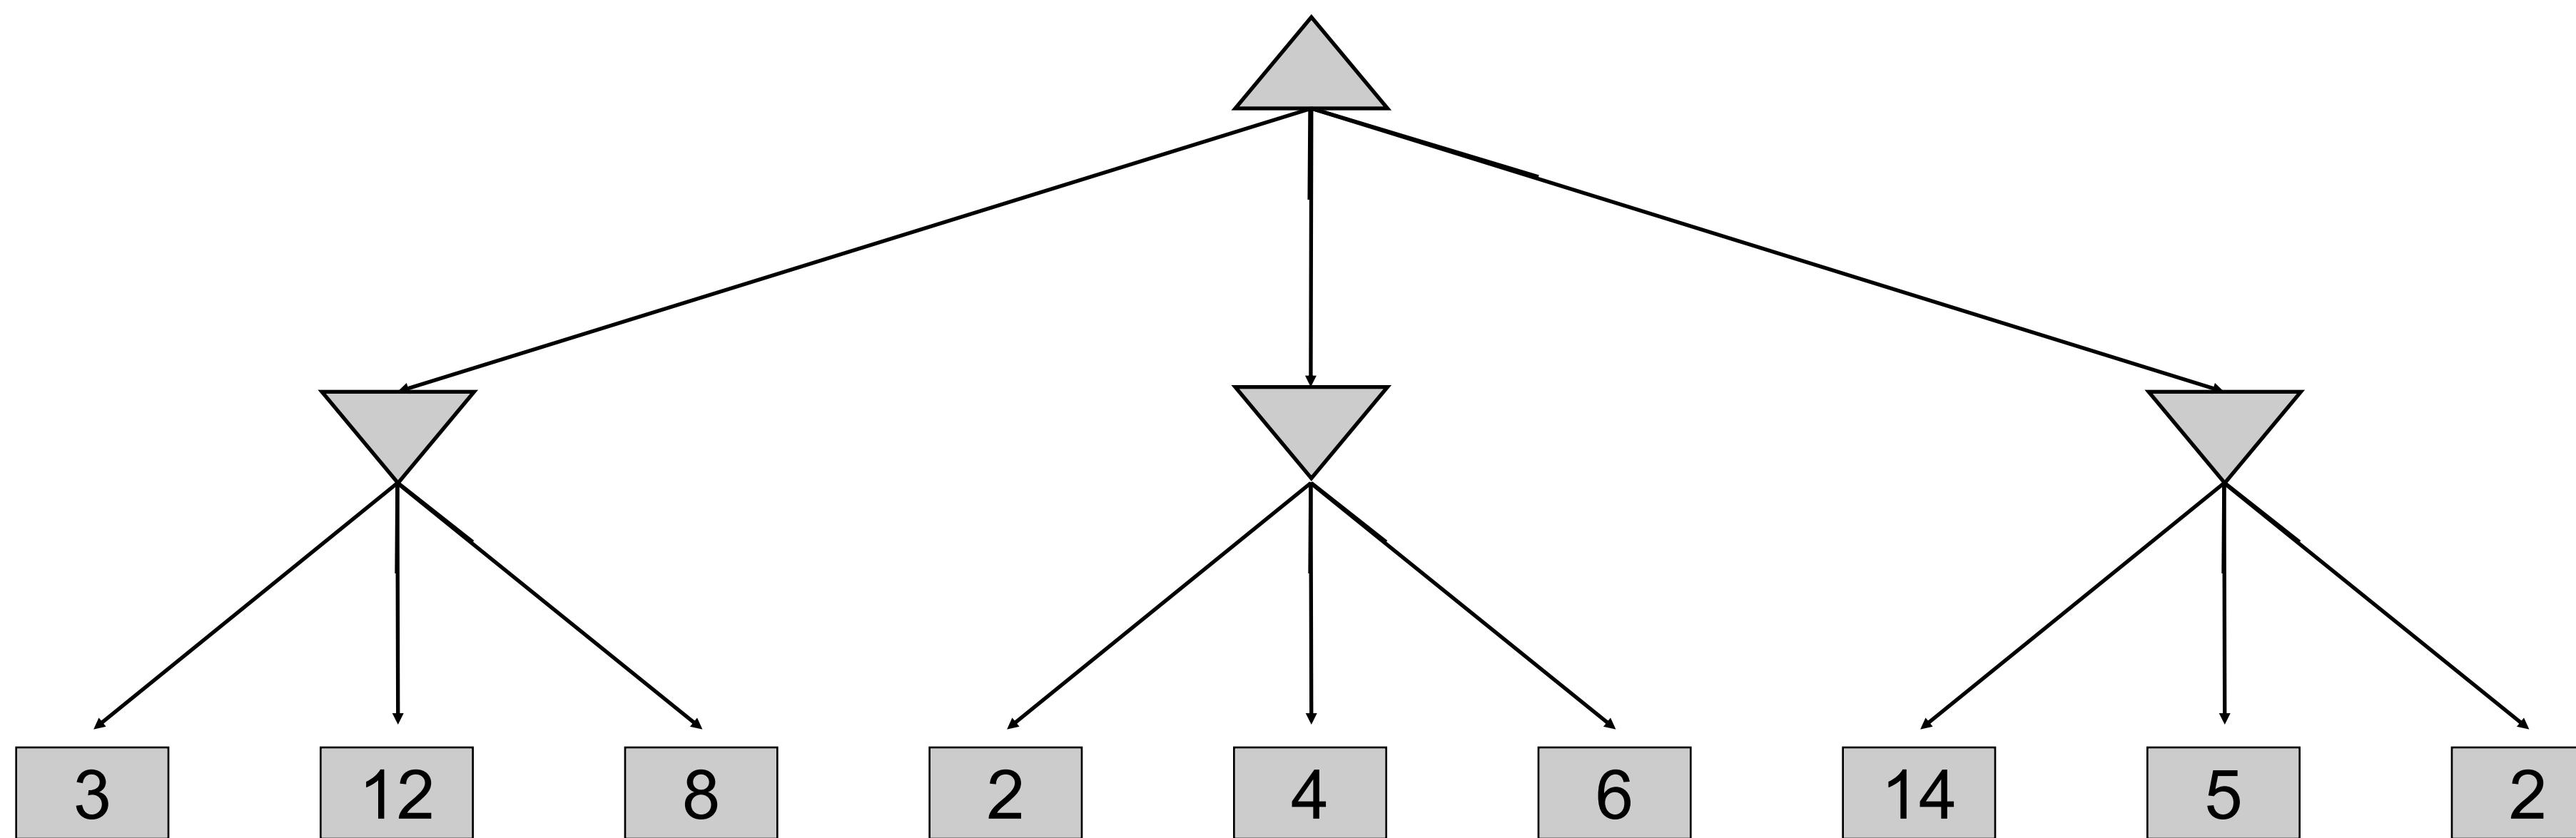
Then it would take **0.1 seconds to search** the entire tic-tac-toe game tree but it would still take  **$10^{18}$  years** to search the full checkers tree.

Do we need to explore the whole tree?

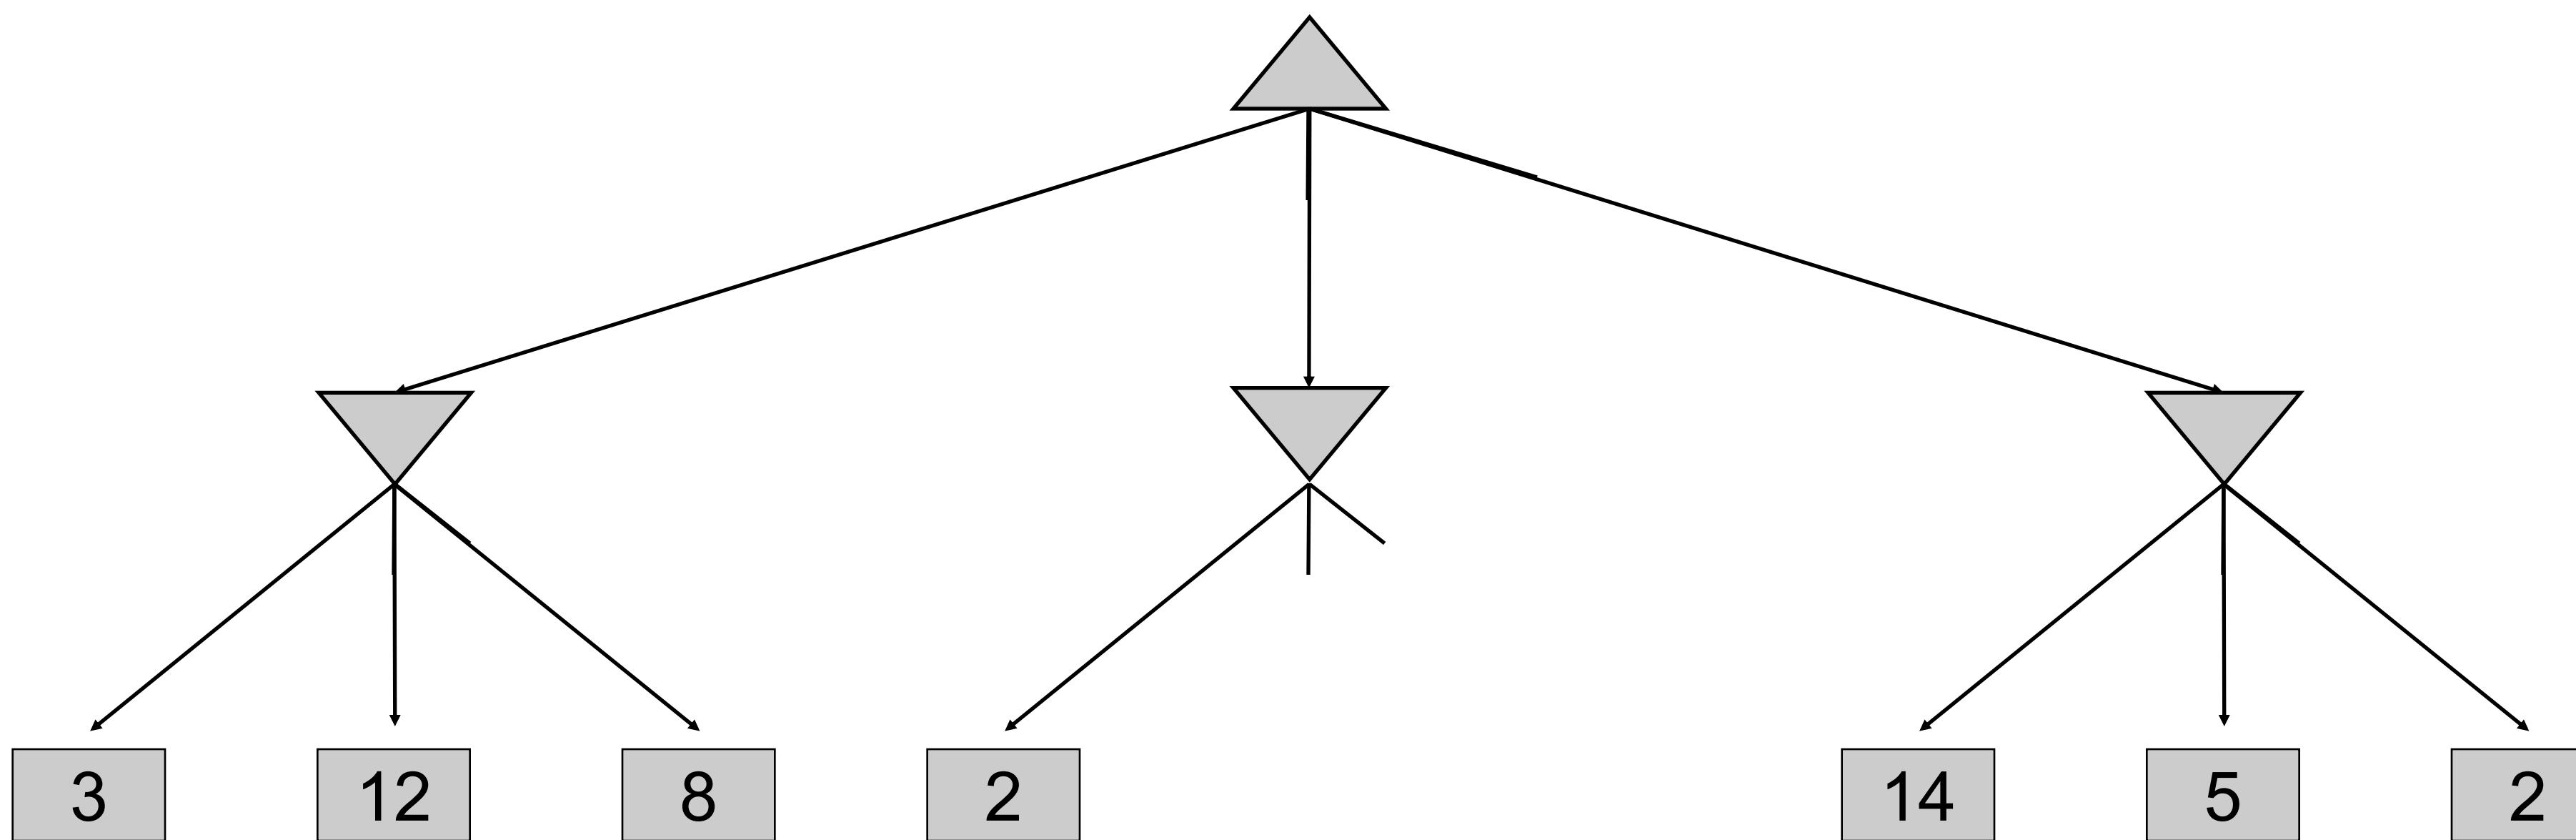
# GAME TREE PRUNING



## MINIMAX EXAMPLE

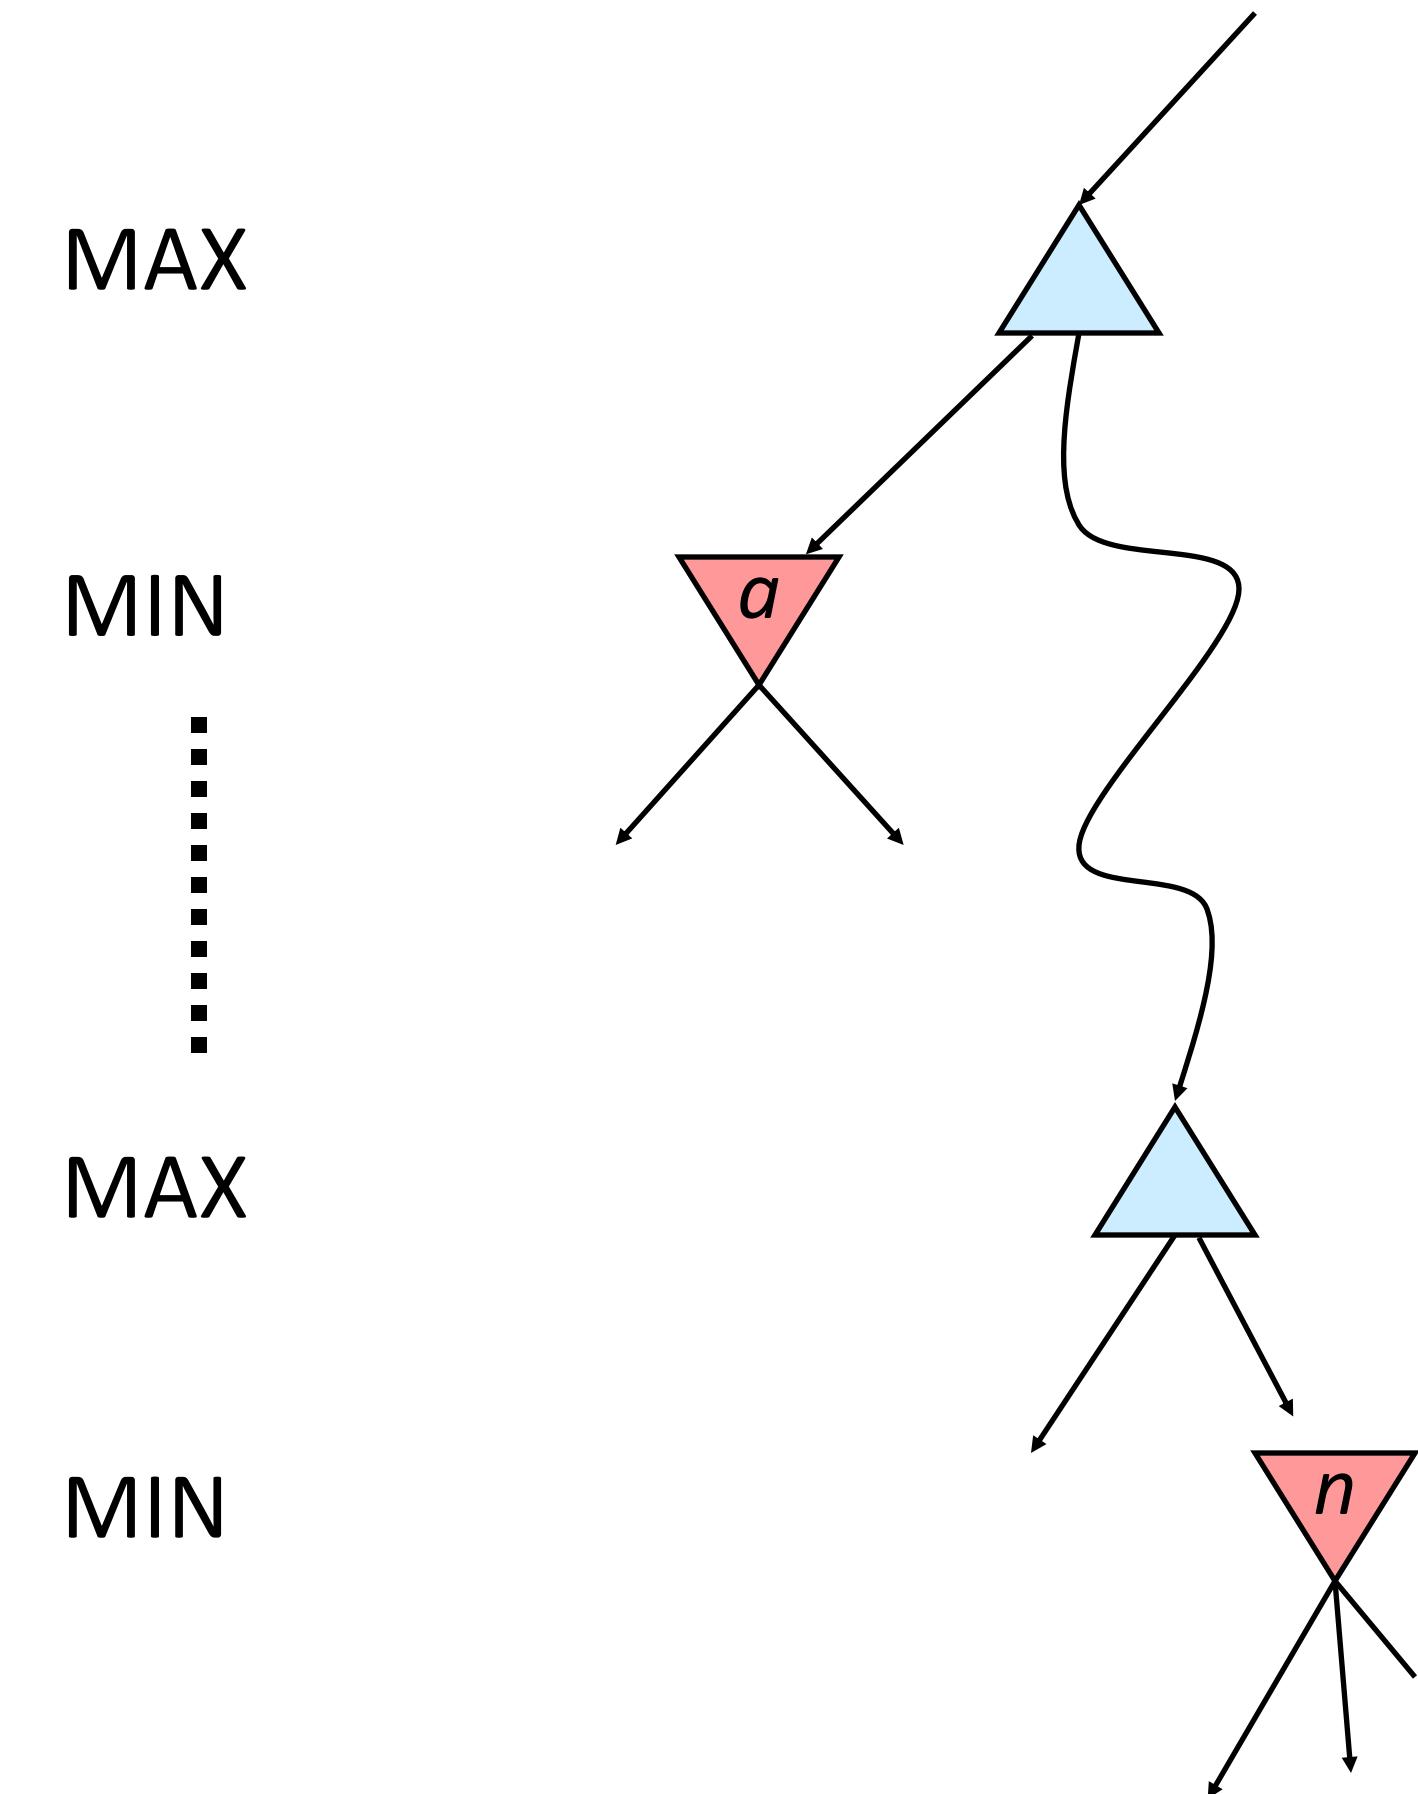


# MINIMAX PRUNING



# ALPHA-BETA PRUNING

- ▶ General configuration (MIN version)
  - ▶ When computing the MIN-VALUE at some node  $n$
  - ▶ We loop over  $n$ 's children
  - ▶  $n$ 's estimate of the children's min is dropping
  - ▶ Who cares about  $n$ 's value? MAX
  - ▶ Let  $\alpha$  be the best value that MAX can get at any choice point along the current path from the root so far
  - ▶ If  $n$  becomes worse than  $\alpha$ , MAX will avoid it, so we can stop considering  $n$ 's other children (it's already bad enough that it won't be played)
- ▶ MAX version is symmetric



# Alpha-Beta Implementation

$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state, α, β):  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, value(successor,α,β))  
        if v ≥ β return v  
        α = max(α, v)  
    return v
```

```
def min-value(state , α, β):  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, value(successor,α,β))  
        if v ≤ α return v  
        β = min(β, v)  
    return v
```

## ALPHA-BETA PRUNING

