

PURDUE CS47100

---

# INTRODUCTION TO AI

## ANNOUNCEMENTS

- ▶ Reminder: Assignment 1 is due on this Friday!
- ▶ Cast your vote on Ed about Assignment 2 due date
  - ▶ Do you want the deadline of Assignment 2 to be extended by 2 days (from Oct 14 to Oct 16)?
  - ▶ Potentially decrease your workload during the fall break, but it also implies that we can at best release the solution to the written part of Assignment 2 on the midterm day (Oct 20)!
  - ▶ Closed by 5pm, Sept 23

---

## RECAP: ADVERSARIAL SEARCH

- ▶ Minimax search:
  - ▶ The max player's minimax value is the maximum of its successors' values
  - ▶ The min player's minimax value is the minimum of its successors' values
- ▶ Alpha-Beta pruning:
  - ▶ If the current estimate for a min player's value is lower than the estimate for some of its ancestor max player's value, prune the remaining branches!
  - ▶ If the current estimate for a max player's value is higher than the estimate for some of its ancestor min player's value, prune the remaining branches!

# Recap: Alpha-Beta Implementation

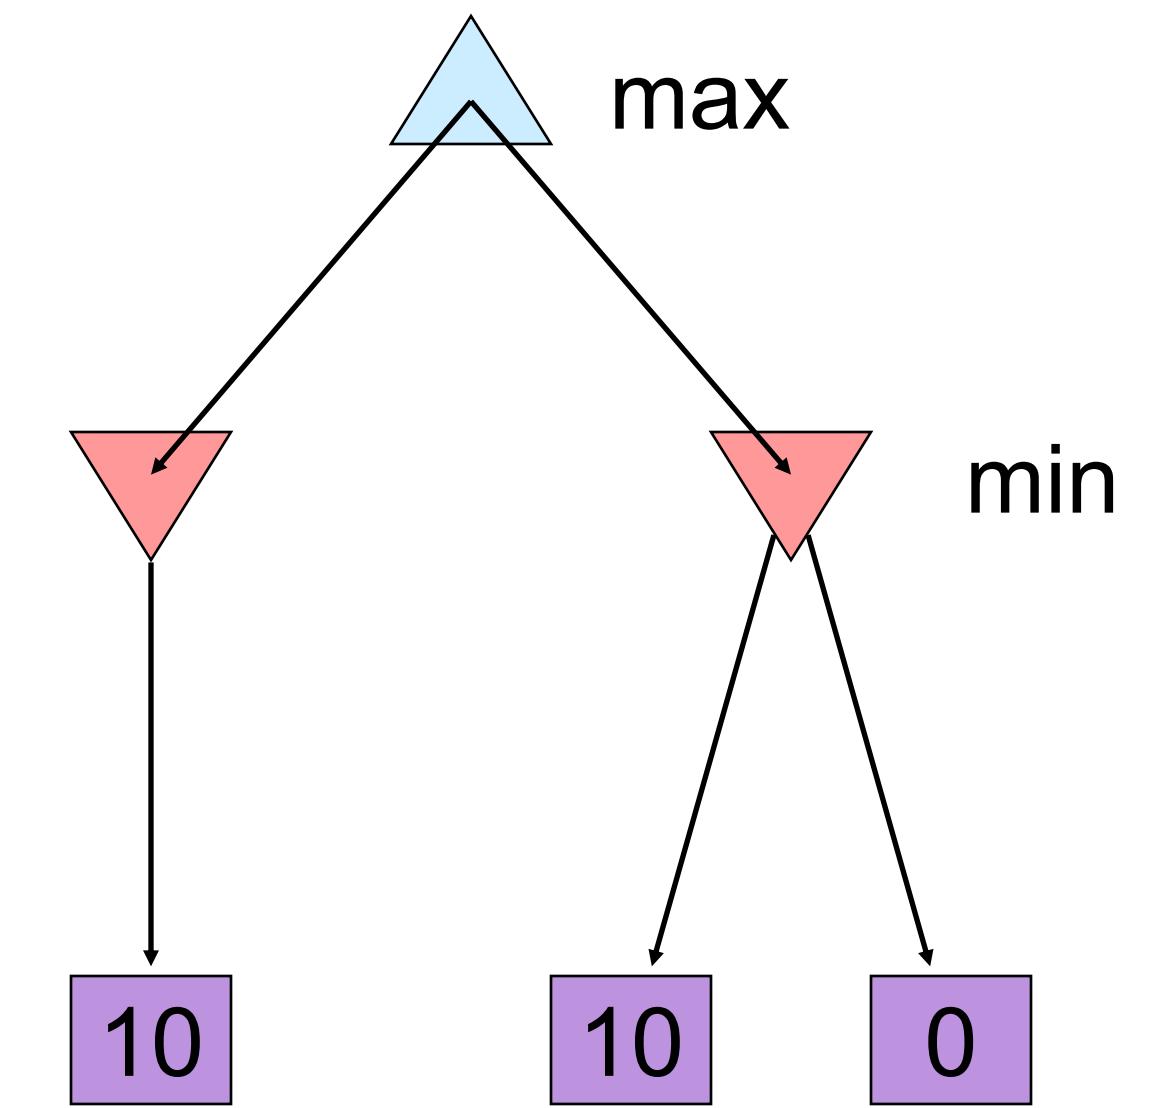
$\alpha$ : MAX's best option on path to root  
 $\beta$ : MIN's best option on path to root

```
def max-value(state, α, β):  
    initialize v = -∞  
    for each successor of state:  
        v = max(v, value(successor,α,β))  
        if v ≥ β return v  
        α = max(α, v)  
    return v
```

```
def min-value(state , α, β):  
    initialize v = +∞  
    for each successor of state:  
        v = min(v, value(successor,α,β))  
        if v ≤ α return v  
        β = min(β, v)  
    return v
```

# ALPHA-BETA PRUNING PROPERTIES

- ▶ This pruning has **no effect** on minimax value computed for the root
- ▶ Values of intermediate nodes might be wrong
  - ▶ Important: children of the root may have the wrong value
  - ▶ So the most naïve version won't let you do action selection
- ▶ Good child ordering improves effectiveness of pruning
- ▶ With "perfect ordering":
  - ▶ Time complexity drops to  $O(b^{m/2})$
  - ▶ Doubles solvable depth
  - ▶ Full search of, e.g. chess, is still hopeless...
- ▶ This is a simple example of **metareasoning** (computing about what to compute)



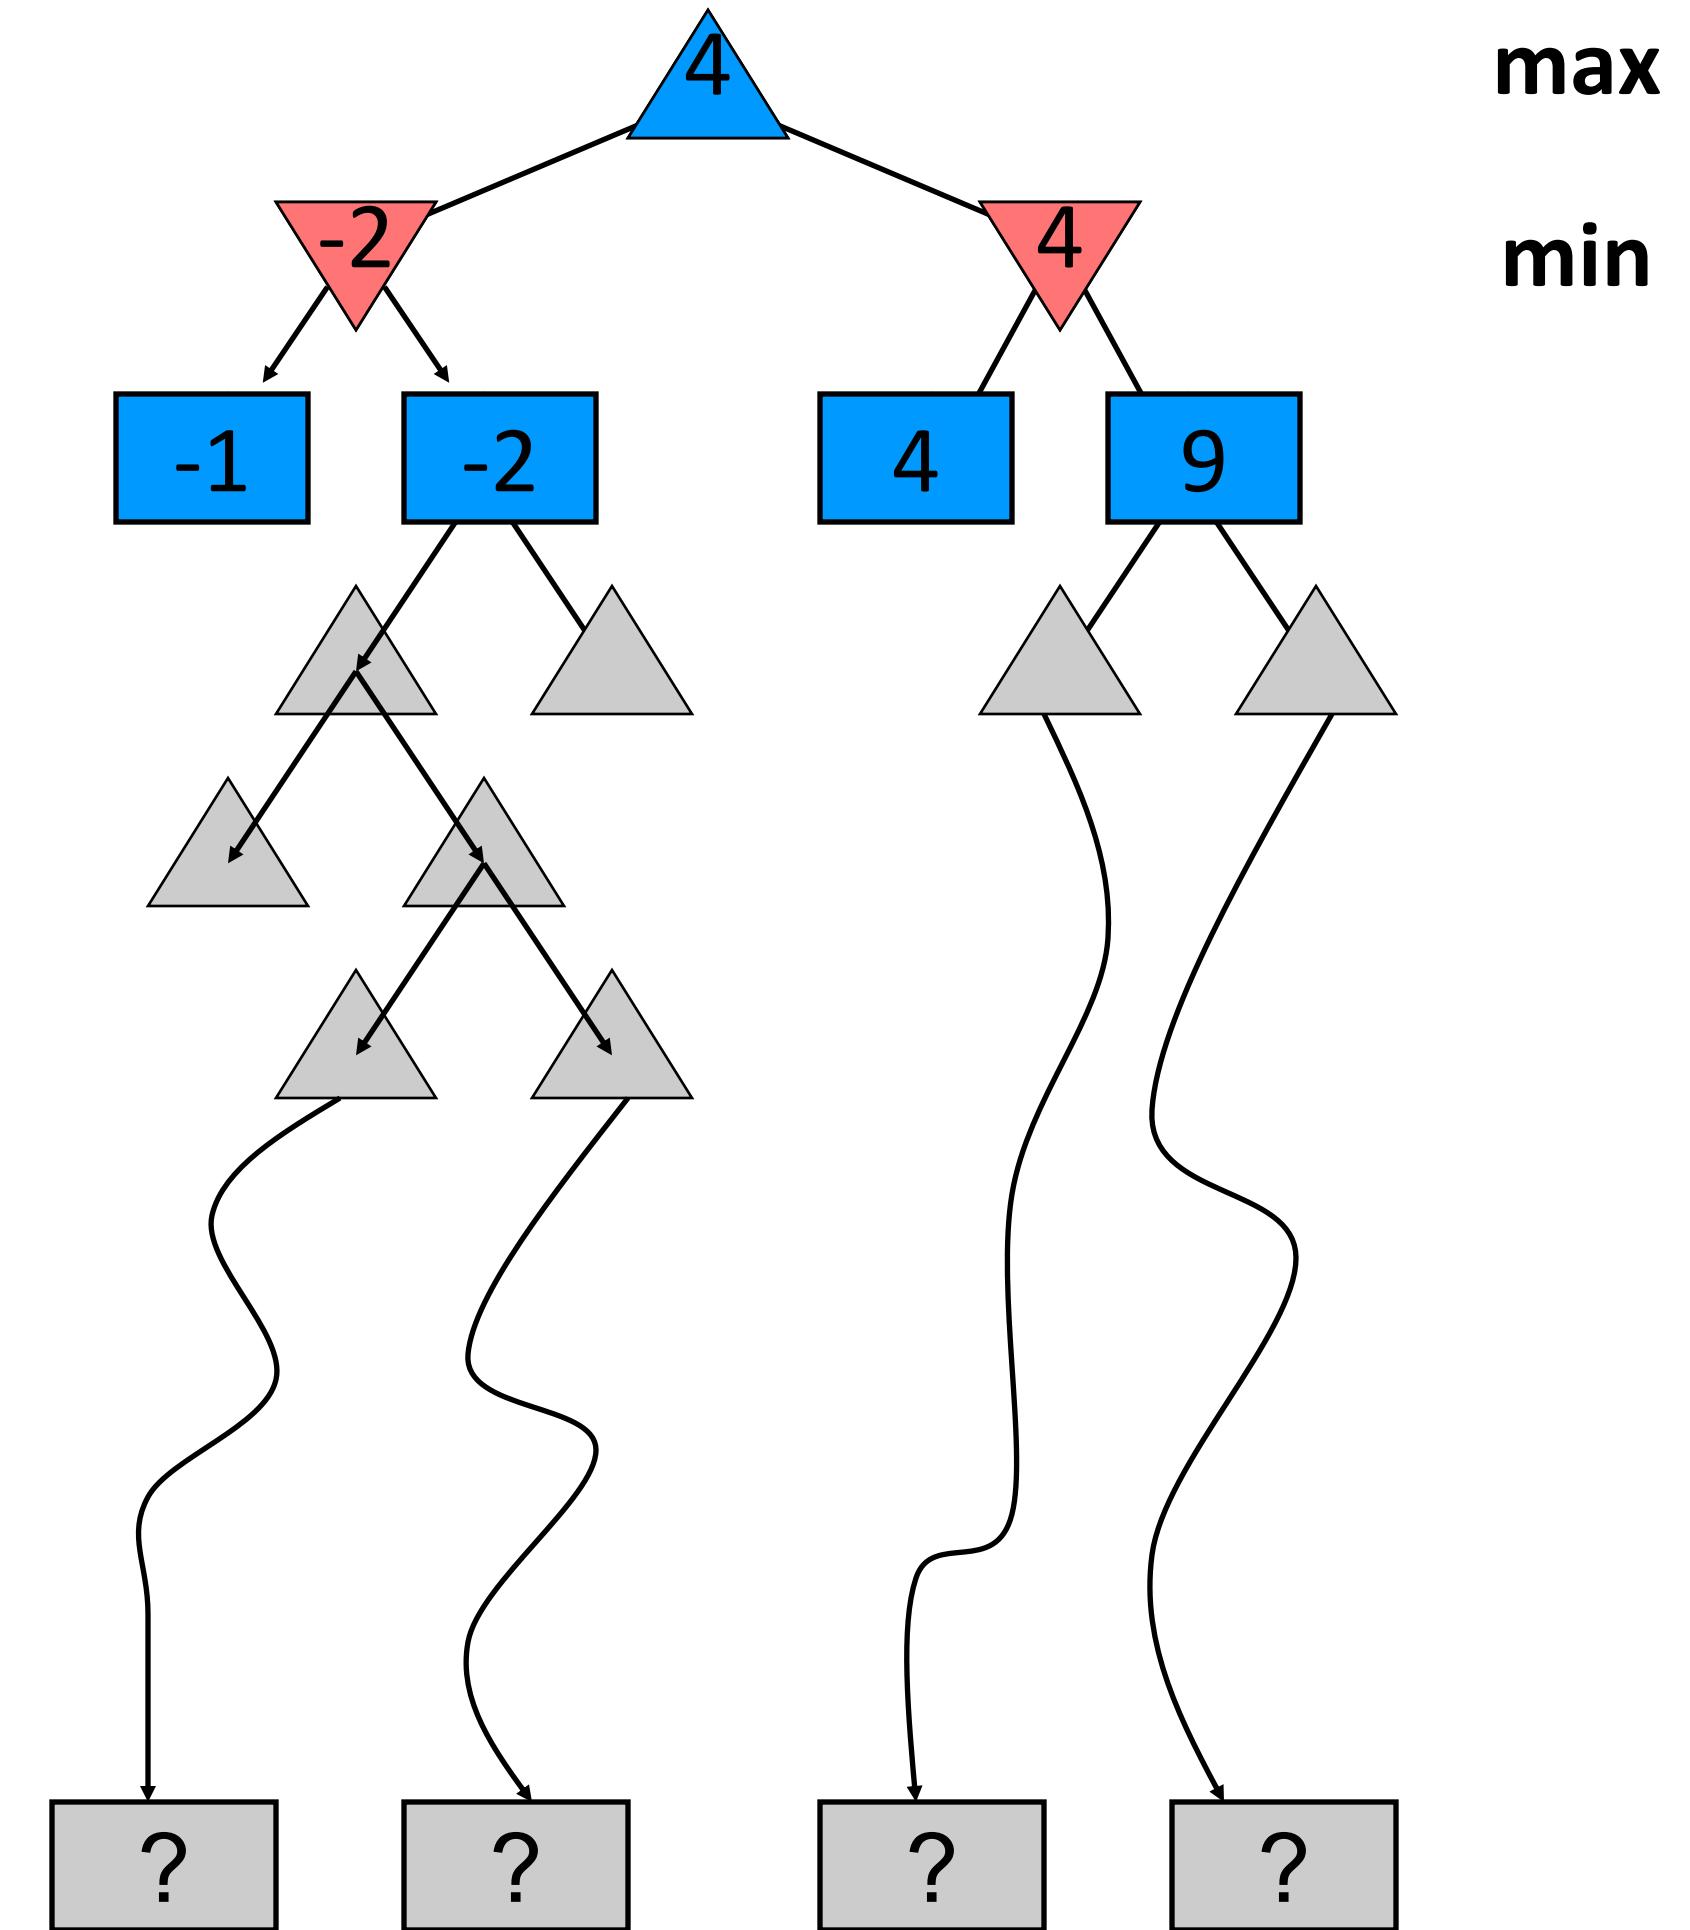
---

# RESOURCE LIMITS

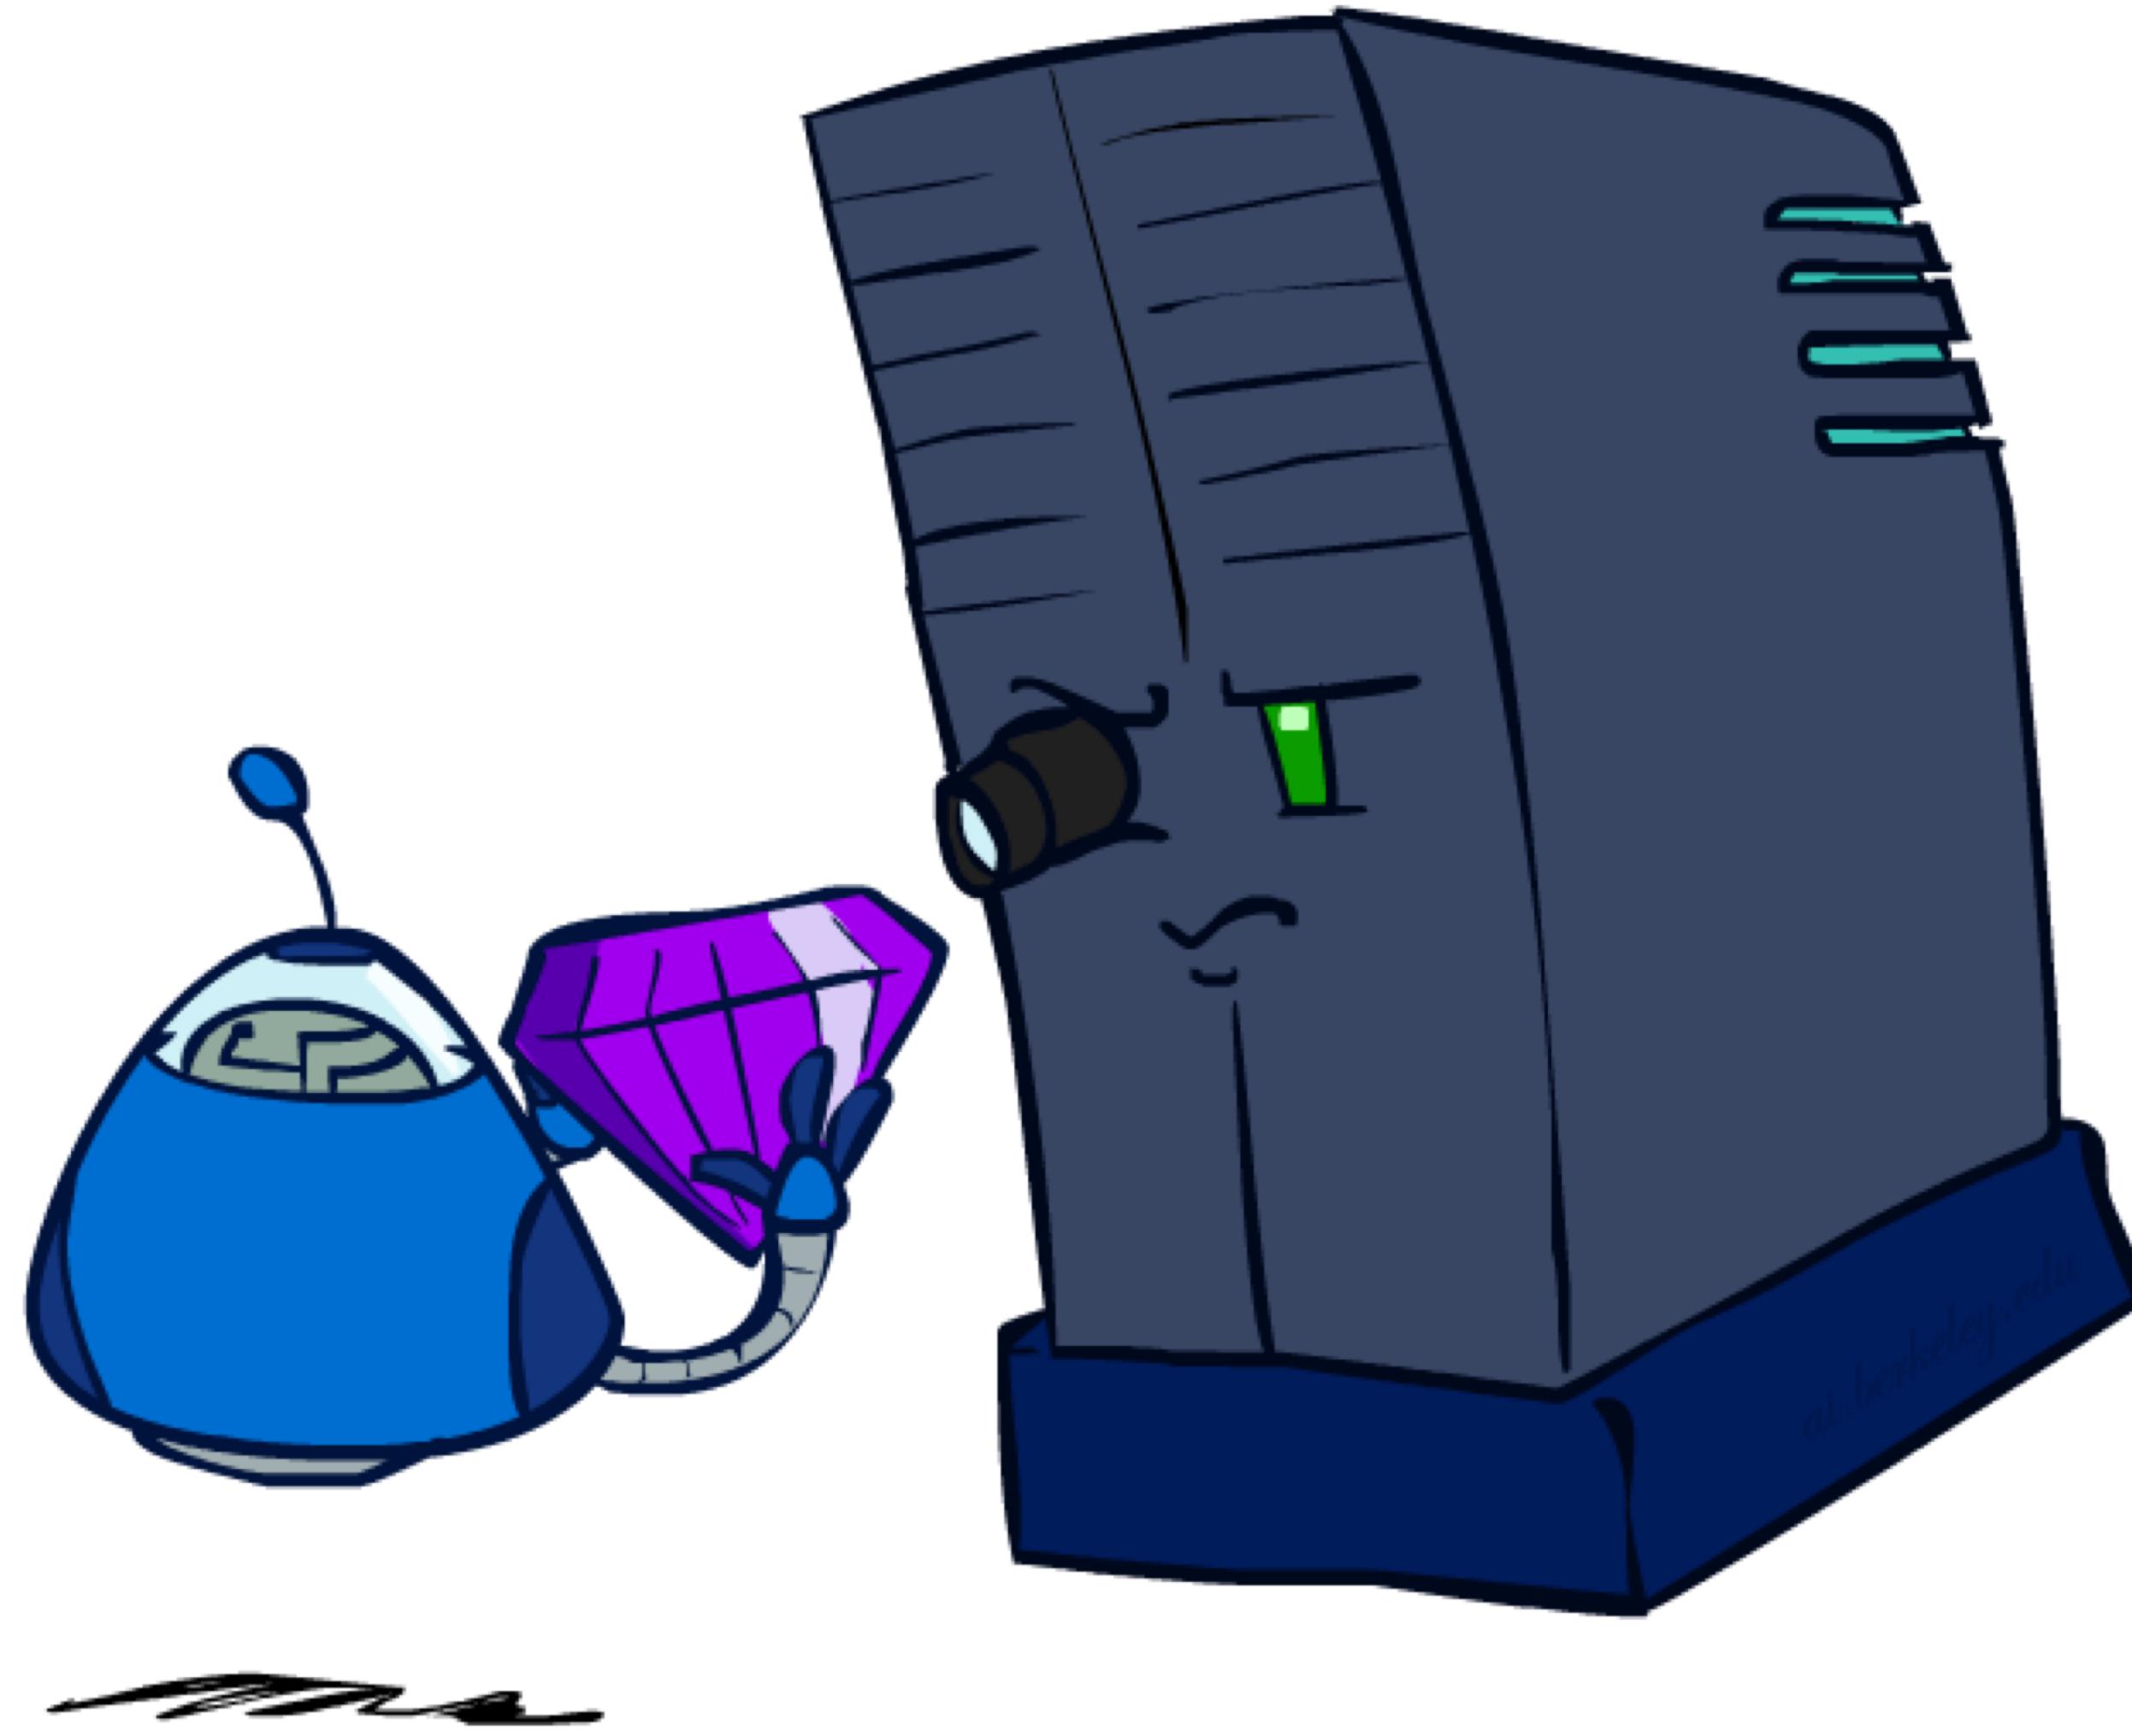


# RESOURCE LIMITS

- ▶ **Problem:** In realistic games, cannot search to leaves
- ▶ **Example:**
  - ▶ Suppose we have 100 seconds, can explore 10K nodes / sec
  - ▶ So can check 1M nodes per move
  - ▶  $\alpha\beta$  reaches about depth 8 - decent chess program
- ▶ **Solution:** Depth-limited search
  - ▶ Instead, search only to a limited depth in the tree
  - ▶ Replace terminal utilities with an **evaluation function** for non-terminal positions
  - ▶ Guarantee of optimal play is gone; More plies makes a BIG difference



# EVALUATION FUNCTIONS



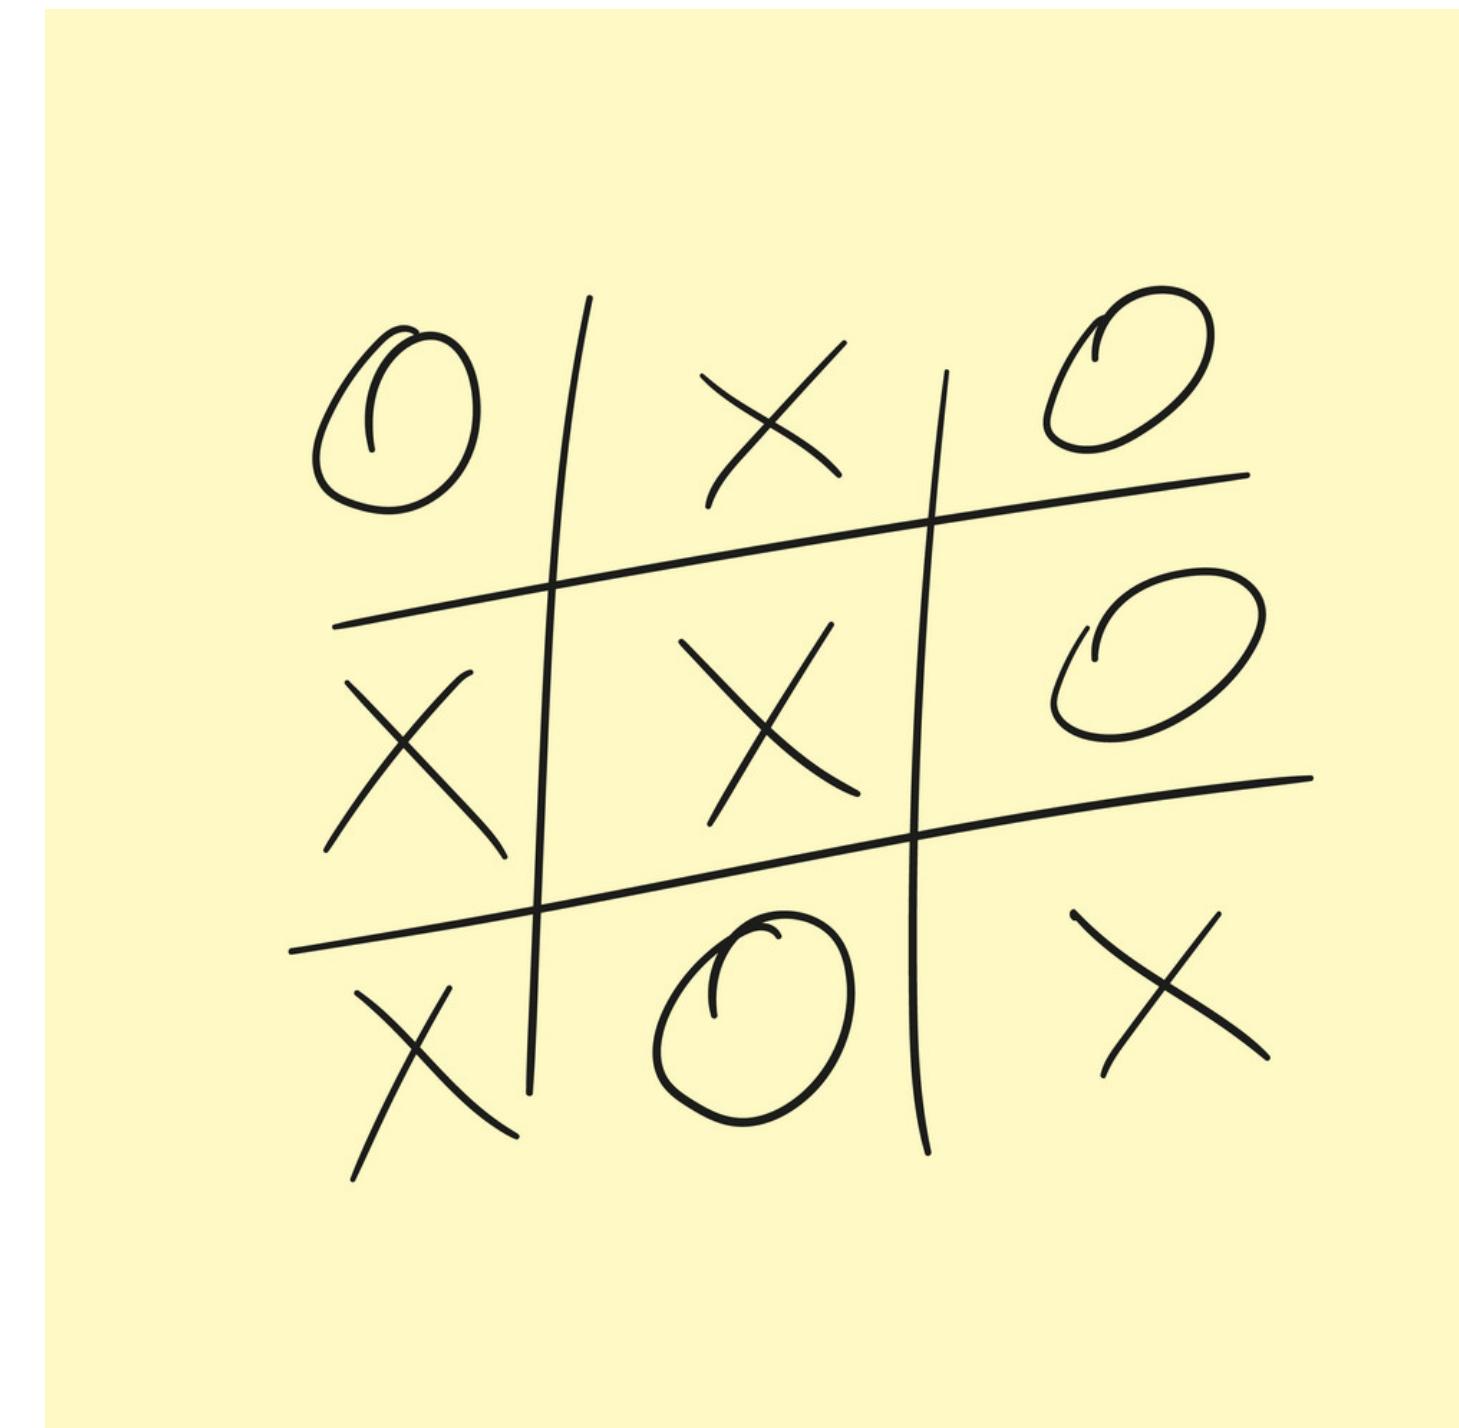
# EVALUATION FUNCTIONS

- ▶ Desirable properties
  - ▶ Order terminal states in same way as true utility function
  - ▶ Strongly correlated with the actual minimax value of the states
  - ▶ Efficient to calculate
- ▶ Typically use **features** – simple characteristics of the game state that are correlated with the *probability of winning*
- ▶ The evaluation function combines feature values to produce a score:

$$Eval(x) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s) = \sum_{i=1}^n w_i f_i(s)$$

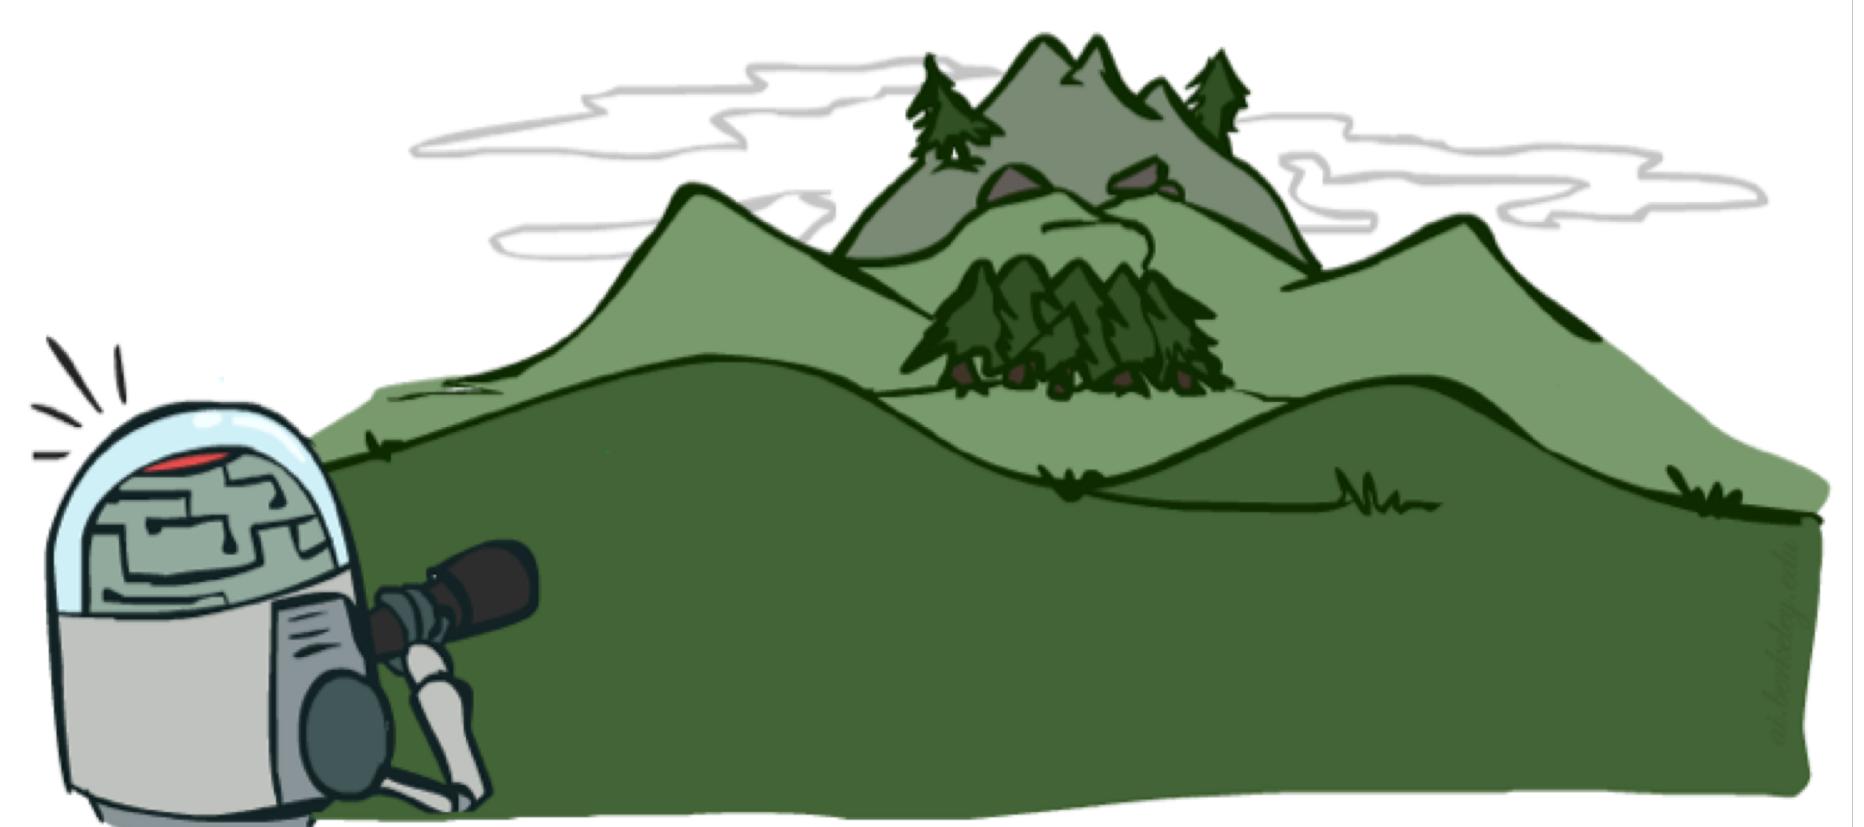
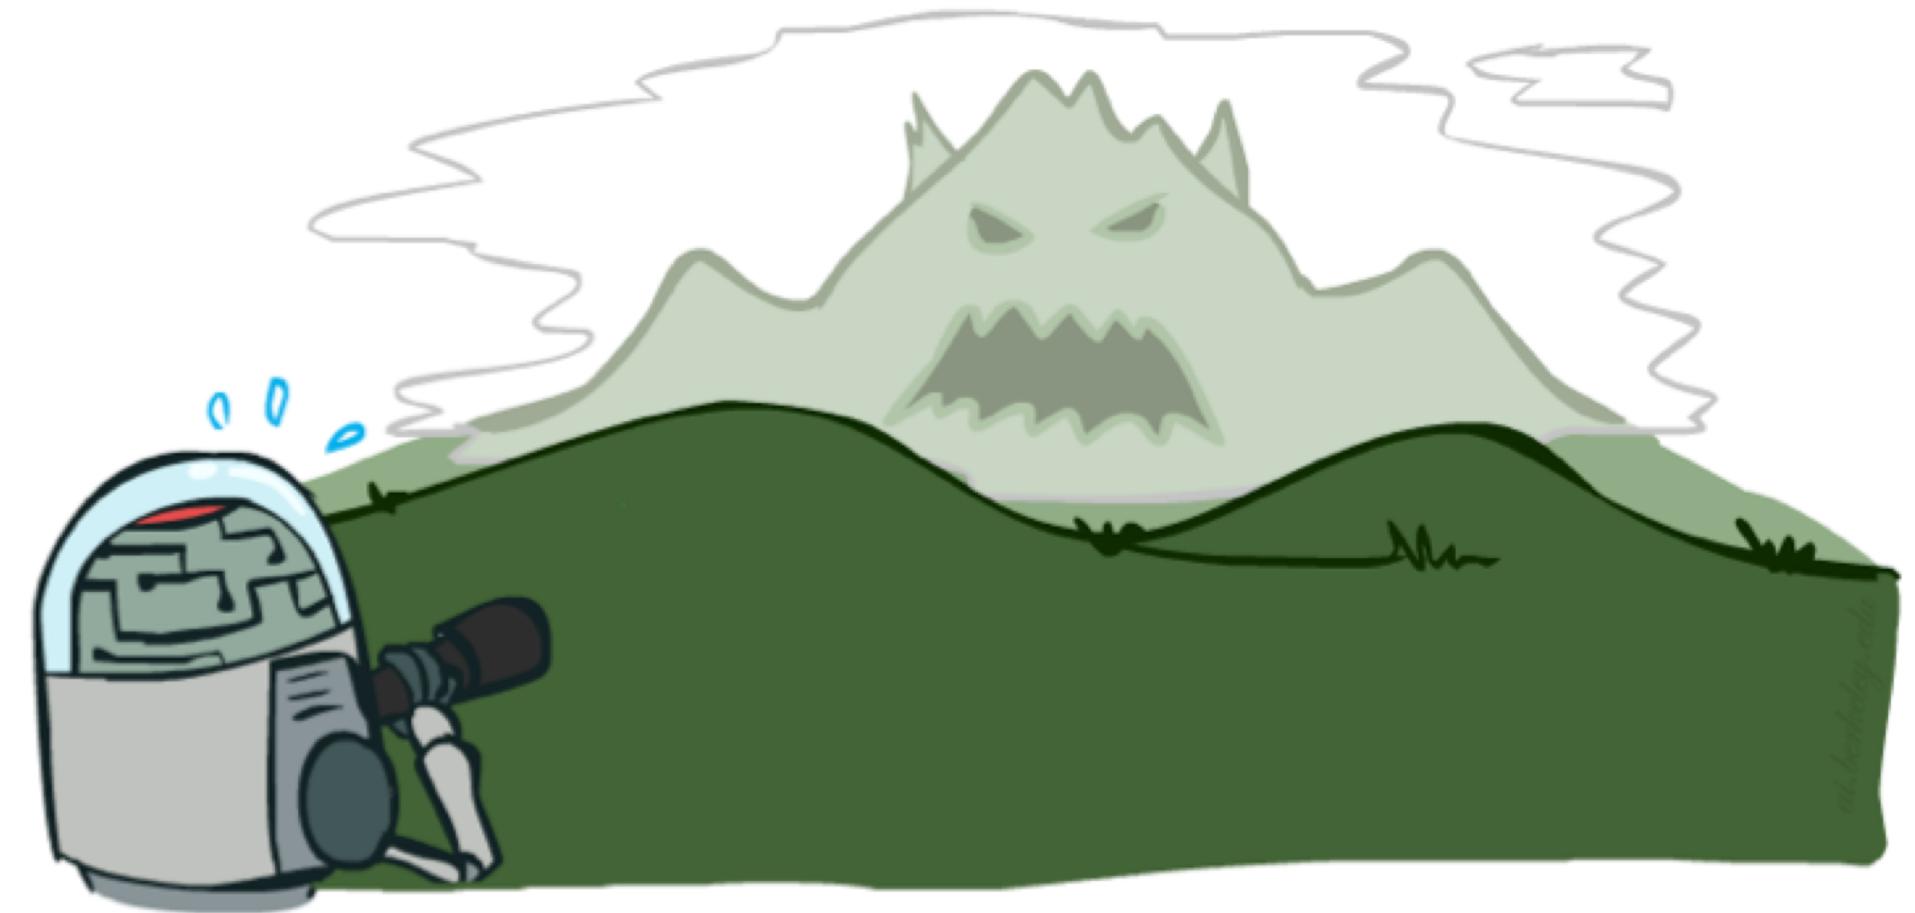
## EXAMPLE FEATURES

- ▶ What would be some useful features for tic-tac-toe?
  - ▶ The number of rows/columns you symbols spread?
  - ▶ Among all rows/columns/diagonals, how many are still reachable by you?
  - ▶ Appearance of winning patterns

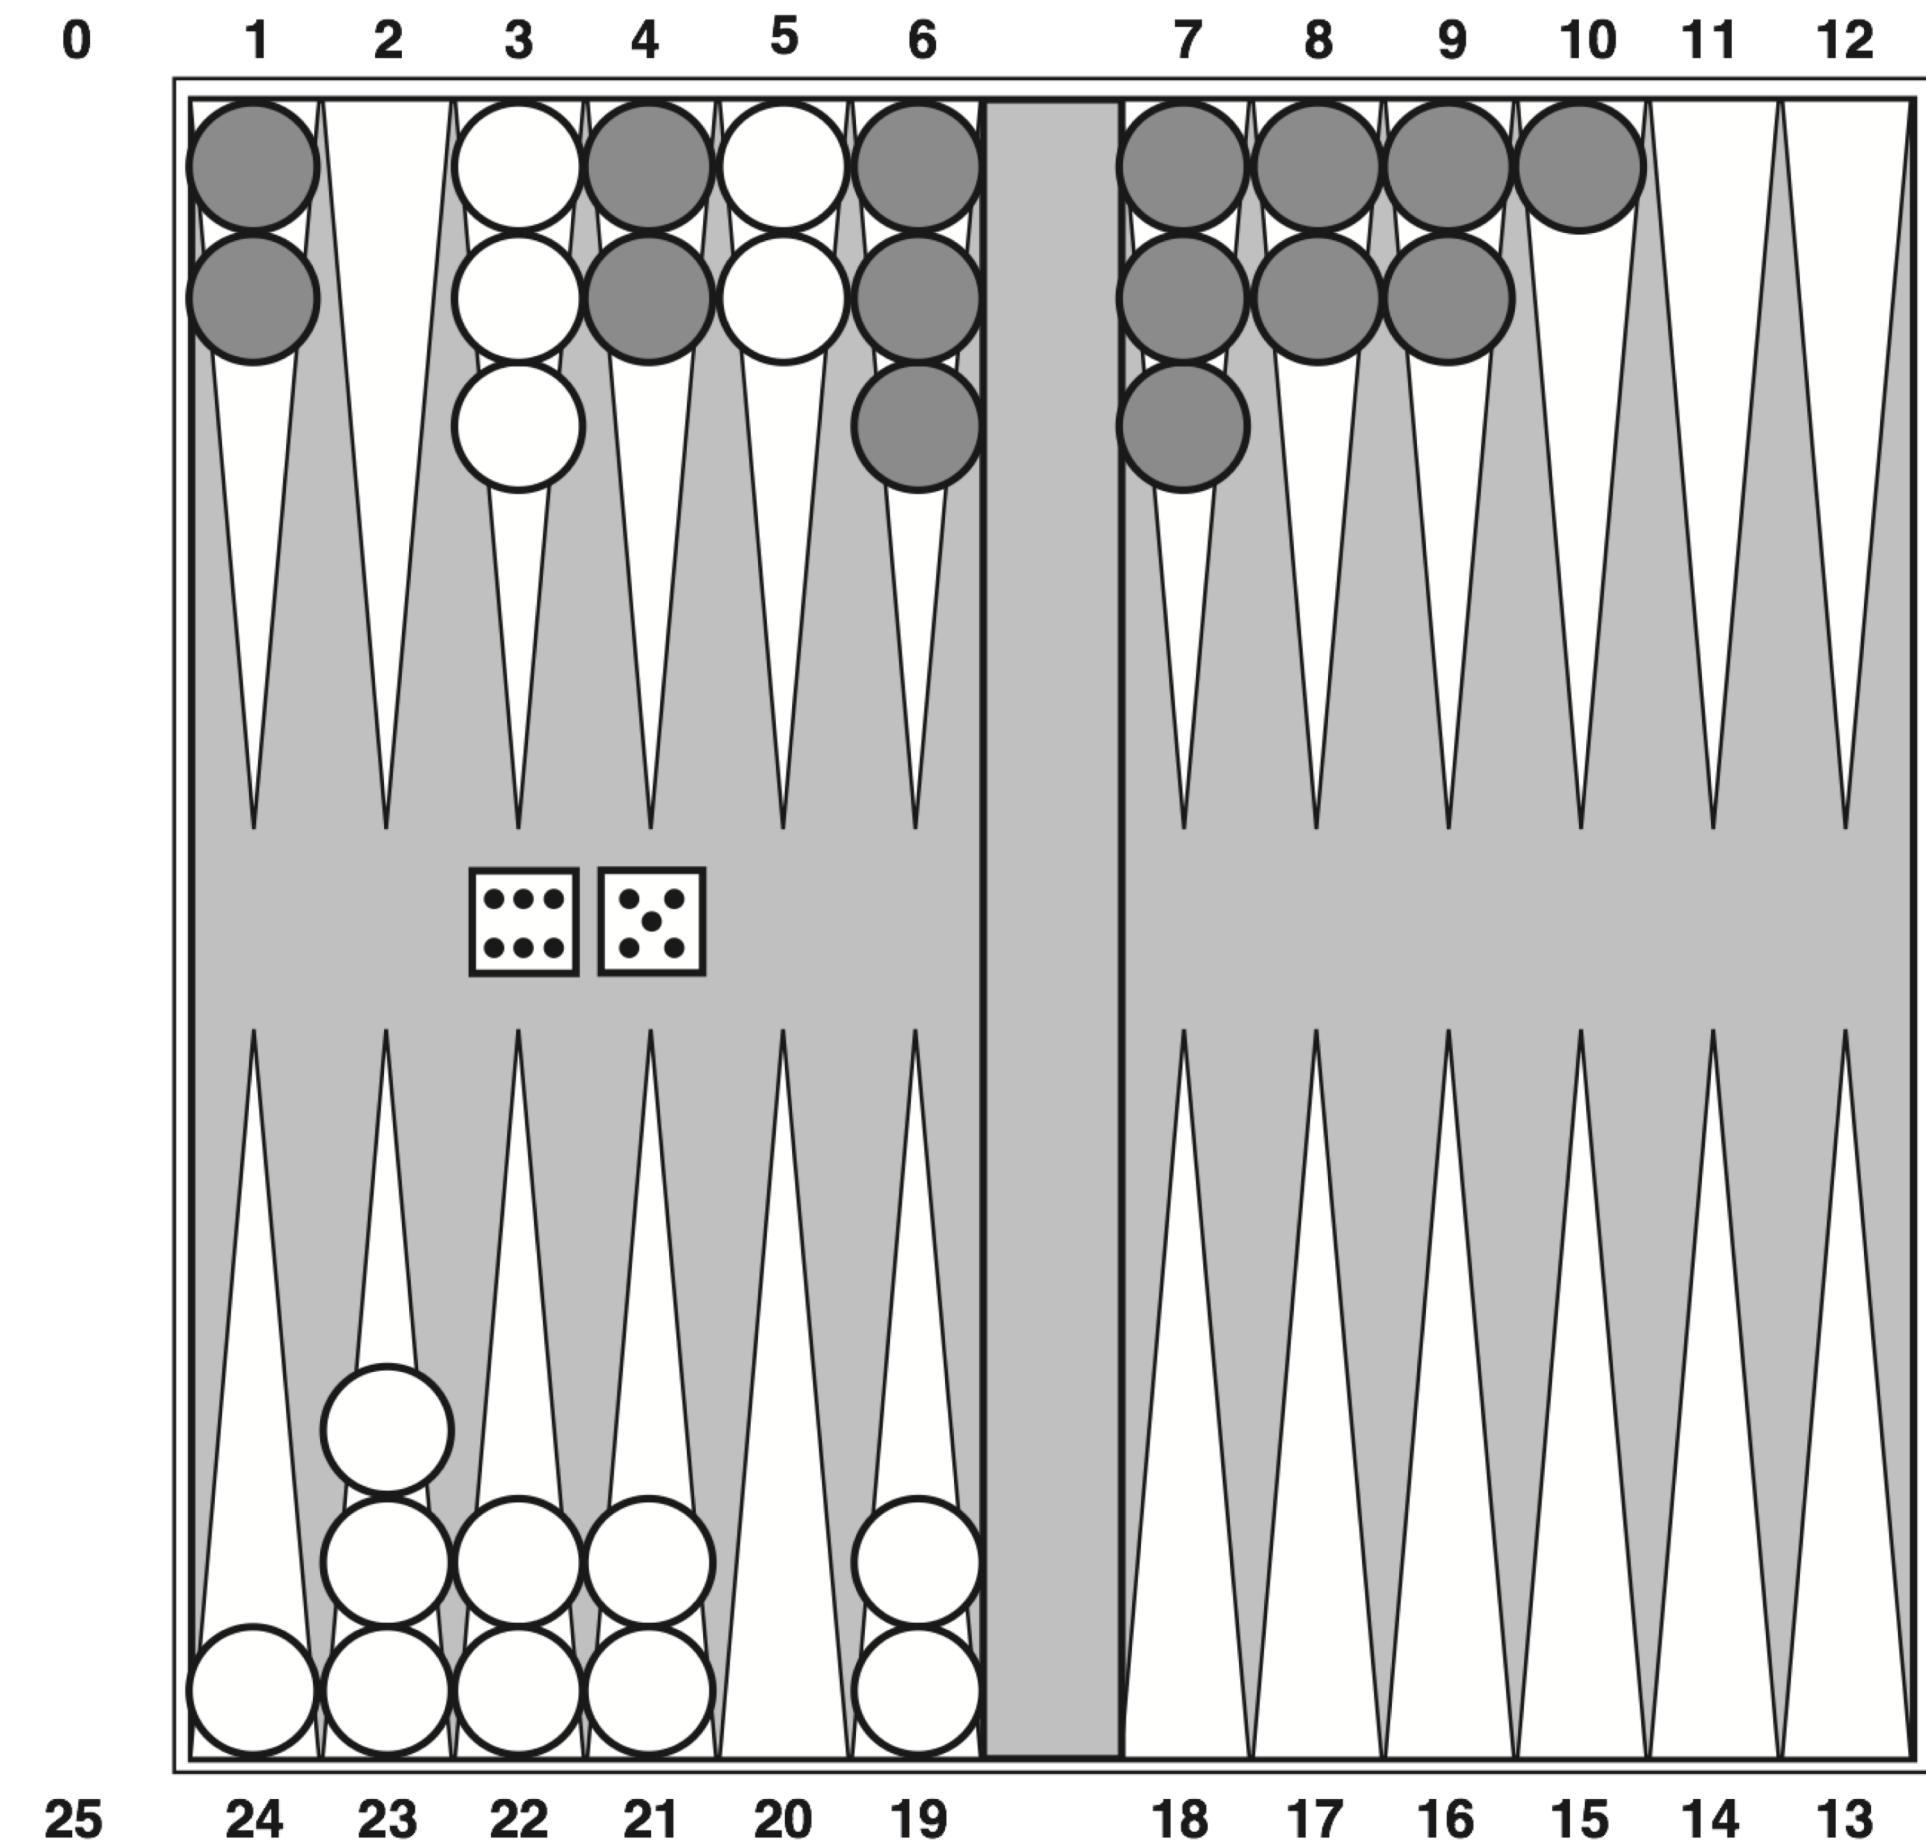


# DEPTH MATTERS

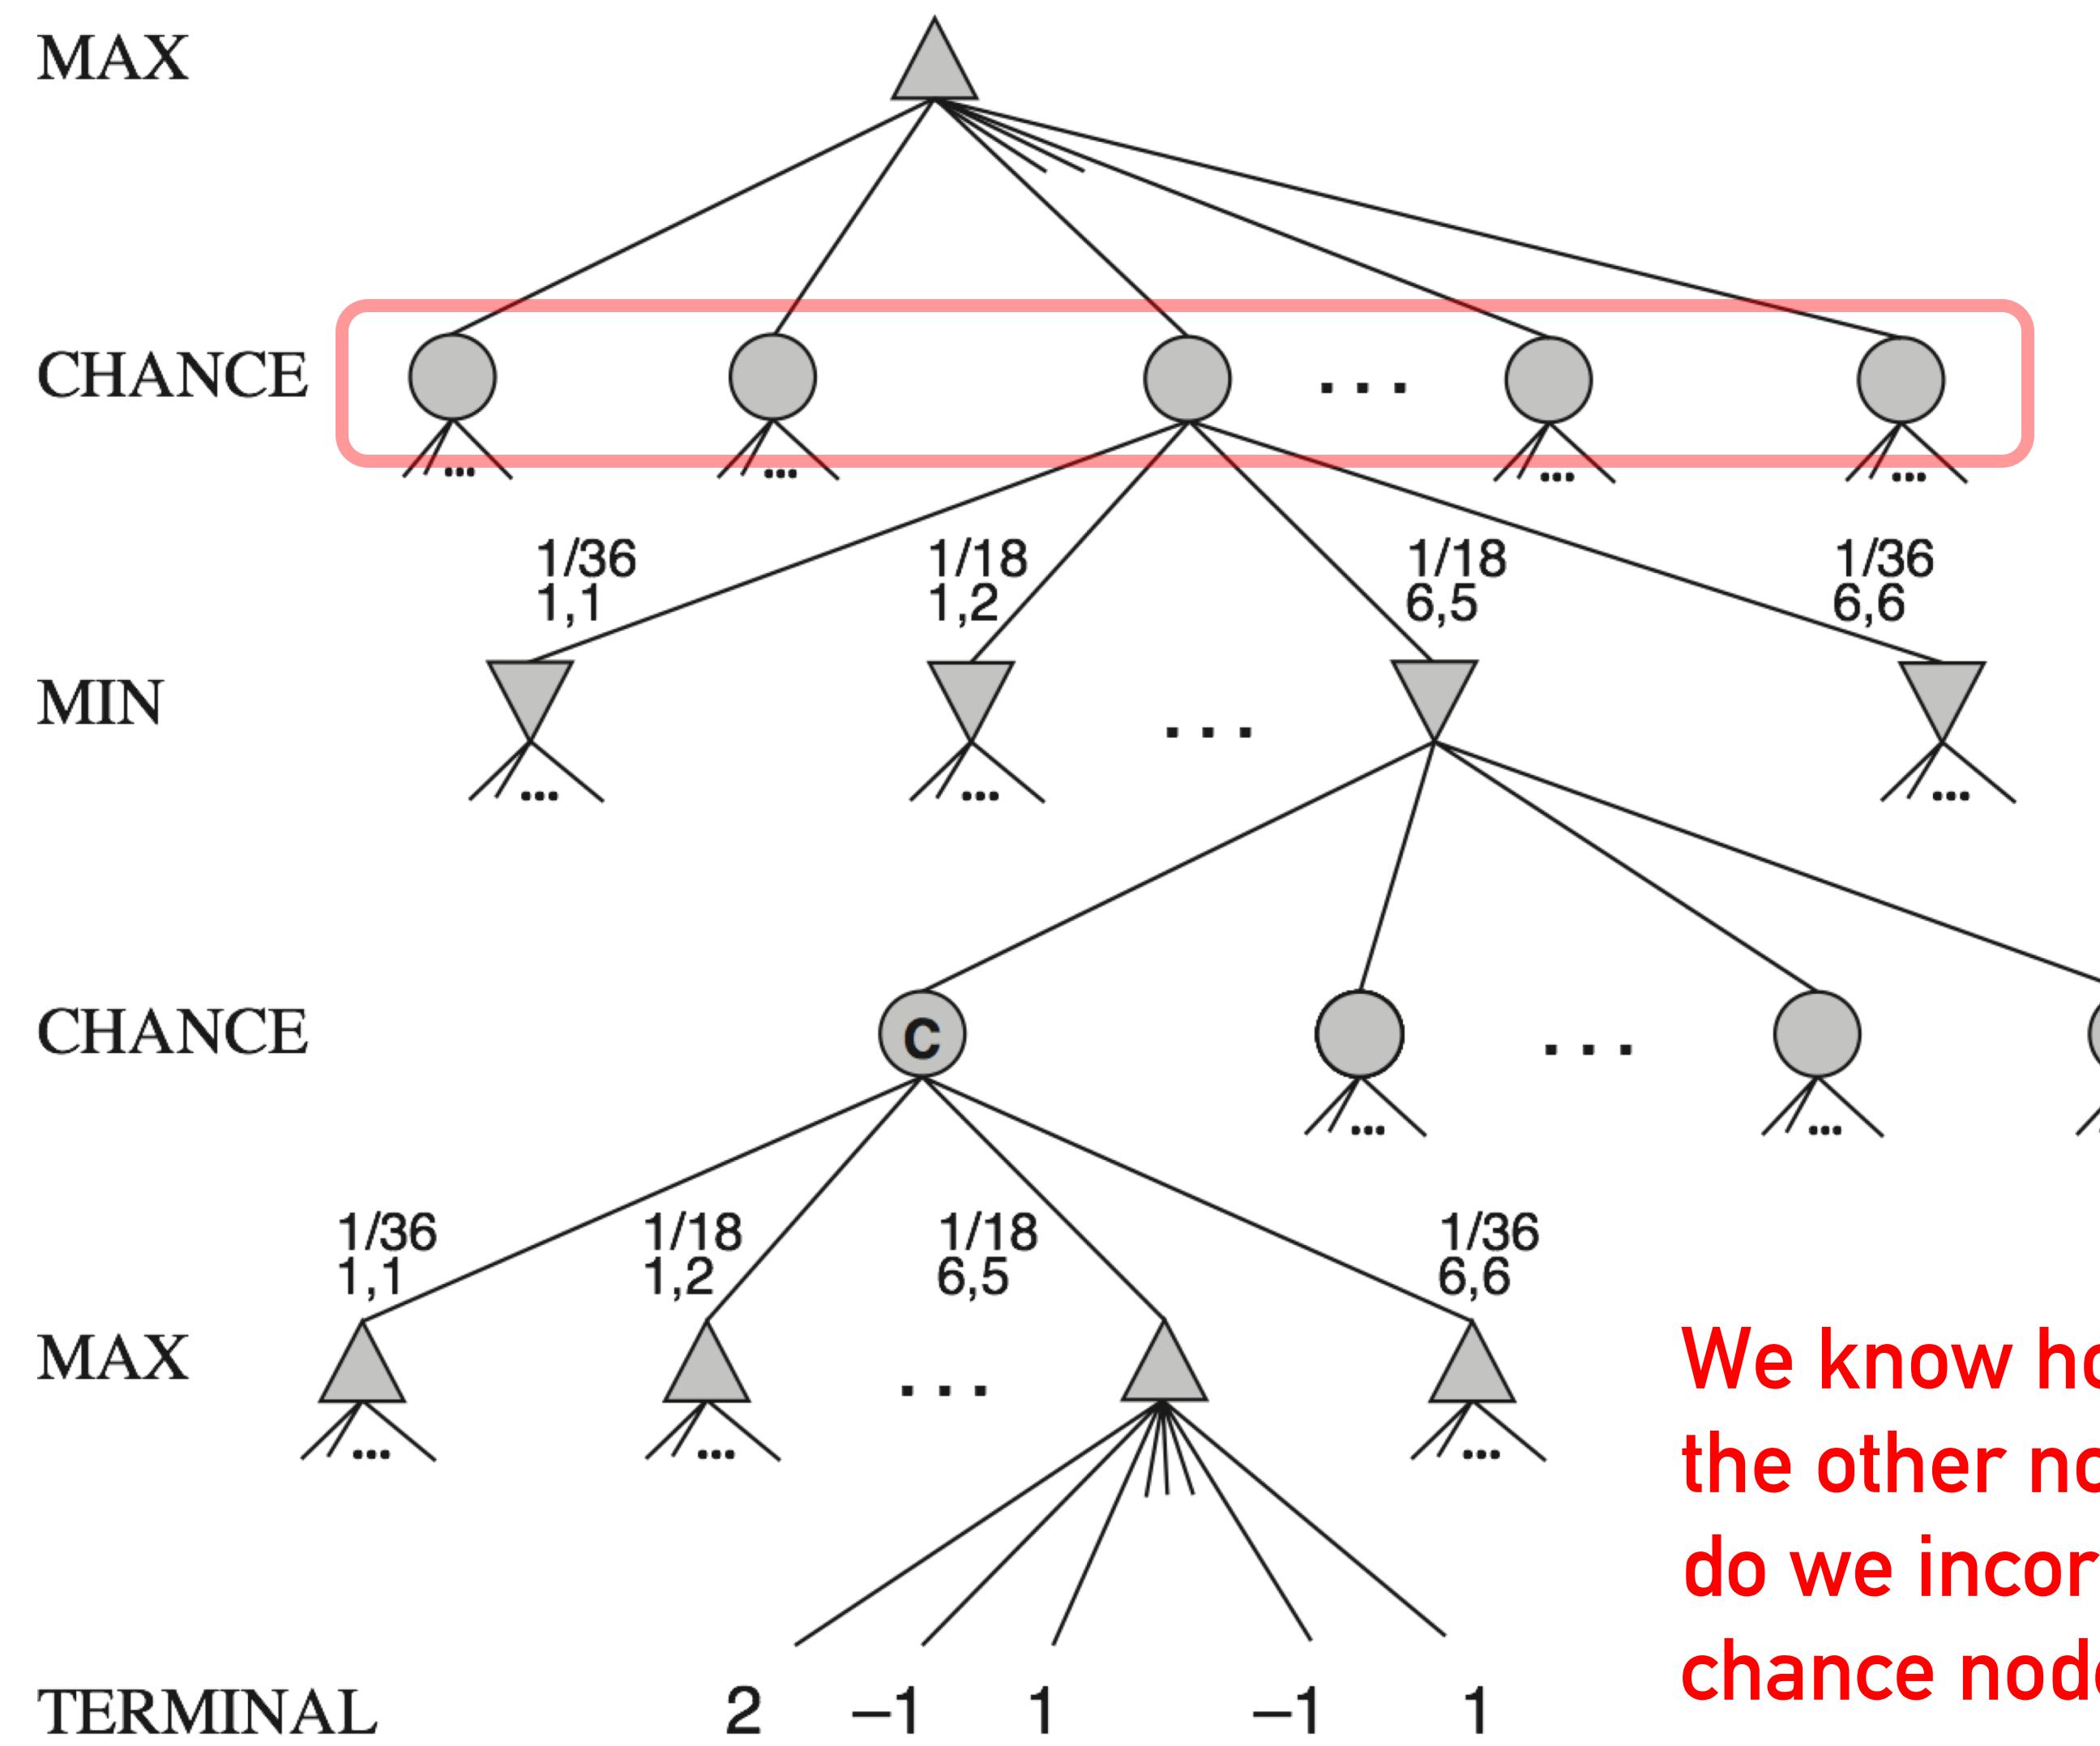
- ▶ Evaluation functions are always imperfect
- ▶ The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters
- ▶ An important example of the tradeoff between complexity of features and complexity of computation



## WHAT IF A GAME HAS A “CHANCE ELEMENT”?



## GAME TREE WITH CHANCE ELEMENT



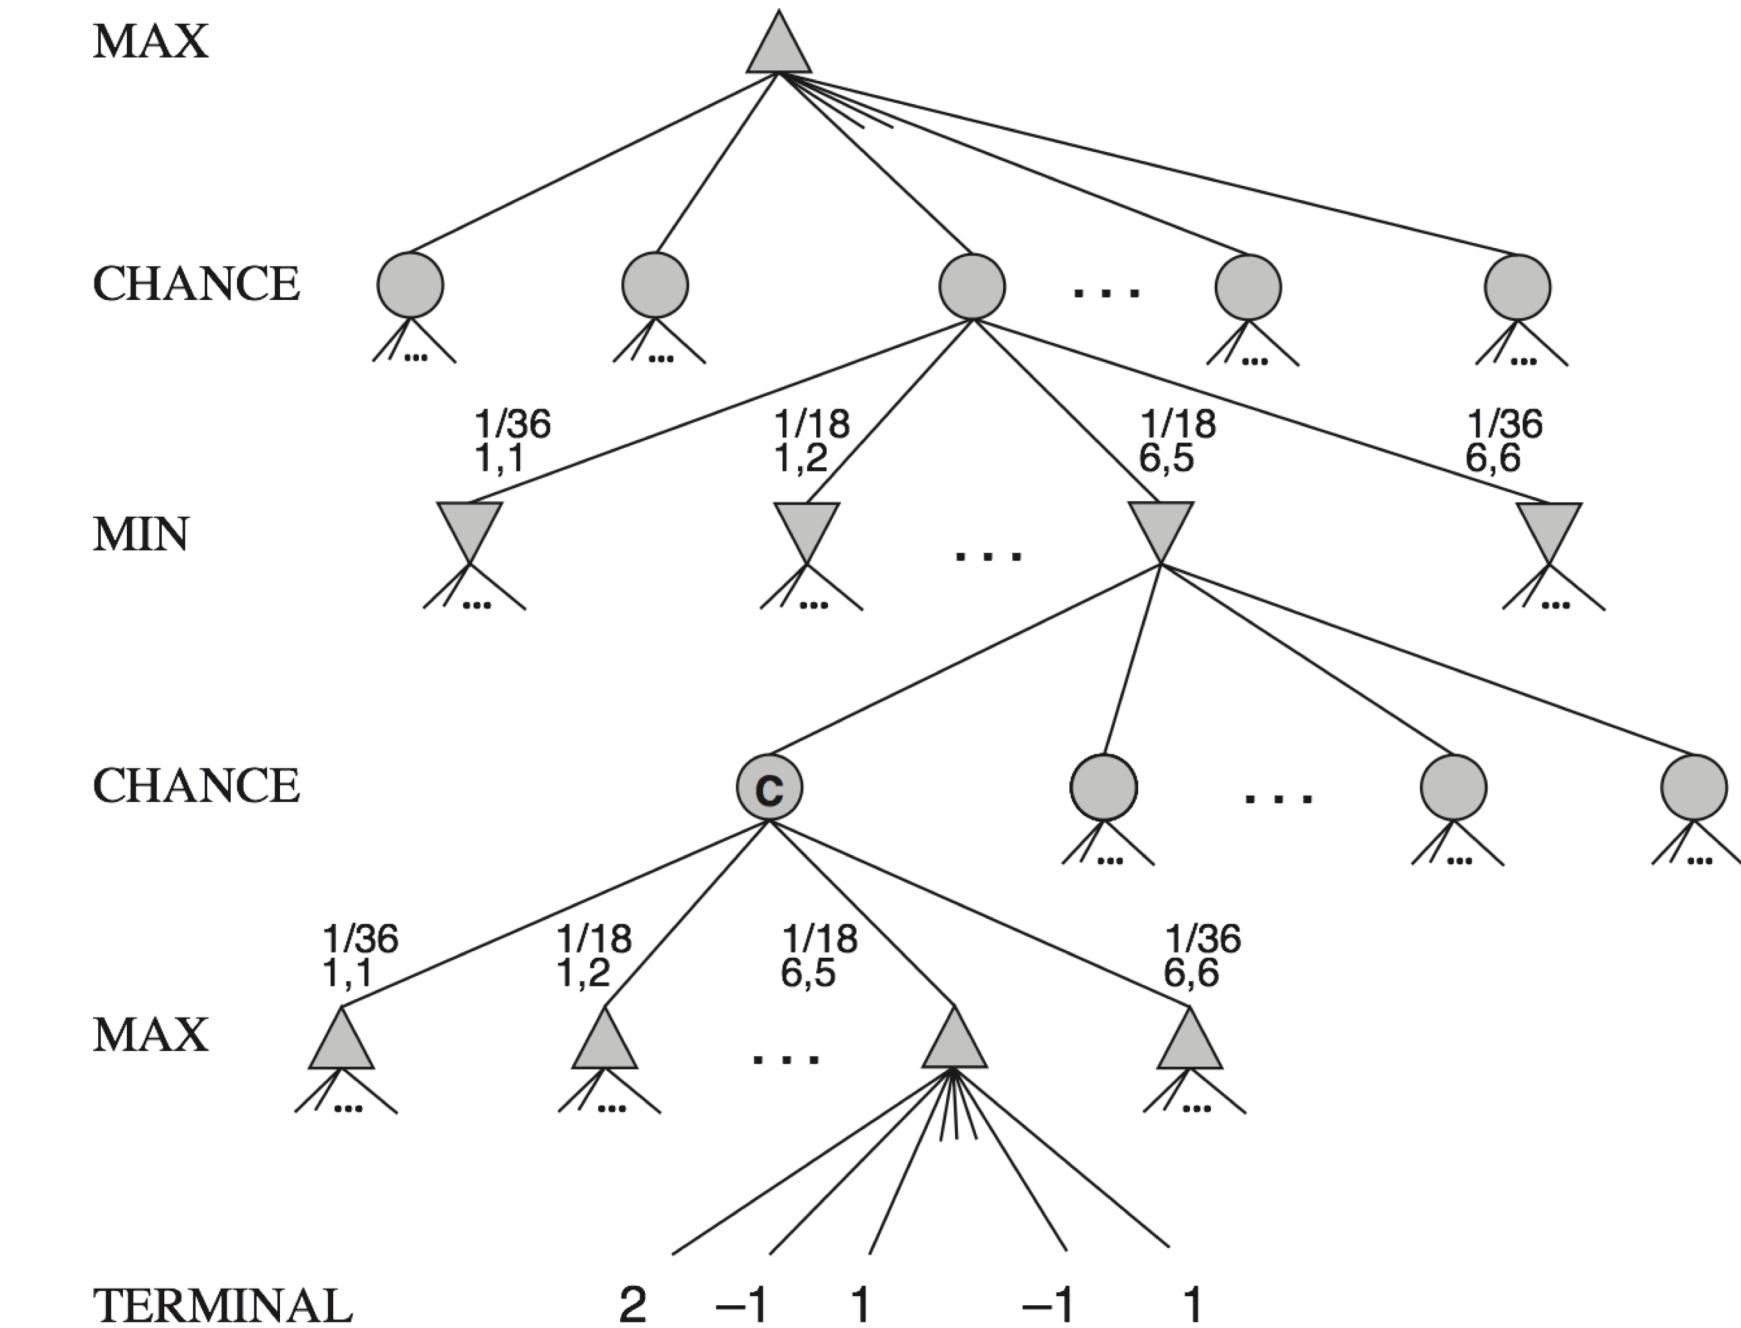
We know how to value the other nodes. How do we incorporate chance nodes?

# EXPECTED MINIMAX VALUE

- The sum of the probability of each possible outcome multiplied by its value:

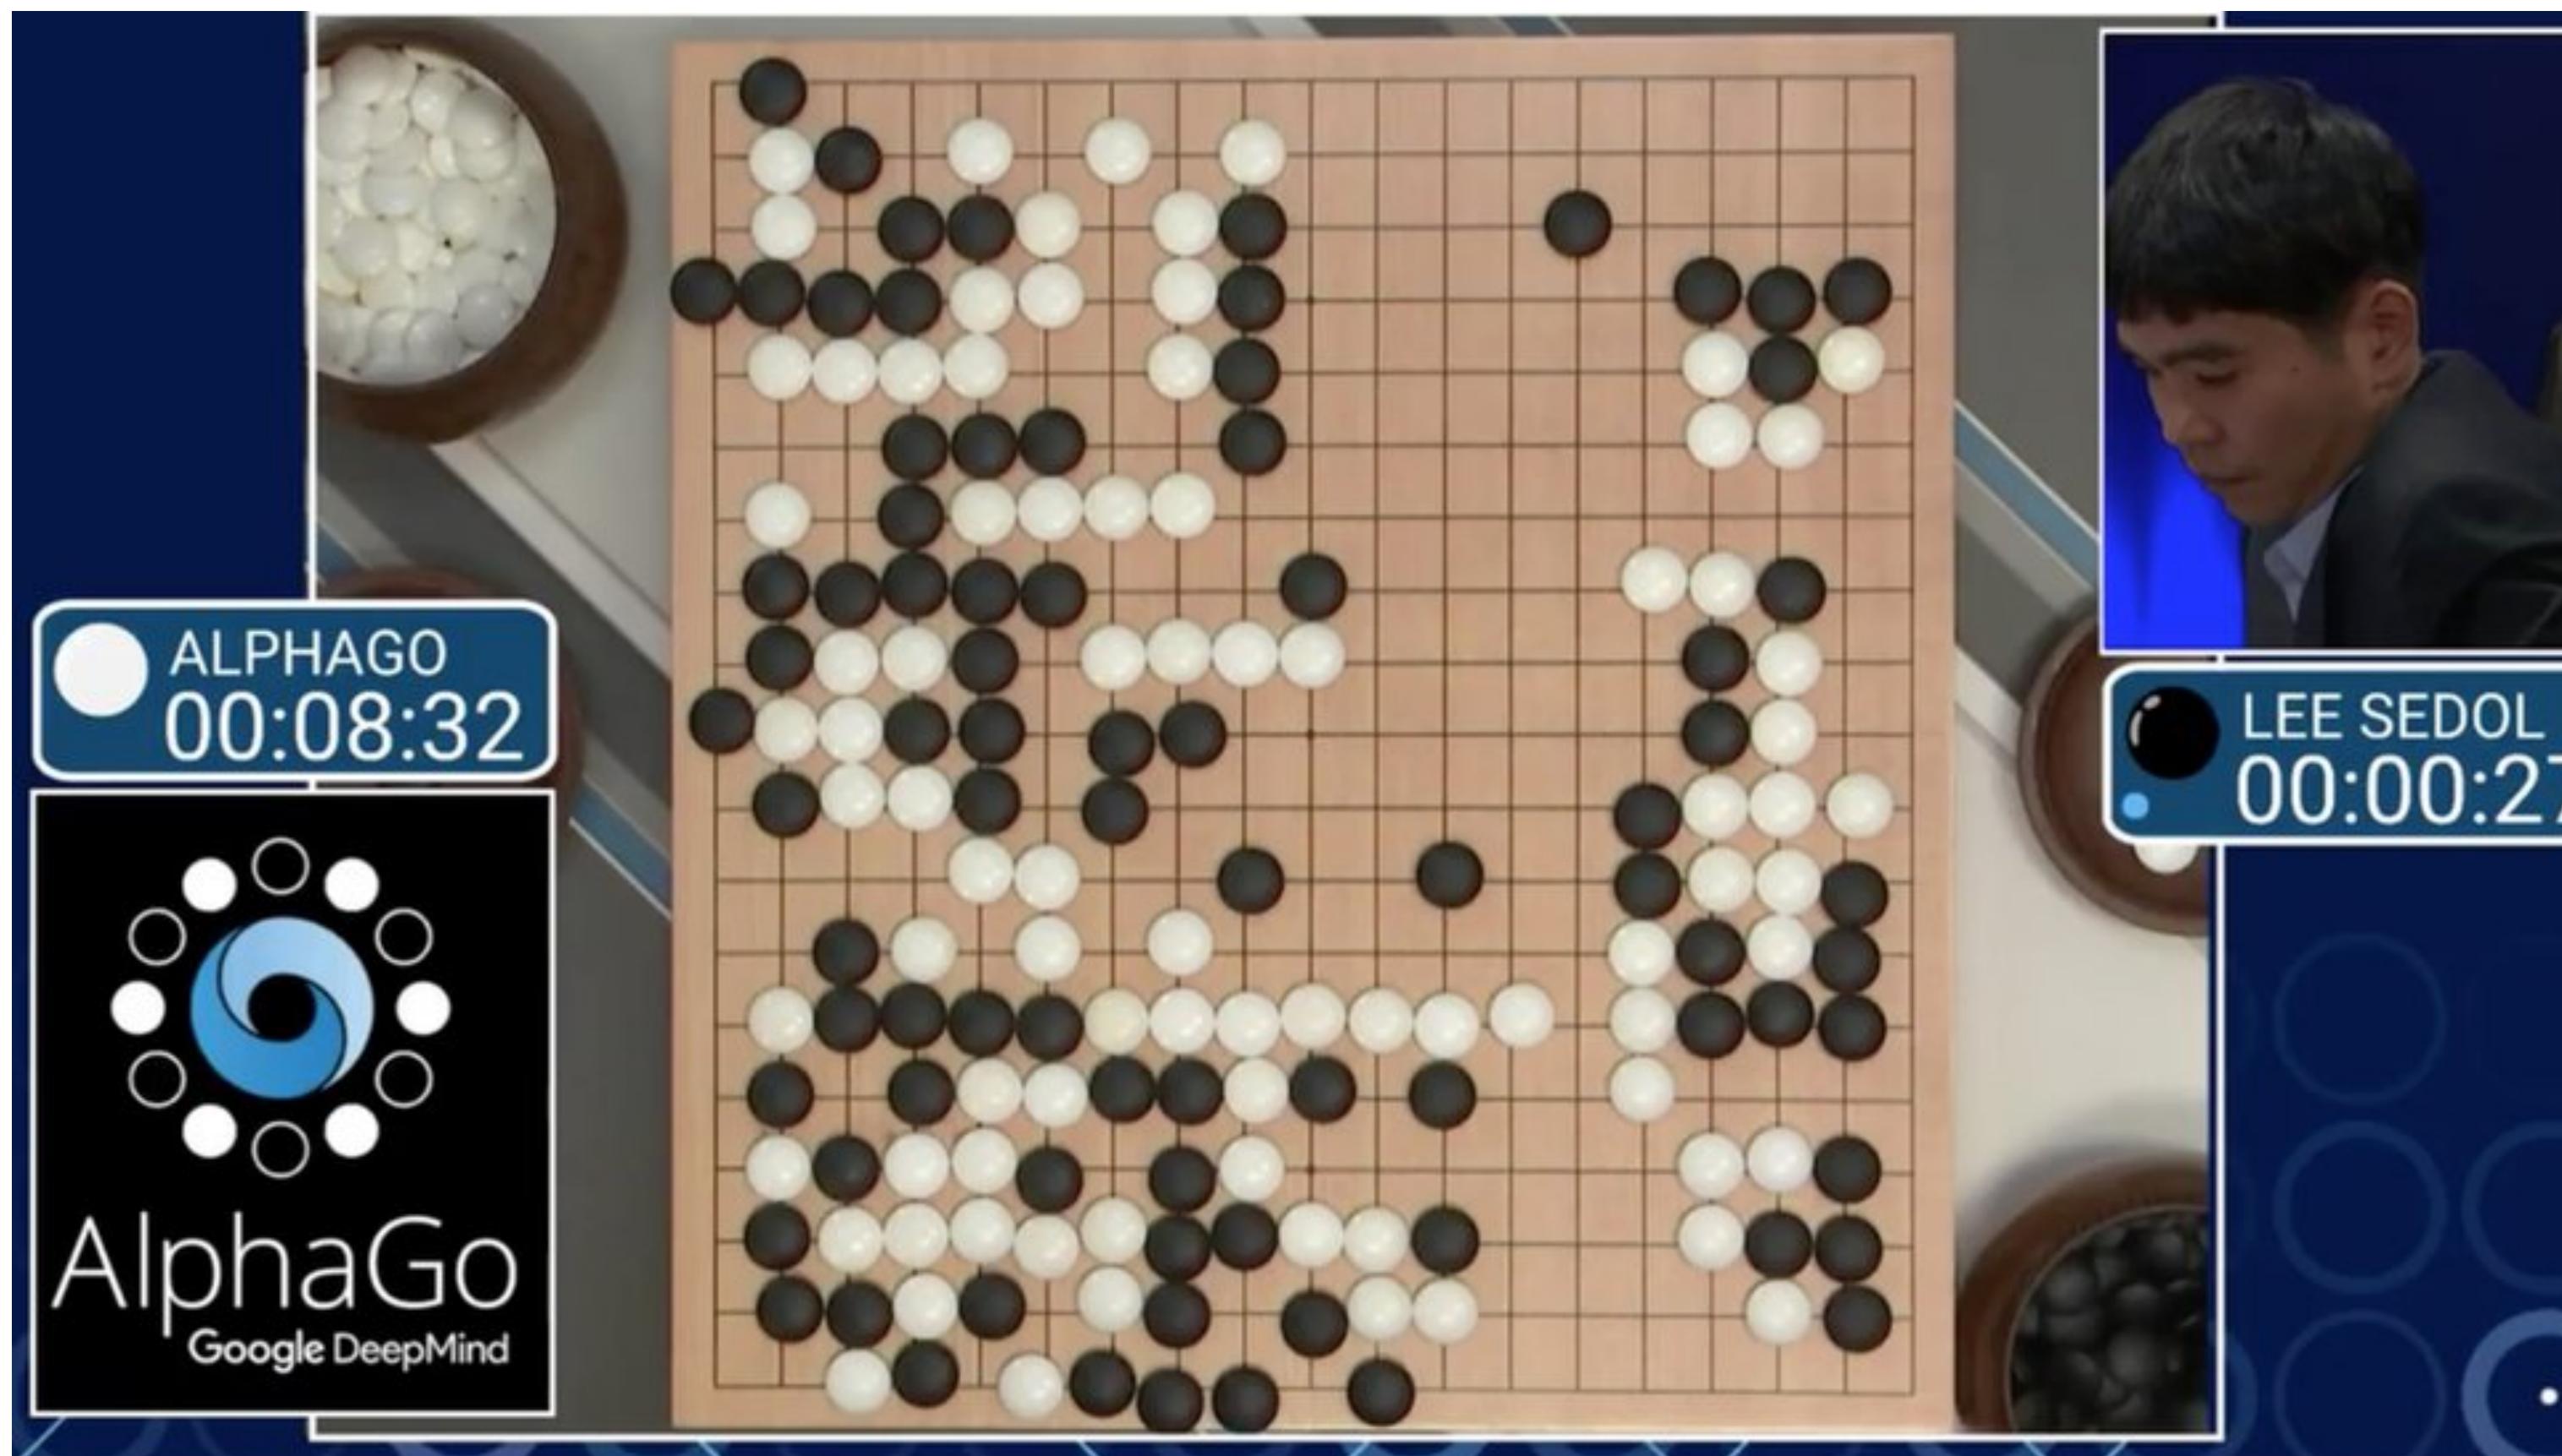
$$E(X) = \sum_i p_i x_i$$

**EXPECTED-MINIMAX-VALUE(n) =**  
 UTILITY(n)  
 $\max_{s \in \text{successors}(n)} \text{EXPECTED-MINIMAX-VALUE}(s)$   
 $\min_{s \in \text{successors}(n)} \text{EXPECTED-MINIMAX-VALUE}(s)$   
 $\sum_{s \in \text{successors}(n)} P(s) \times \text{EXPECTED-MINIMAX-VALUE}(s)$



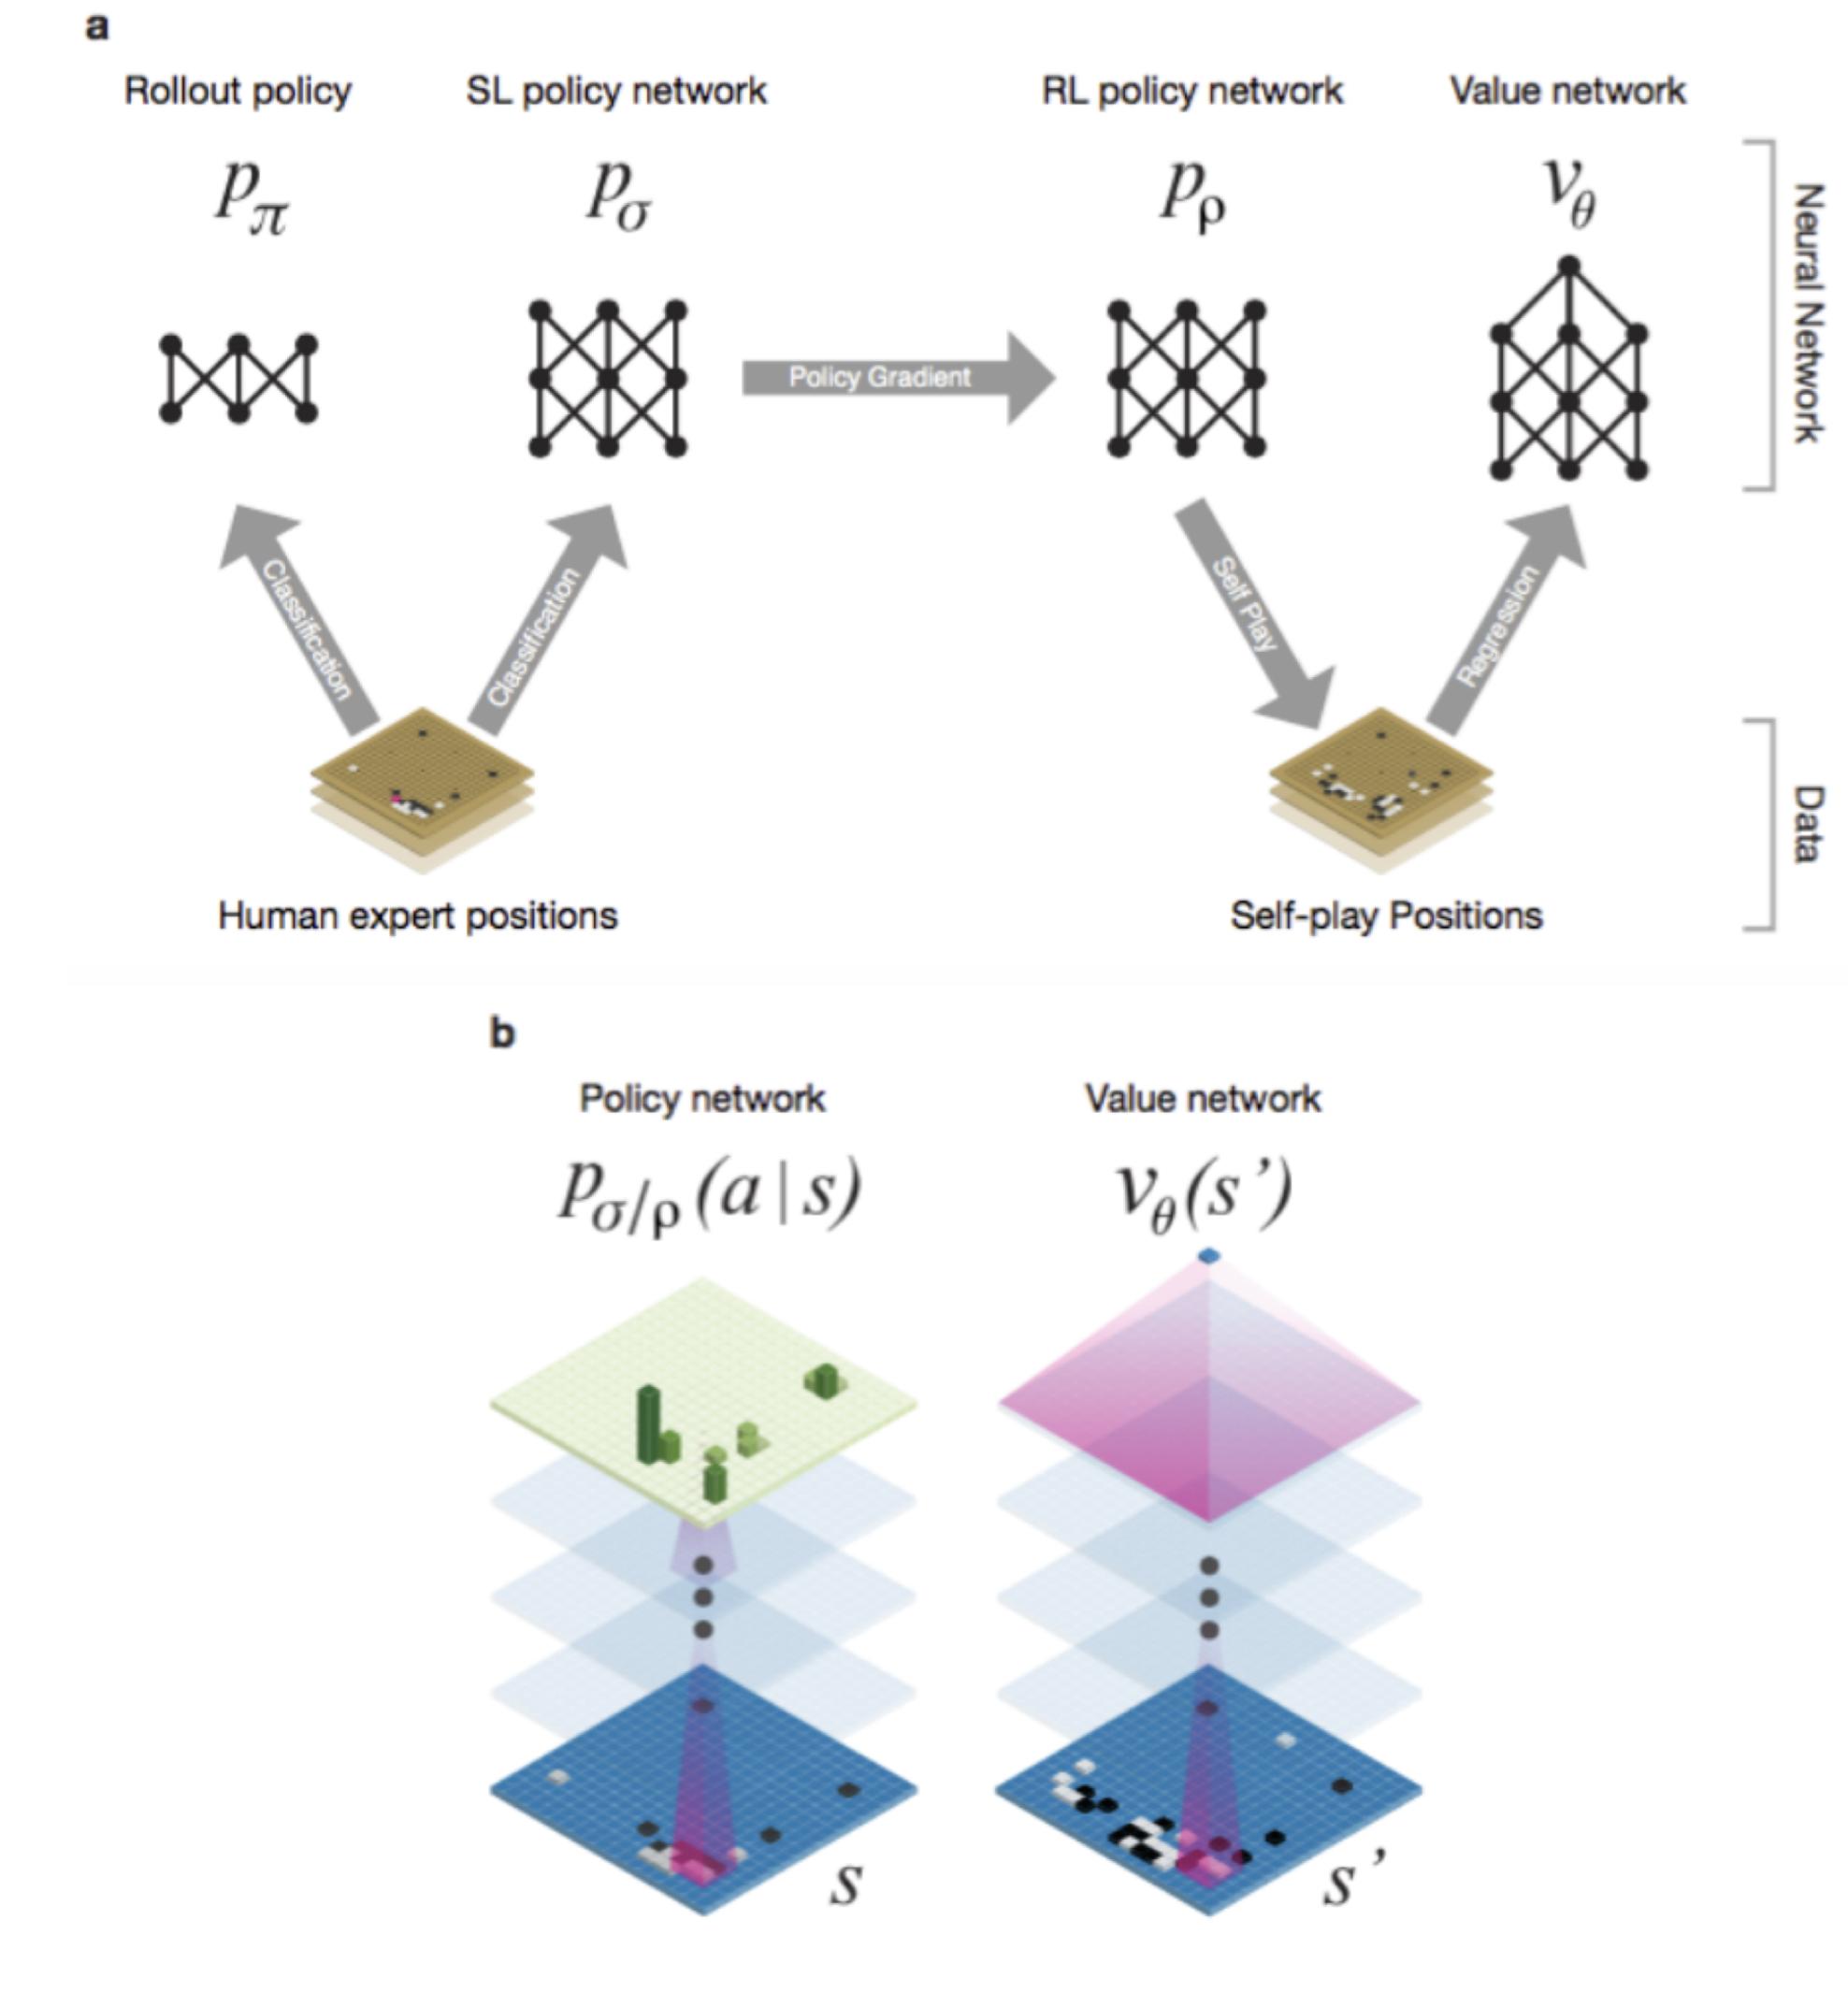
*if terminal node*  
*if MAX node*  
*if MIN node*  
*if CHANCE node*

# APPLICATION: ALPHAGO



# HOW DOES ALPHAGO WORK?

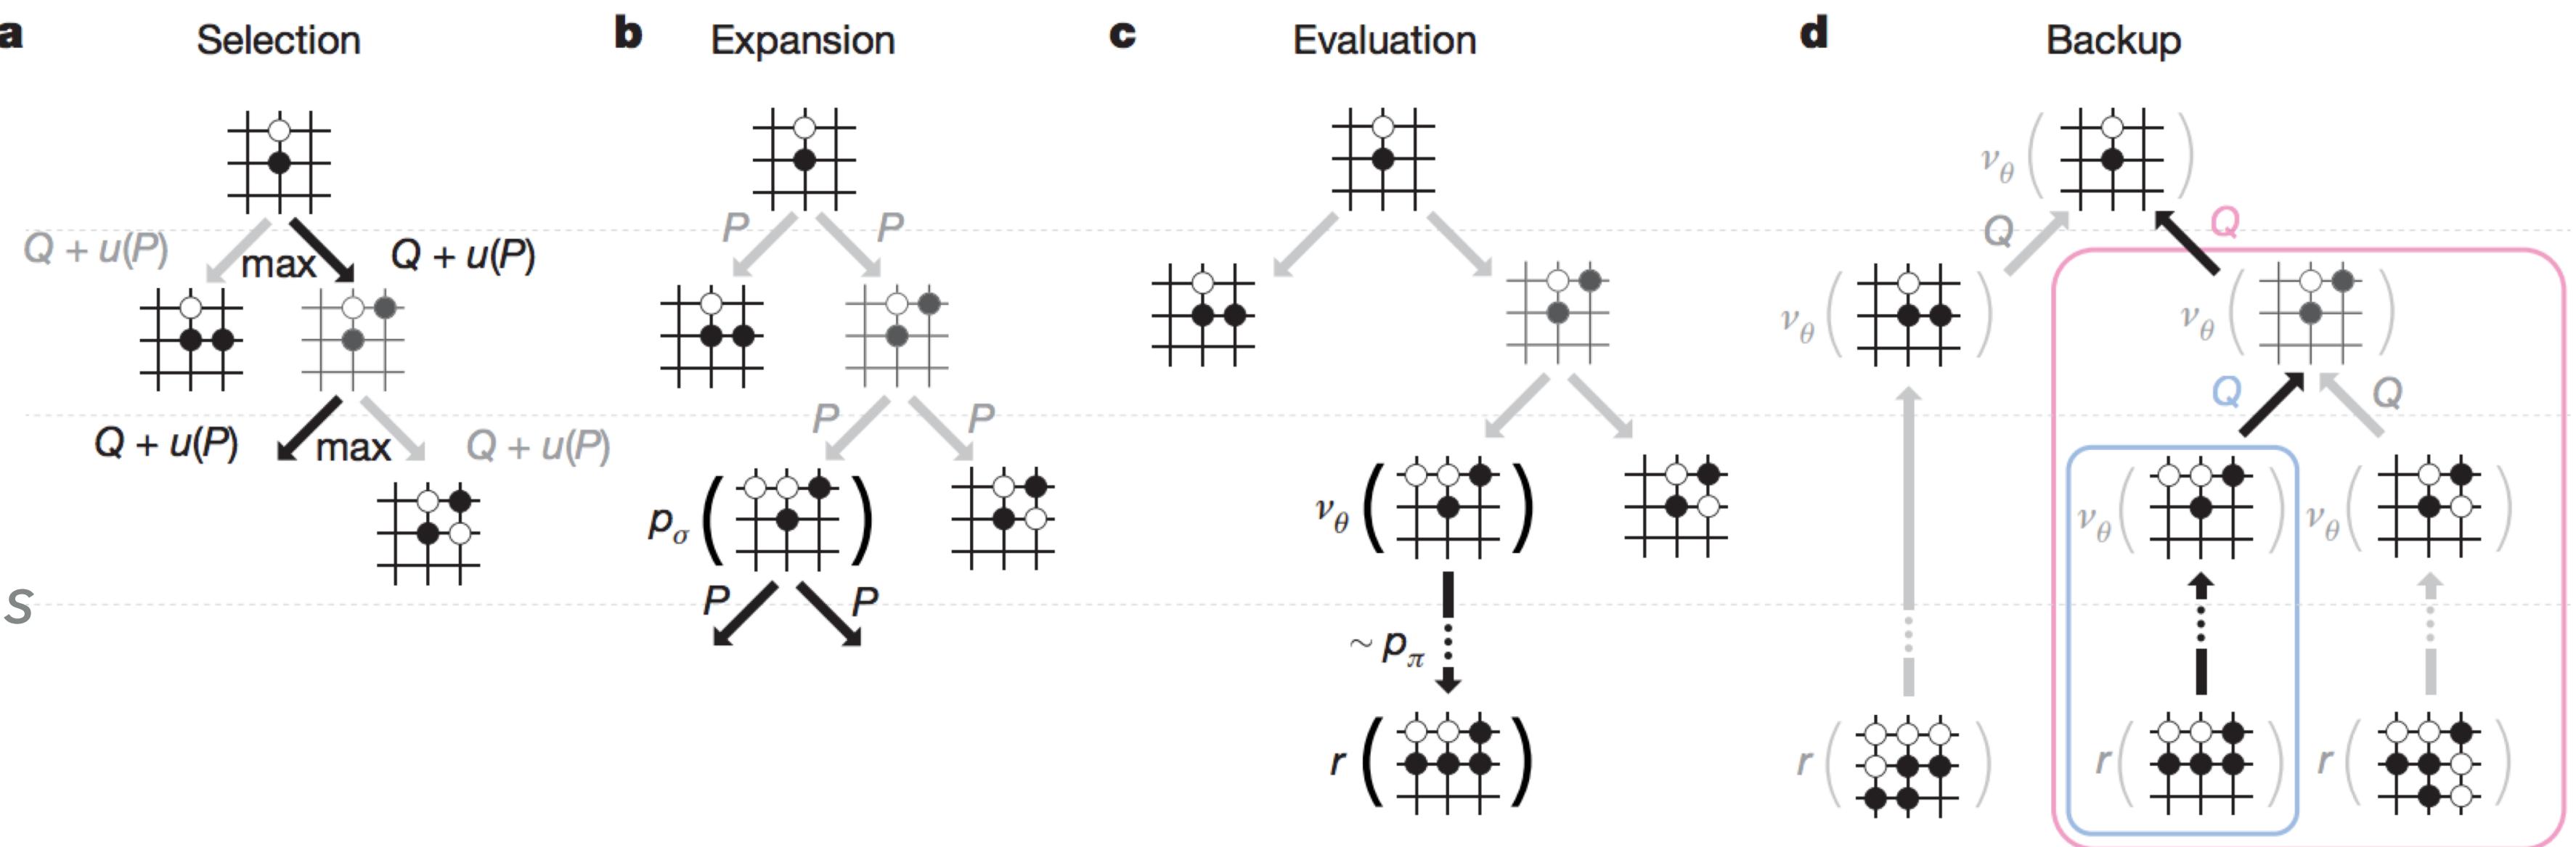
- ▶ Challenge 1: How to evaluate the goodness of a board (i.e., evaluation function of state  $s$ ) **when players play optimally?**
- ▶ Use 30 million positions from human expert games to predict  $p_\sigma(a|s)$ , i.e., the distributions of human moves at a state  $s$ , as a neural network
- ▶ Use reinforcement learning to improve the policy to maximize the chance of winning via self-playing; resulting in a better policy  $p_\rho(a|s)$
- ▶ Learn value of each state  $V_\theta(s)$  as another neural network by starting from 30 million positions and self-playing using policy  $p_\rho(a|s)$  and record the end result (i.e., win/lose)



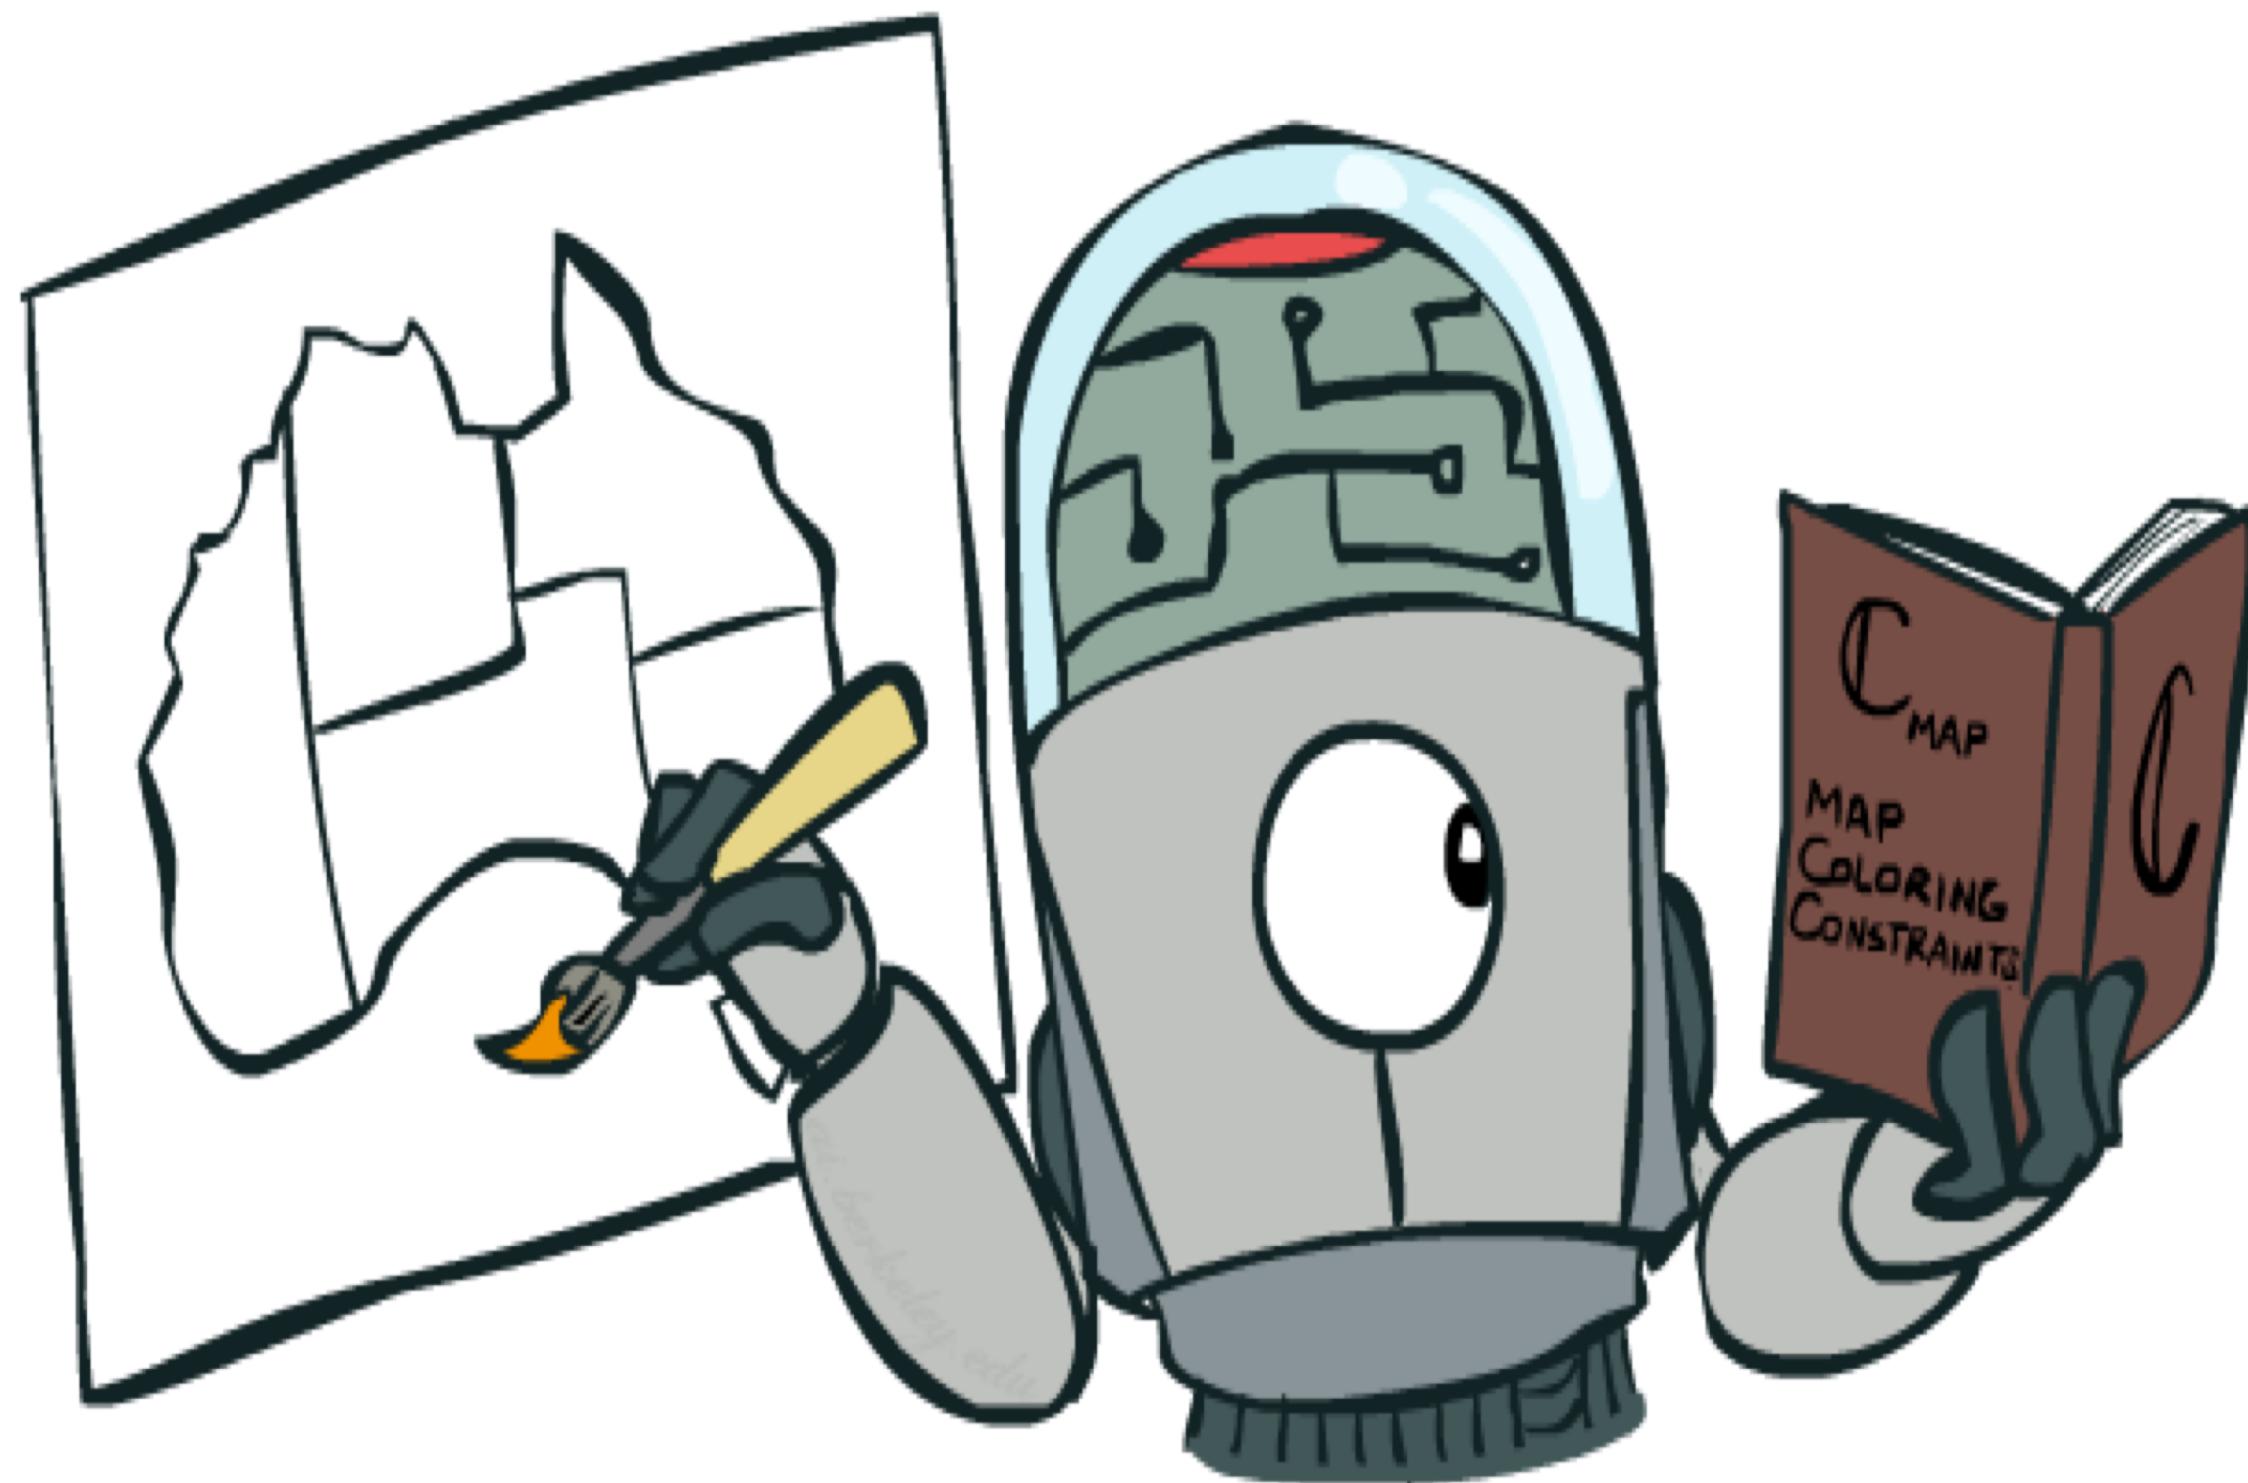
# HOW DOES ALPHAGO WORK?

- ▶ Challenge 2: The branching factor of Go is huge; even a search with several steps is computationally expensive
- ▶ Use Monte Carlo Tree Search (MCTS) to only focus on the “promising” part!

$u(s,a)$  is initialized with  $p_\sigma(a|s)$  and decrease with  $N(s,a)$ ;  
 $Q(s,a)$  reflects the estimated value of taking action  $a$  at state  $s$



# CONSTRAINT SATISFACTION PROBLEMS



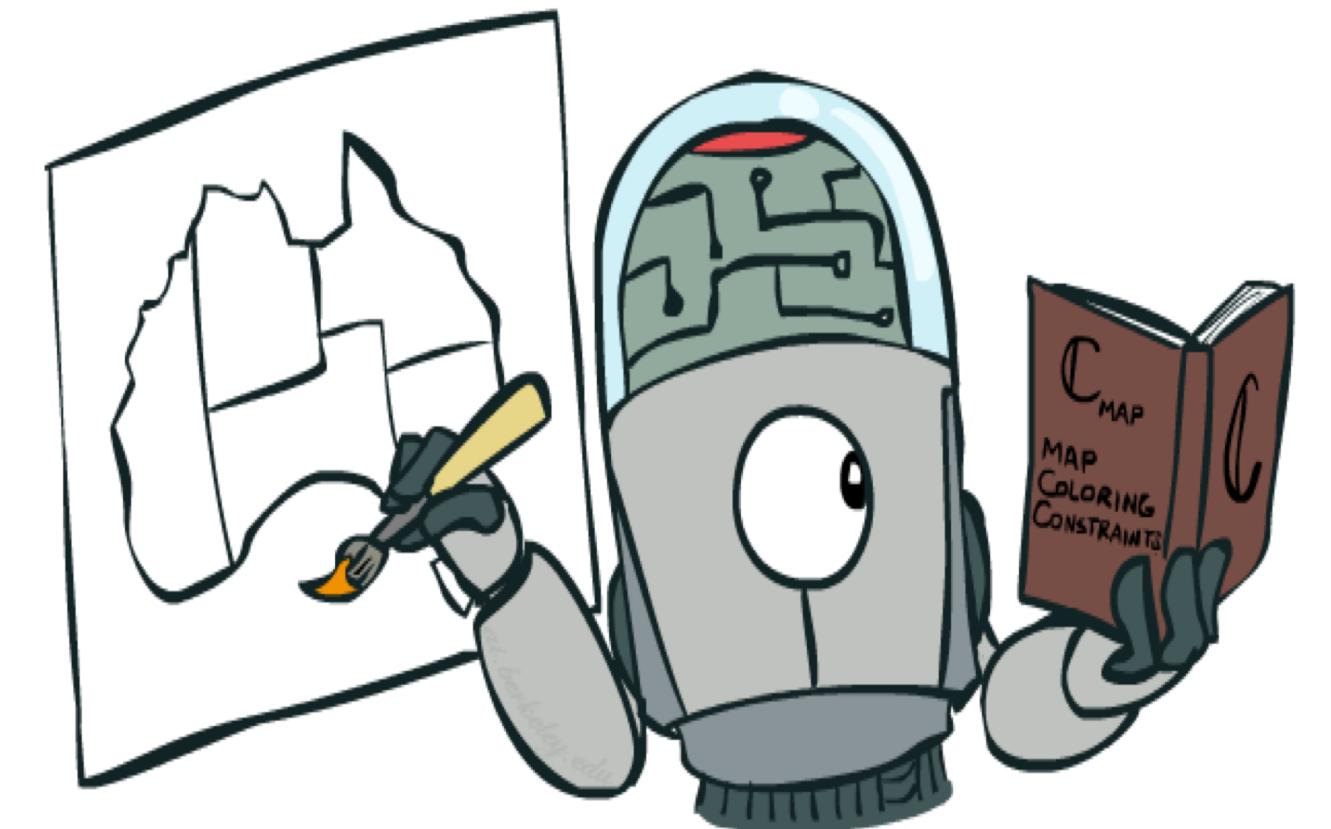
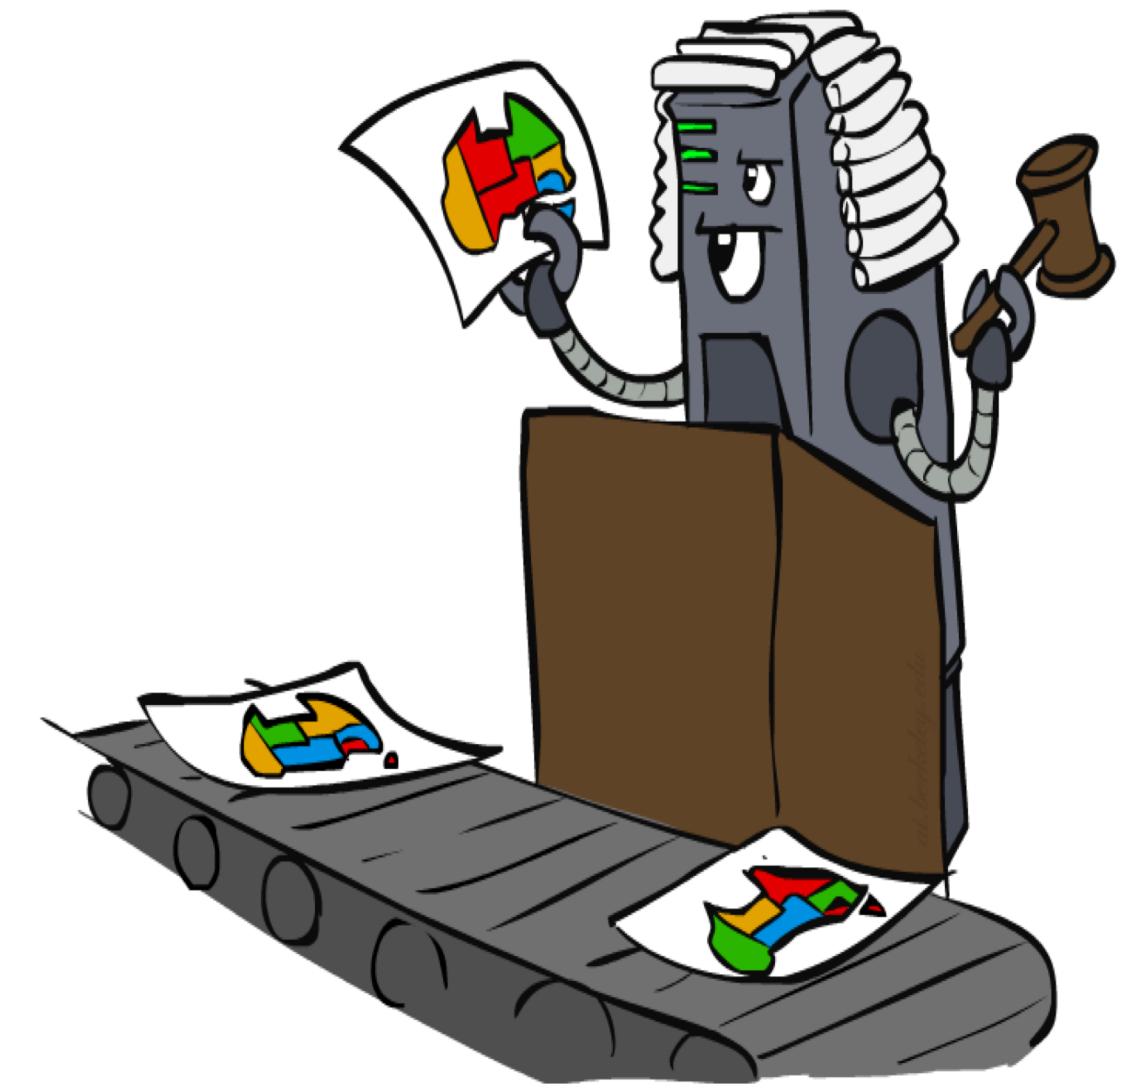
# WHAT IS SEARCH FOR?

- ▶ **Search/Planning:** sequences of actions
  - ▶ The path to the goal is the important thing
  - ▶ Paths have various costs, depths
  - ▶ Heuristics give problem-specific guidance
  
- ▶ **Identification:** assignments to variables
  - ▶ The goal itself is important, not the path
  - ▶ All paths at the same depth (for some formulations)

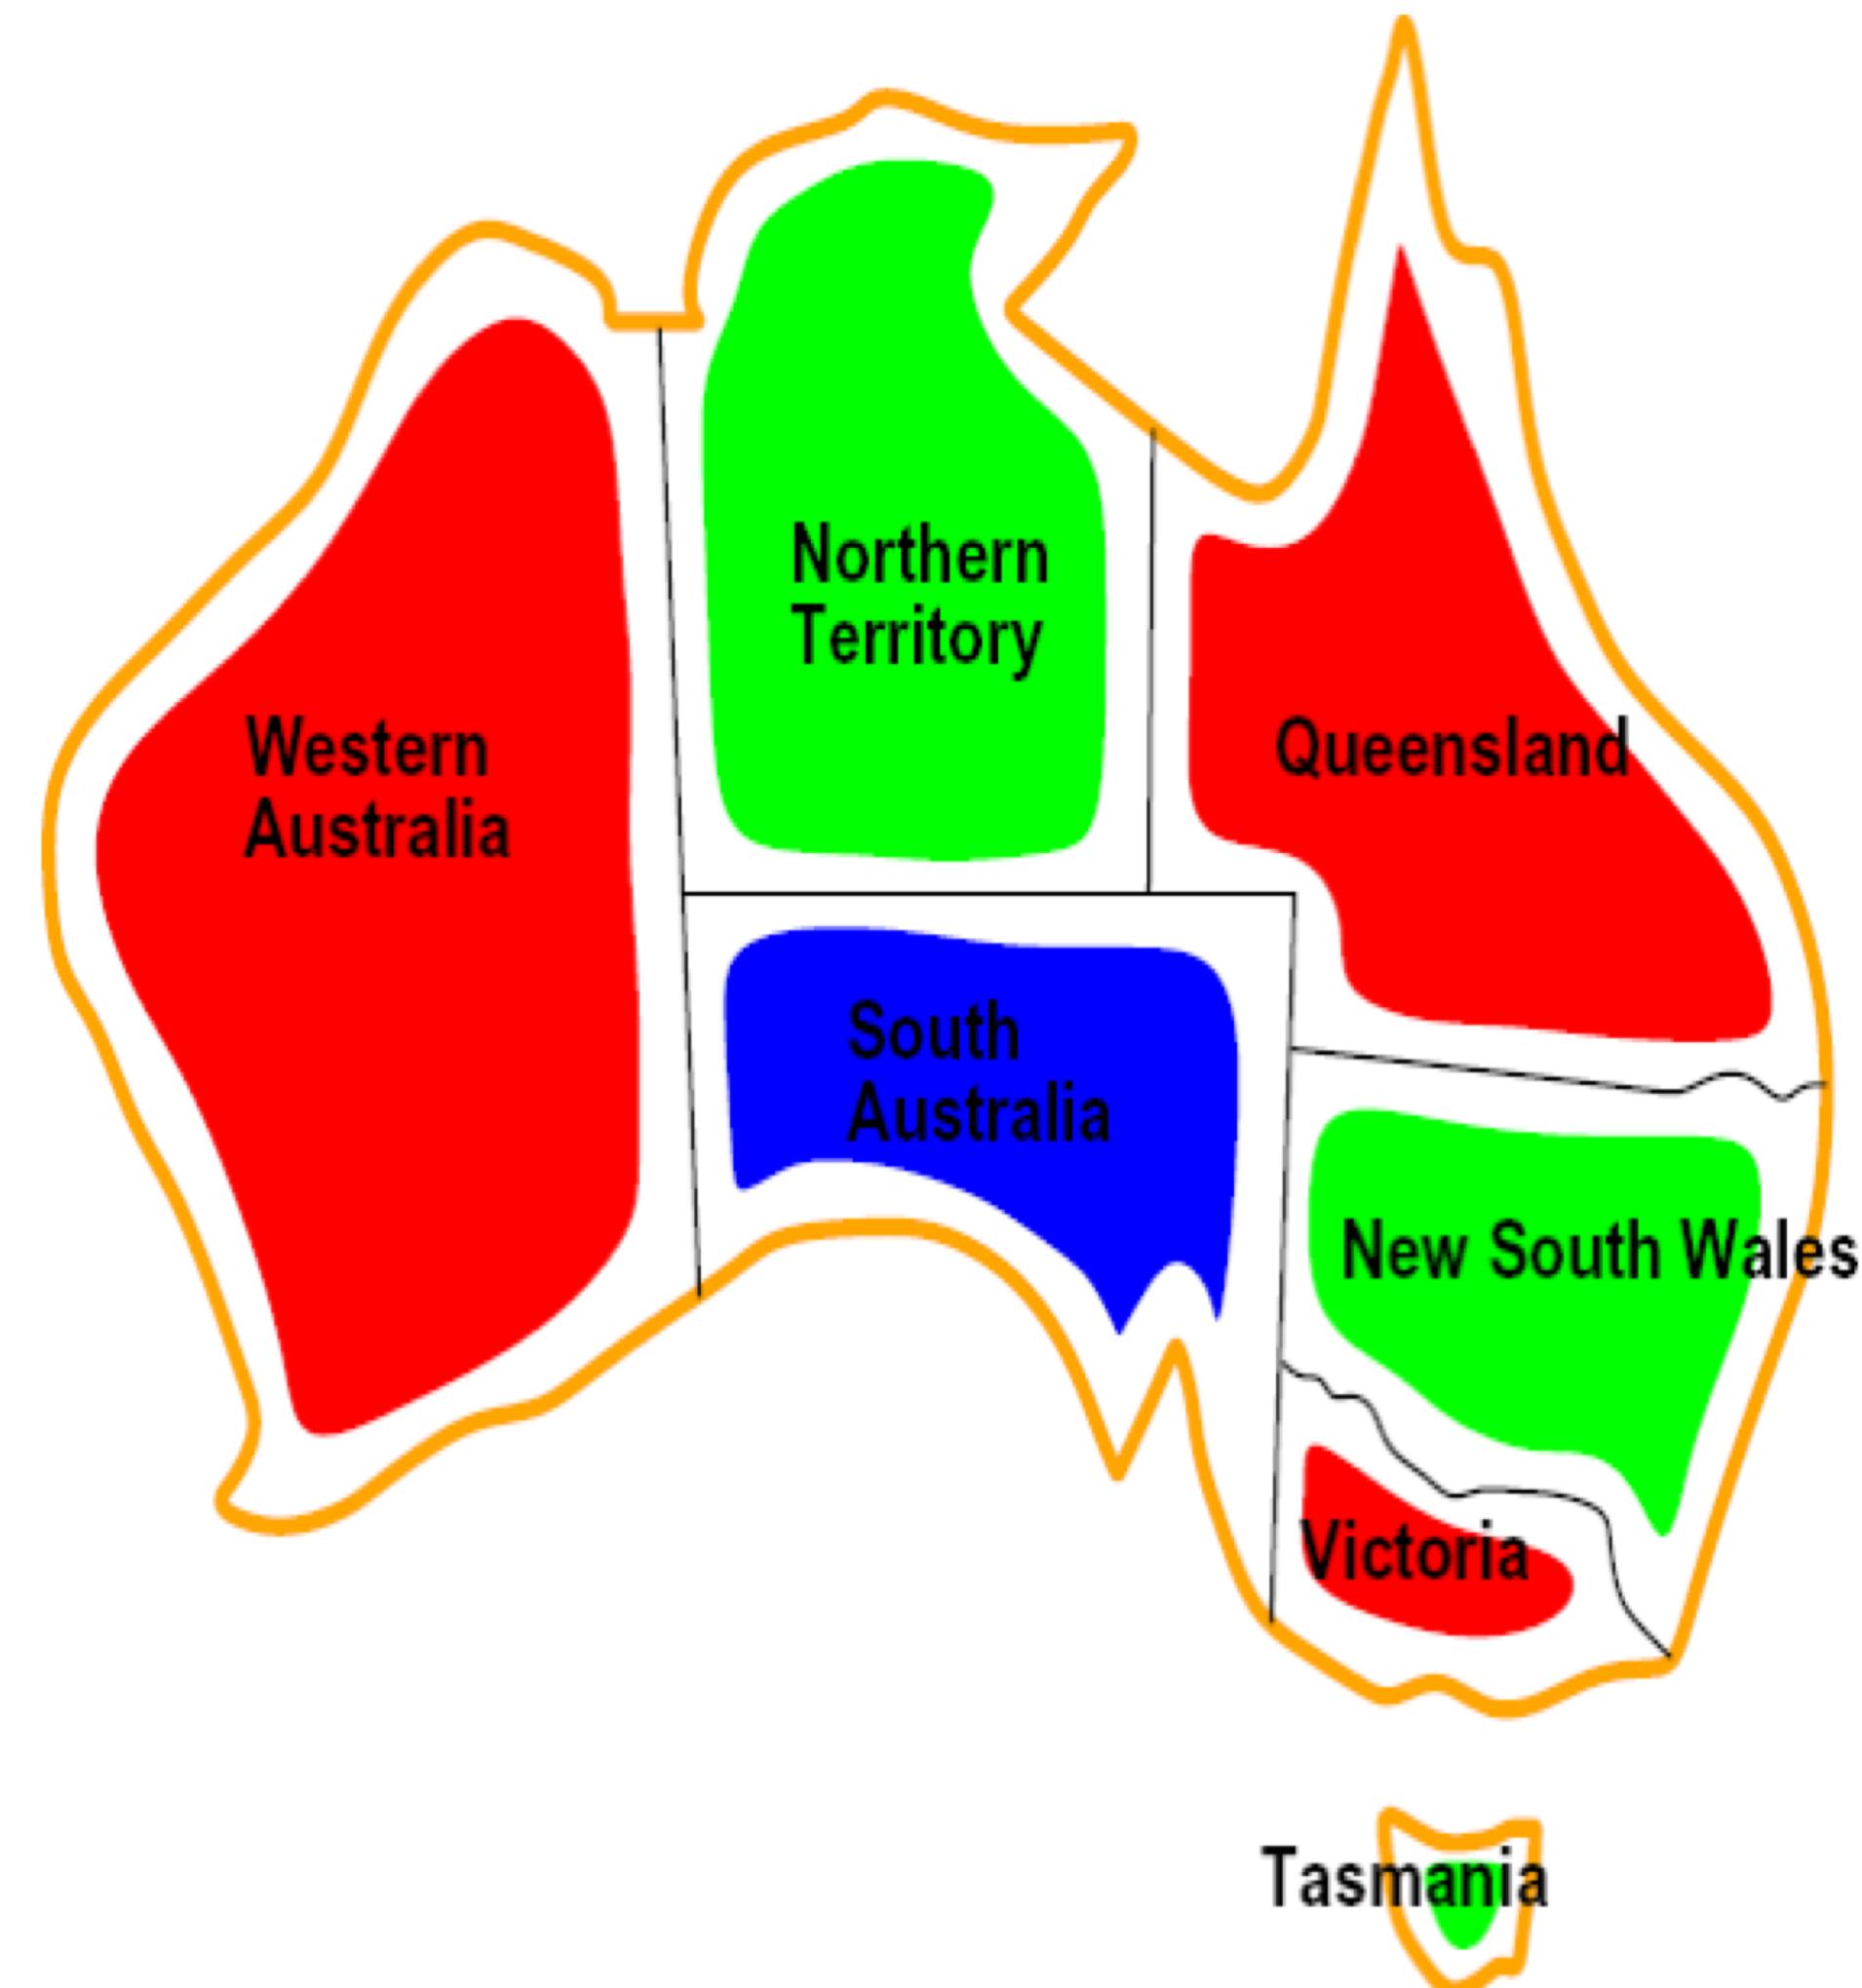


# CONSTRAINT SATISFACTION PROBLEMS

- ▶ Standard search problems:
  - ▶ State is a “black box”: arbitrary data structure
  - ▶ Goal test can be any function over states
- ▶ Constraint satisfaction problems (CSPs) – a special subset of search problems
  - ▶ State is defined by **variables  $X_i$**  with values from a **domain D**
  - ▶ Goal test is a set of constraints specifying allowable combinations of values for subsets of variables
  - ▶ Allows useful general-purpose algorithms with more power than standard search algorithms



## CSP EXAMPLES



## EXAMPLE: MAP COLORING

- ▶ Variables: WA, NT, Q, NSW, V, SA, T
- ▶ Domains:  $D = \{\text{red, green, blue}\}$
- ▶ Constraints: adjacent regions must have different colors

Implicit:  $WA \neq NT$

Explicit:  $(WA, NT) \in \{(red, green), (red, blue), \dots\}$

- ▶ Solutions are assignments satisfying all constraints, e.g.:
- $\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green}, V=\text{red}, SA=\text{blue}, T=\text{green}\}$

